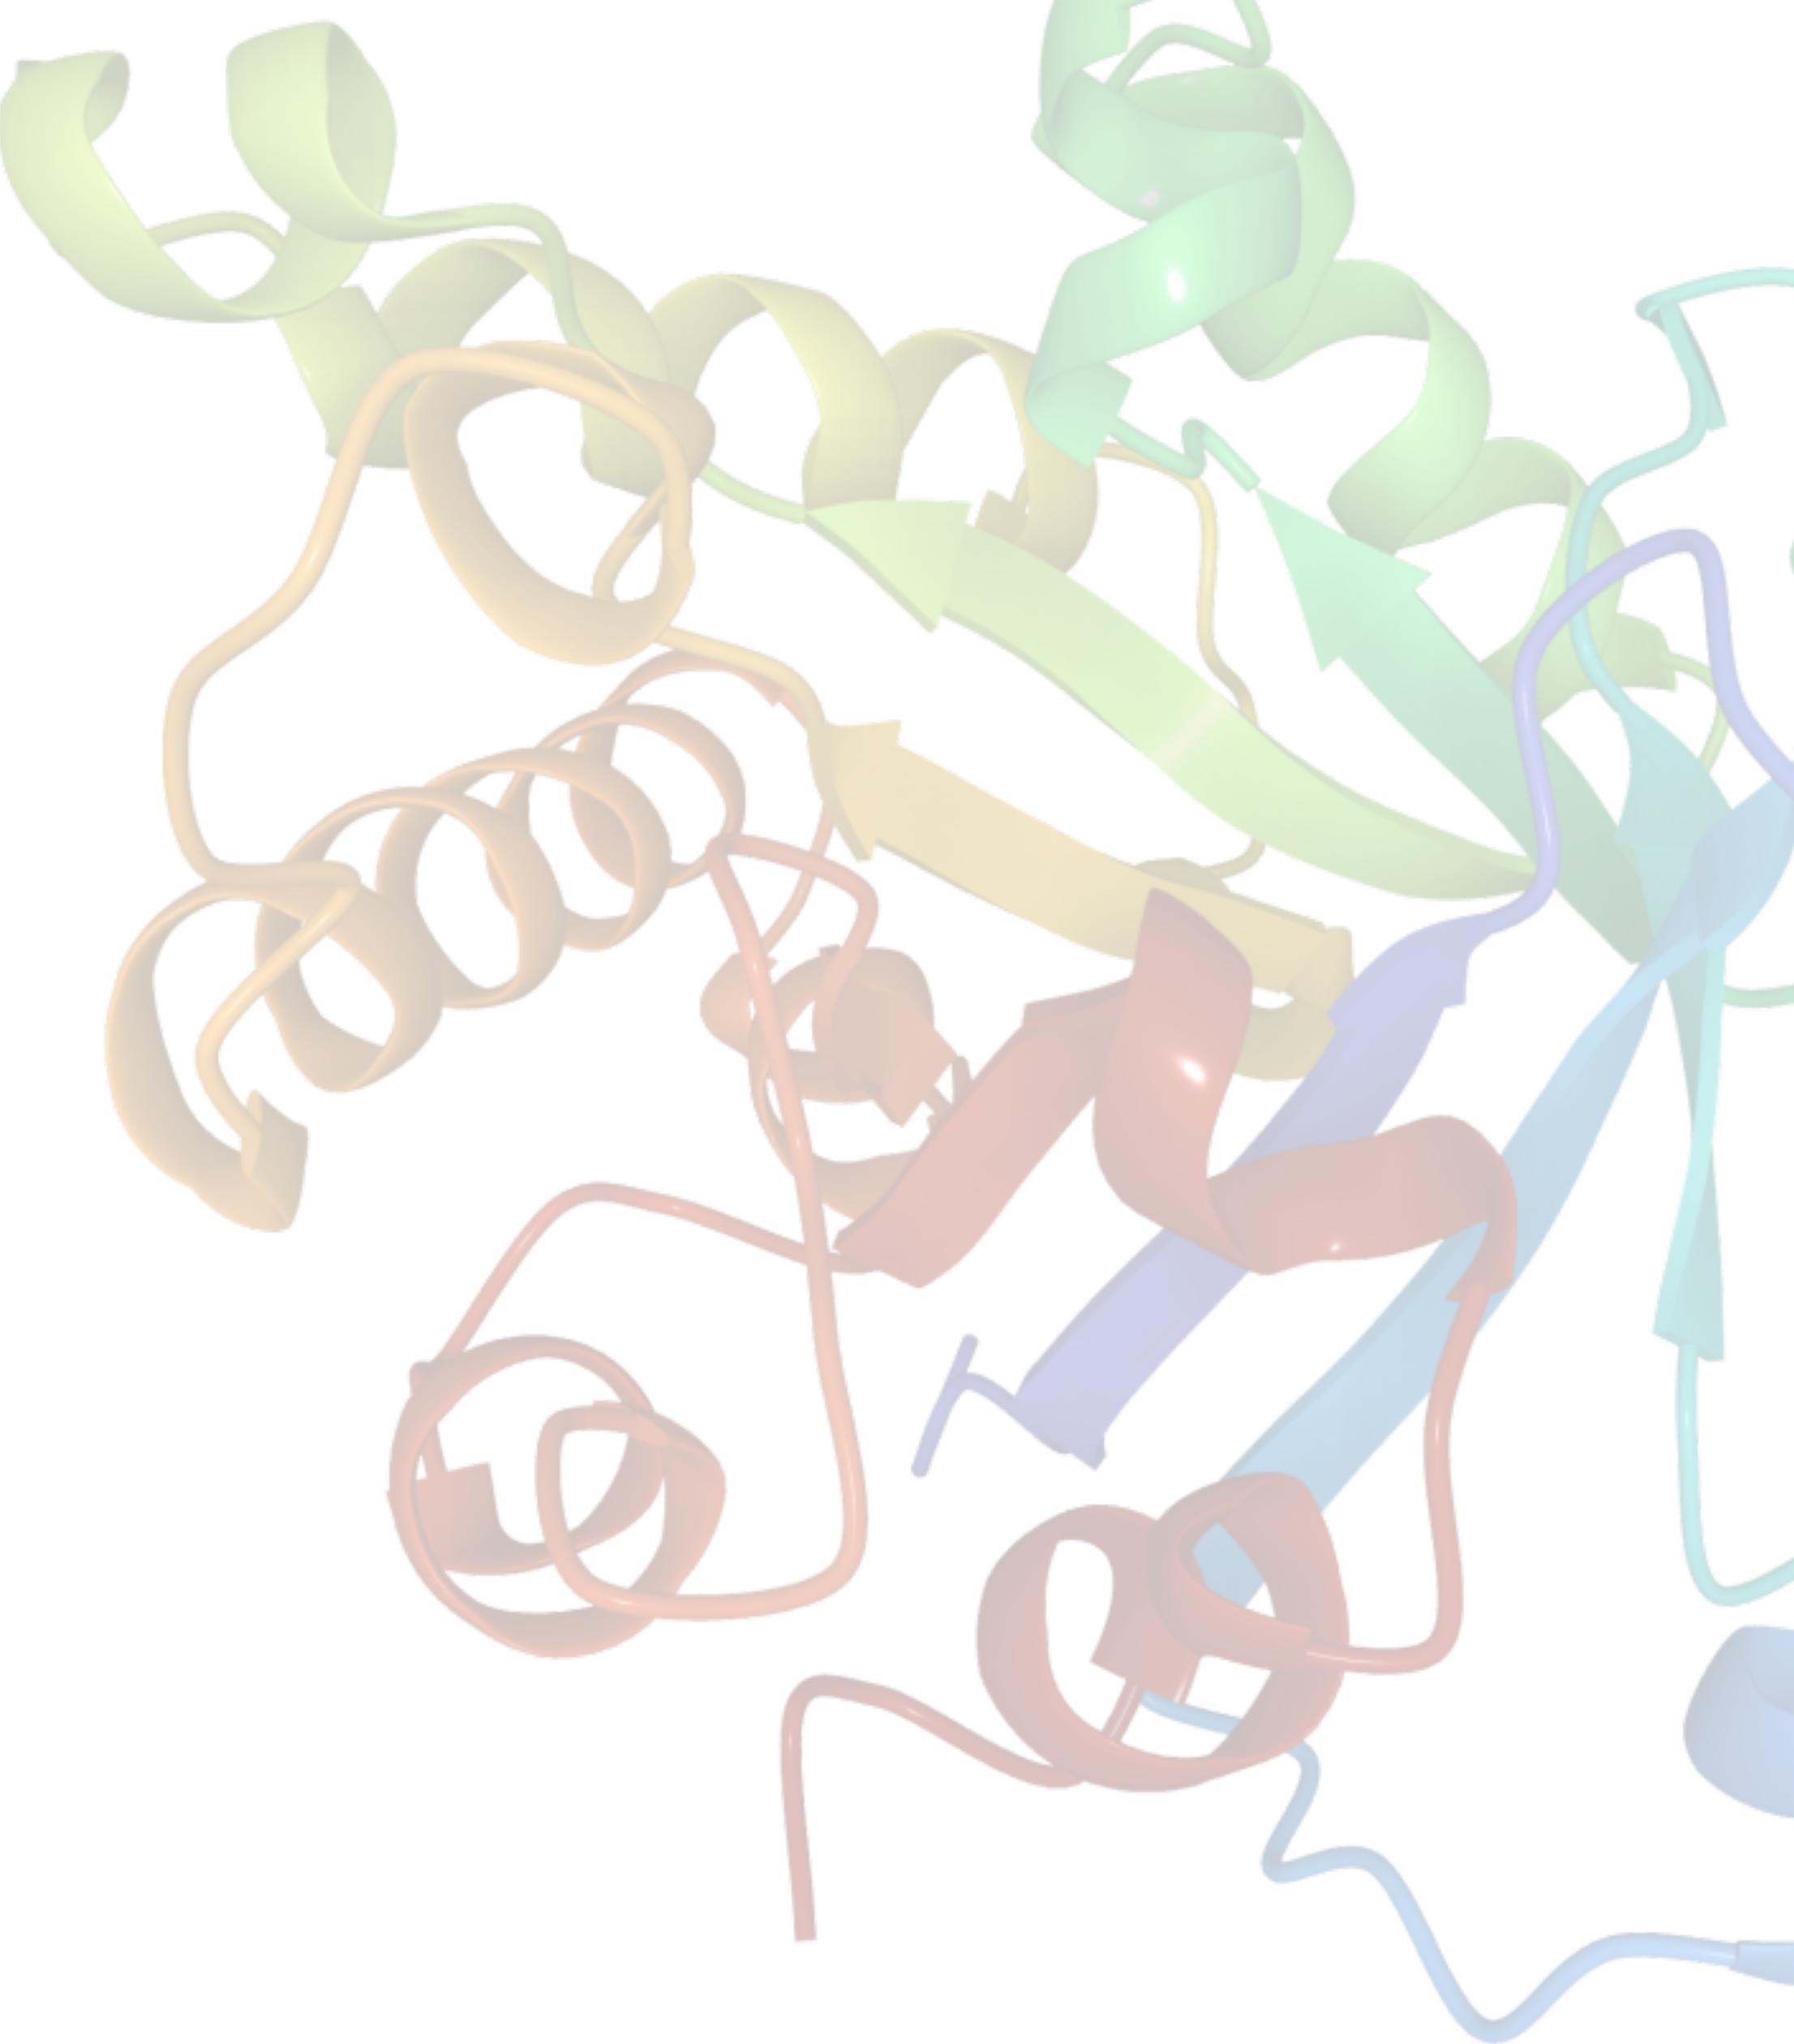


Flow Matching

Alex Tong

Probabilistic AI Summer School 2025

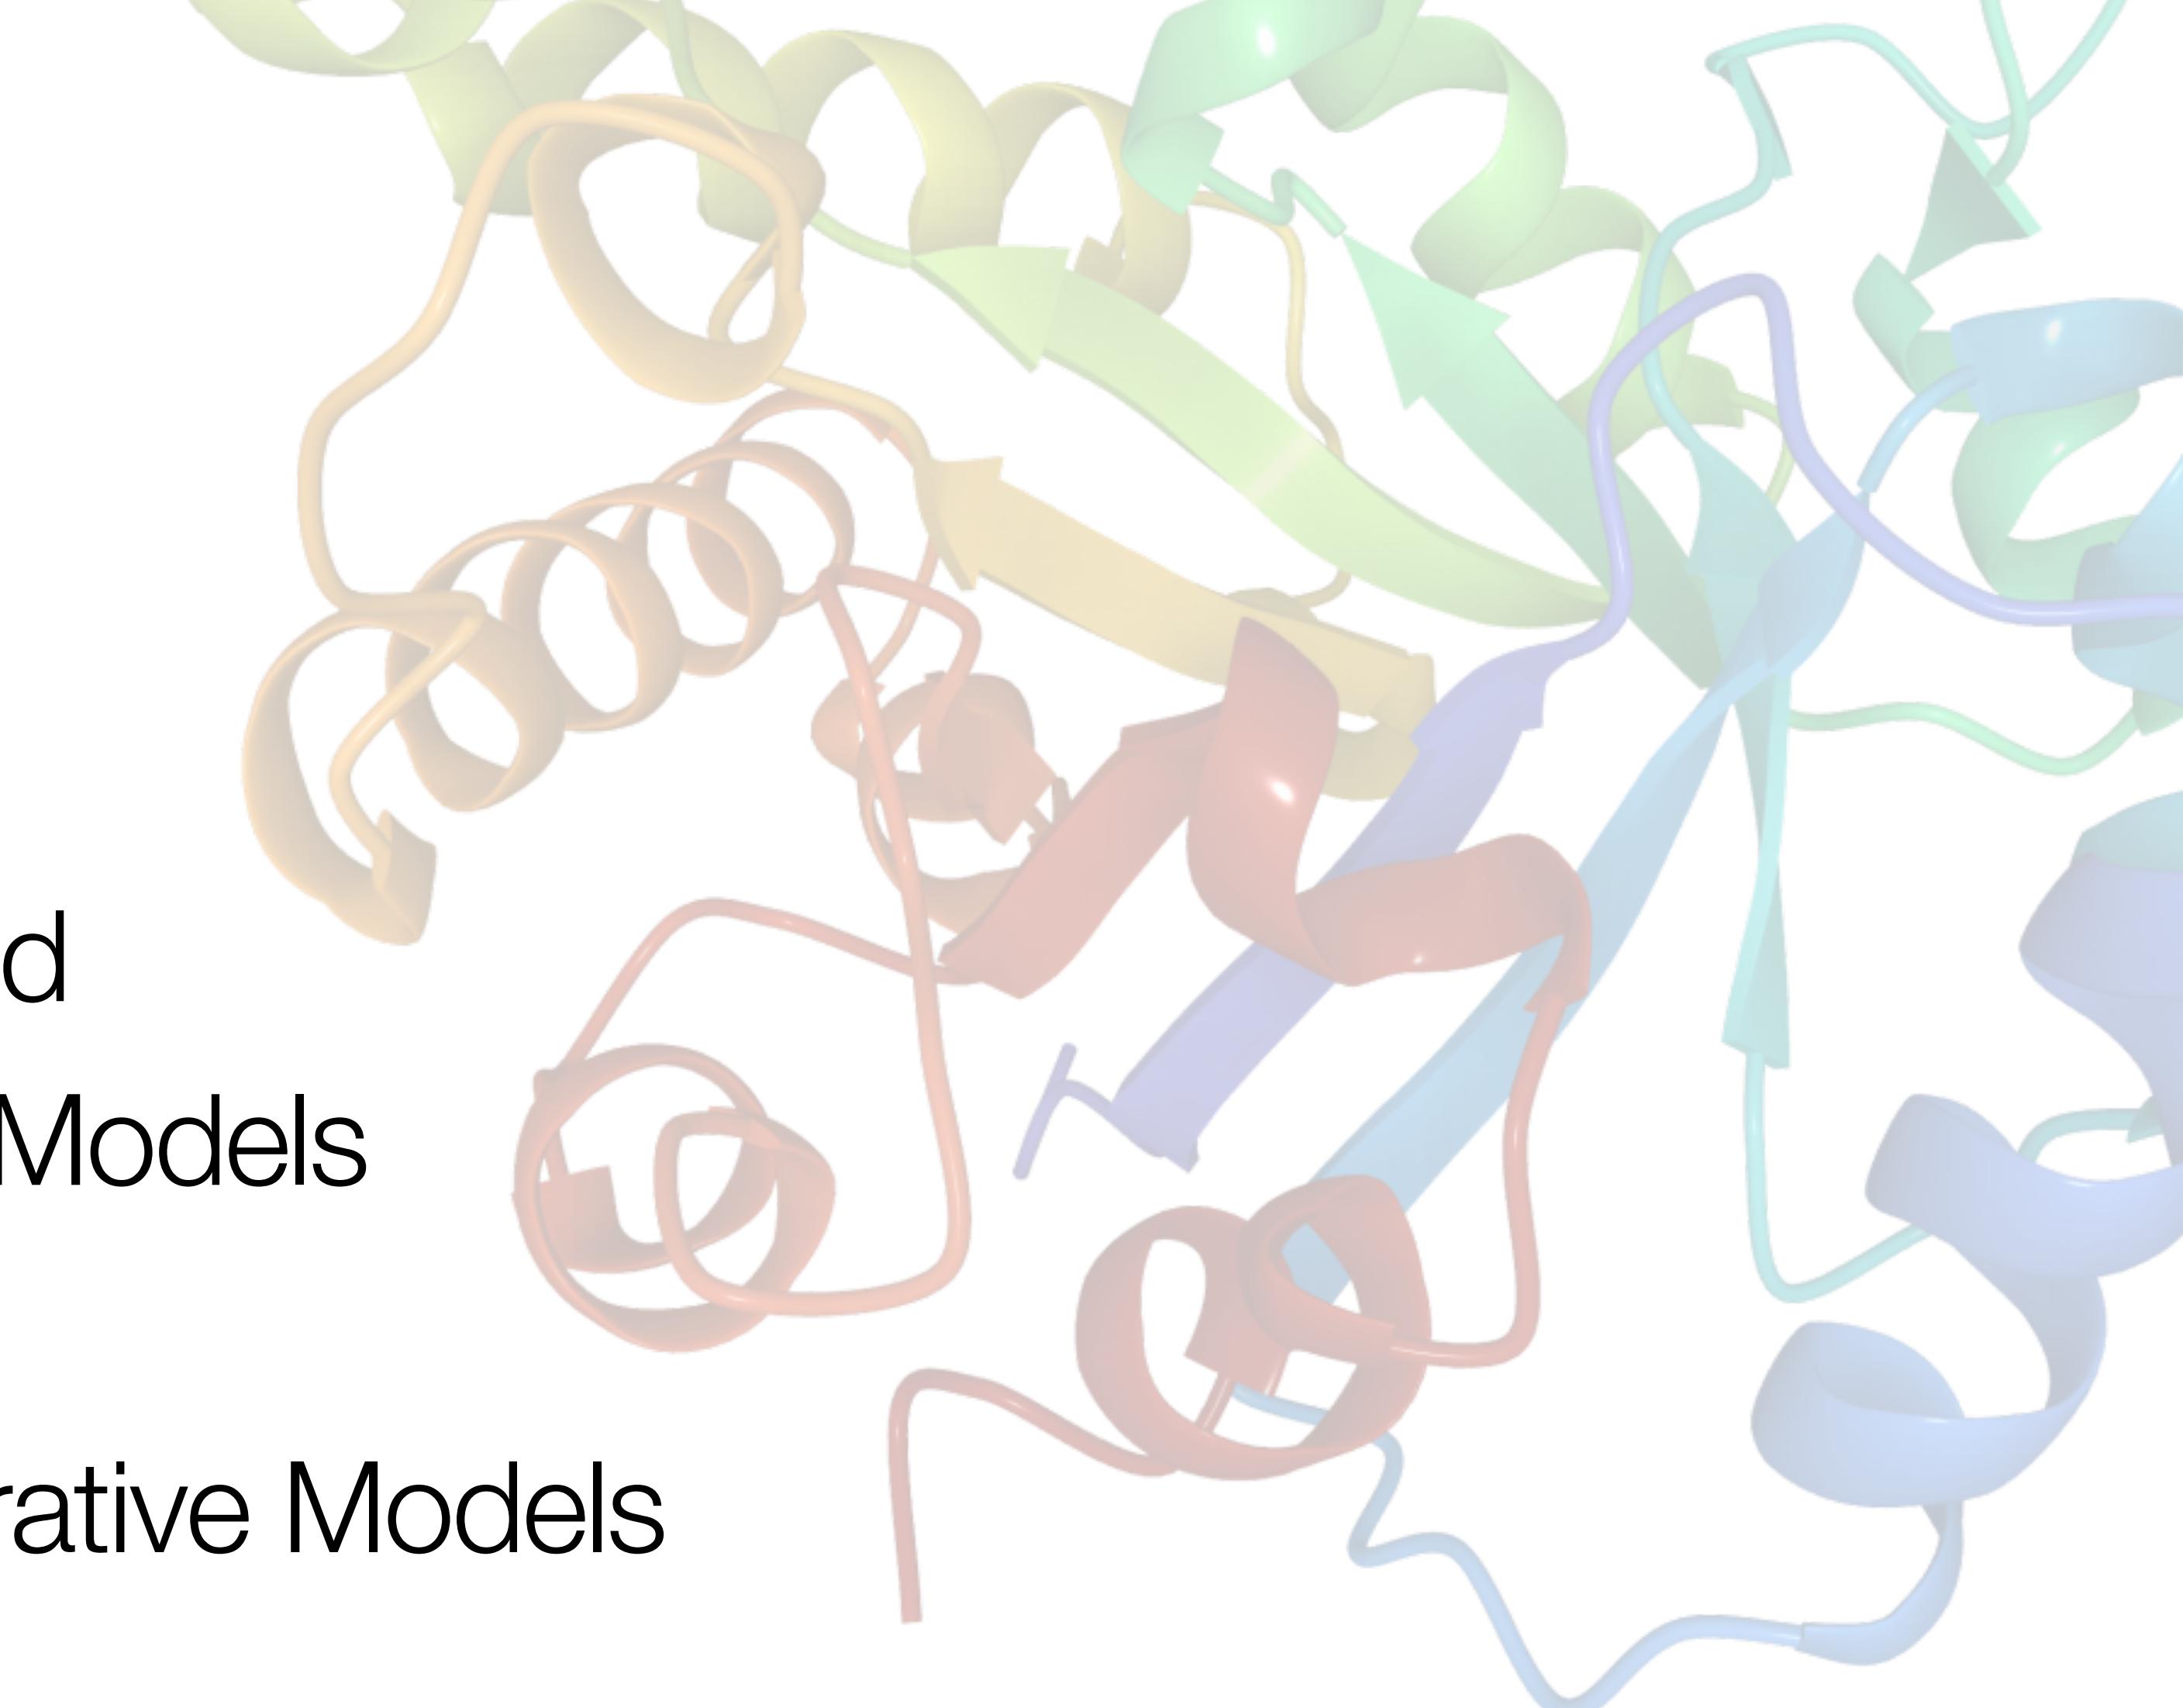


Outline:

Part I: Dynamical Systems and
Simulation-based Generative Models

Part II: Simulation-Free Generative Models

Part III: Extensions

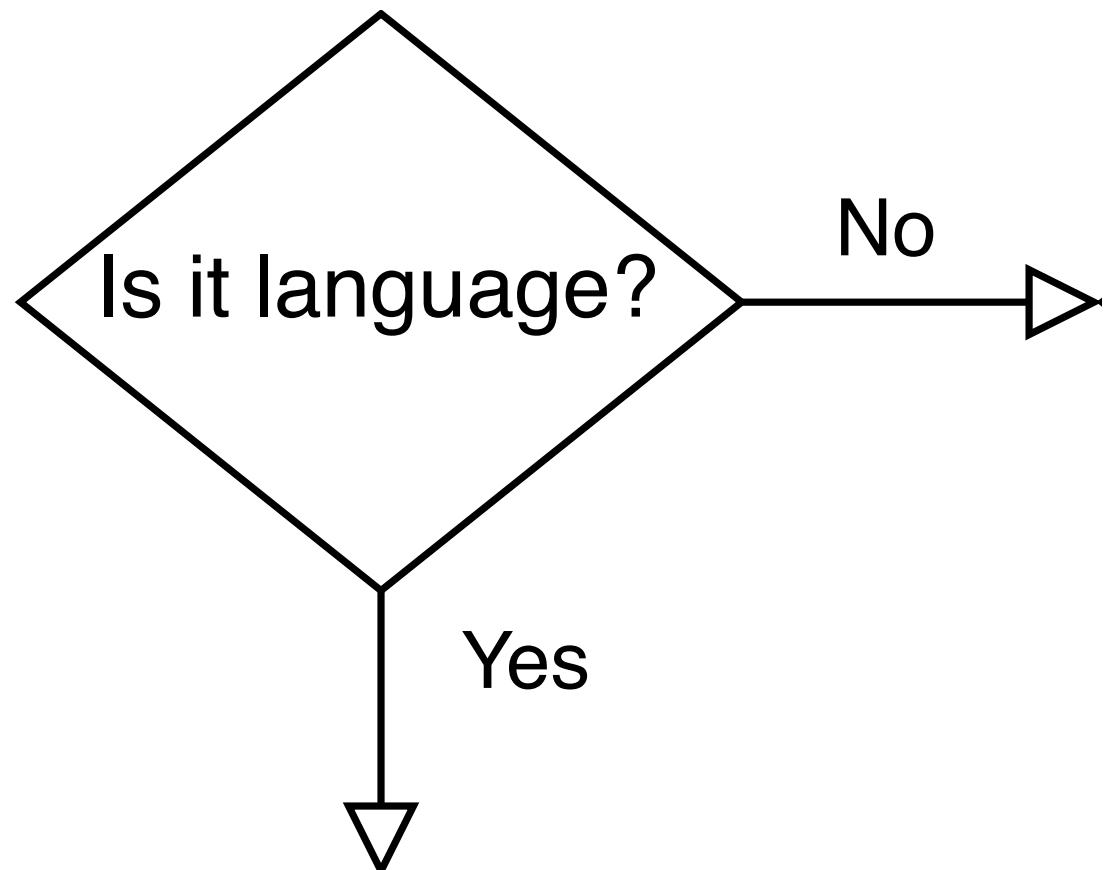


What do deep learning people mean by generative models?

Models trained to generate samples from the data

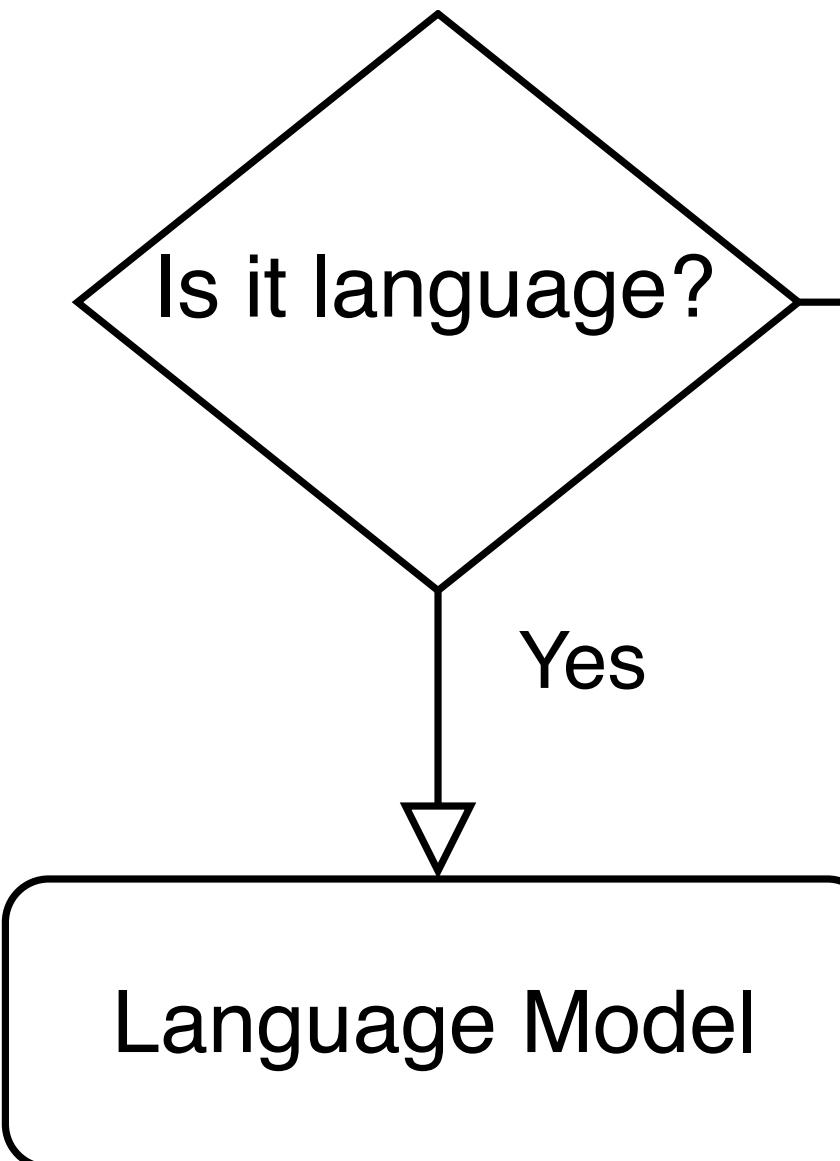
What do deep learning people mean by generative models?

Models trained to generate samples from the data



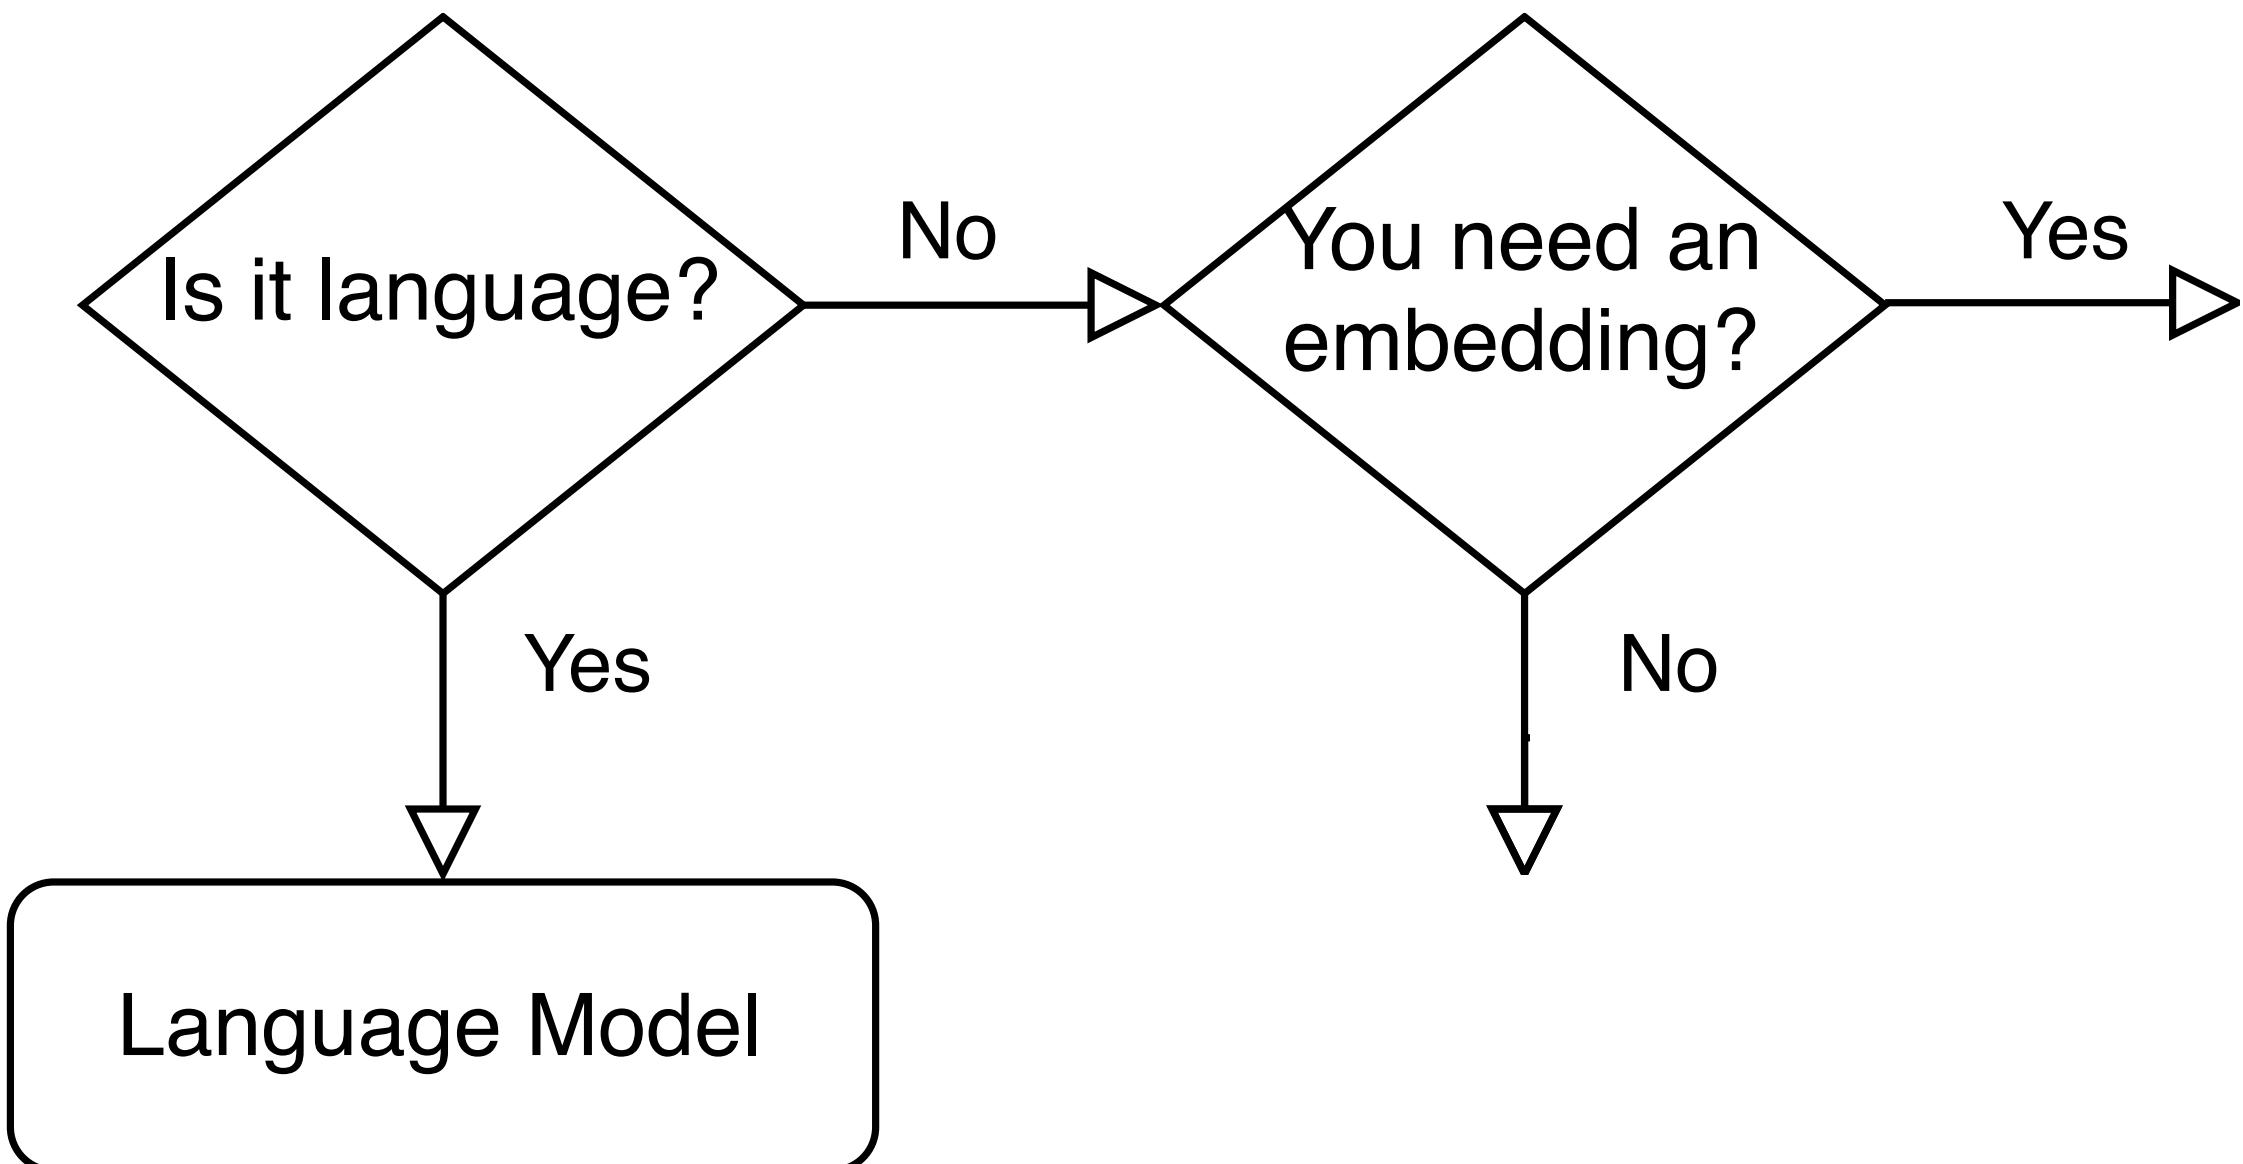
What do deep learning people mean by generative models?

Models trained to generate samples from the data



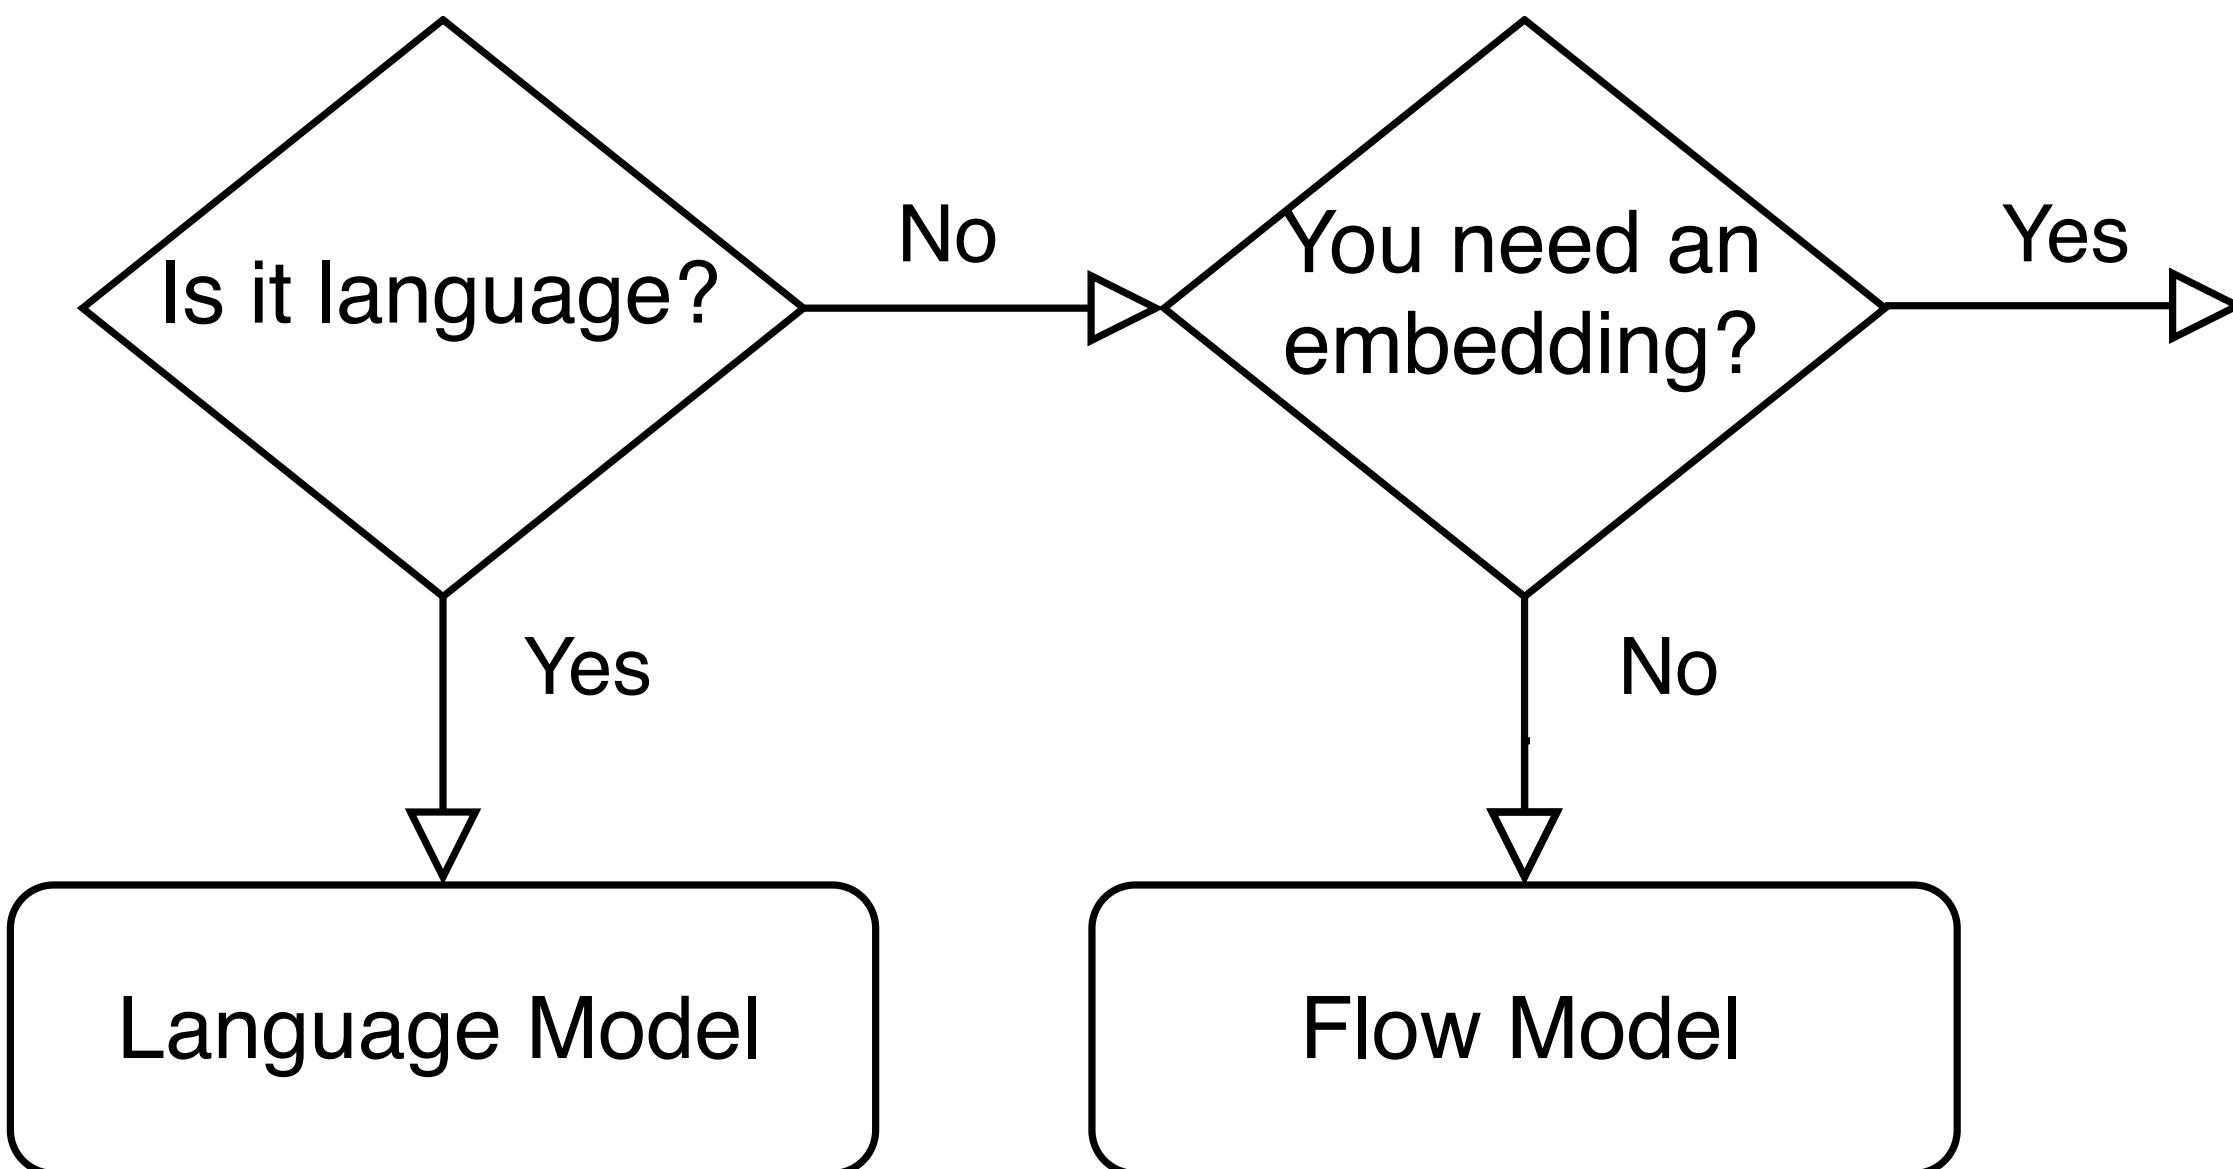
What do deep learning people mean by generative models?

Models trained to generate samples from the data



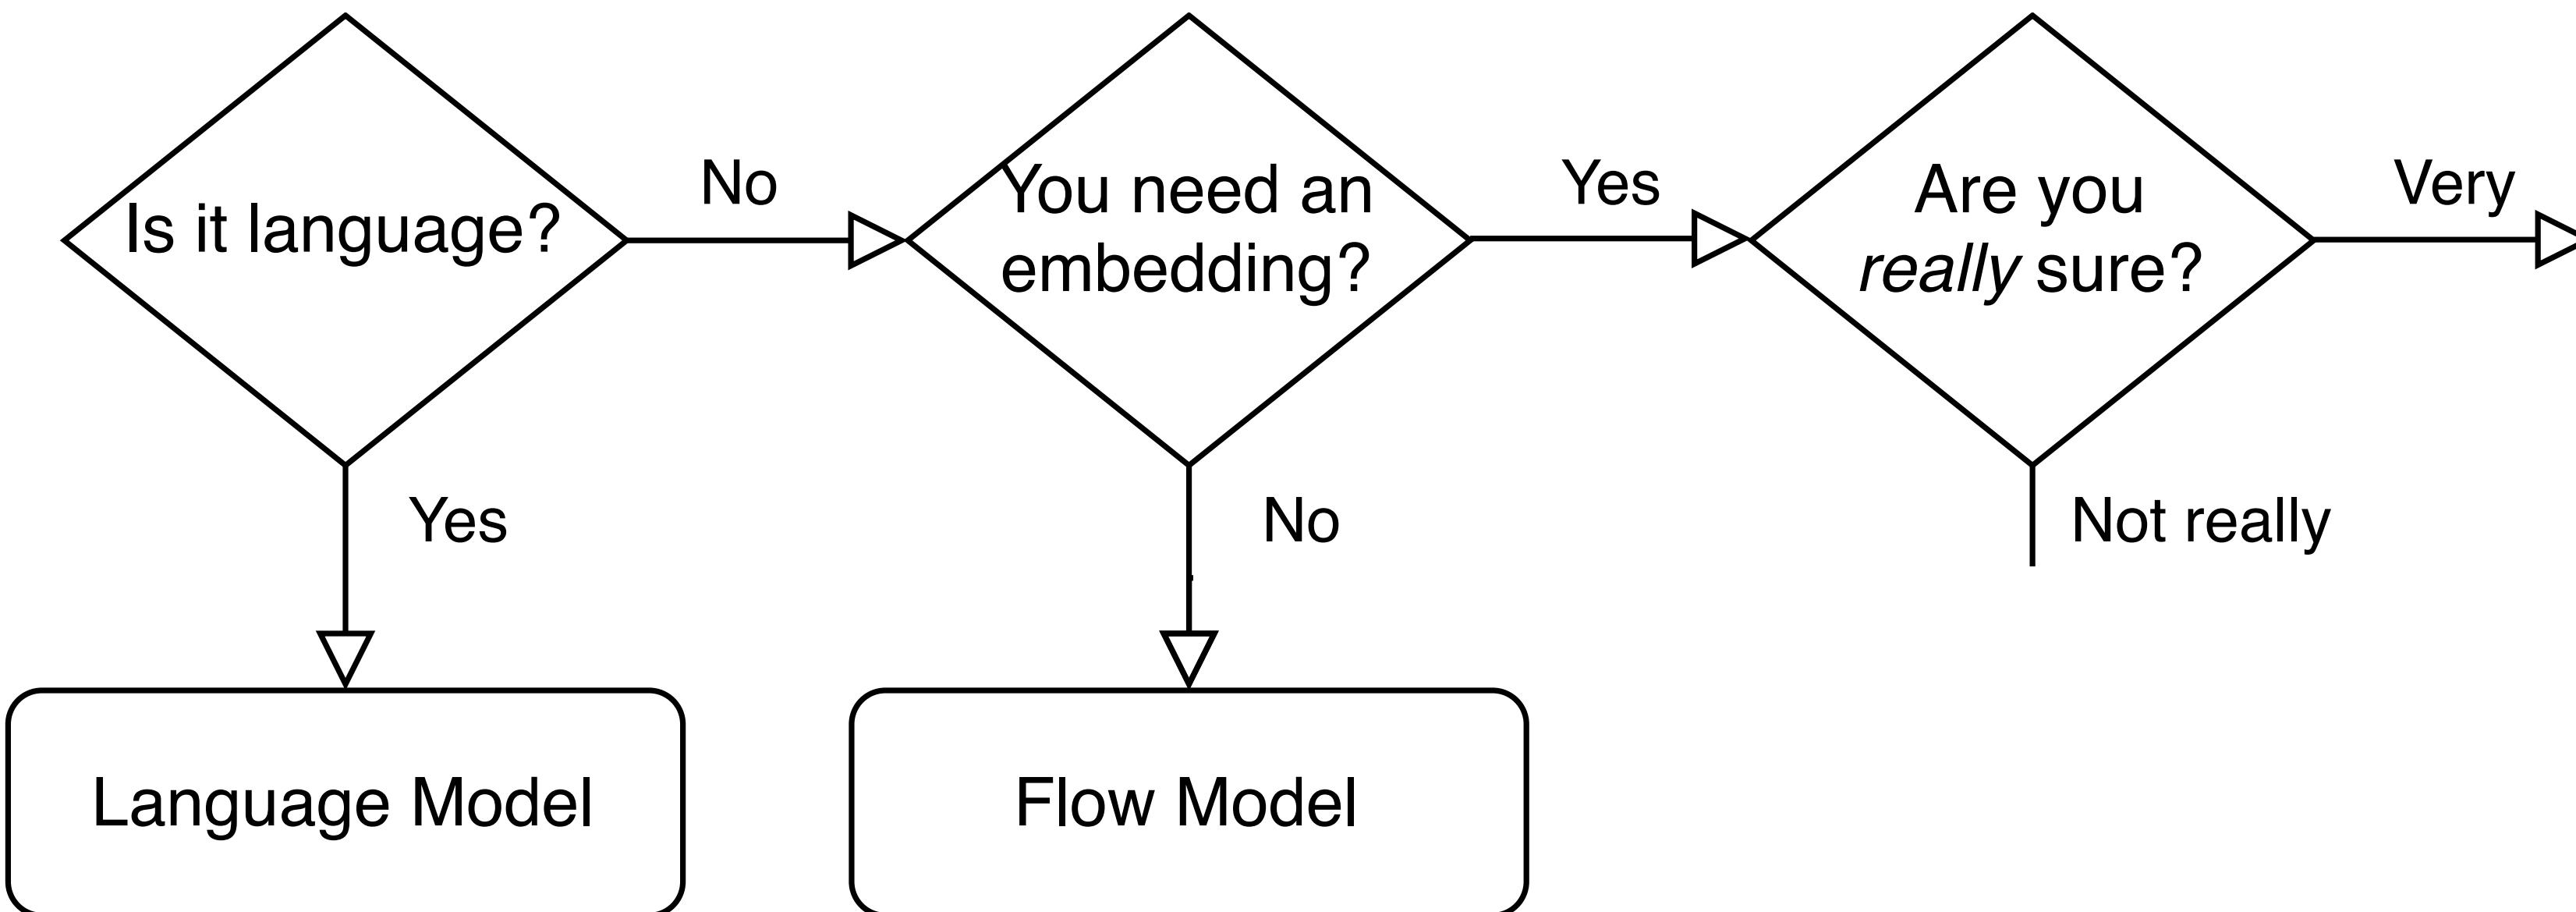
What do deep learning people mean by generative models?

Models trained to generate samples from the data



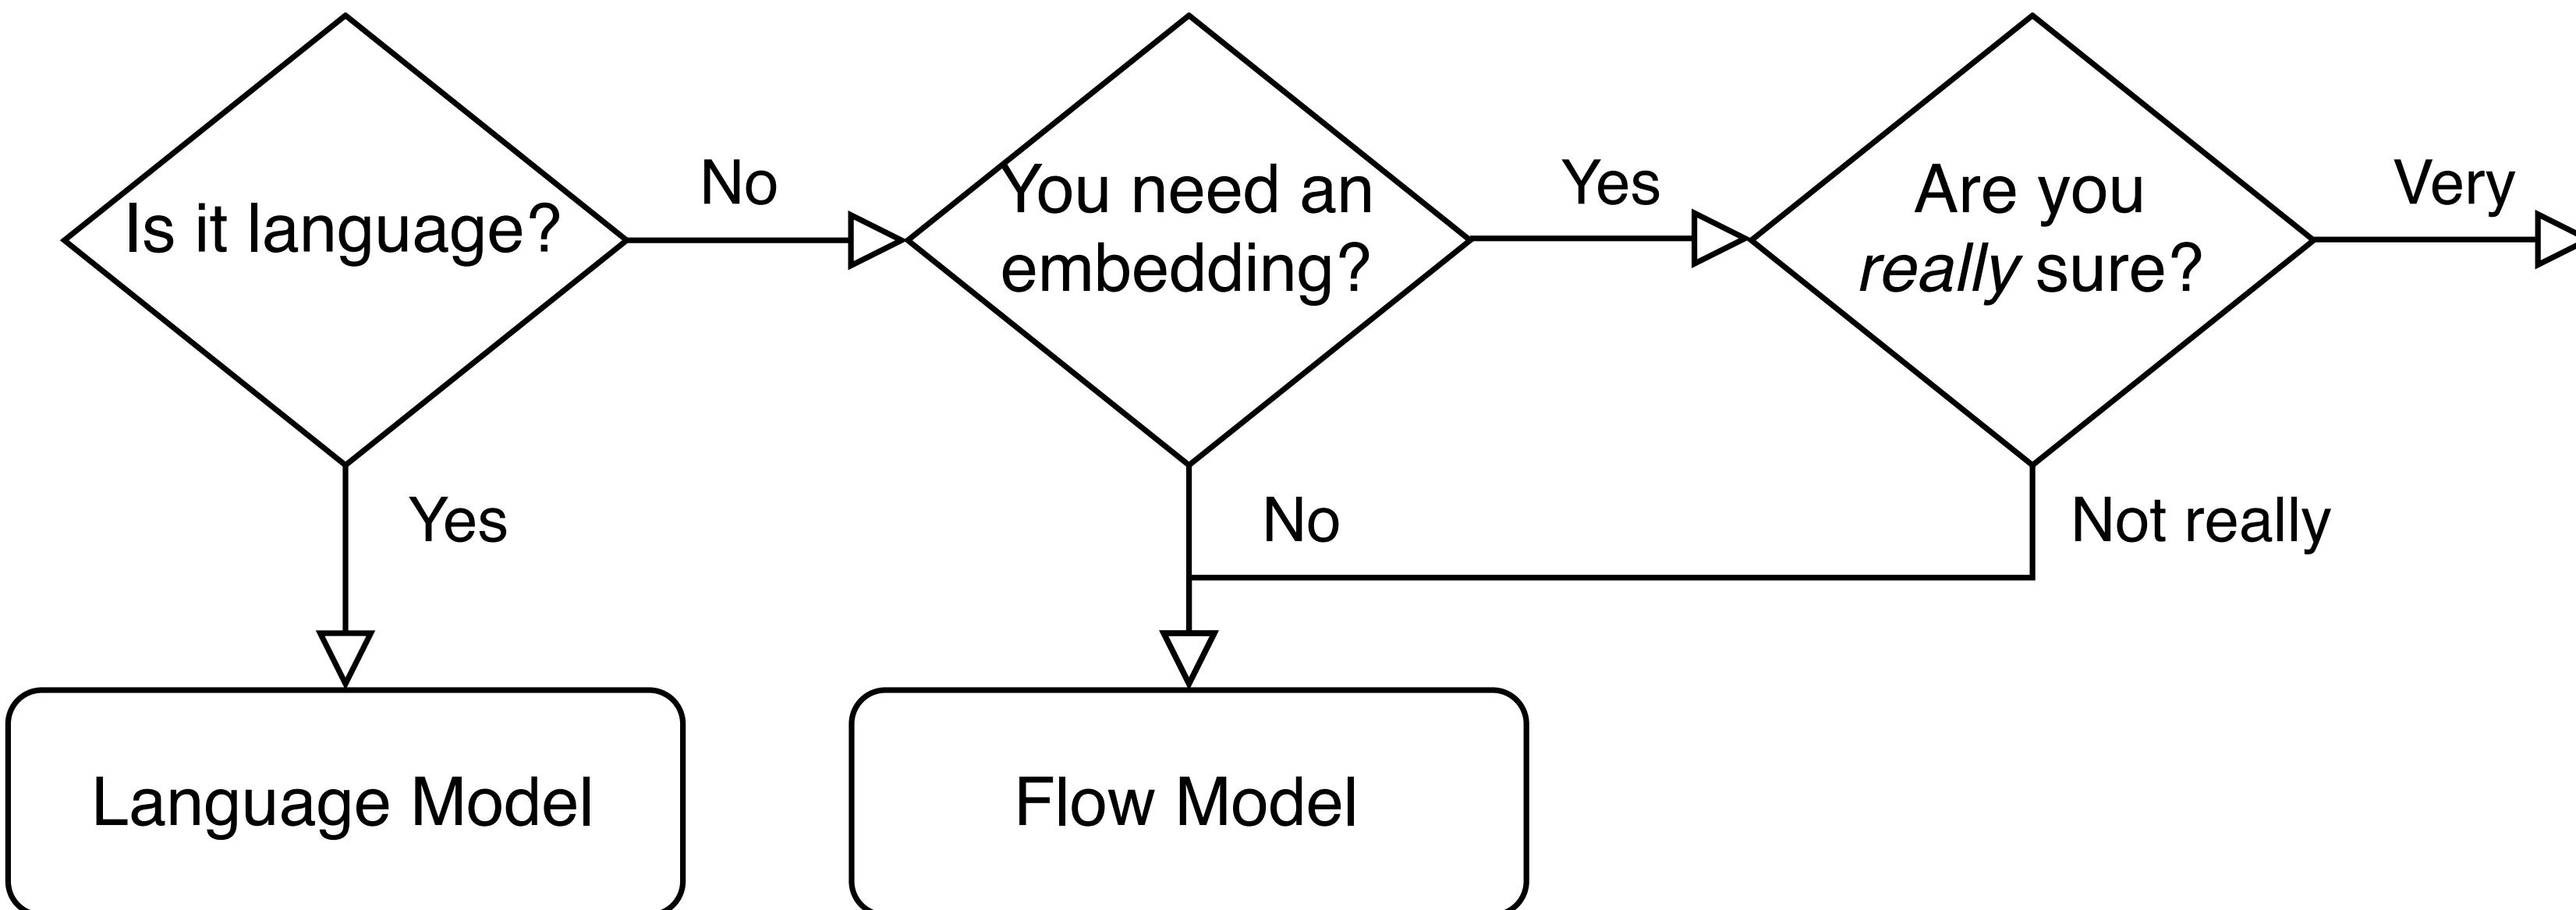
What do deep learning people mean by generative models?

Models trained to generate samples from the data



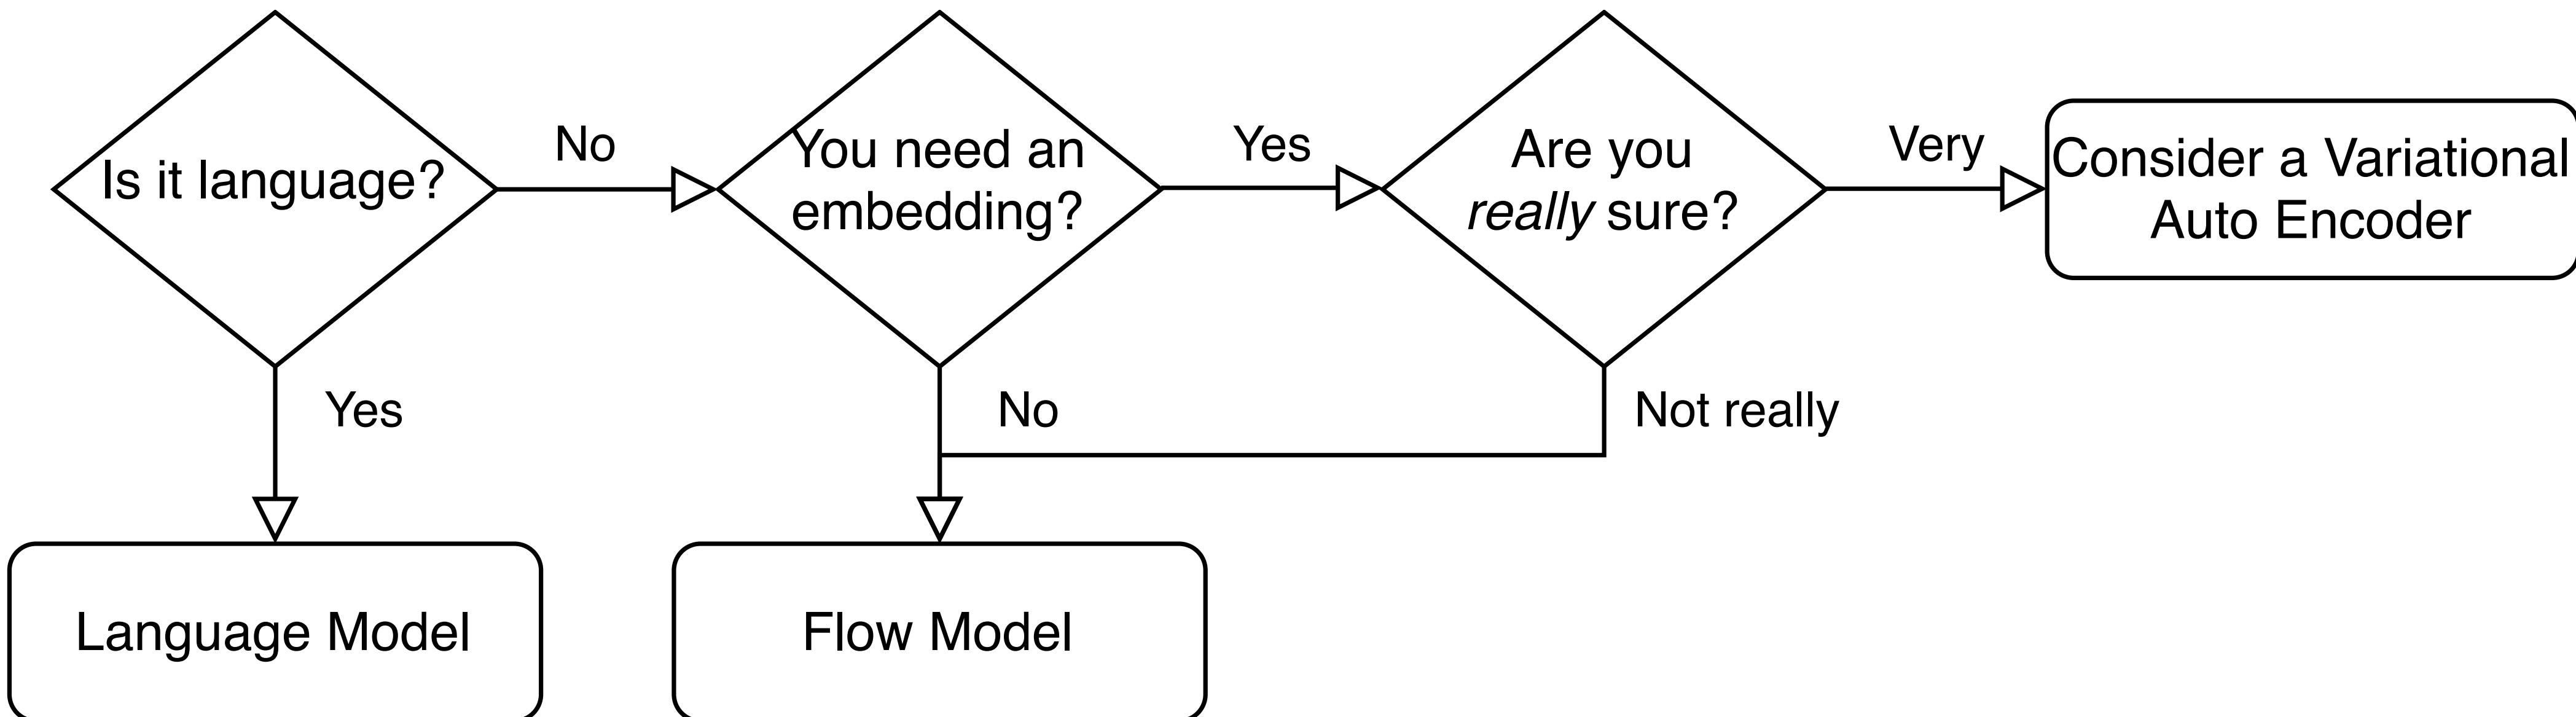
What do deep learning people mean by generative models?

Models trained to generate samples from the data



What do deep learning people mean by generative models?

Models trained to generate samples from the data



Flow Matching

- Image generation
- Video generation
- Text to speech
- Audio generation
- Molecule generation
- Protein design



[atong01/conditional-flow-matching](#)

Public

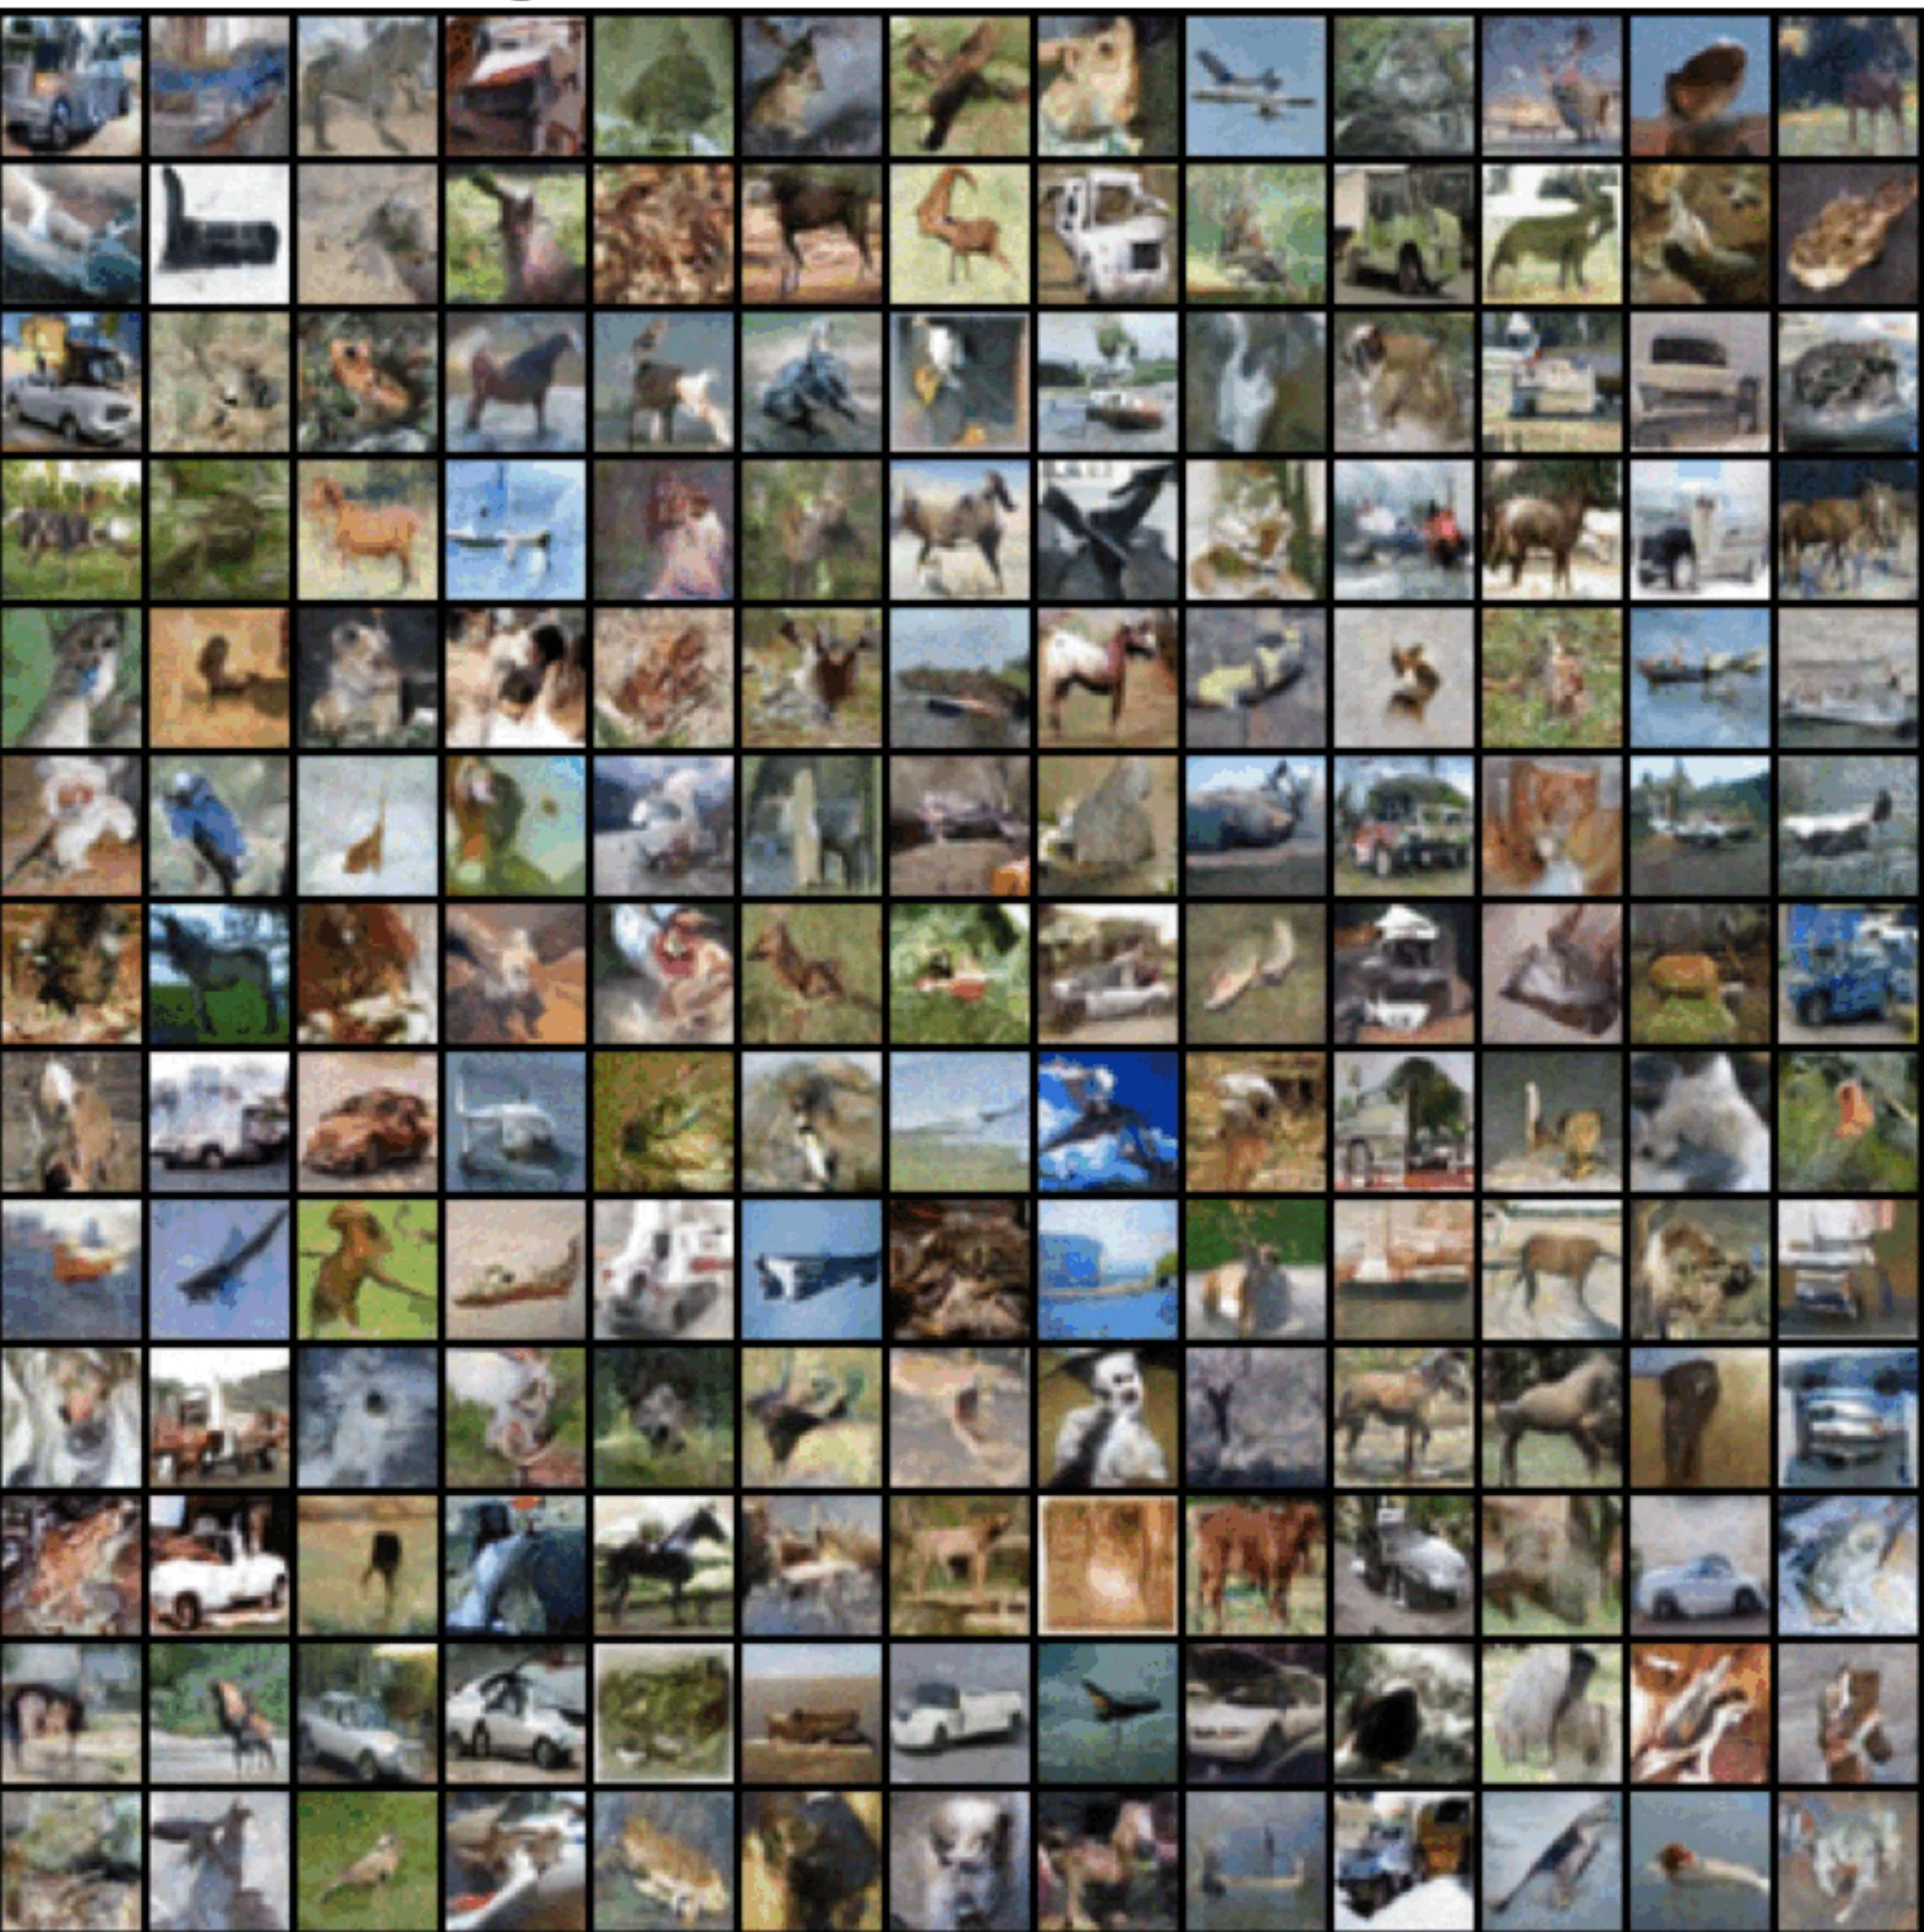
TorchCFM: a Conditional Flow Matching library

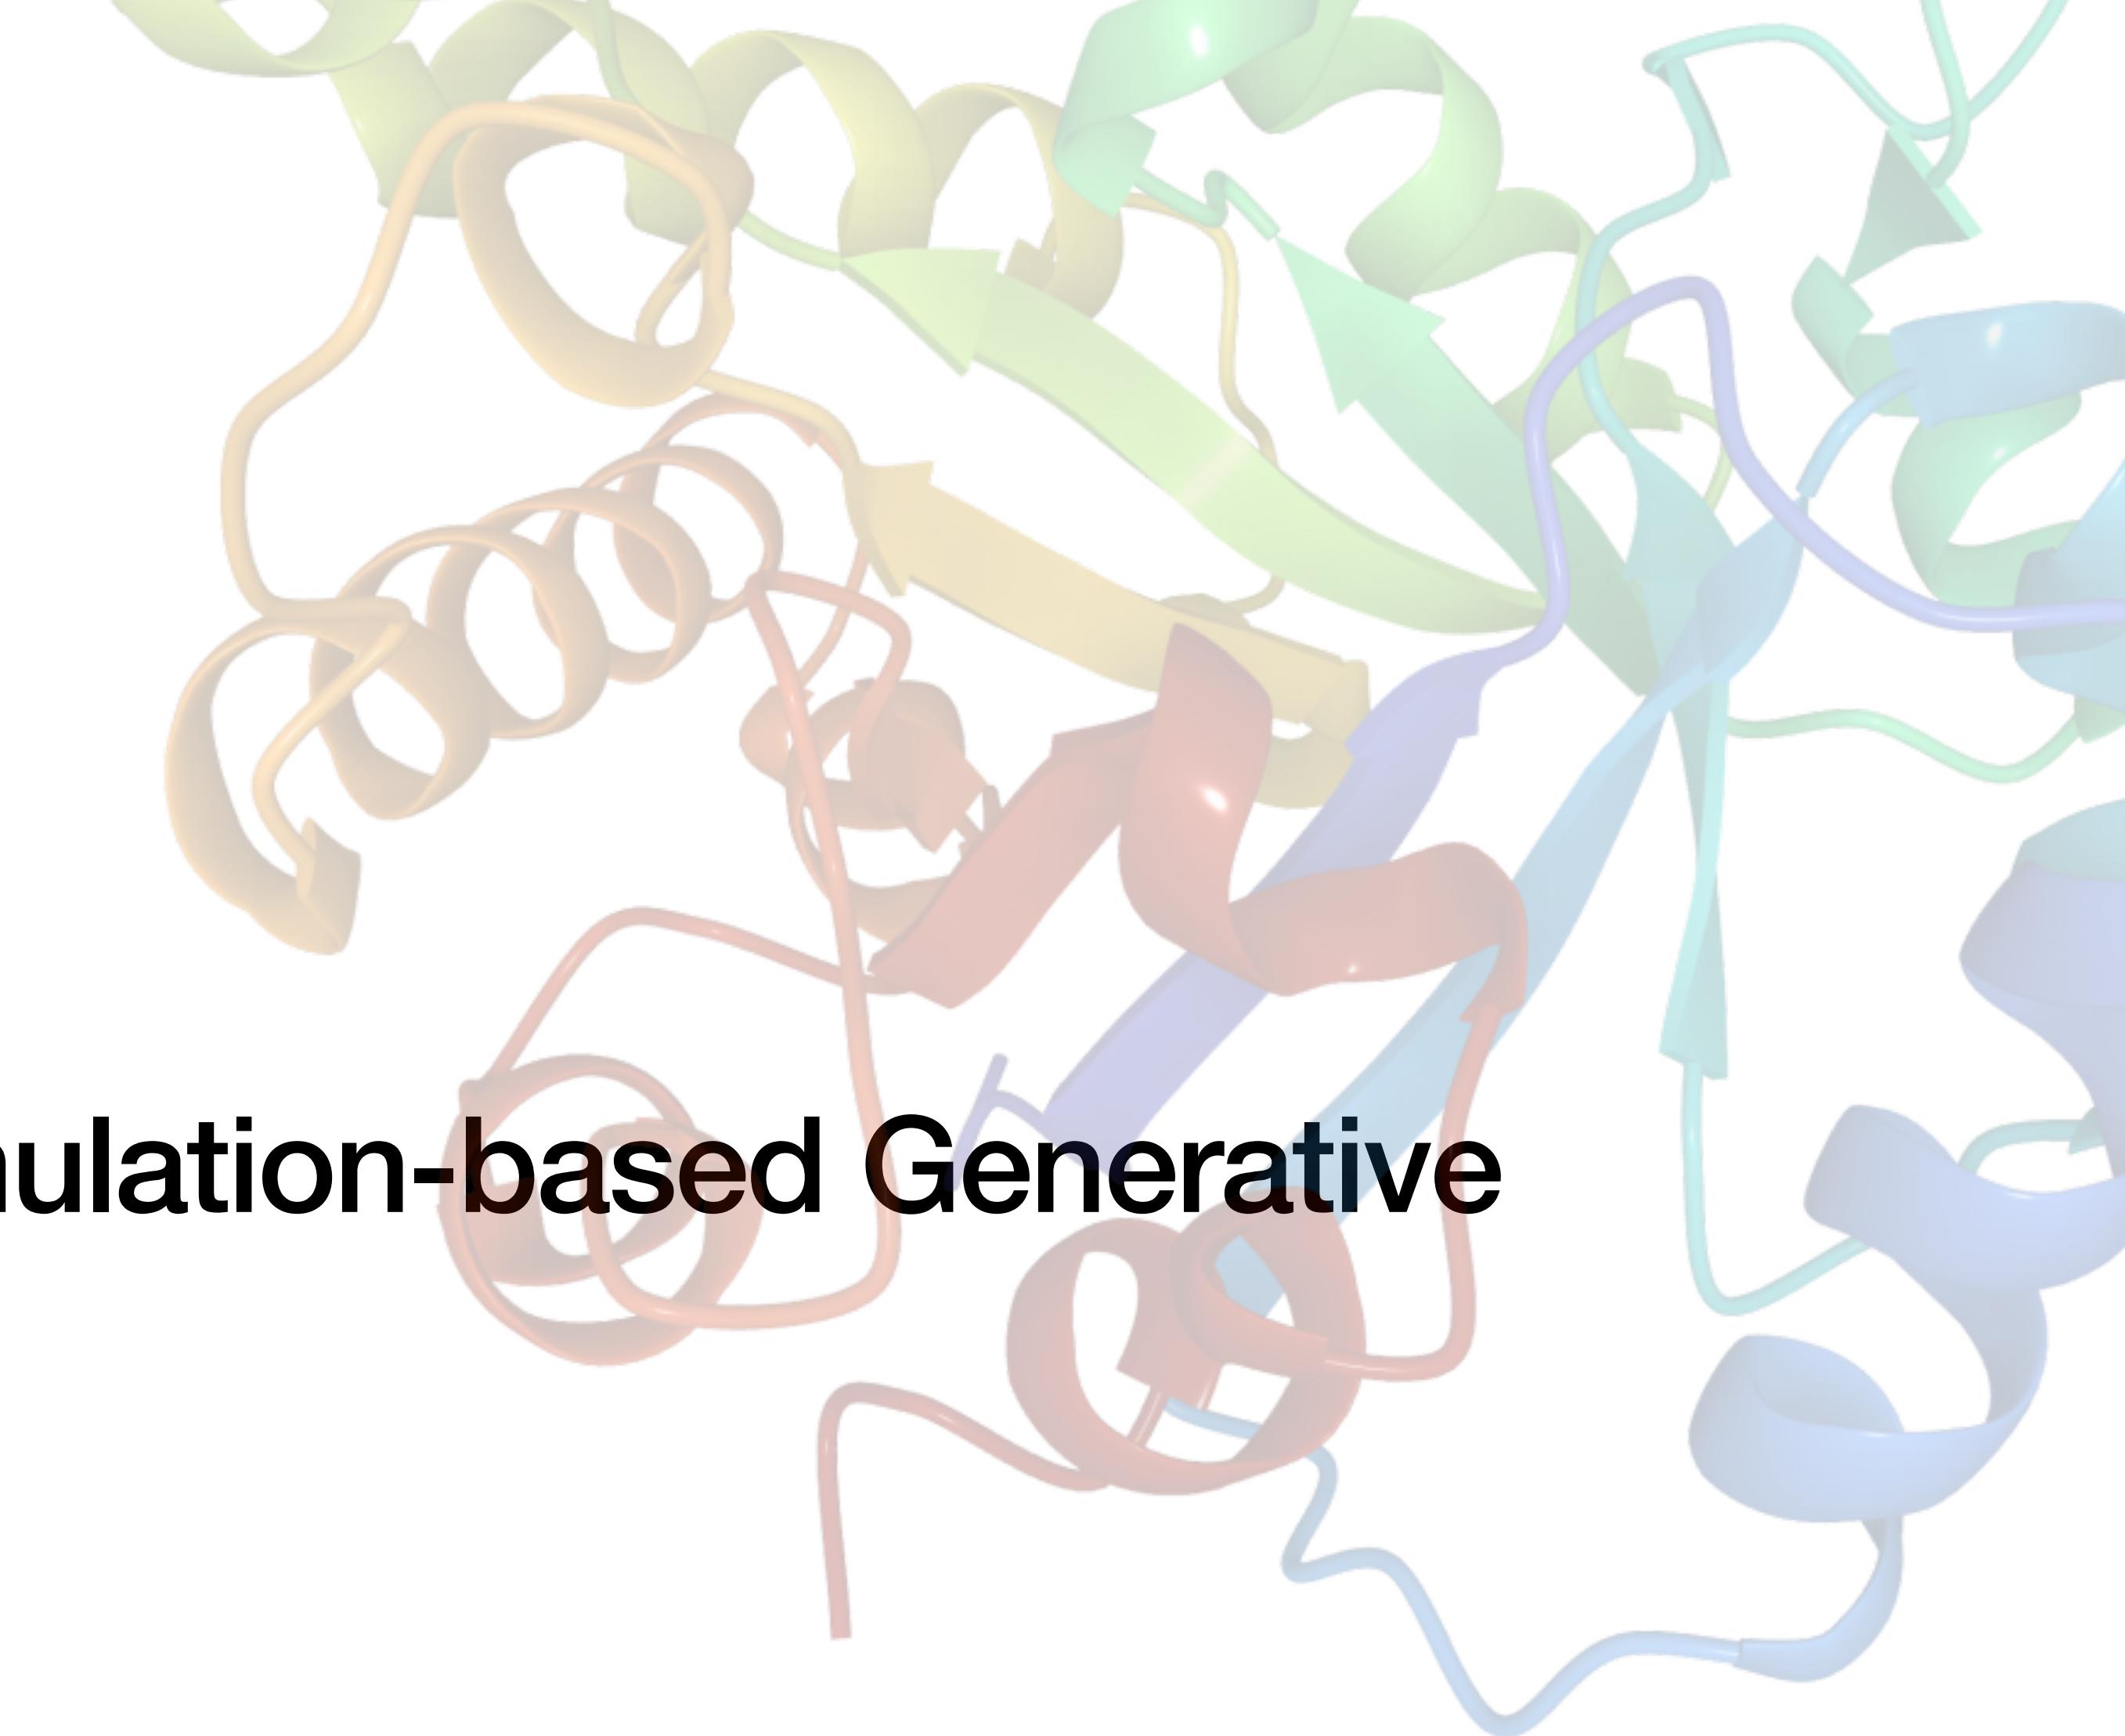
Python

1.8k

141

Generated images with OT-CFM at iteration 40k (FID: 47.4)





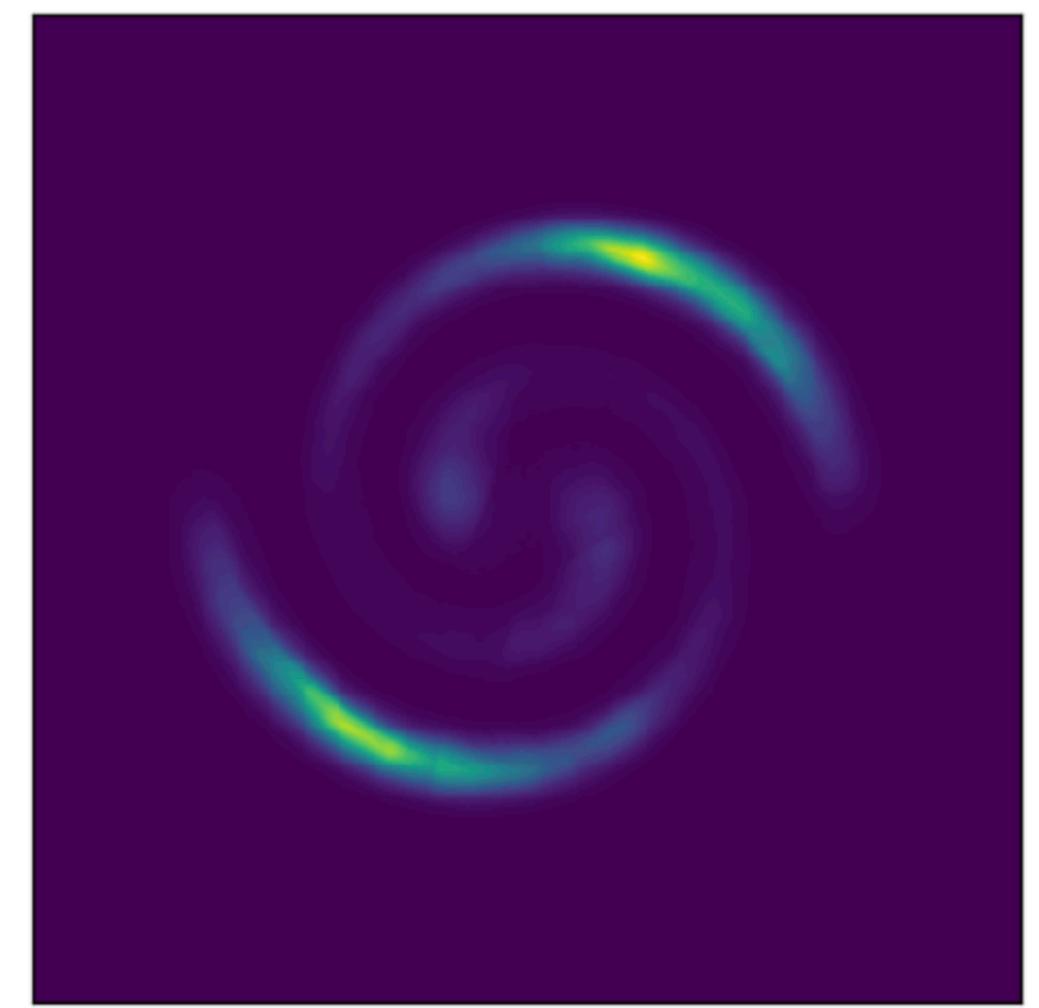
Part I:

Dynamical Systems and Simulation-based Generative

Models

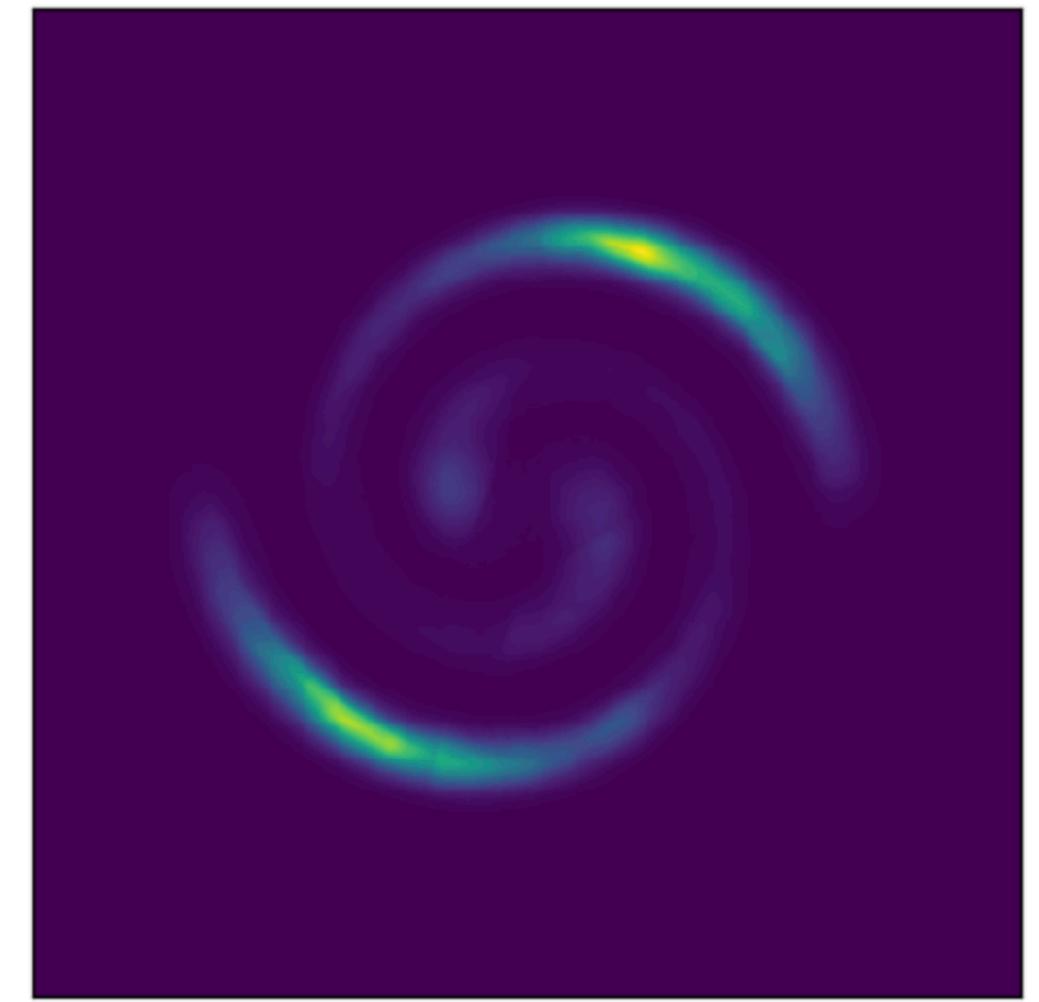
Problem Setting: Generative Modeling

- Unknown: data distribution q



Problem Setting: Generative Modeling

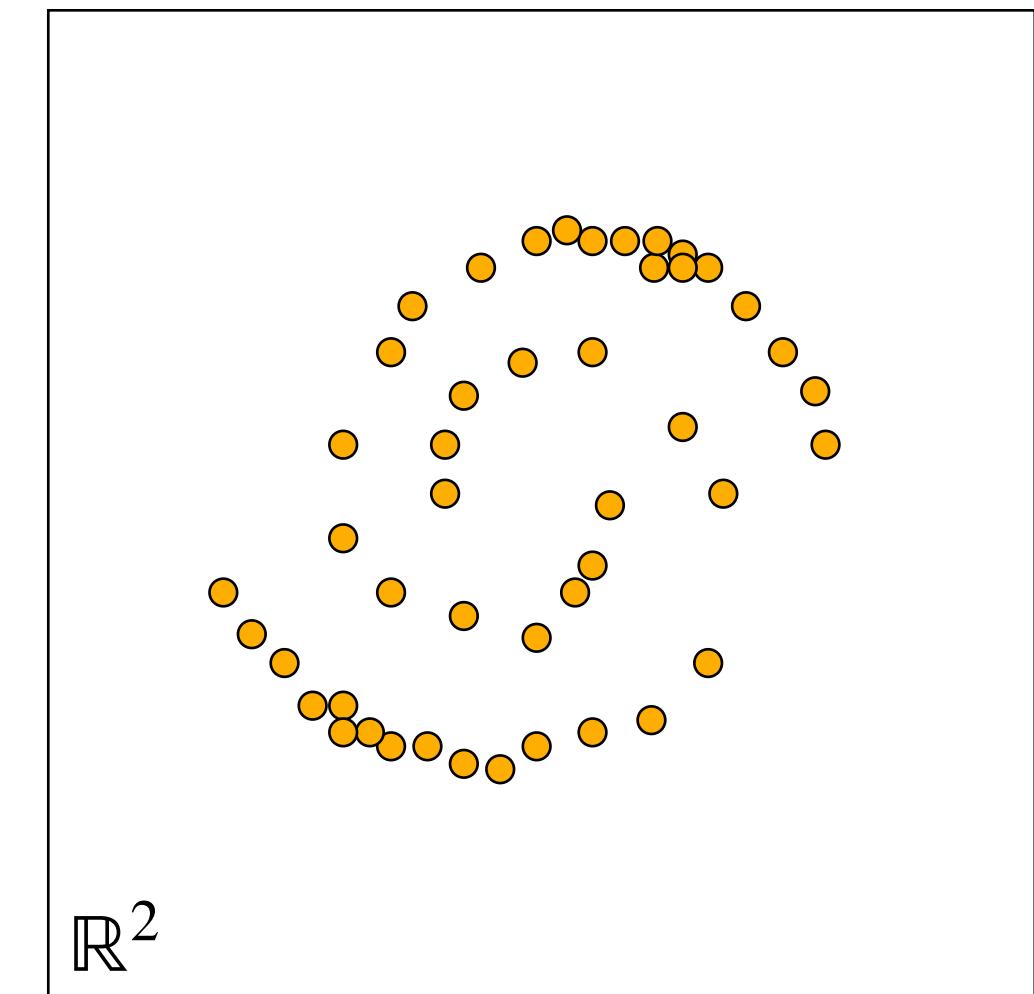
- **Unknown:** data distribution q
- **Given:** samples $x_1 \sim q$



Problem Setting: Generative Modeling

- Unknown: data distribution q
- Given: samples $x_1 \sim q$

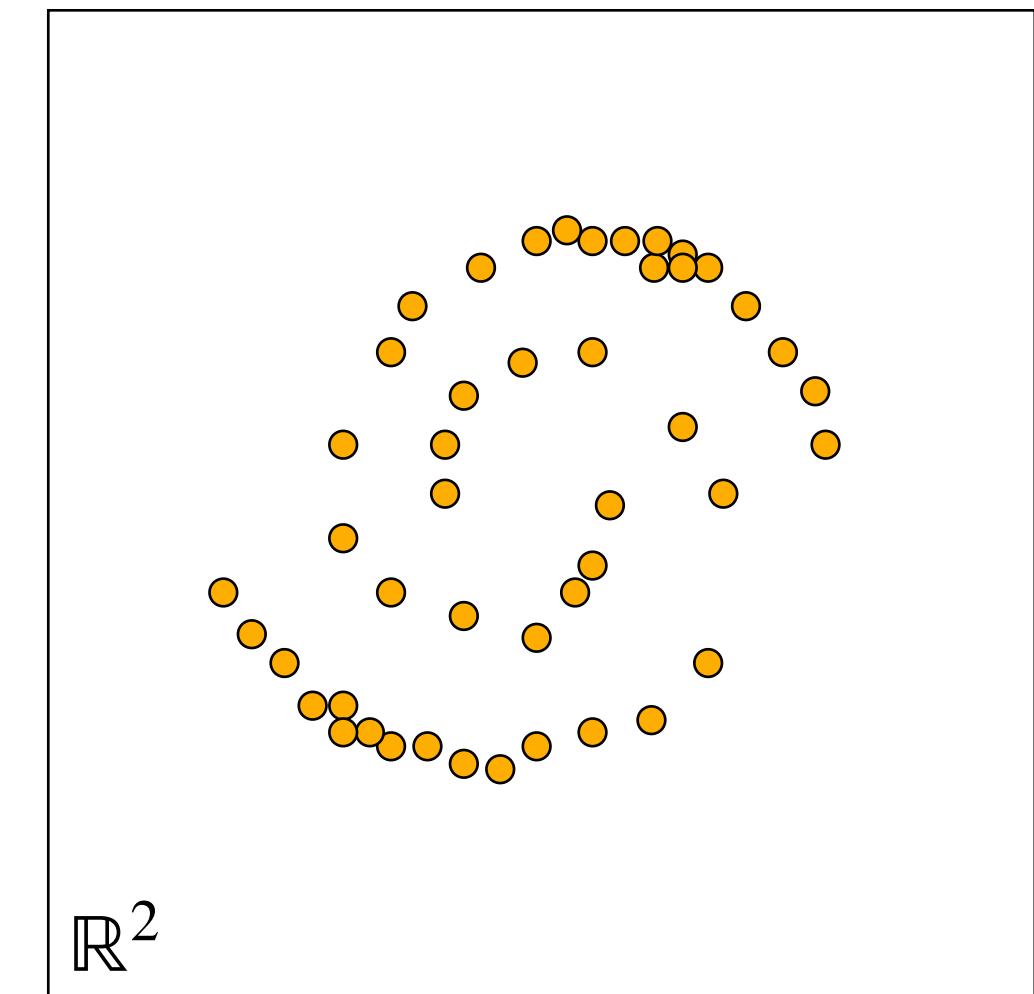
$$x_1 \sim q$$



Problem Setting: Generative Modeling

- Unknown: data distribution q
- Given: samples $x_1 \sim q$

$$x_1 \sim q$$

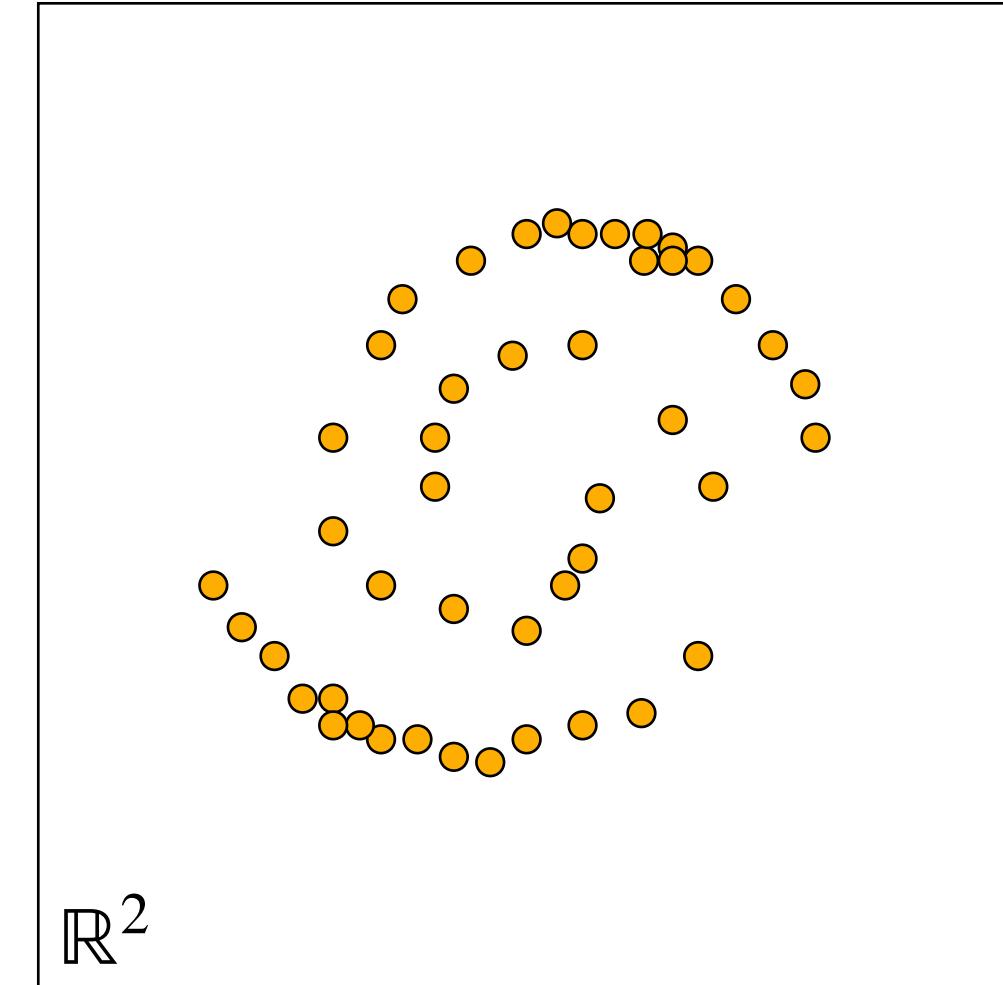


Goal: learn *a sampler* from the unknown q

Deep Generative Modeling

- Unknown: data distribution q
- Given: samples $x_1 \sim q$

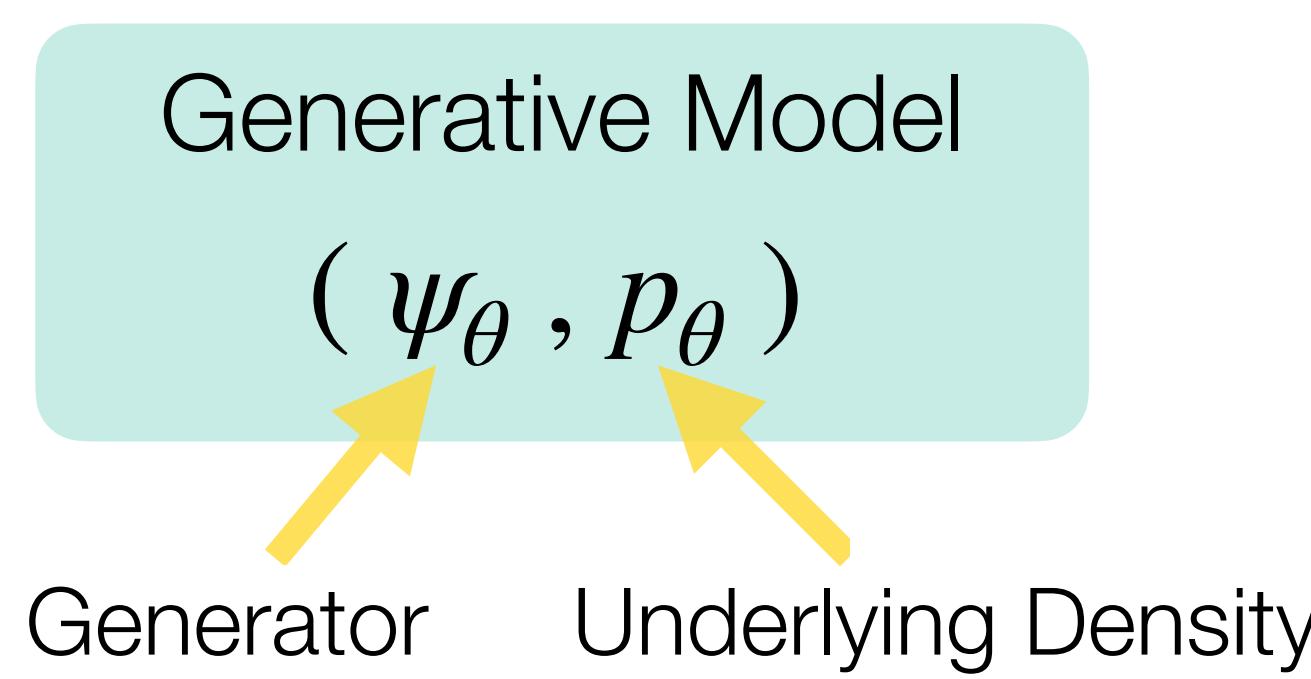
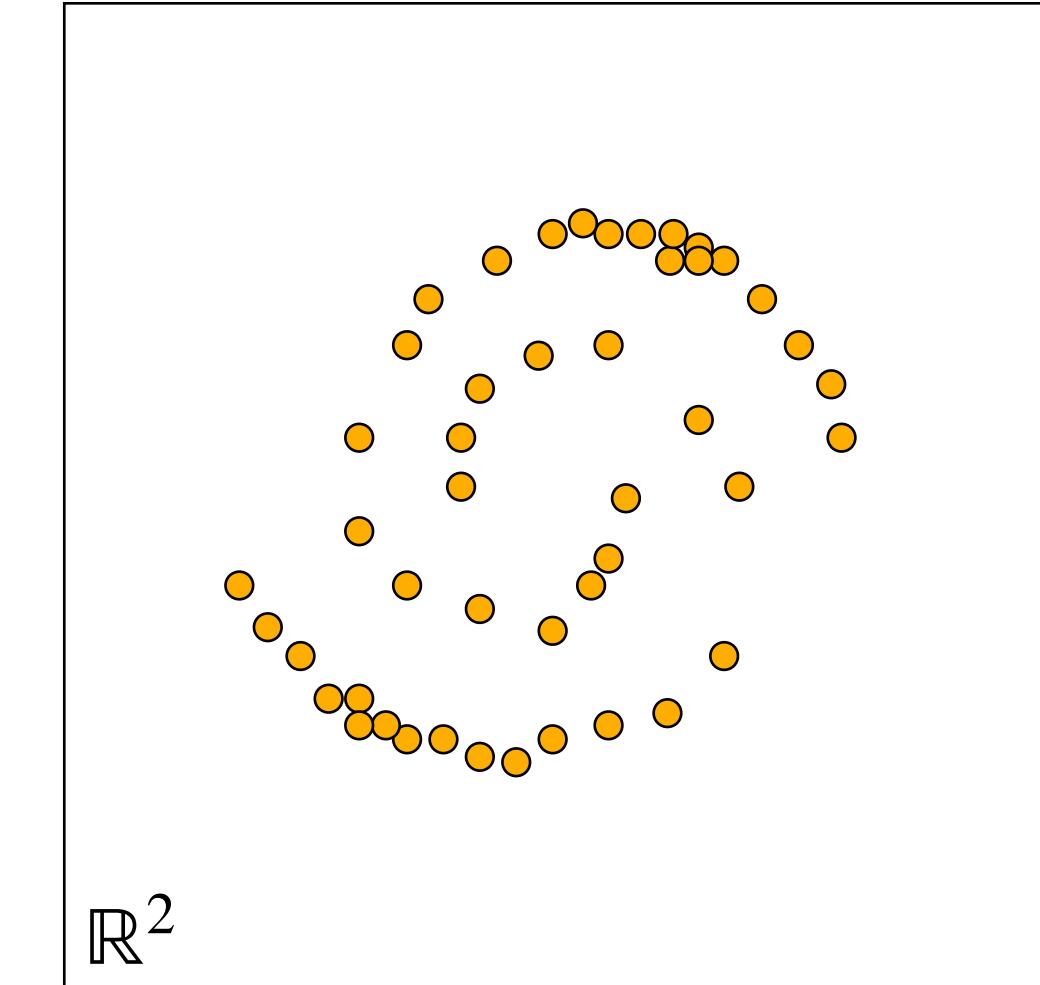
$$x_1 \sim q$$



Deep Generative Modeling

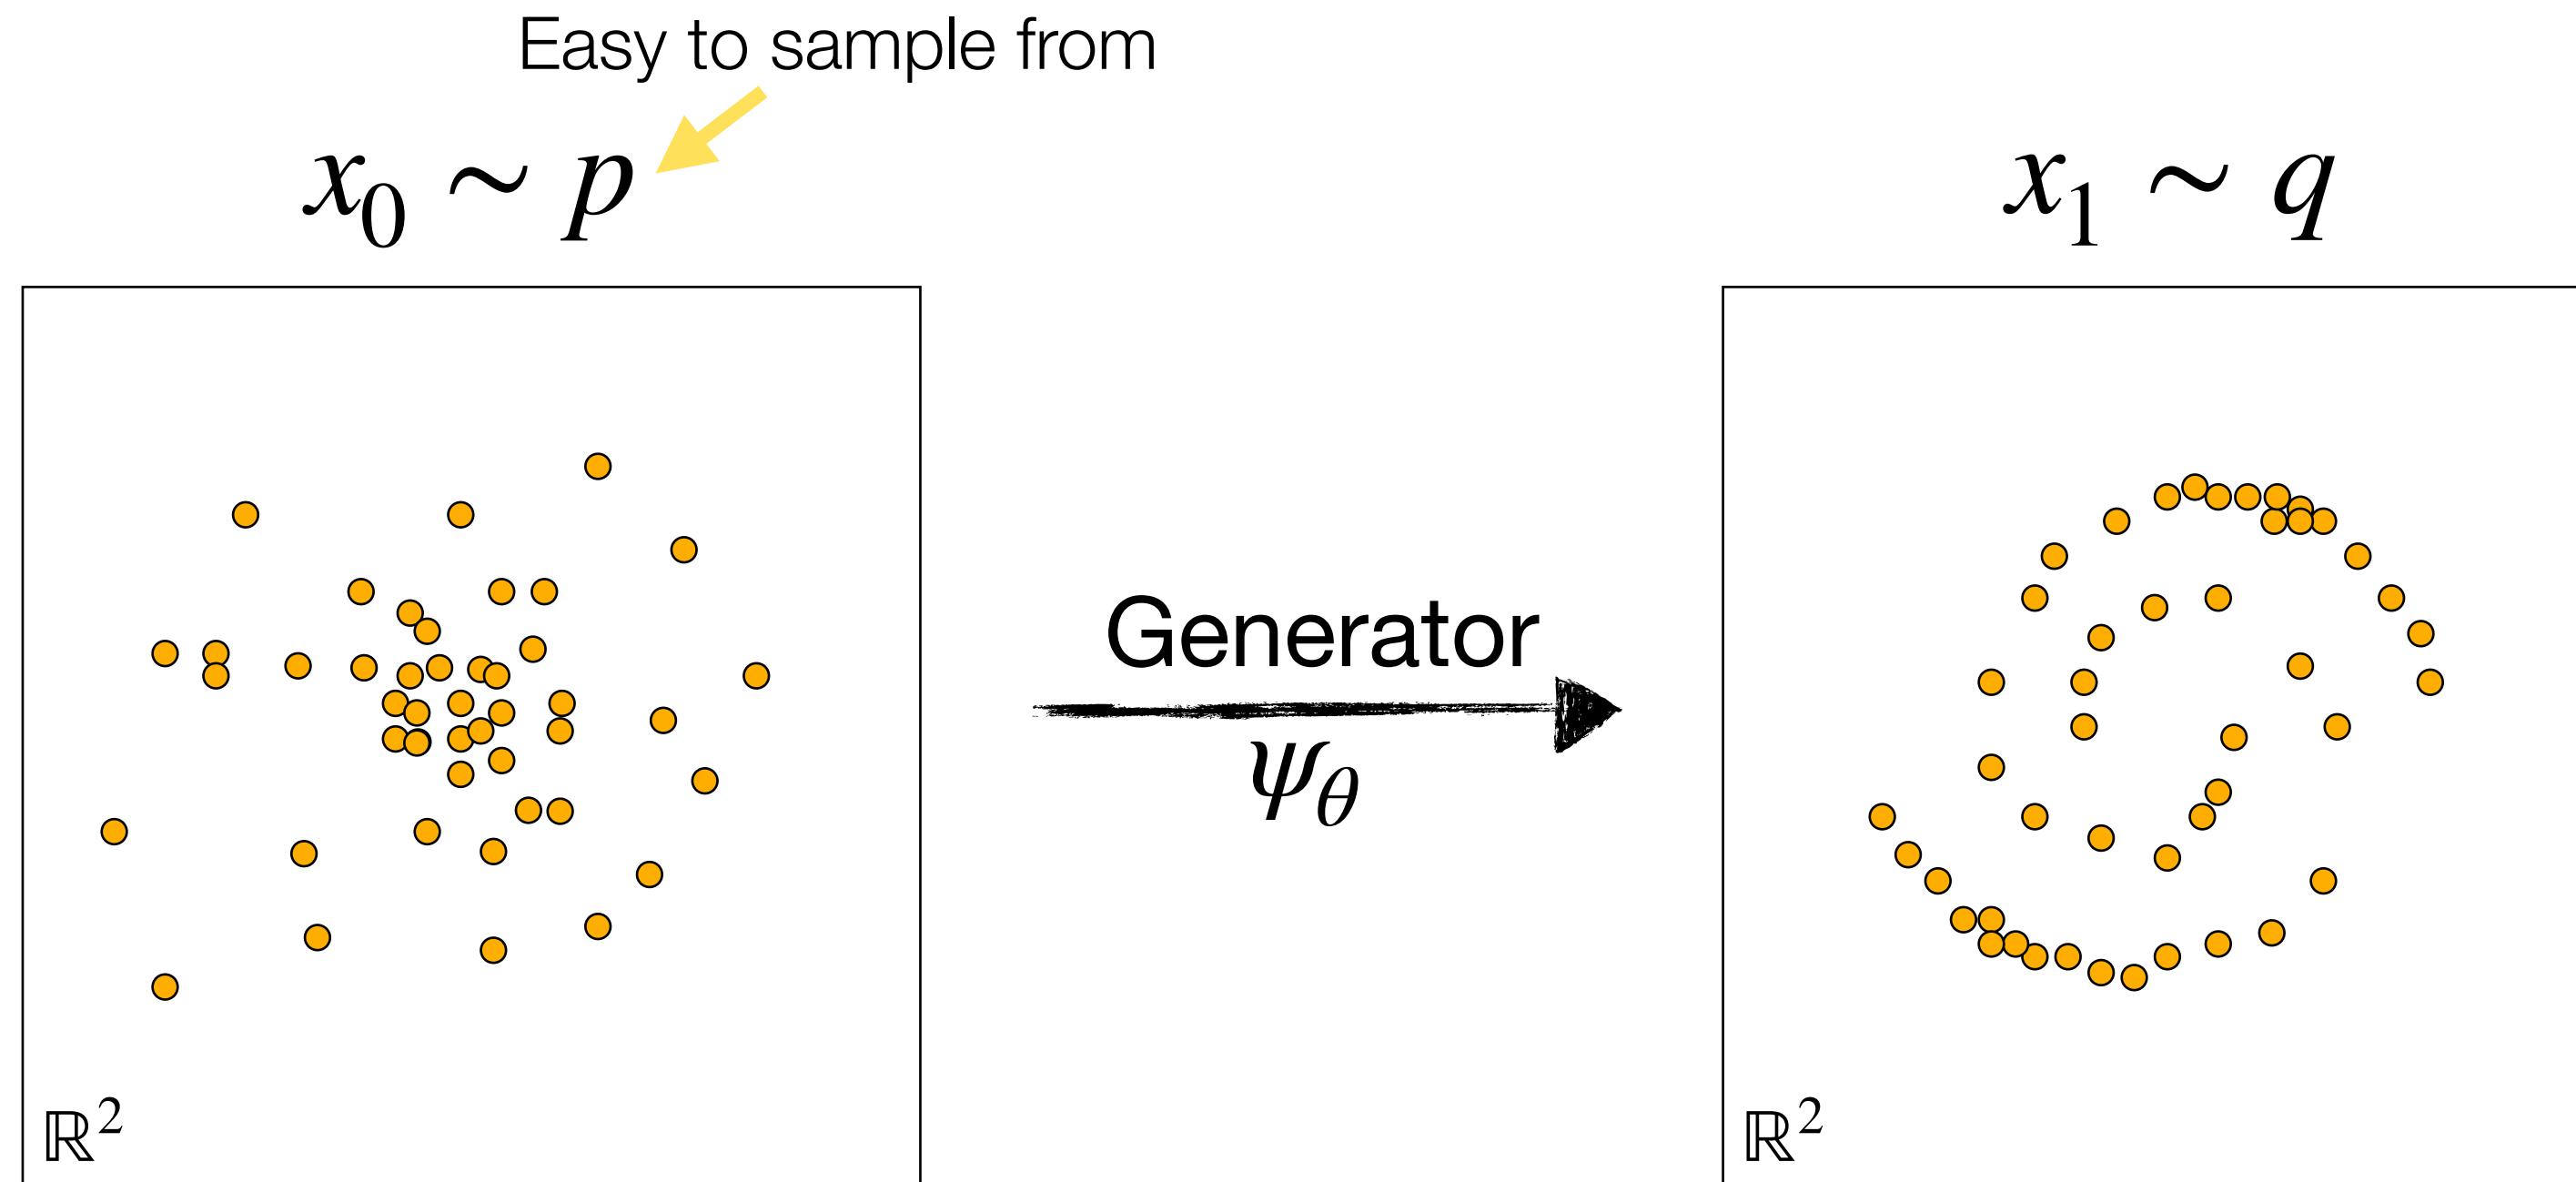
- Unknown: data distribution q
- Given: samples $x_1 \sim q$
- Learn: neural network with parameters θ

$$x_1 \sim q$$



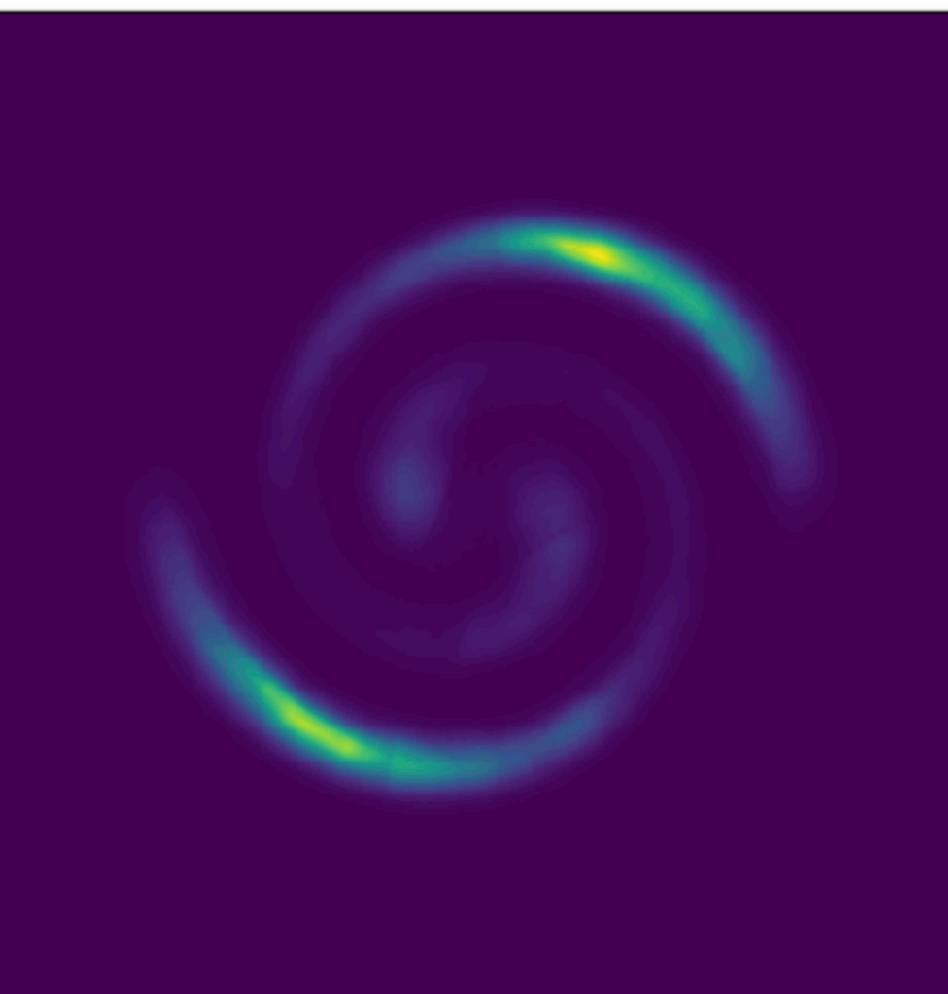
Goal: find parameters θ s.t. $p_\theta \approx q$

Deep Generative Modeling



Sampling
 $x_0 \sim p$
 $\psi_\theta(x_0) \sim q$

How to model ψ_θ ?



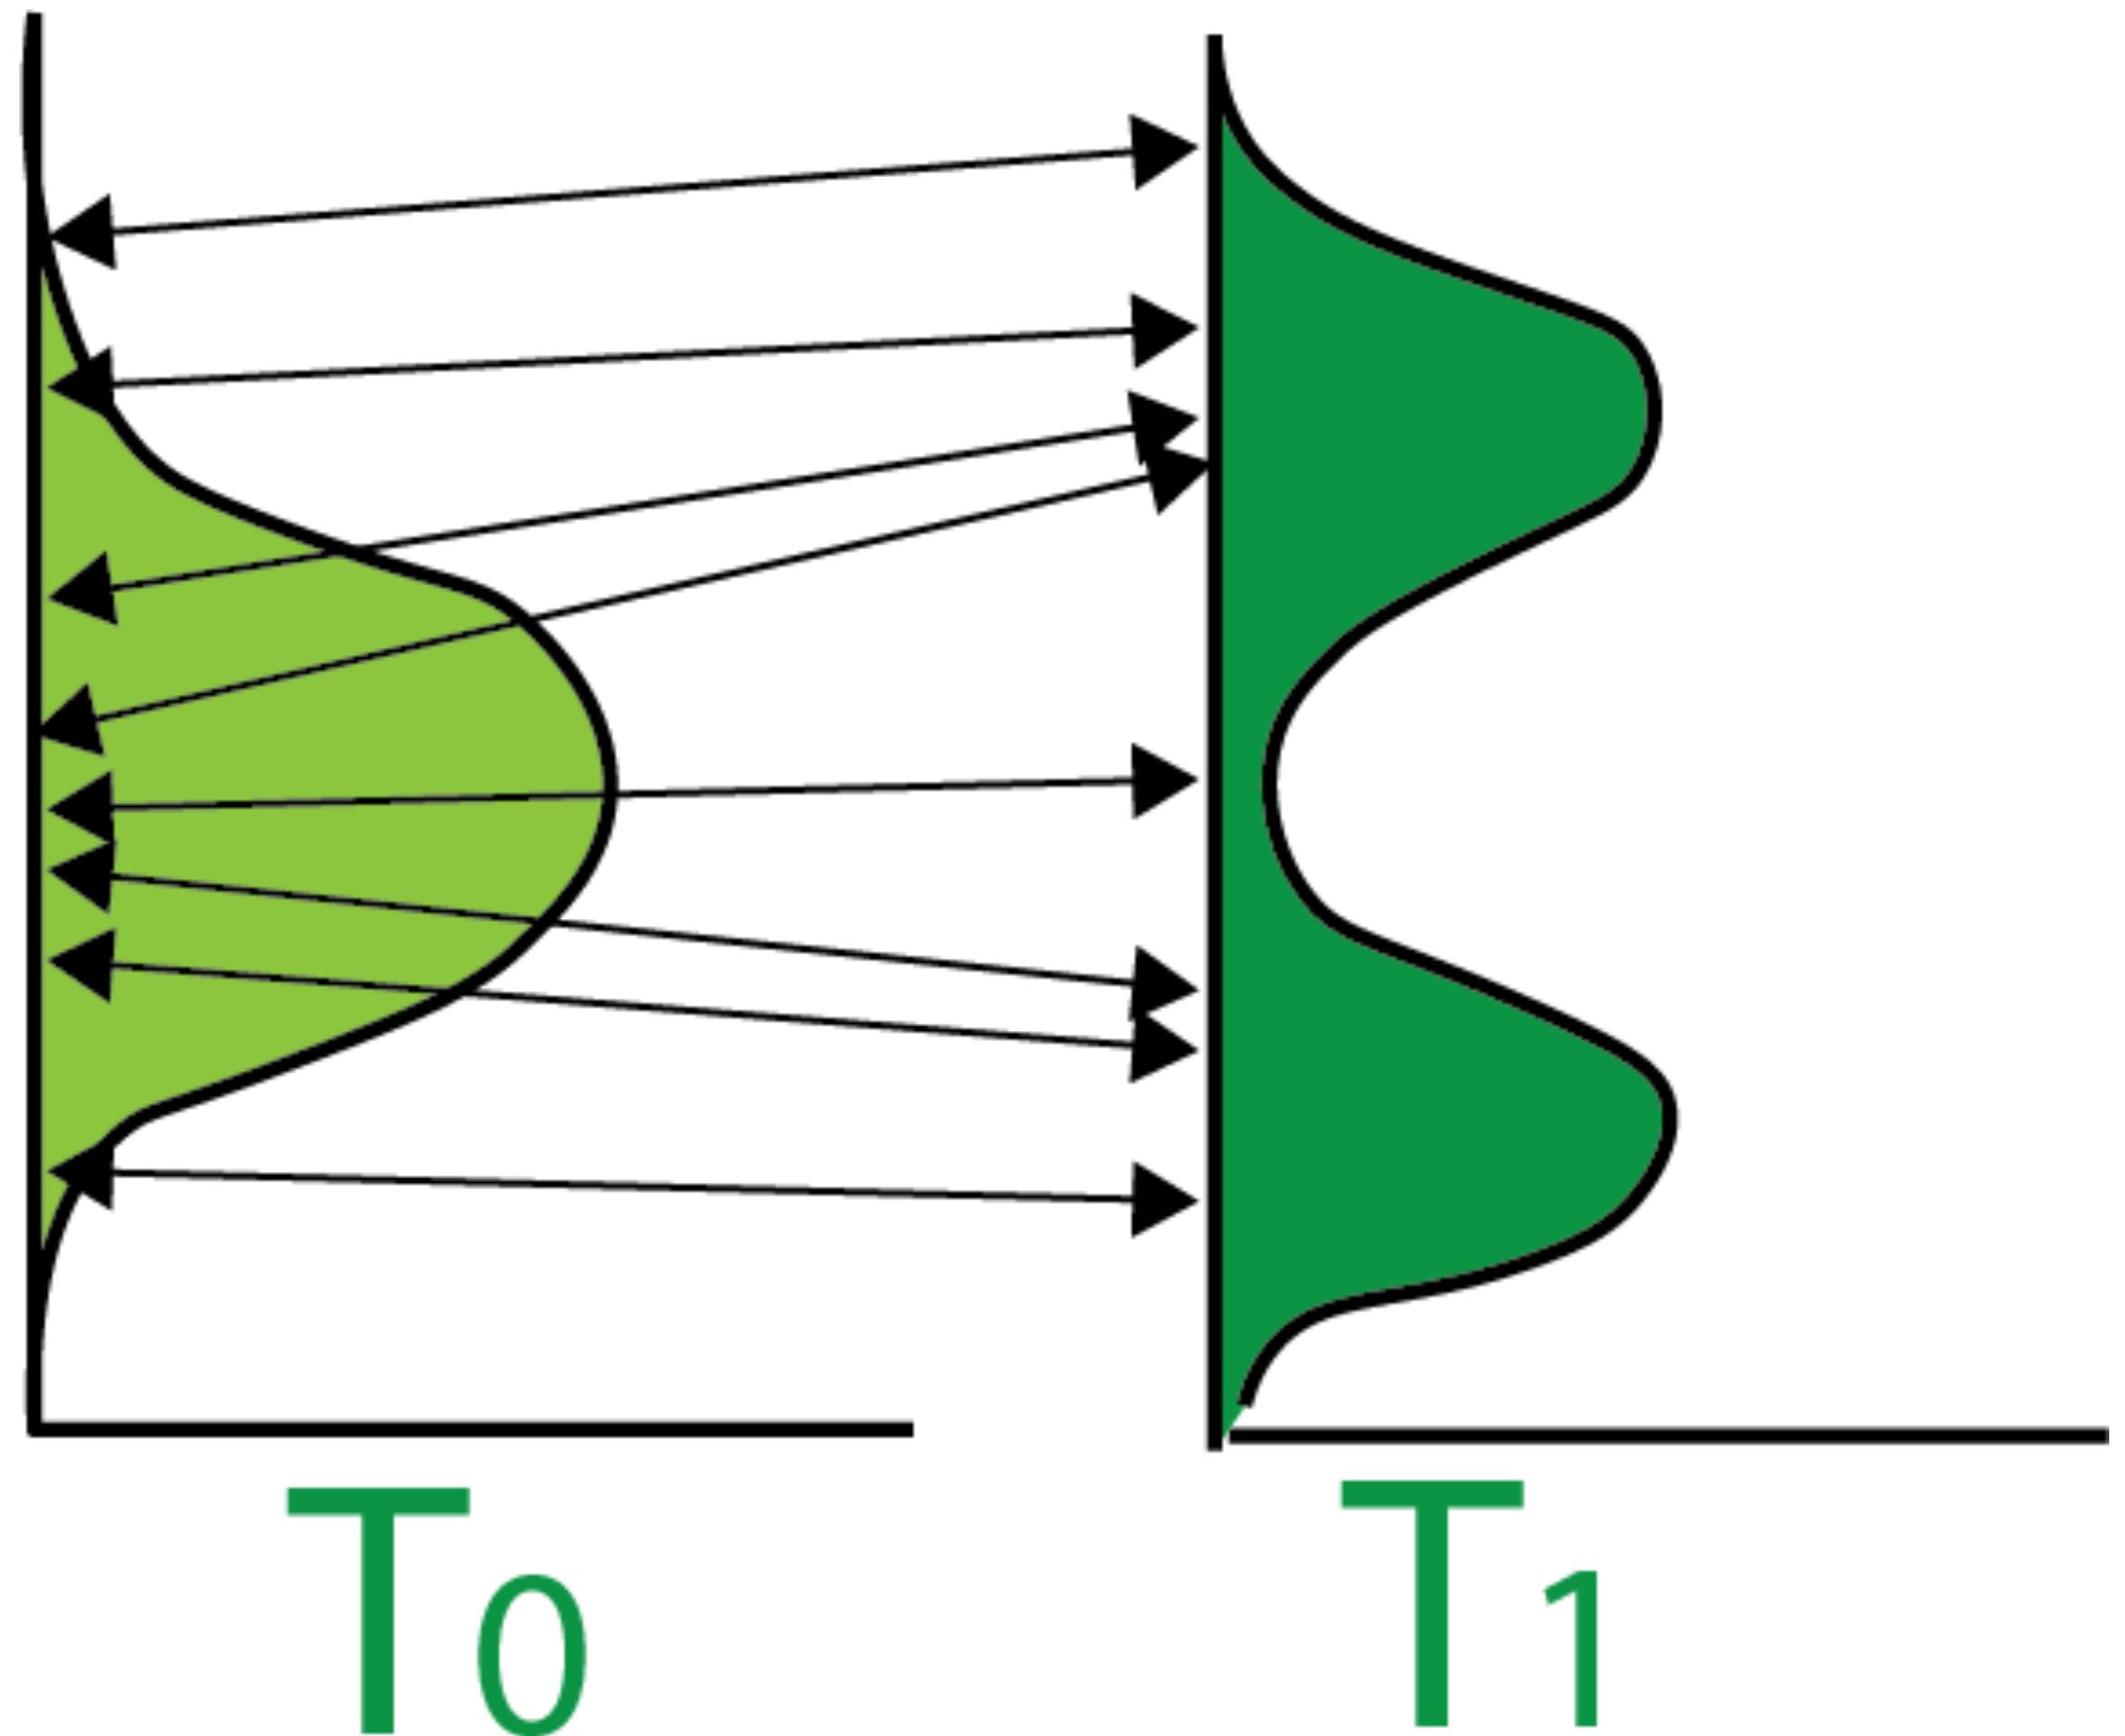
Density Estimation
 $p_\theta \approx q$

Normalizing Flows (NFs)

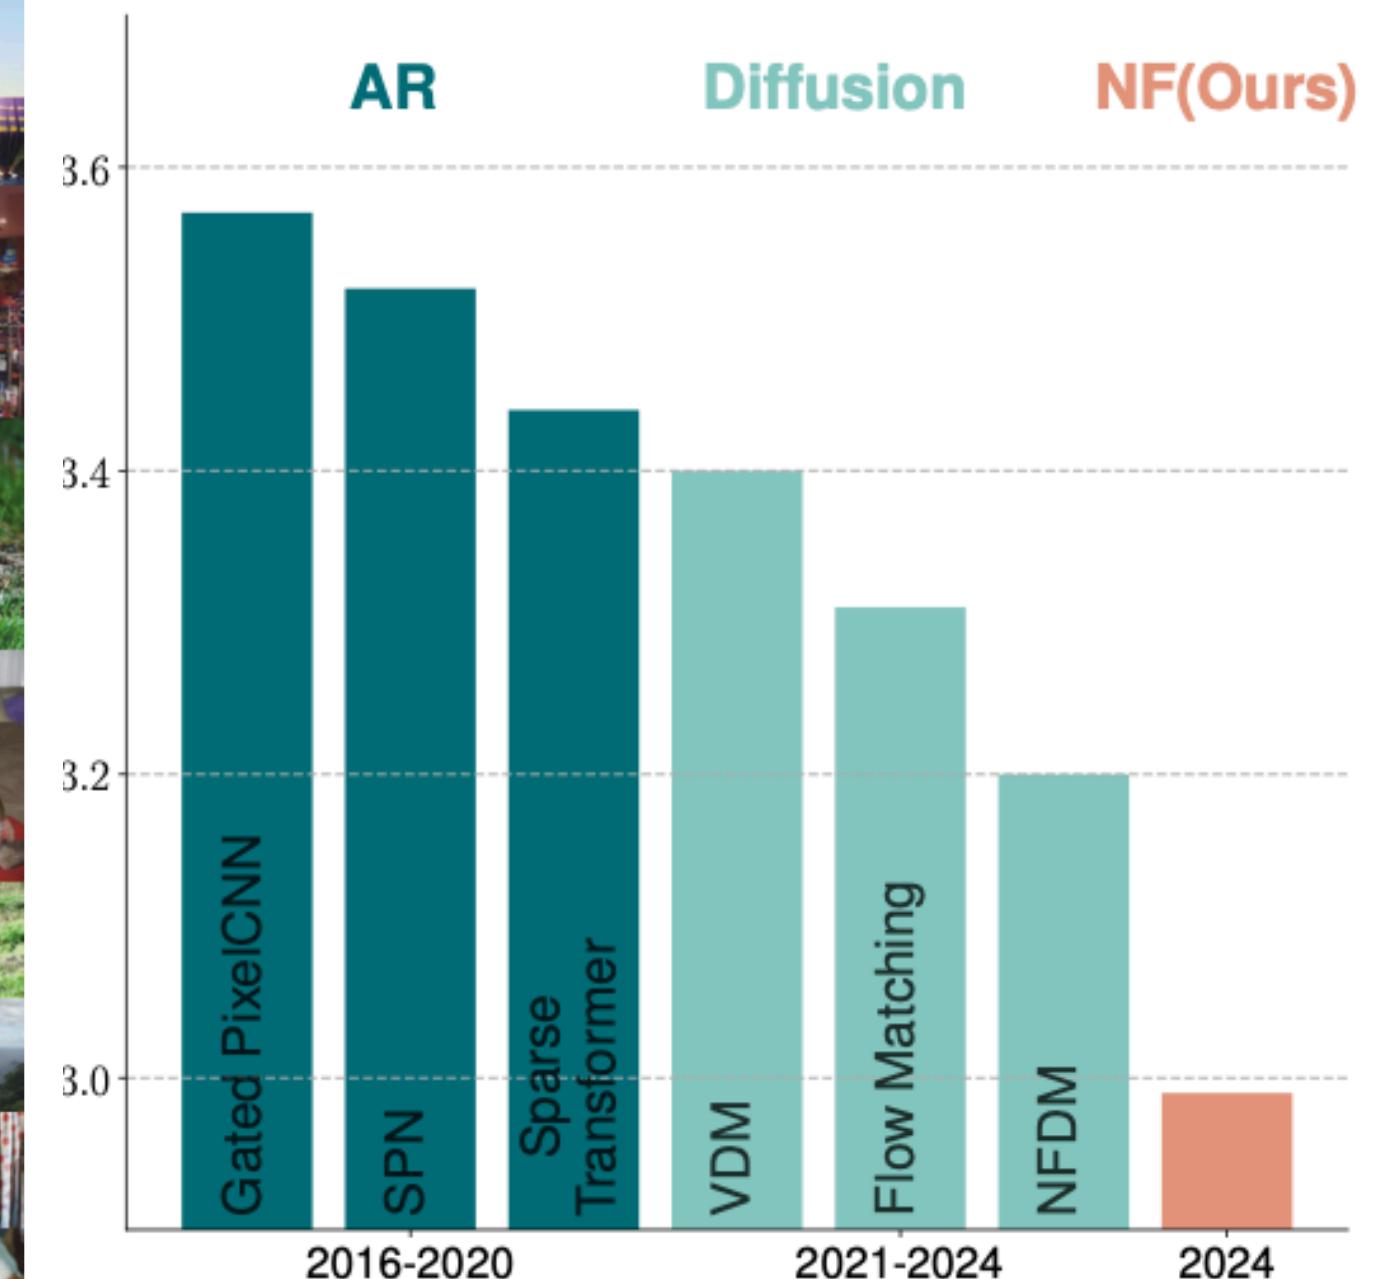
Sample from some complicated distribution by sampling from a simple distribution then applying U

- Begin with a simple distribution

$$p_{t_0}(x) \sim \mathcal{N}(0, 1)$$



Normalizing Flows (NFs)



Normalizing Flows are Capable Generative Models

Shuangfei Zhai¹ Ruixiang Zhang¹ Preetum Nakkiran¹ David Berthelot¹ Jiatao Gu¹ Huangjie Zheng¹
Tianrong Chen¹ Miguel Angel Bautista¹ Navdeep Jaitly¹ Josh Susskind¹



Normalizing Flows (NFs)

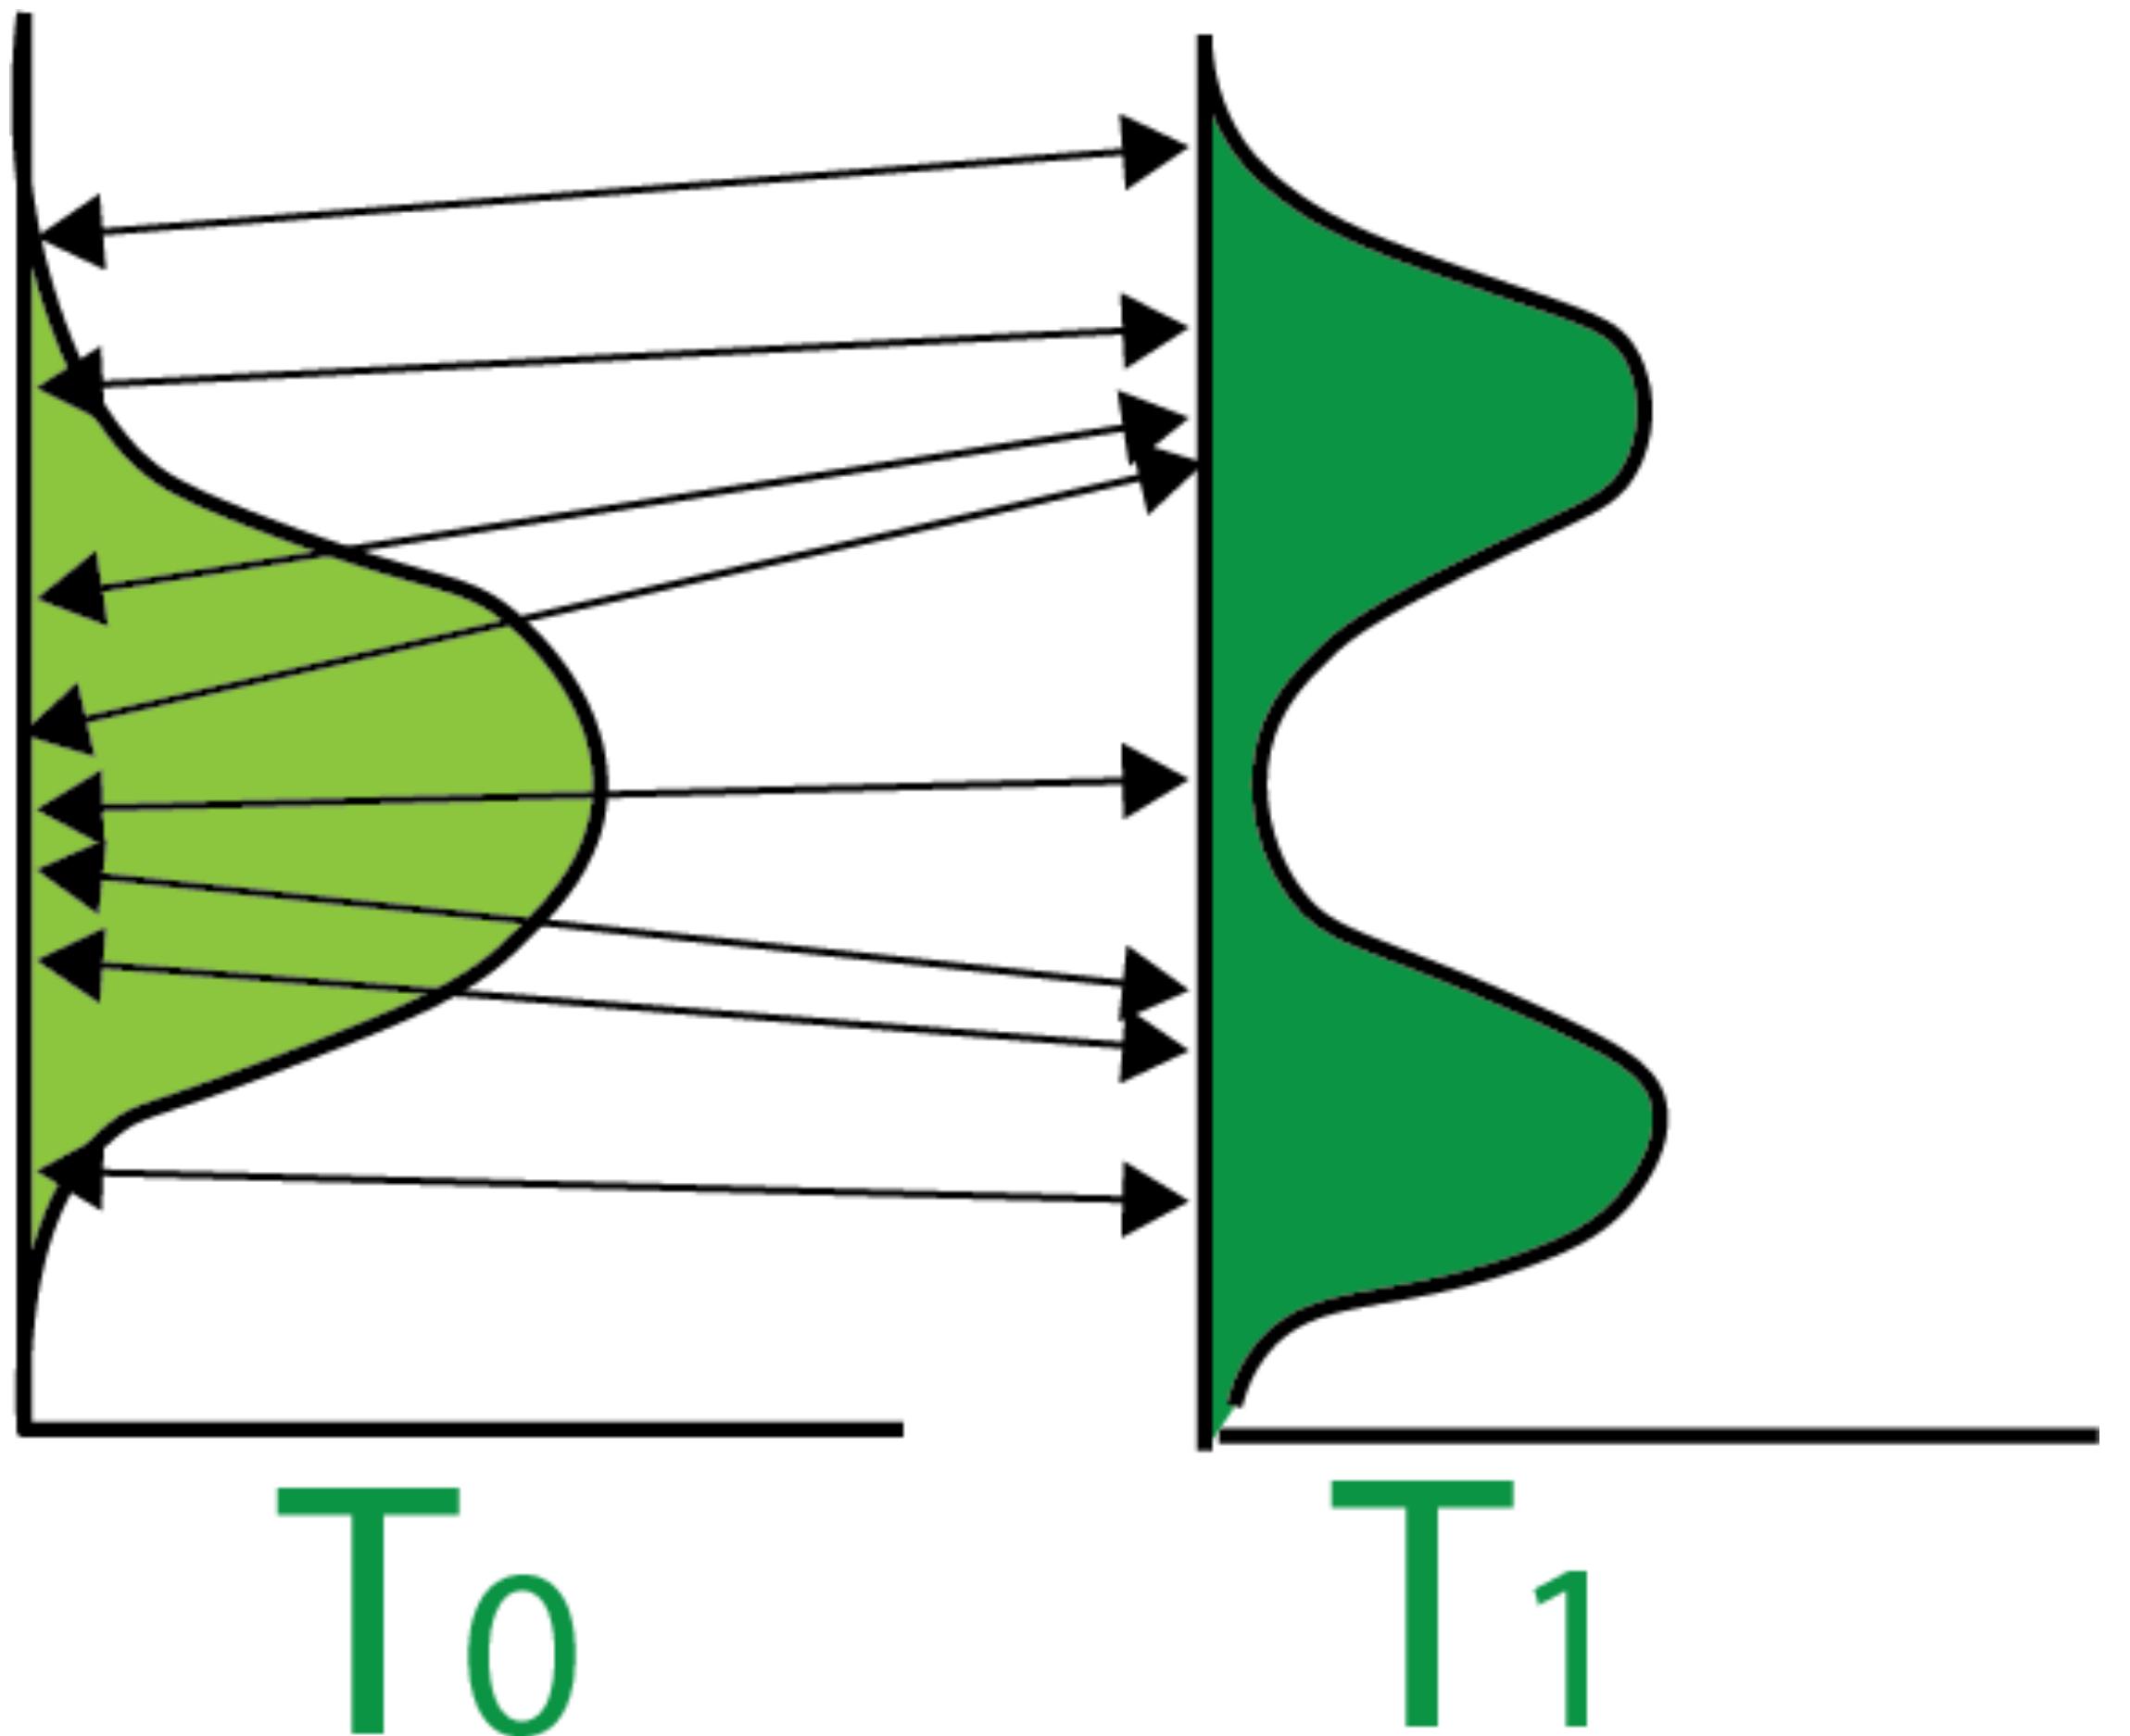
Sample from some complicated distribution by sampling from a simple distribution then applying U

- Begin with a simple distribution

$$p_{t_0}(x) \sim \mathcal{N}(0, 1)$$

- Apply an invertible transformation(s)

$$x_{t_1} = U(x_{t_0})$$



Normalizing Flows (NFs)

Sample from some complicated distribution by sampling from a simple distribution then applying U

- Begin with a simple distribution

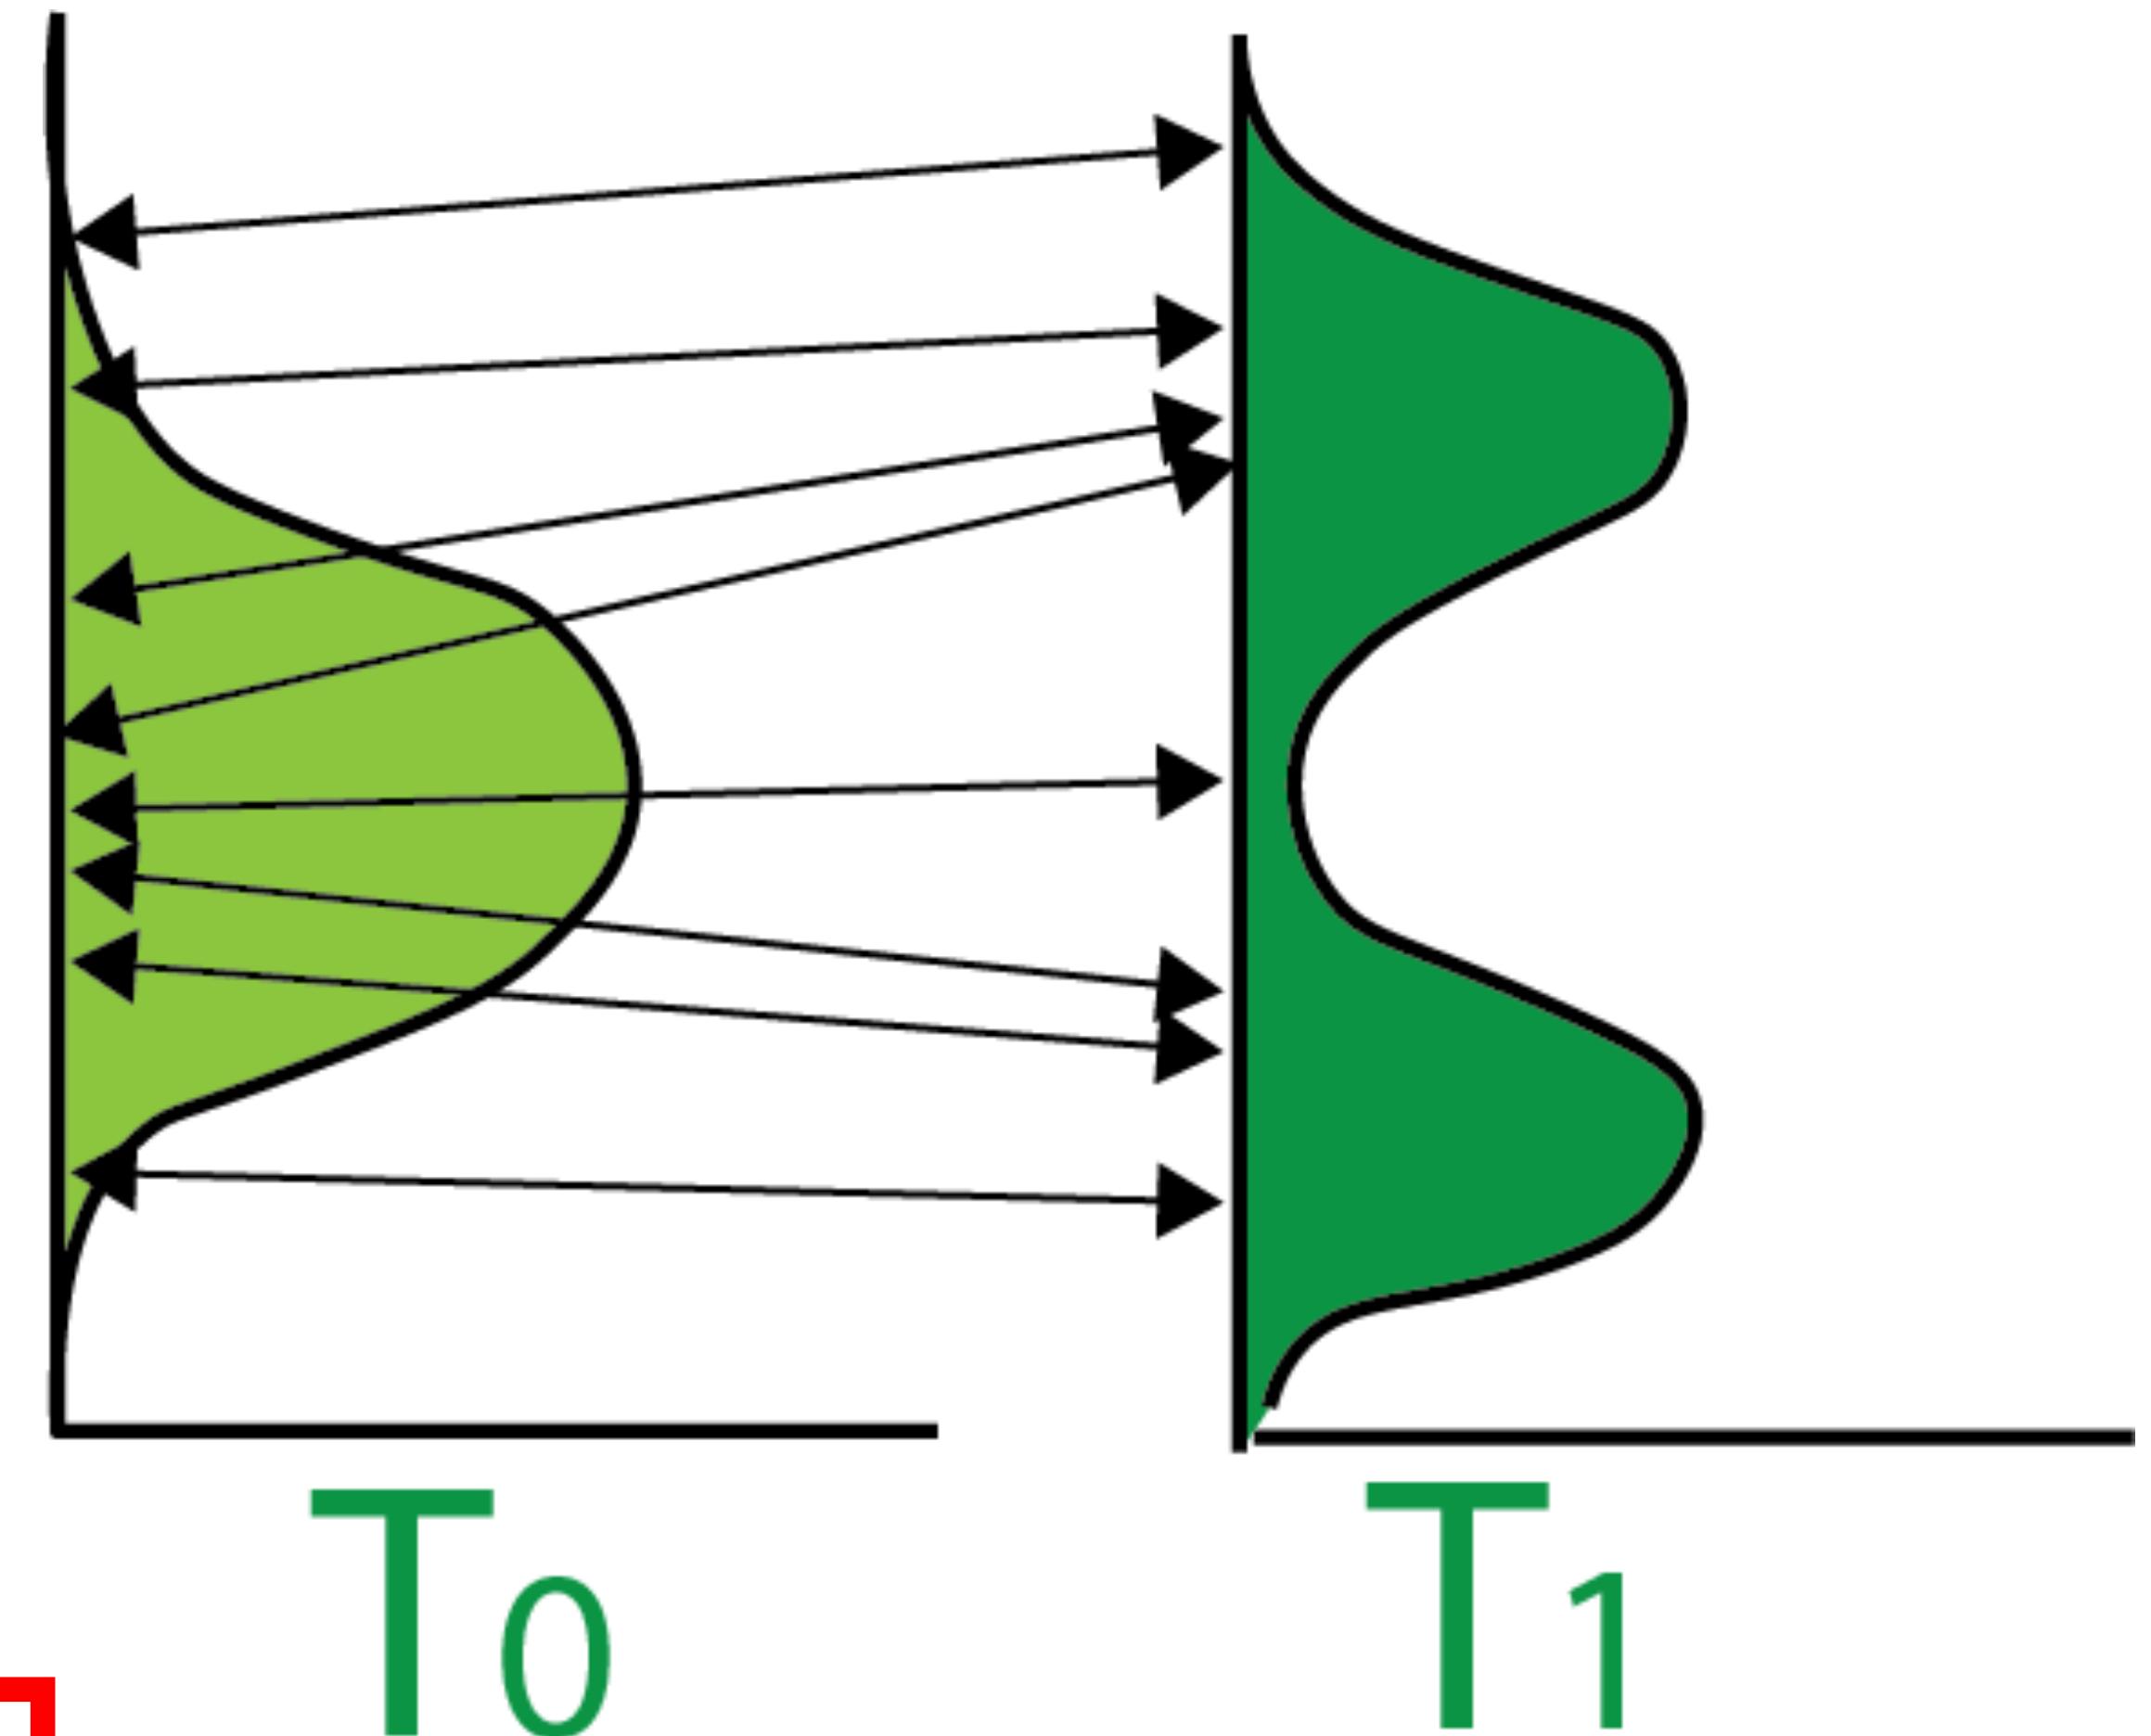
$$p_{t_0}(x) \sim \mathcal{N}(0, 1)$$

- Apply an invertible transformation(s)

$$x_{t_1} = U(x_{t_0})$$

- Use change of variables to calculate probability

$$\log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \log \det \left| \frac{\partial U}{\partial x_{t_0}} \right|$$



Deep Normalizing Flows (NFs)

Apply a series of transformations

$$x_{t_1} = U(x_{t_0}) \rightarrow x_{t_N} = u_N \circ u_{N-1} \circ \cdots \circ u_1(x_{t_0})$$

Use change of variables to calculate probability

$$\log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \log \det \left| \frac{\partial U}{\partial x_{t_0}} \right| \rightarrow \log p_{t_N}(x_{t_N}) = \log p_{t_0}(x_{t_0}) - \sum_{n=1}^N \log \det \left| \frac{\partial u_n}{\partial x_{t_{n-1}}} \right|$$

Continuous Normalizing Flows

Apply a series of transformations

$$x_{t_N} = u_N \circ u_{N-1} \circ \cdots \circ u_1(x_{t_0}) \rightarrow x_{t_1} = U(x_{t_0}) = \int_{t_0}^{t_1} u(x(t), t) dt$$

Use change of variables to calculate probability

$$\log p_{t_N}(x_{t_N}) = \log p_{t_0}(x_{t_0}) - \sum_{n=1}^N \log \det \left| \frac{\partial u_n}{\partial x_{t_{n-1}}} \right| \rightarrow \log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial u}{\partial x(t)} \right) dt$$

Continuous Normalizing Flows

Apply a series of transformations

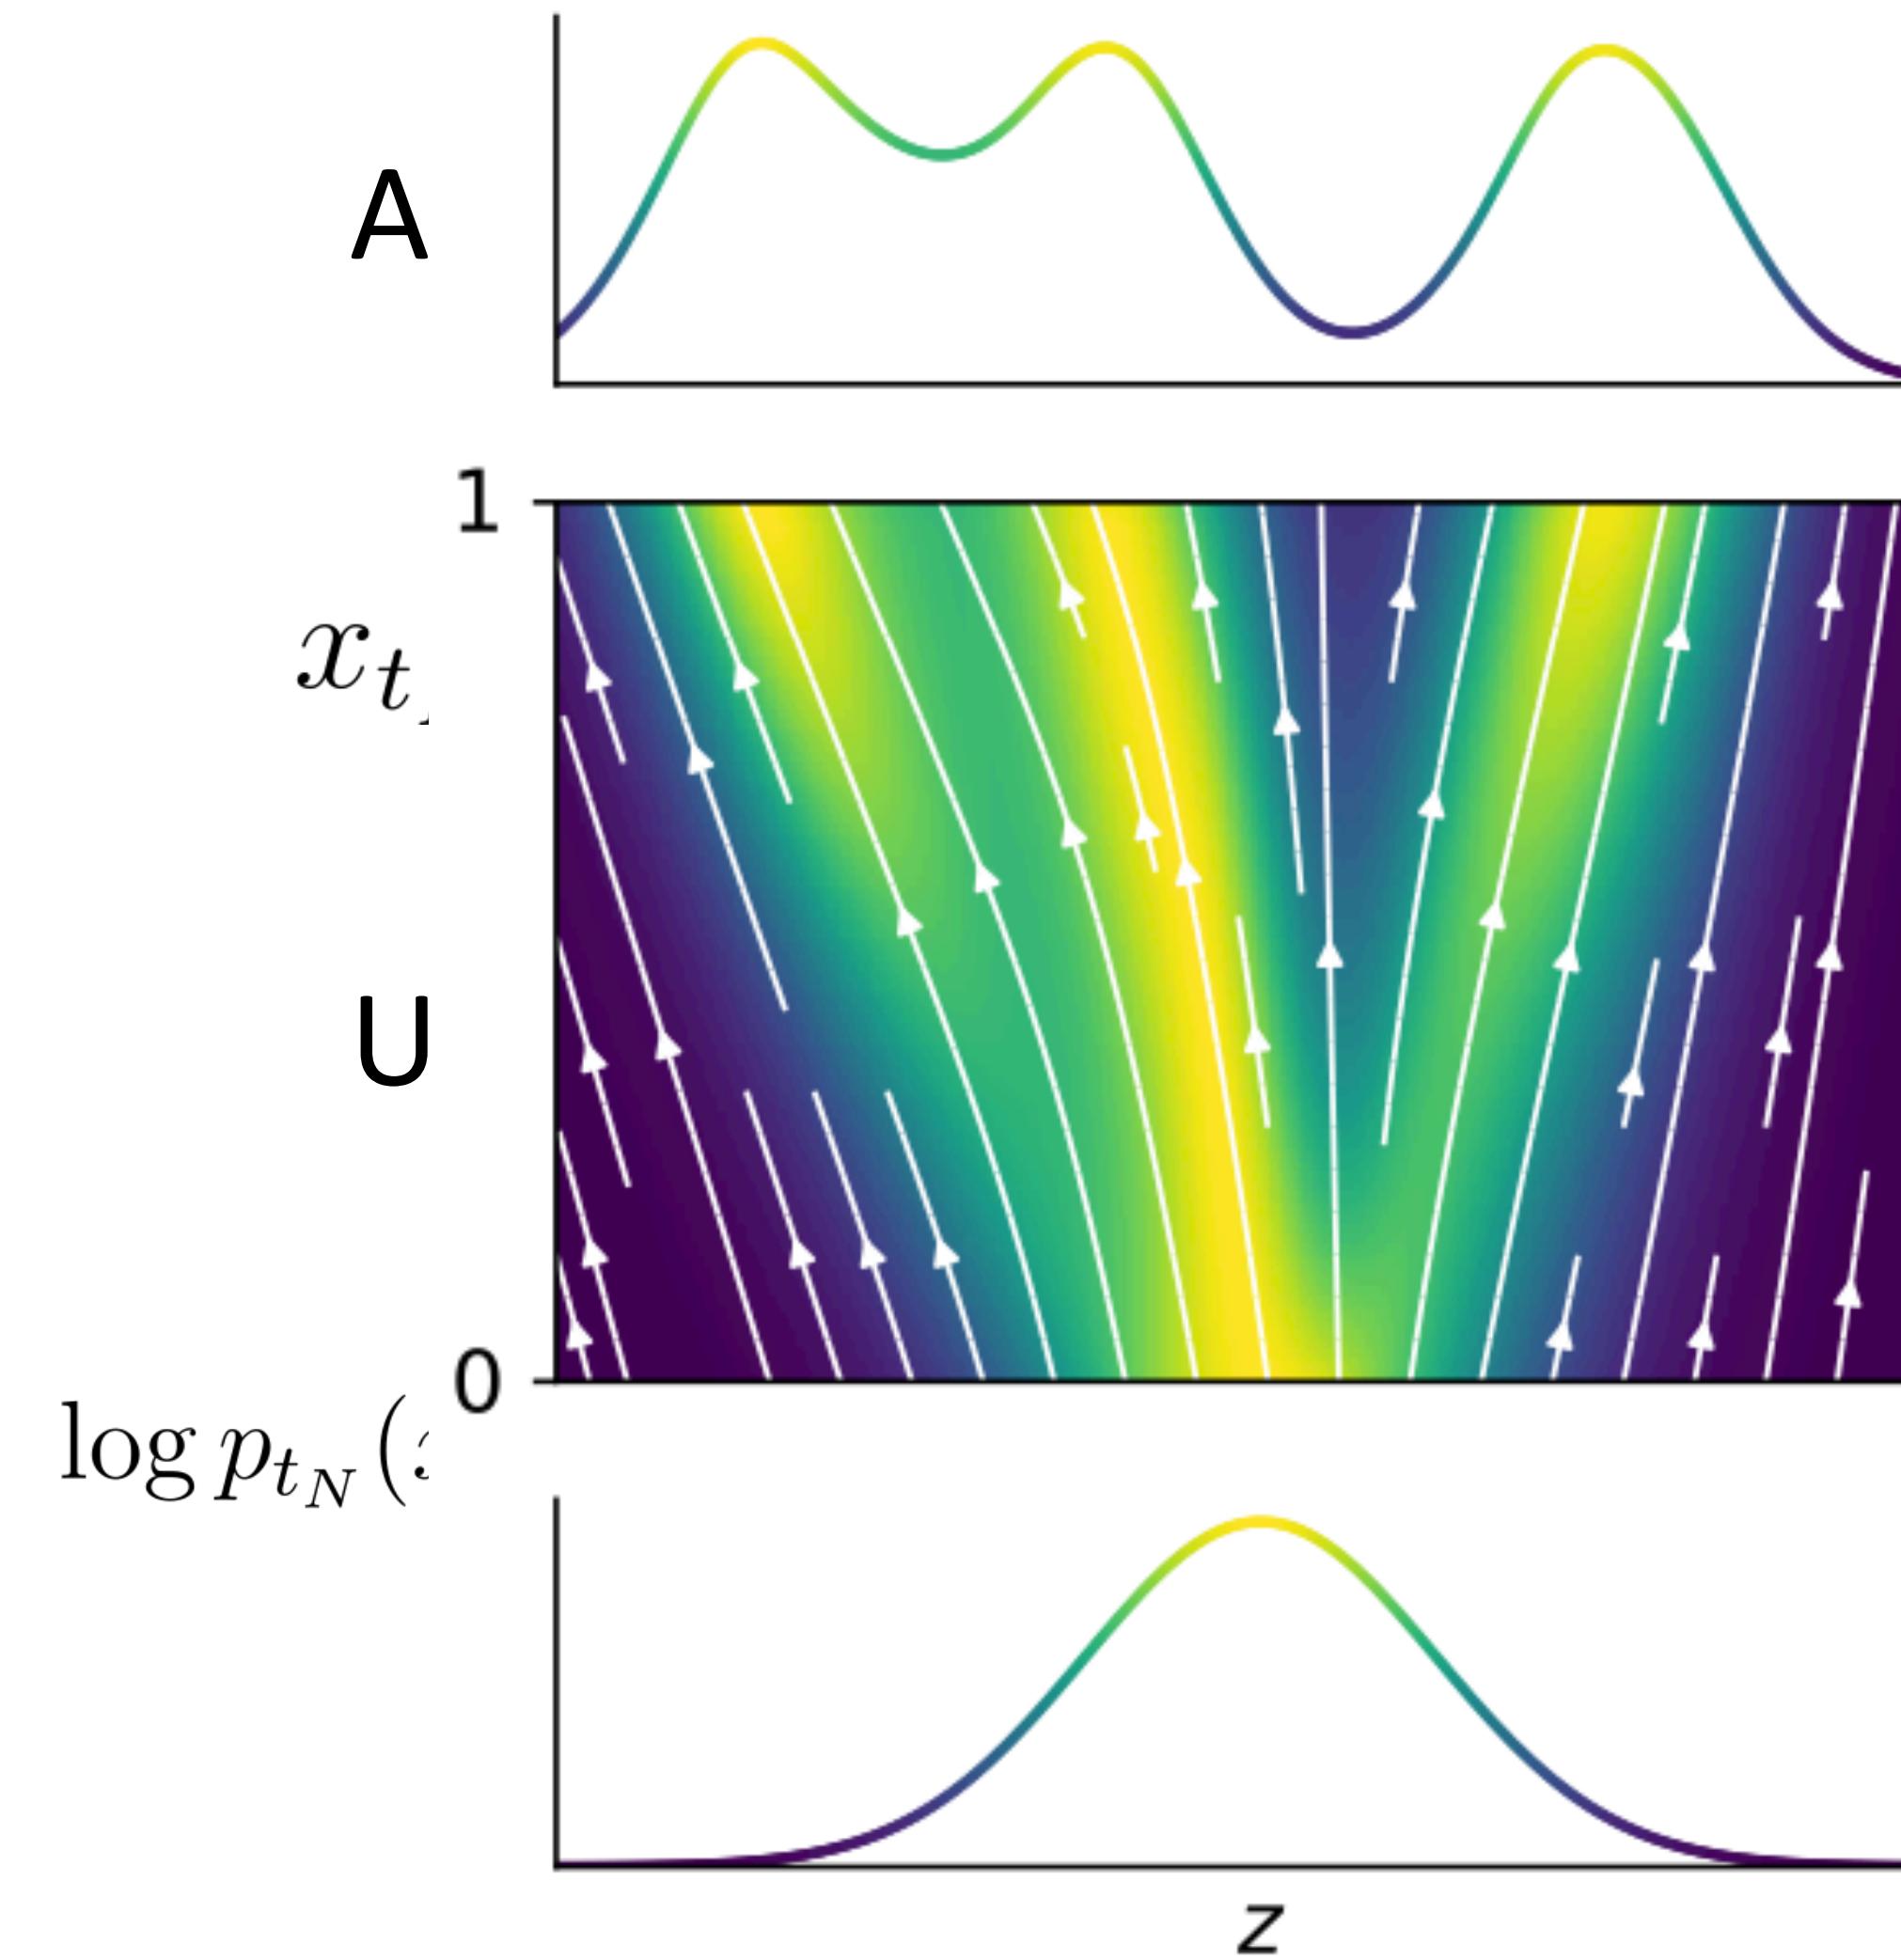
$$x_{t_N} = u_N \circ u_{N-1} \circ \cdots \circ u_1(x_{t_0}) \rightarrow x_{t_1} = U(x_{t_0}) = \int_{t_0}^{t_1} u(x(t), t) dt$$

Use change of variables to calculate probability

$$\log p_{t_N}(x_{t_N}) = \log p_{t_0}(x_{t_0}) - \sum_{n=1}^N \log \det \left| \frac{\partial u_n}{\partial x_{t_{n-1}}} \right| \rightarrow \log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial u}{\partial x(t)} \right) dt$$

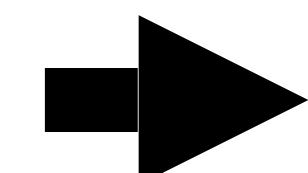
Constructs an invertible
transformation ψ for any f !

Continuous Normalizing Flows



ons

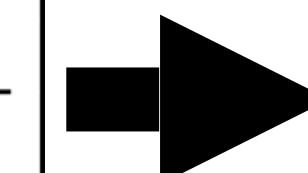
x_{t_0}



$$x_{t_1} = U(x_{t_0}) = \int_{t_0}^{t_1} u(x(t), t) dt$$

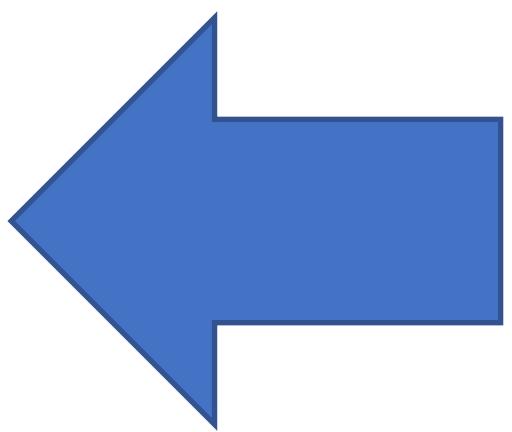
culate probability

$$\frac{\partial u_n}{r_{t_{n-1}}}$$



$$\log p_{t_1}(x_{t_1}) = \log p_{t_0}(x_{t_0}) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial u}{\partial x(t)} \right) dt$$

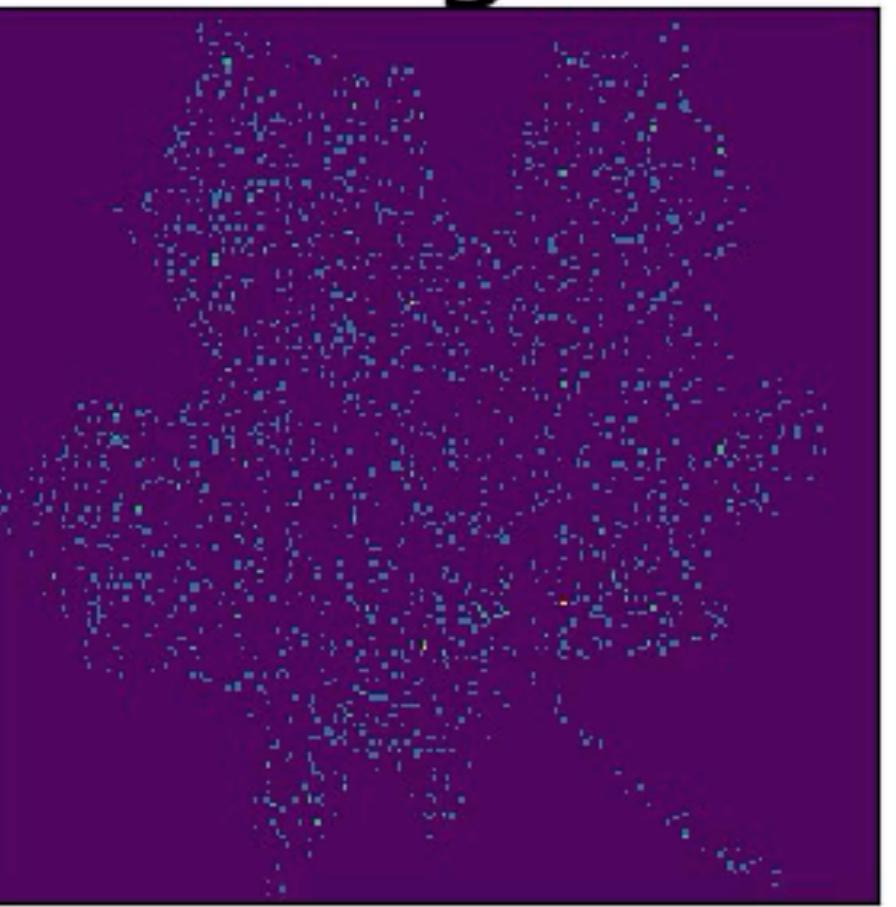
What does this look like?



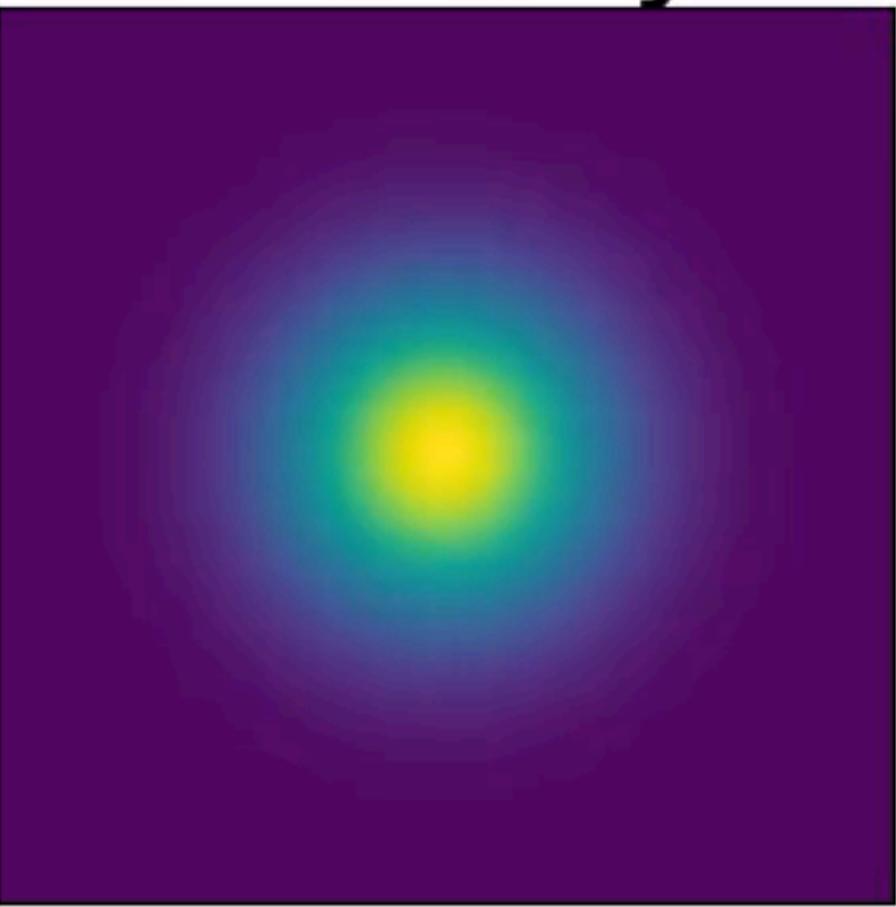
What does this look like?



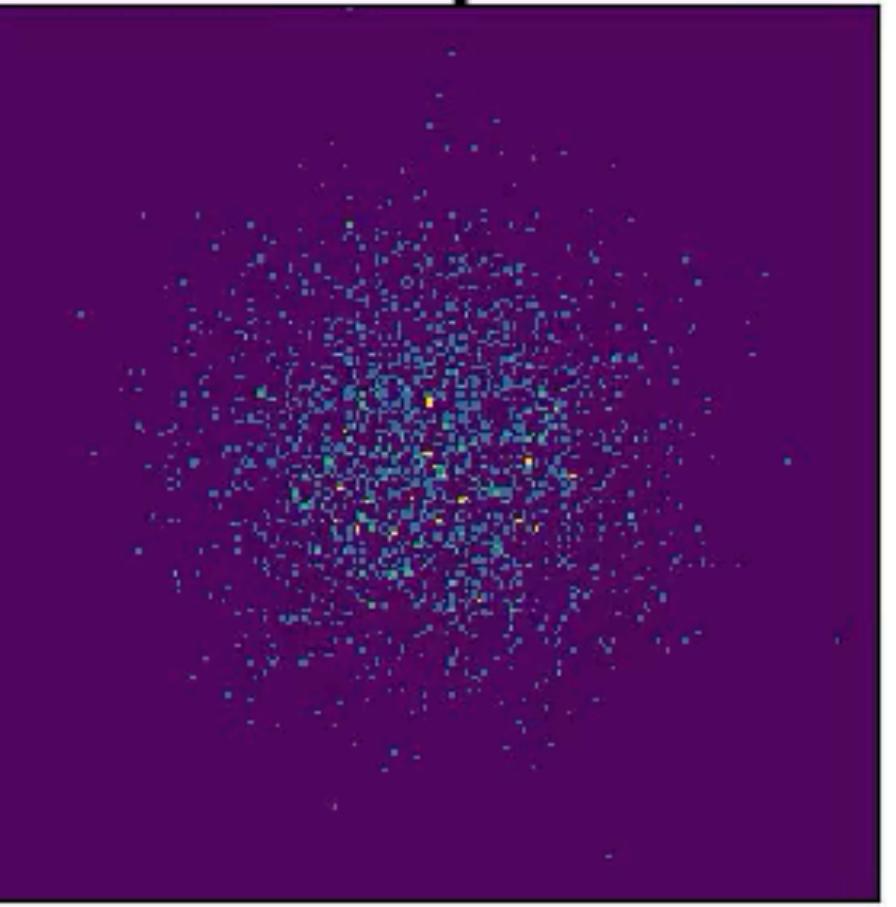
Target



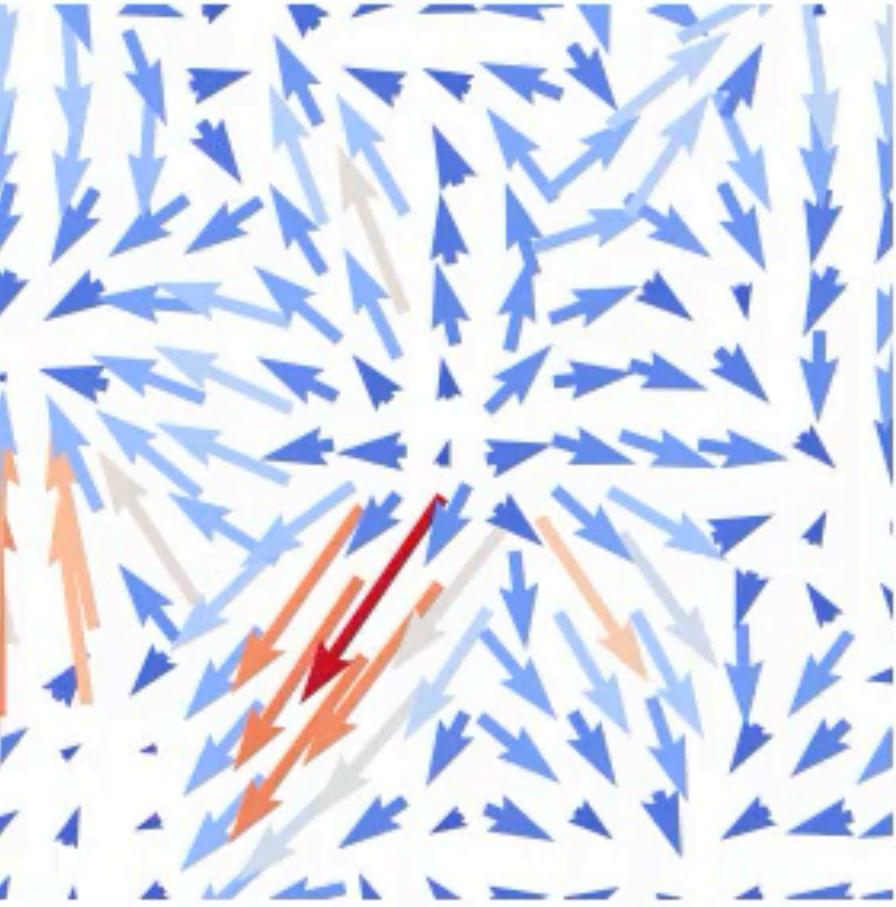
Density



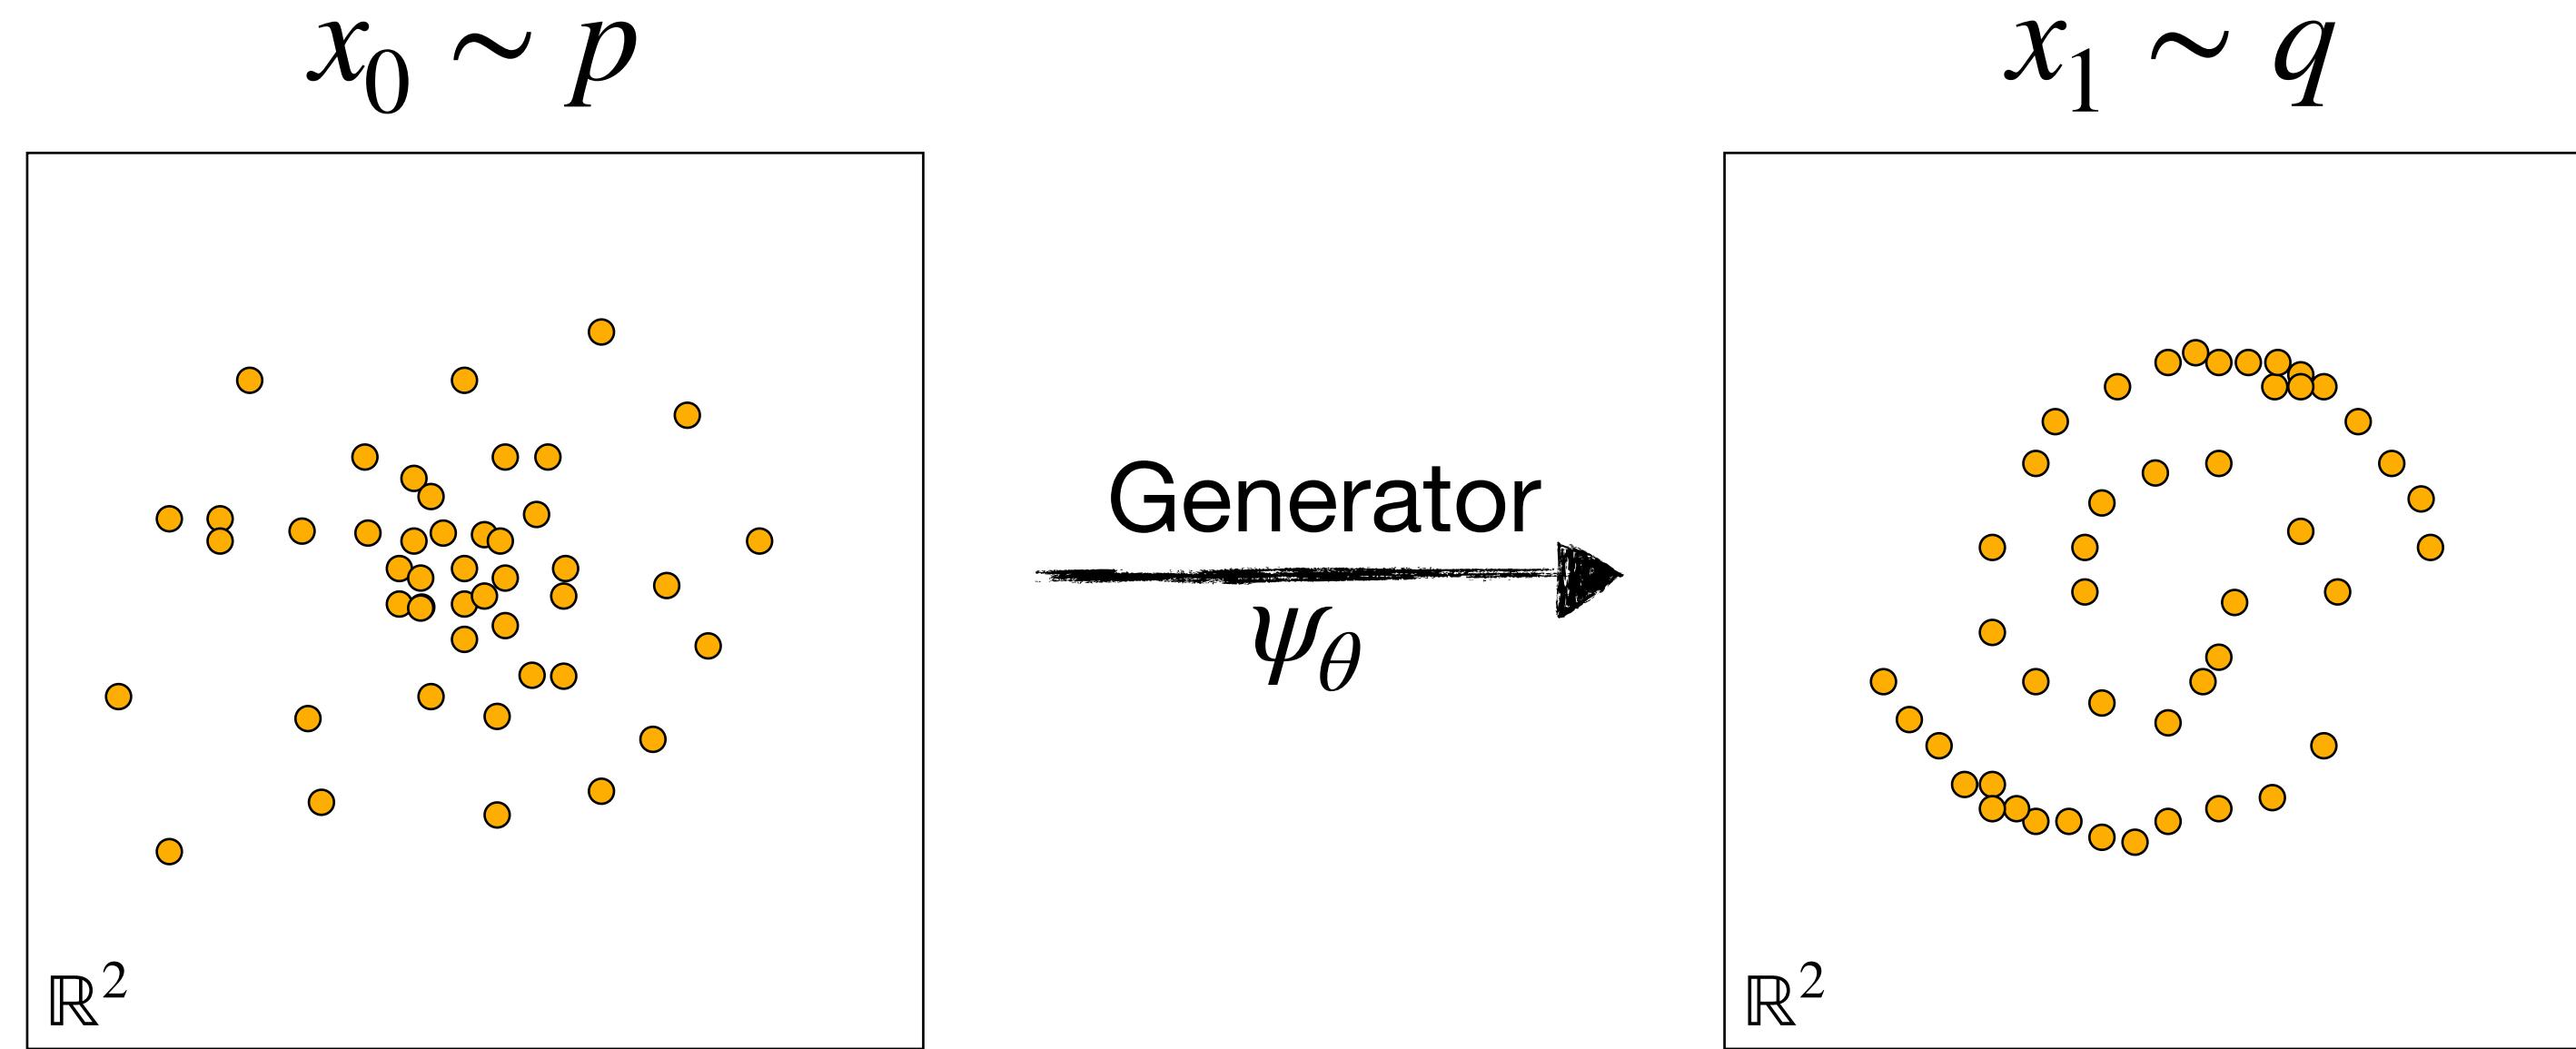
Samples



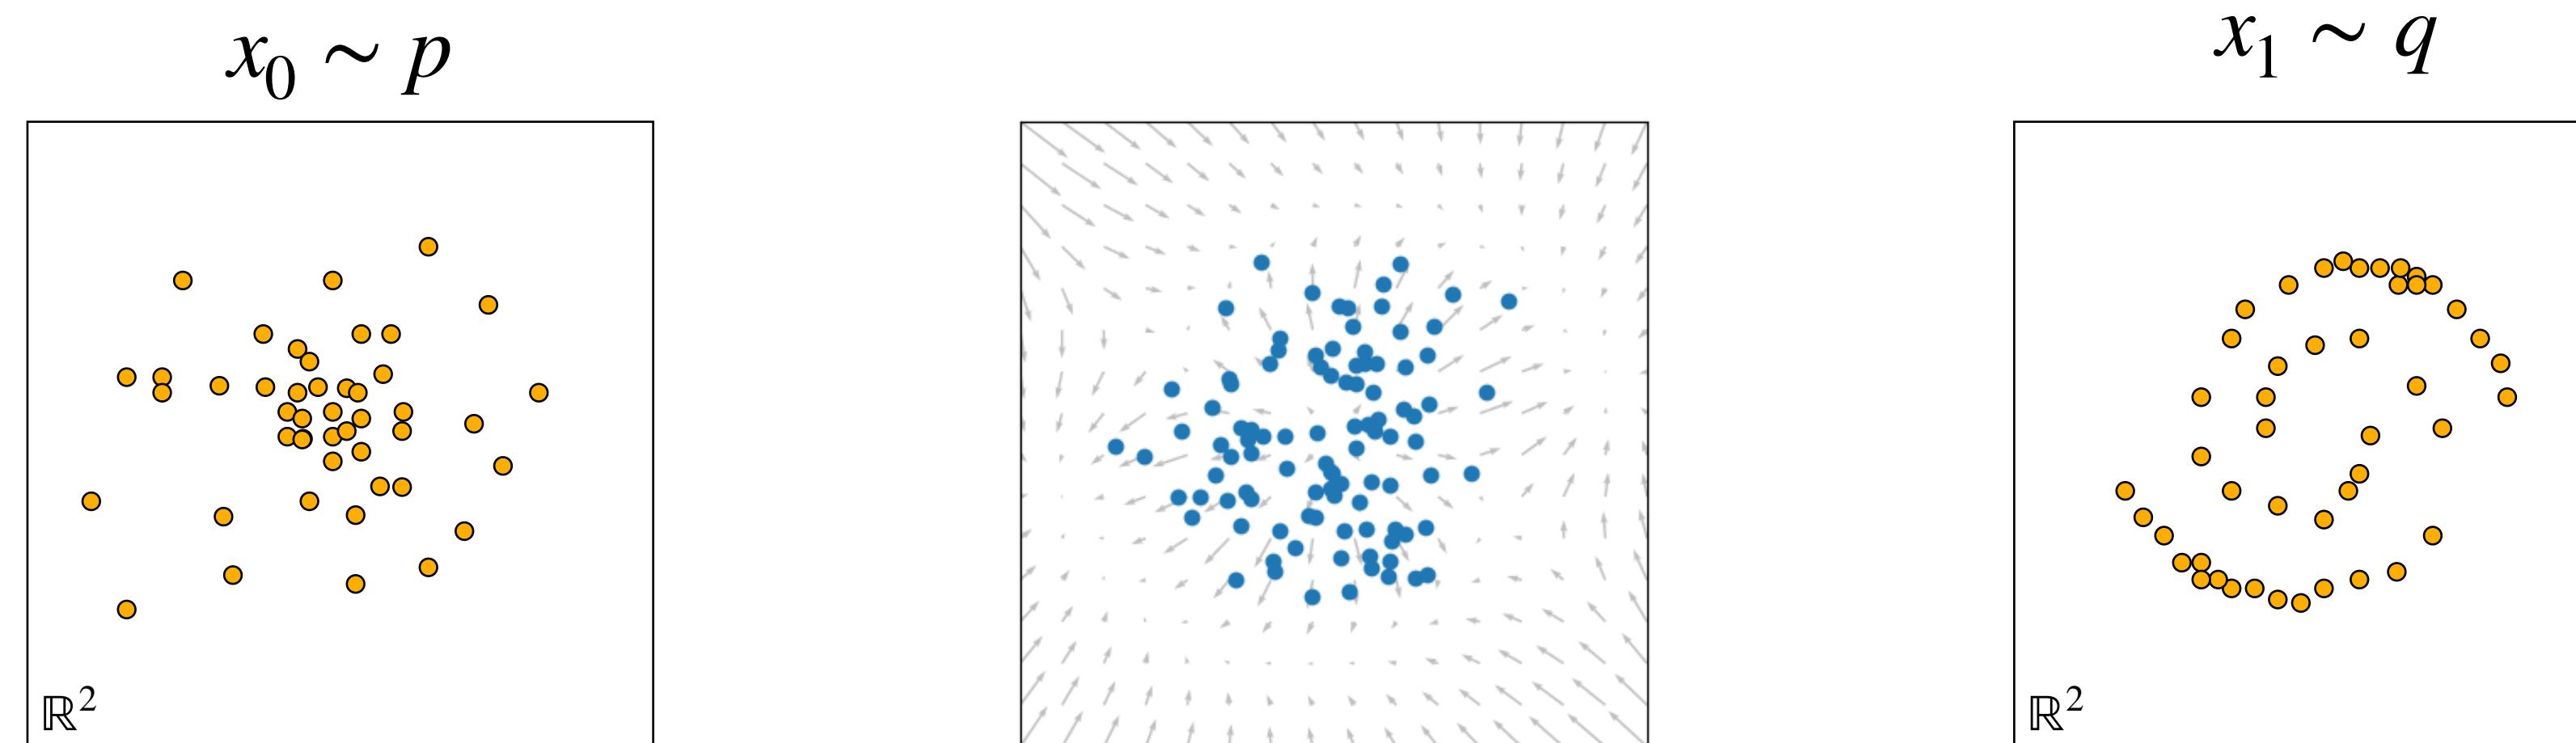
Vector Field



Focus: Generative Models as Dynamical Systems



Focus: Dynamical Systems as Generative Models



$$\psi_t : [0,1] \times \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Time-Dependent Generator

Sampling = Simulating

$$x_0 \sim p \quad \xrightarrow{\text{simulate}} \quad \text{simulate}(x_0, t) = \psi_t(x_0)$$

Deterministic and Stochastic Dynamics

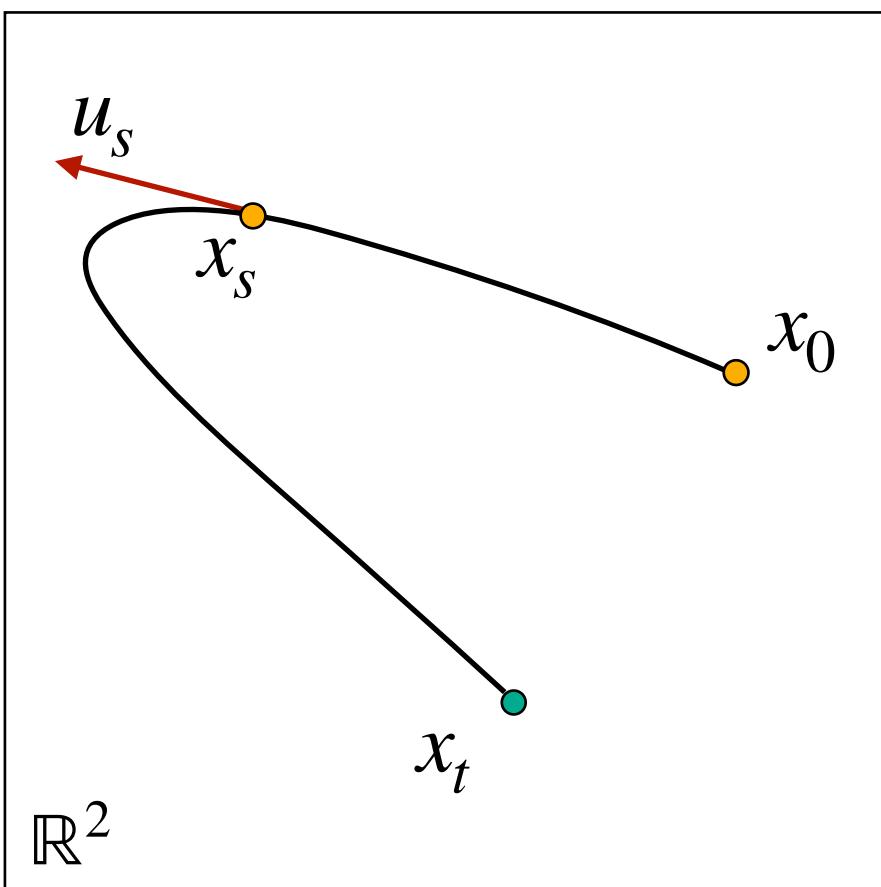
<p>Flows</p> <p>ODE</p> $dx_t = u_t(x_t)dt$ <p>Velocity field</p>	<p>Diffusion</p> <p>SDE</p> $dx_t = f_t(x_t)dt + g_t dw_t$ <p>Drift</p> <p>Diffusion Coefficient</p> <p>Brownian Motion</p>
---	---

Deterministic and Stochastic Dynamics

Flows

ODE
$$dx_t = u_t(x_t)dt$$

Velocity field



Deterministic

$$x_t = x_0 + \int_0^t u_s(x_s)ds$$

Diffusion

SDE
$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

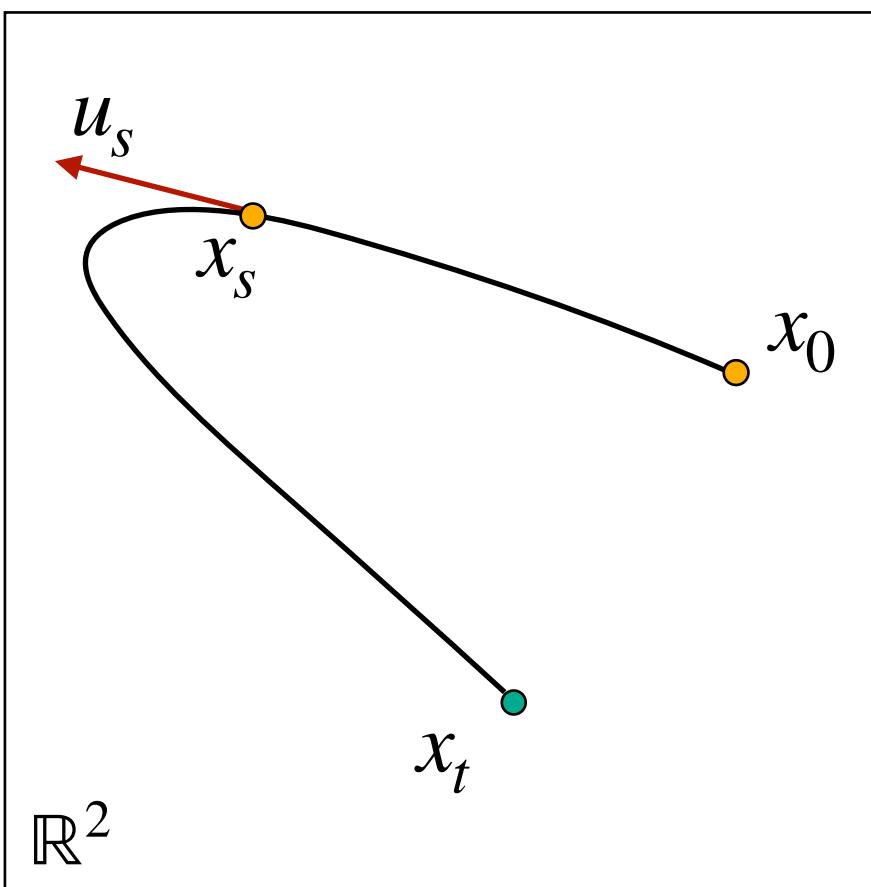
Brownian
Motion

Deterministic and Stochastic Dynamics

Flows

ODE
$$dx_t = u_t(x_t)dt$$

Velocity field



Deterministic

$$x_t = x_0 + \int_0^t u_s(x_s)ds$$

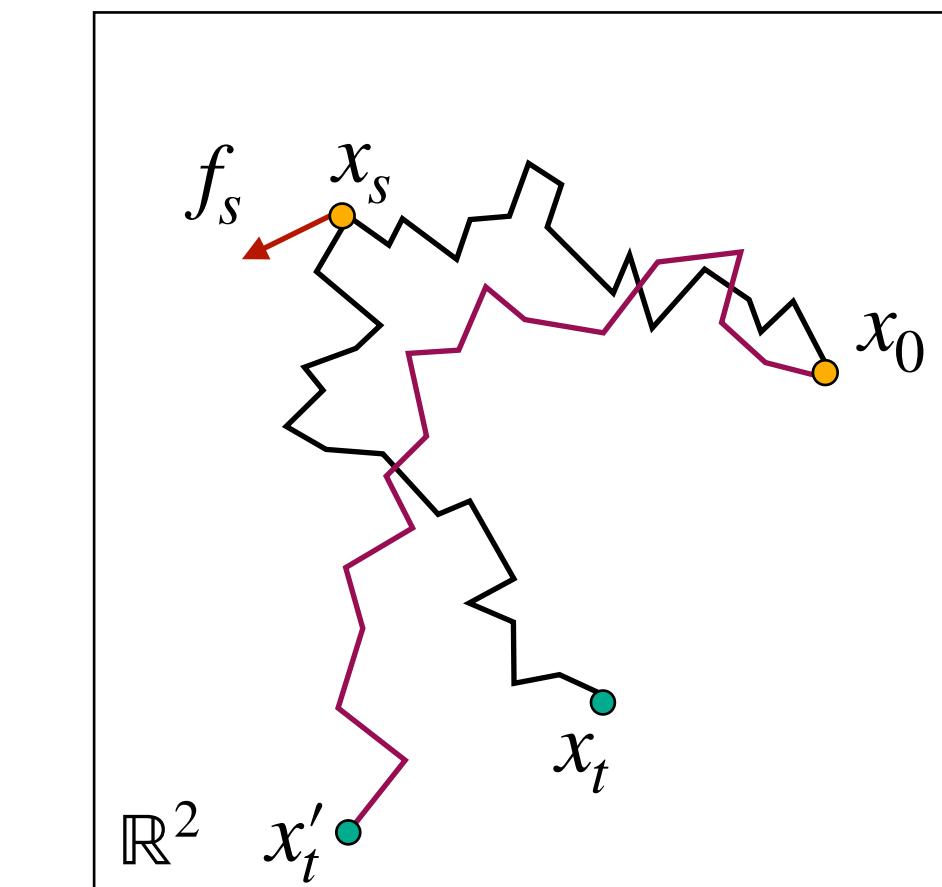
Diffusion

SDE
$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

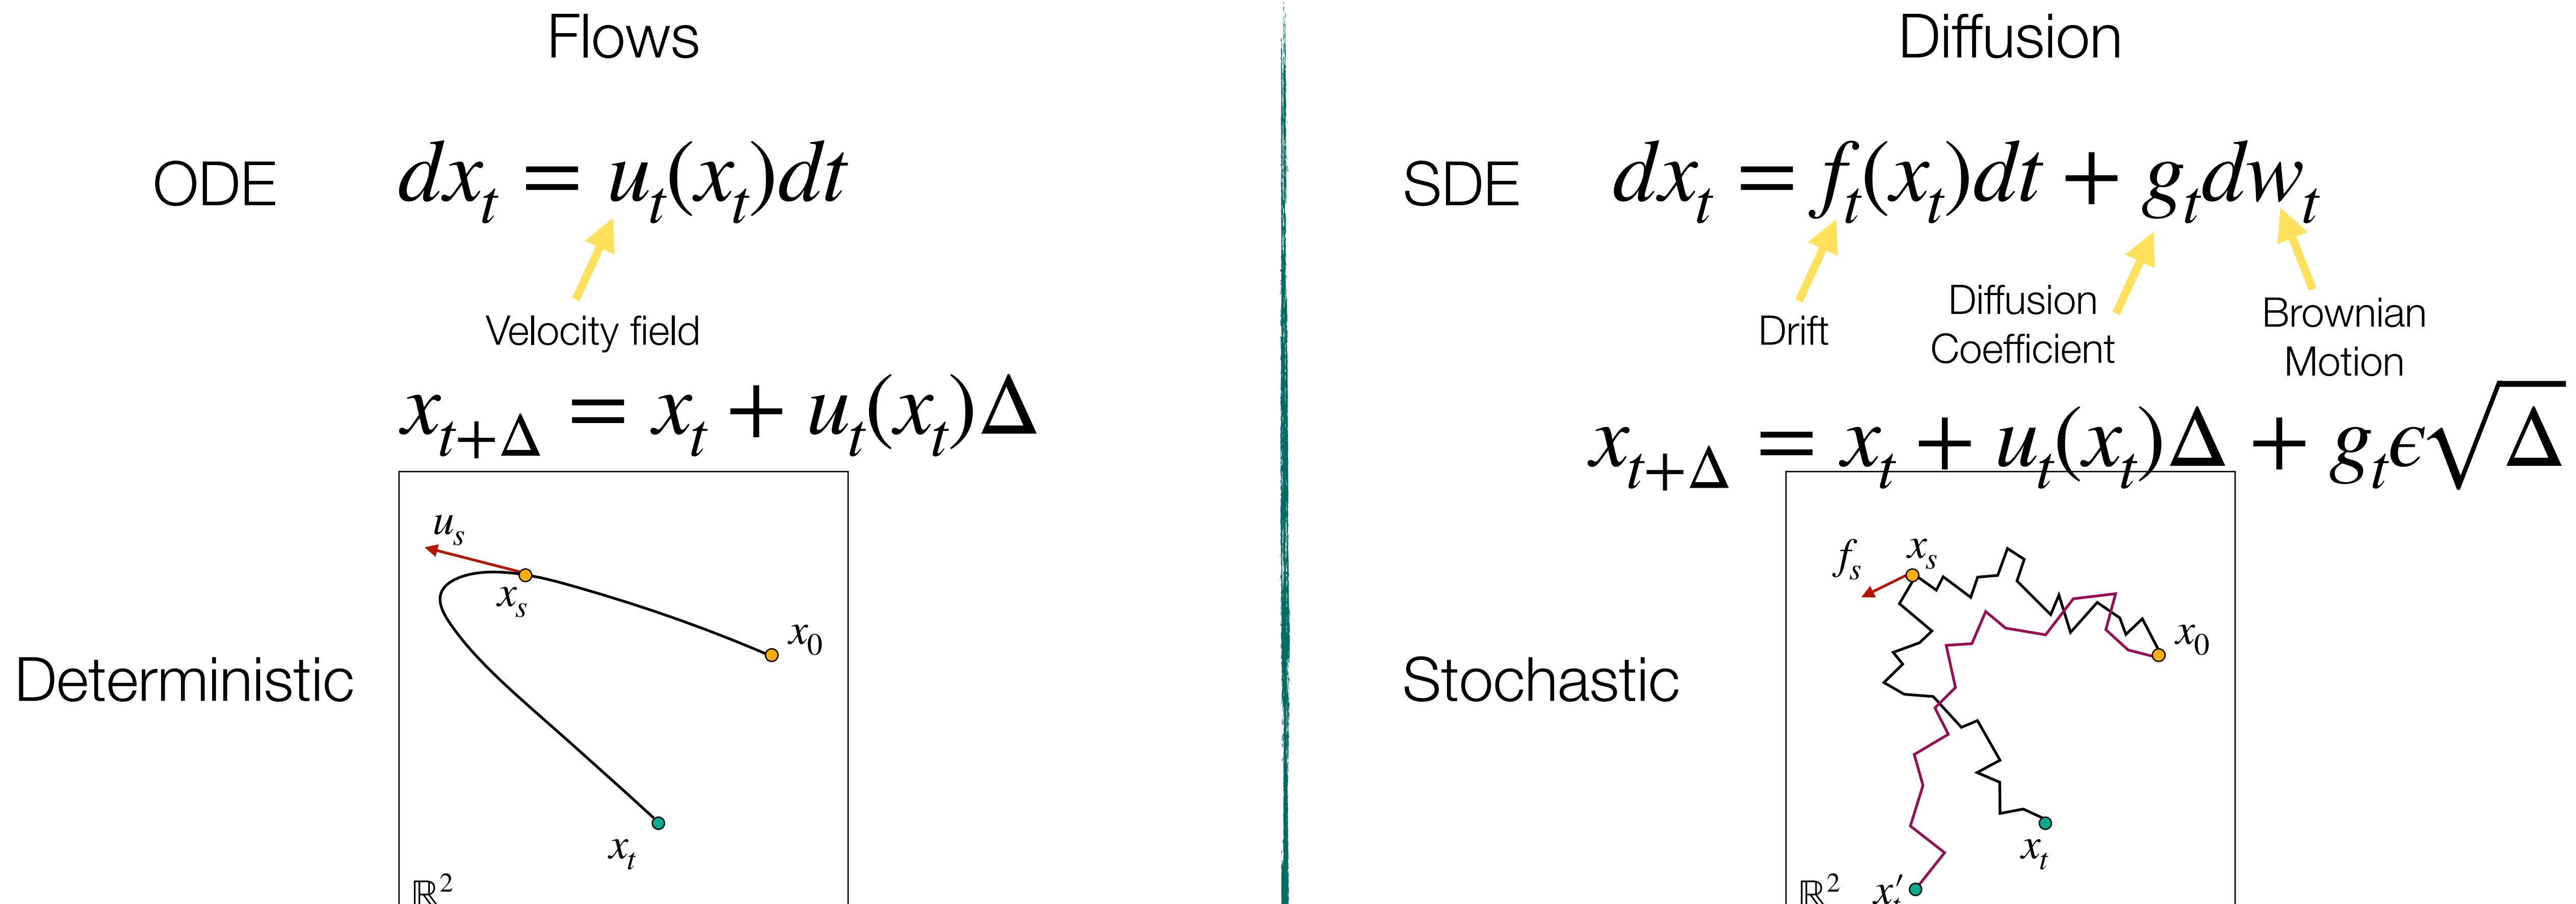
Brownian
Motion



Stochastic

$$x_t = x_0 + \int_0^t f_s(x_s)ds + \int_0^t g_s dw_s$$

Deterministic and Stochastic Dynamics



$$x_t = x_0 + \int_0^t u_s(x_s)ds$$

$$x_t = x_0 + \int_0^t f_s(x_s)ds + \int_0^t g_s dw_s$$

Where are the probabilities?

Flows

$$\text{ODE} \quad dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$

Diffusion

$$\text{SDE} \quad dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t f_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Where are the probabilities?

Flows

$$\text{ODE} \quad dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{u}_t)$$

Yes!

Can we build a generative model with these?

Diffusion

$$\text{SDE} \quad dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion
Coefficient

Brownian
Motion

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{f}_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing...



Diffusion and Score Models

SDE
$$dx_t = f_t(x_t)dt + g_t dw_t$$

The diagram shows three yellow arrows pointing from labels to specific terms in the SDE equation. The first arrow points from 'Drift' to $f_t(x_t)dt$. The second arrow points from 'Diffusion Coefficient' to g_t . The third arrow points from 'Brownian Motion' to dw_t .

Drift Diffusion Coefficient Brownian Motion

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{f}_t) + \frac{1}{2} \mathbf{g}_t^2 \nabla^2 p_t$$

Need one more thing...

Diffusion and Score Models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data → Noise

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{f}_t) + \frac{1}{2} \mathbf{g}_t^2 \nabla^2 p_t$$

Need one more thing...

Diffusion and Score Models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data → Noise

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t$$

Noise → Data

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{f}_t) + \frac{1}{2} g_t^2 \nabla^2 p_t$$

Need one more thing... The Score!

Diffusion and Score Models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t \quad \text{Noise} \rightarrow \text{Data}$$

Learn the score by regressing to conditional scores:

$$\min_{\theta} \mathbb{E}_{p_{data}, p_t(x|x_{data})} [\|s_t^\theta(x) - \nabla \log p_t(x|x_{data})\|^2]$$

Simulation-free

Known SDEs: Variance Exploding
Variance Preserving

Where are the probabilities?

Flows

ODE
$$dx_t = u_t(x_t)dt$$

Velocity field

The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{u}_t)$$

Yes!

Diffusion

SDE
$$dx_t = f_t(x_t)dt + g_t dw_t$$

Drift

Diffusion Coefficient

Brownian Motion

The Fokker-Planck Equation

$$\partial_t p_t = - \operatorname{div}(p_t \mathbf{f}_t) + \frac{1}{2} \mathbf{g}_t^2 \nabla^2 p_t$$

Learn: score $\nabla \log p_t$

- Only Gaussian source
- Solution asymptotically reaches source

Where are the probabilities?

Flows

ODE

$$dx_t = u_t(x_t)dt$$

Velocity field

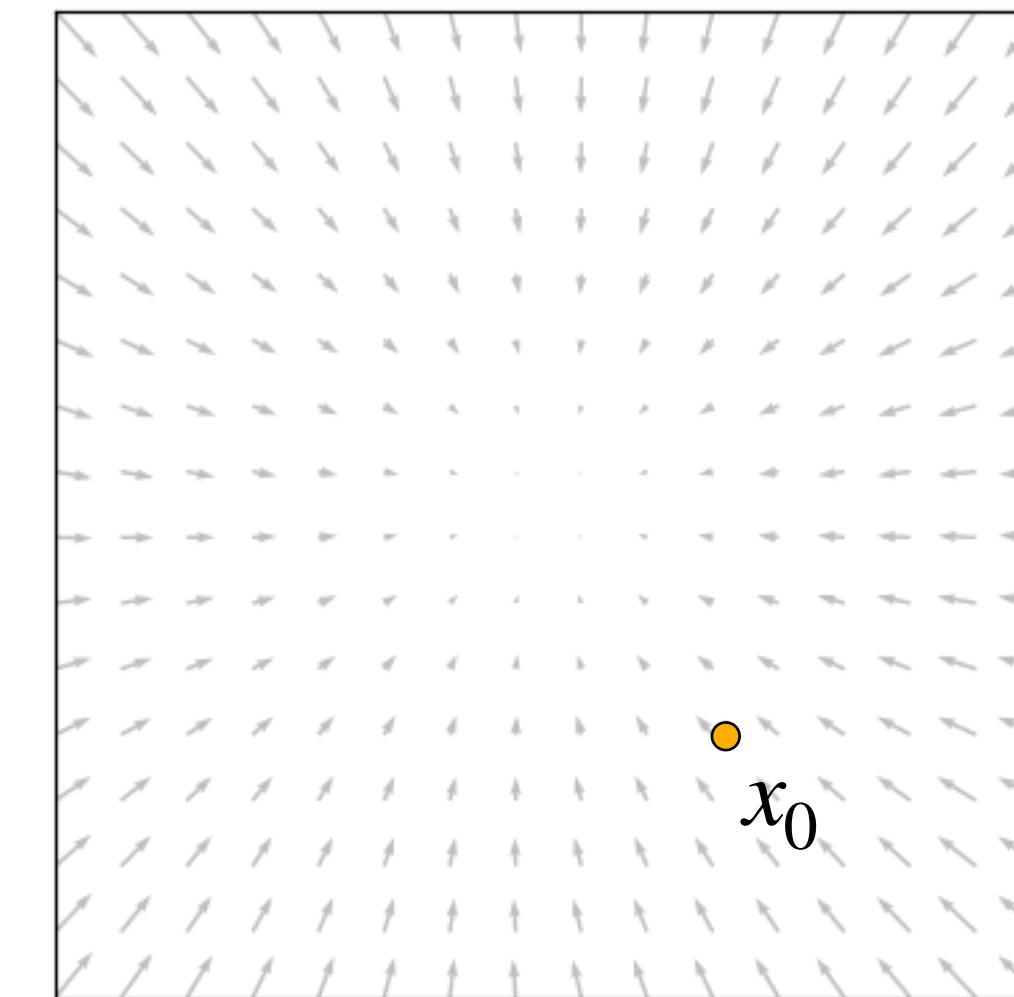
The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t \textcolor{orange}{u}_t)$$

Yes!

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$



$\psi_t(x)$ is smooth with smooth
inverse defined by $-u_t(x)$

Where are the probabilities?

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$

The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$

Learn: velocity field u_t

- Universal transformation between densities
- Defined on finite time interval

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Diffusion

Drift

Diffusion
Coefficient

Brownian
Motion

The Liouville Equation

$$\partial_t p_t = - \operatorname{div}\left(p_t\left(f_t - \frac{1}{2}g_t^2 \nabla \log p_t\right)\right)$$

Learn: score $\nabla \log p_t$

- Only Gaussian source
- Solution asymptotically reaches source

Where are the probabilities?

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$

The Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$

Learn: velocity field u_t

- Universal transformation between densities
- Defined on finite time interval

SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Diffusion

Drift

Brownian Motion

Diffusion Coefficient

The Liouville Equation

$$\partial_t p_t = - \operatorname{div}\left(p_t\left(f_t - \frac{1}{2}g_t^2 \nabla \log p_t\right)\right)$$

Learn: score $\nabla \log p_t$

- Only Gaussian source
- Solution asymptotically reaches source

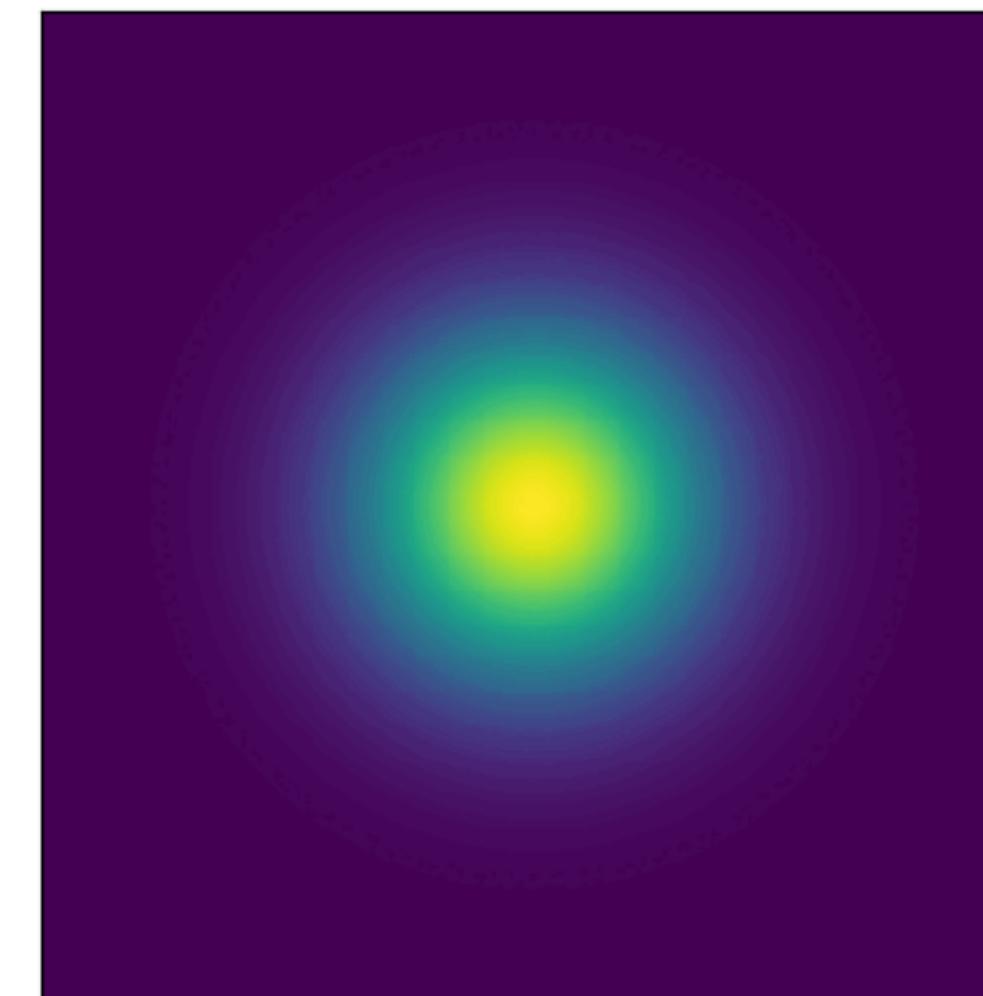
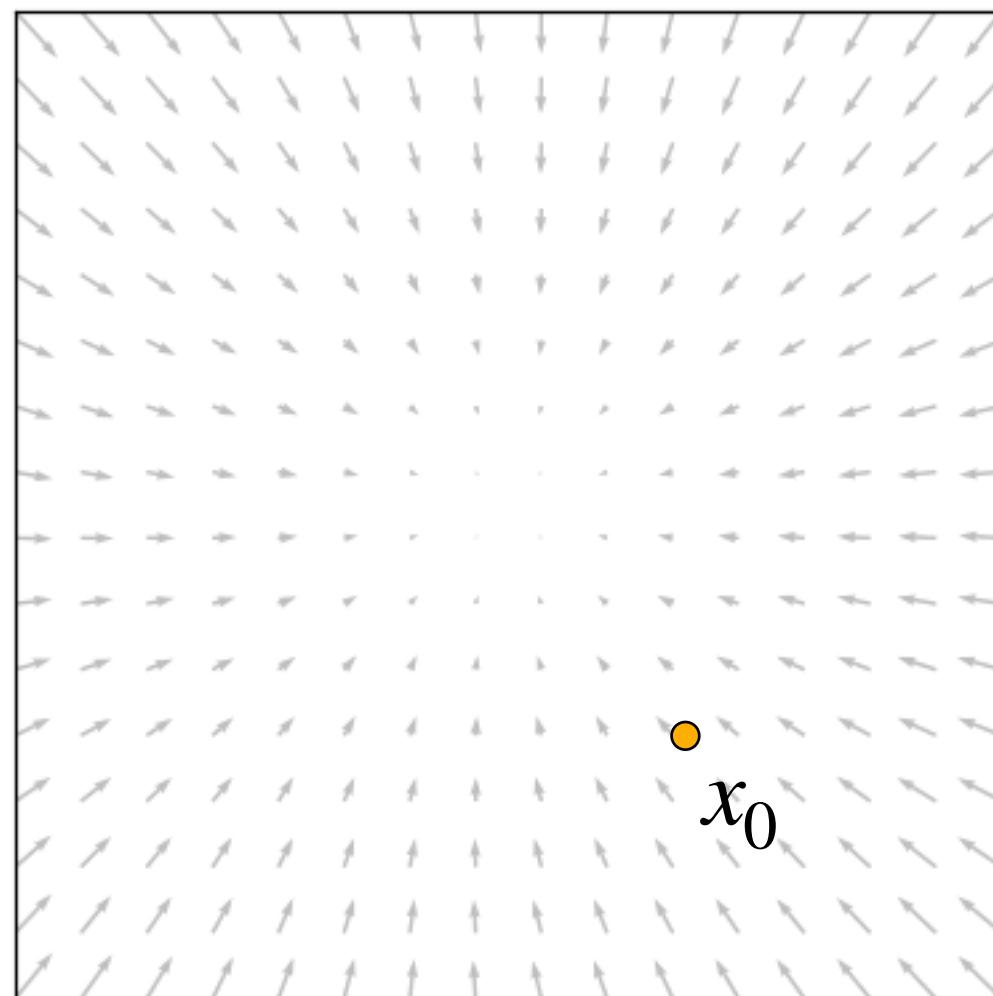
Flows as Generative Models

Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$

Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$



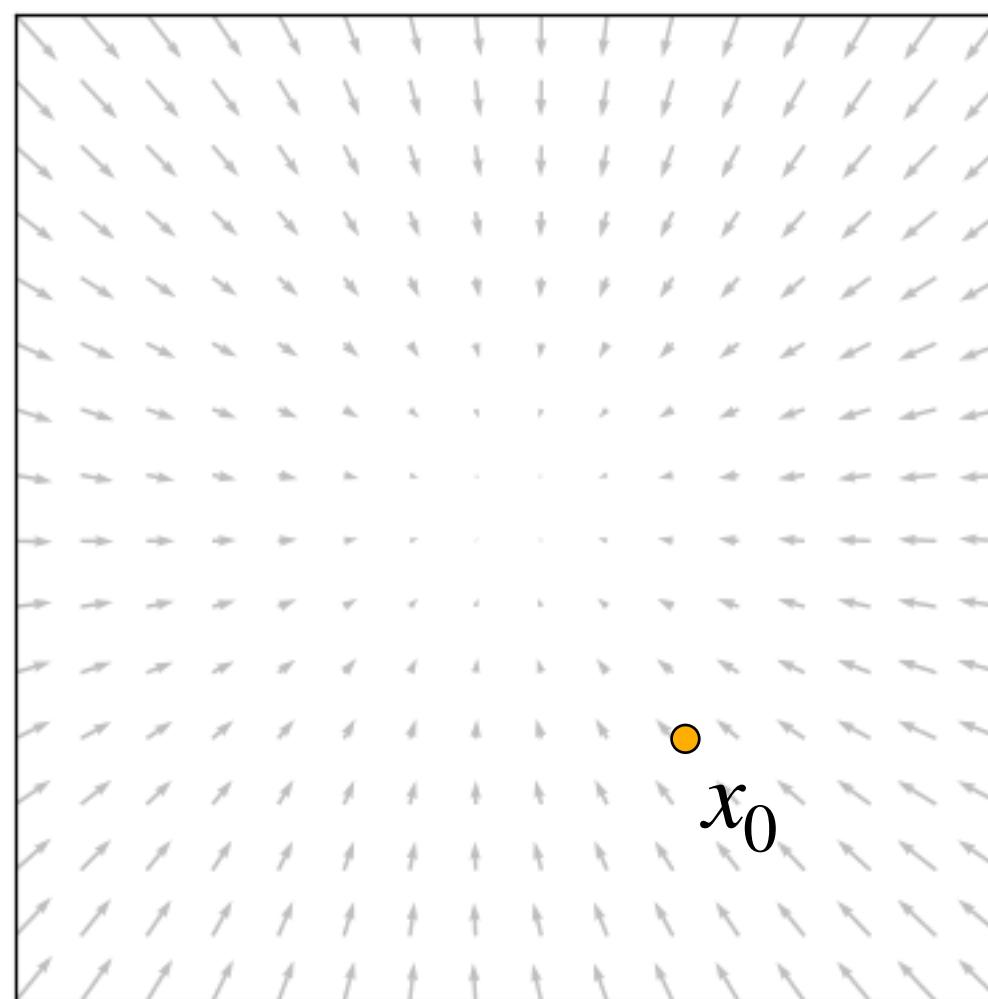
Learn: velocity field u_t

Goal: find velocity field u_t s.t. $p_1 \approx q$

Training with Simulation

Flow ODE

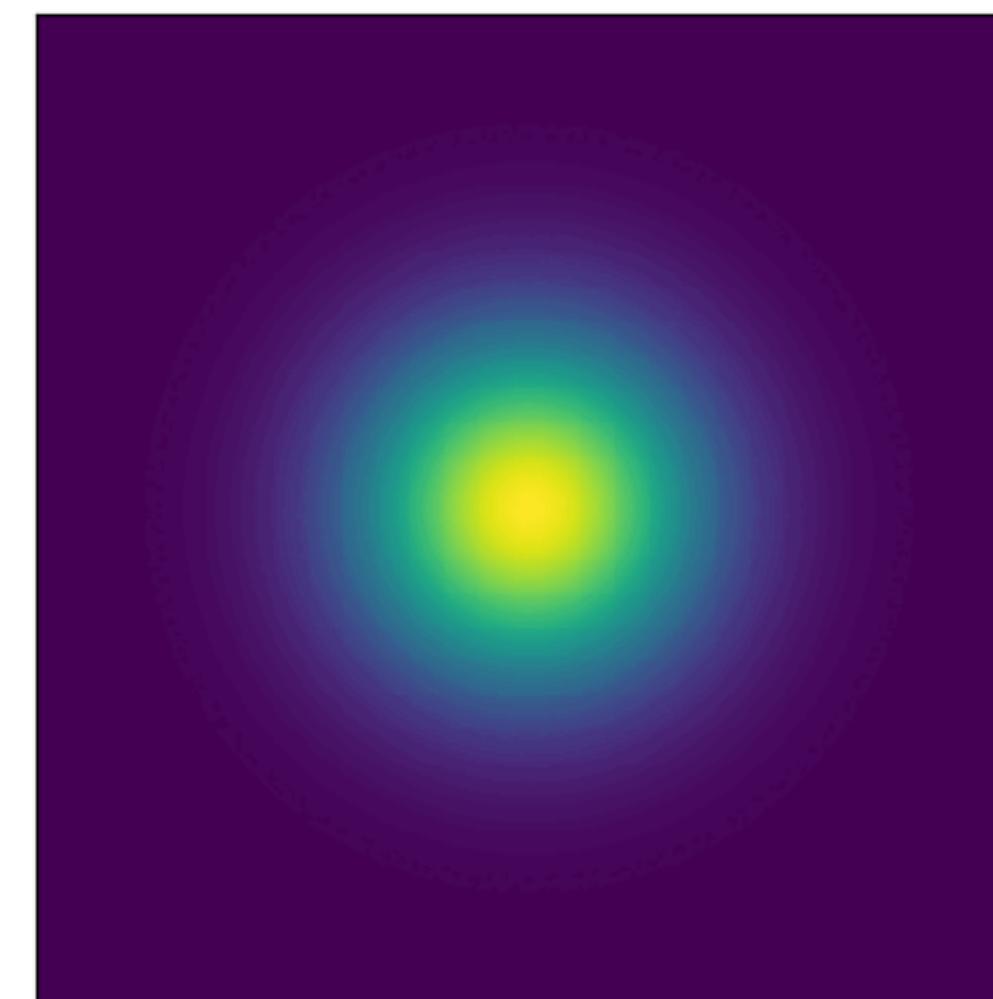
$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$



Learn: velocity field u_t

Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$



Log-likelihood computation

$$\log p_1(x_1) = \log p(x_0) + \int_1^0 \operatorname{div}(u_t(x_t)) dt$$

$$x_t = x_1 + \int_1^t u_s(x_s) ds$$

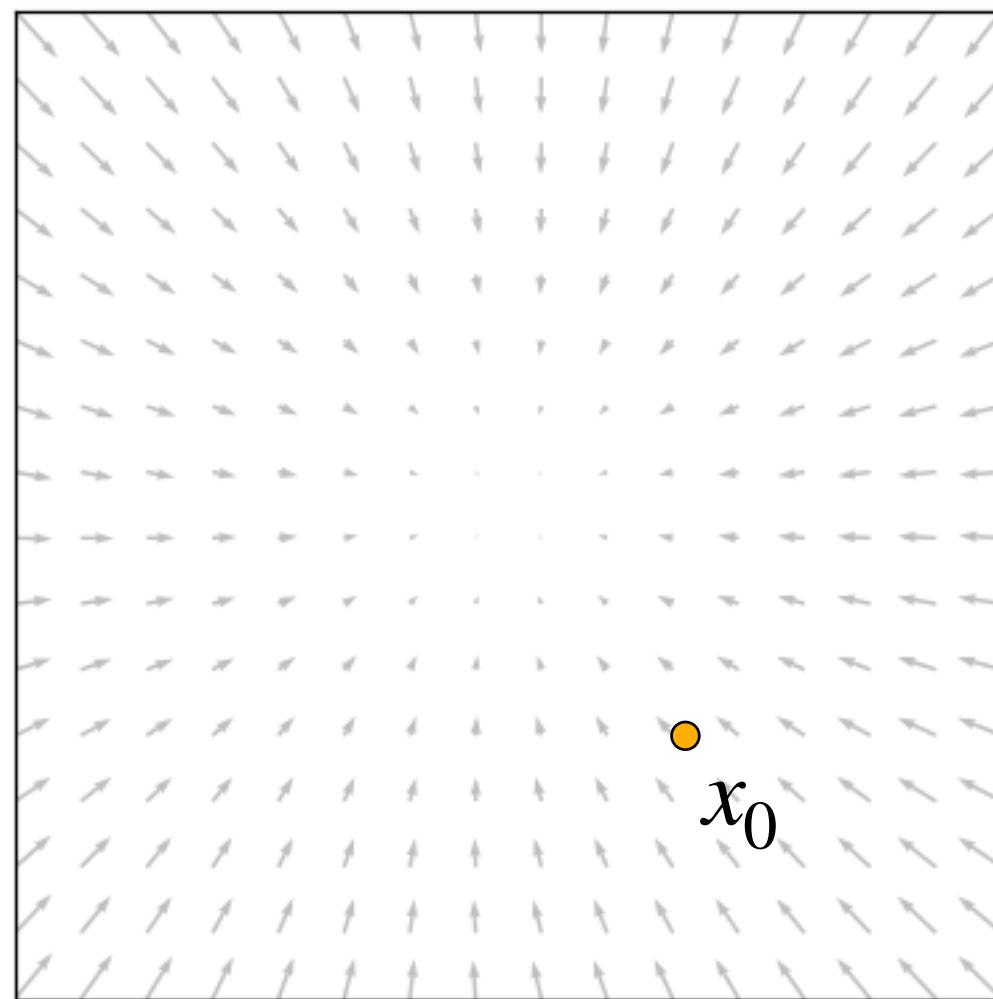
Maximum Likelihood Objective

$$D_{\text{KL}}(q \parallel p_1) = - \mathbb{E}_{x \sim q} \log p_1(x) + c$$

Training with Simulation

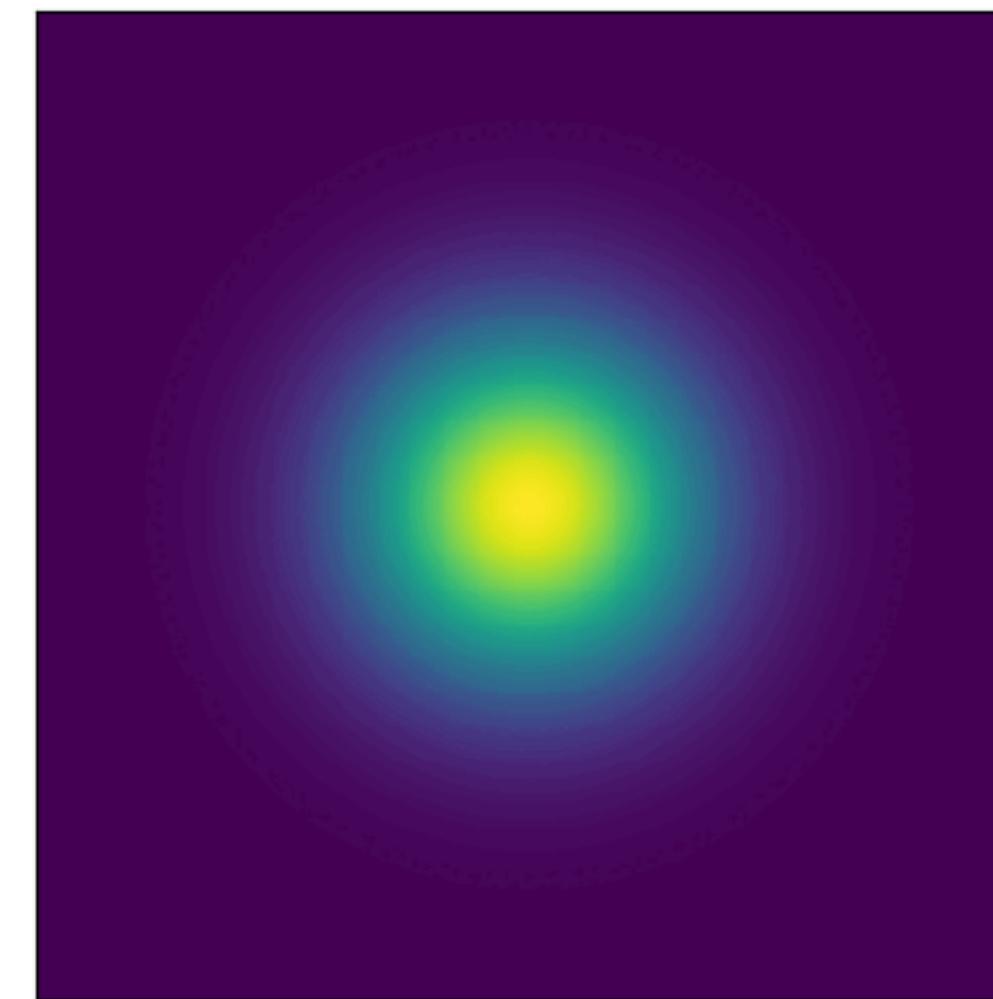
Flow ODE

$$\dot{\psi}_t(x_0) = u_t(\psi_t(x_0))$$



Continuity Equation

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$



Log-likelihood computation

$$\log p_1(x_1) = \log p(x_0) + \int_1^0 \operatorname{div}(u_t(x_t)) dt$$

$$x_t = x_1 + \int_1^t u_s(x_s) ds$$

Maximum Likelihood Objective

Vs. Normalizing Flows:

- Avoids invertible architectures!
- Requires computing divergence

$$D_{\text{KL}}(q \parallel p_1) = - \mathbb{E}_{x \sim q} \log p_1(x) + c$$

Training with Simulation

ODE $dx_t = u_t(x_t)dt$

Flows

Velocity field

Log-likelihood computation

$$\log p_1(x_1) = \log p(x_0) + \int_1^0 \operatorname{div}(u_t(x_t))dt$$
$$x_t = x_1 + \int_1^t u_s(x_s)ds$$

The Continuity Equation

$$\partial_t p_t = -\operatorname{div}(p_t u_t)$$

Requires:

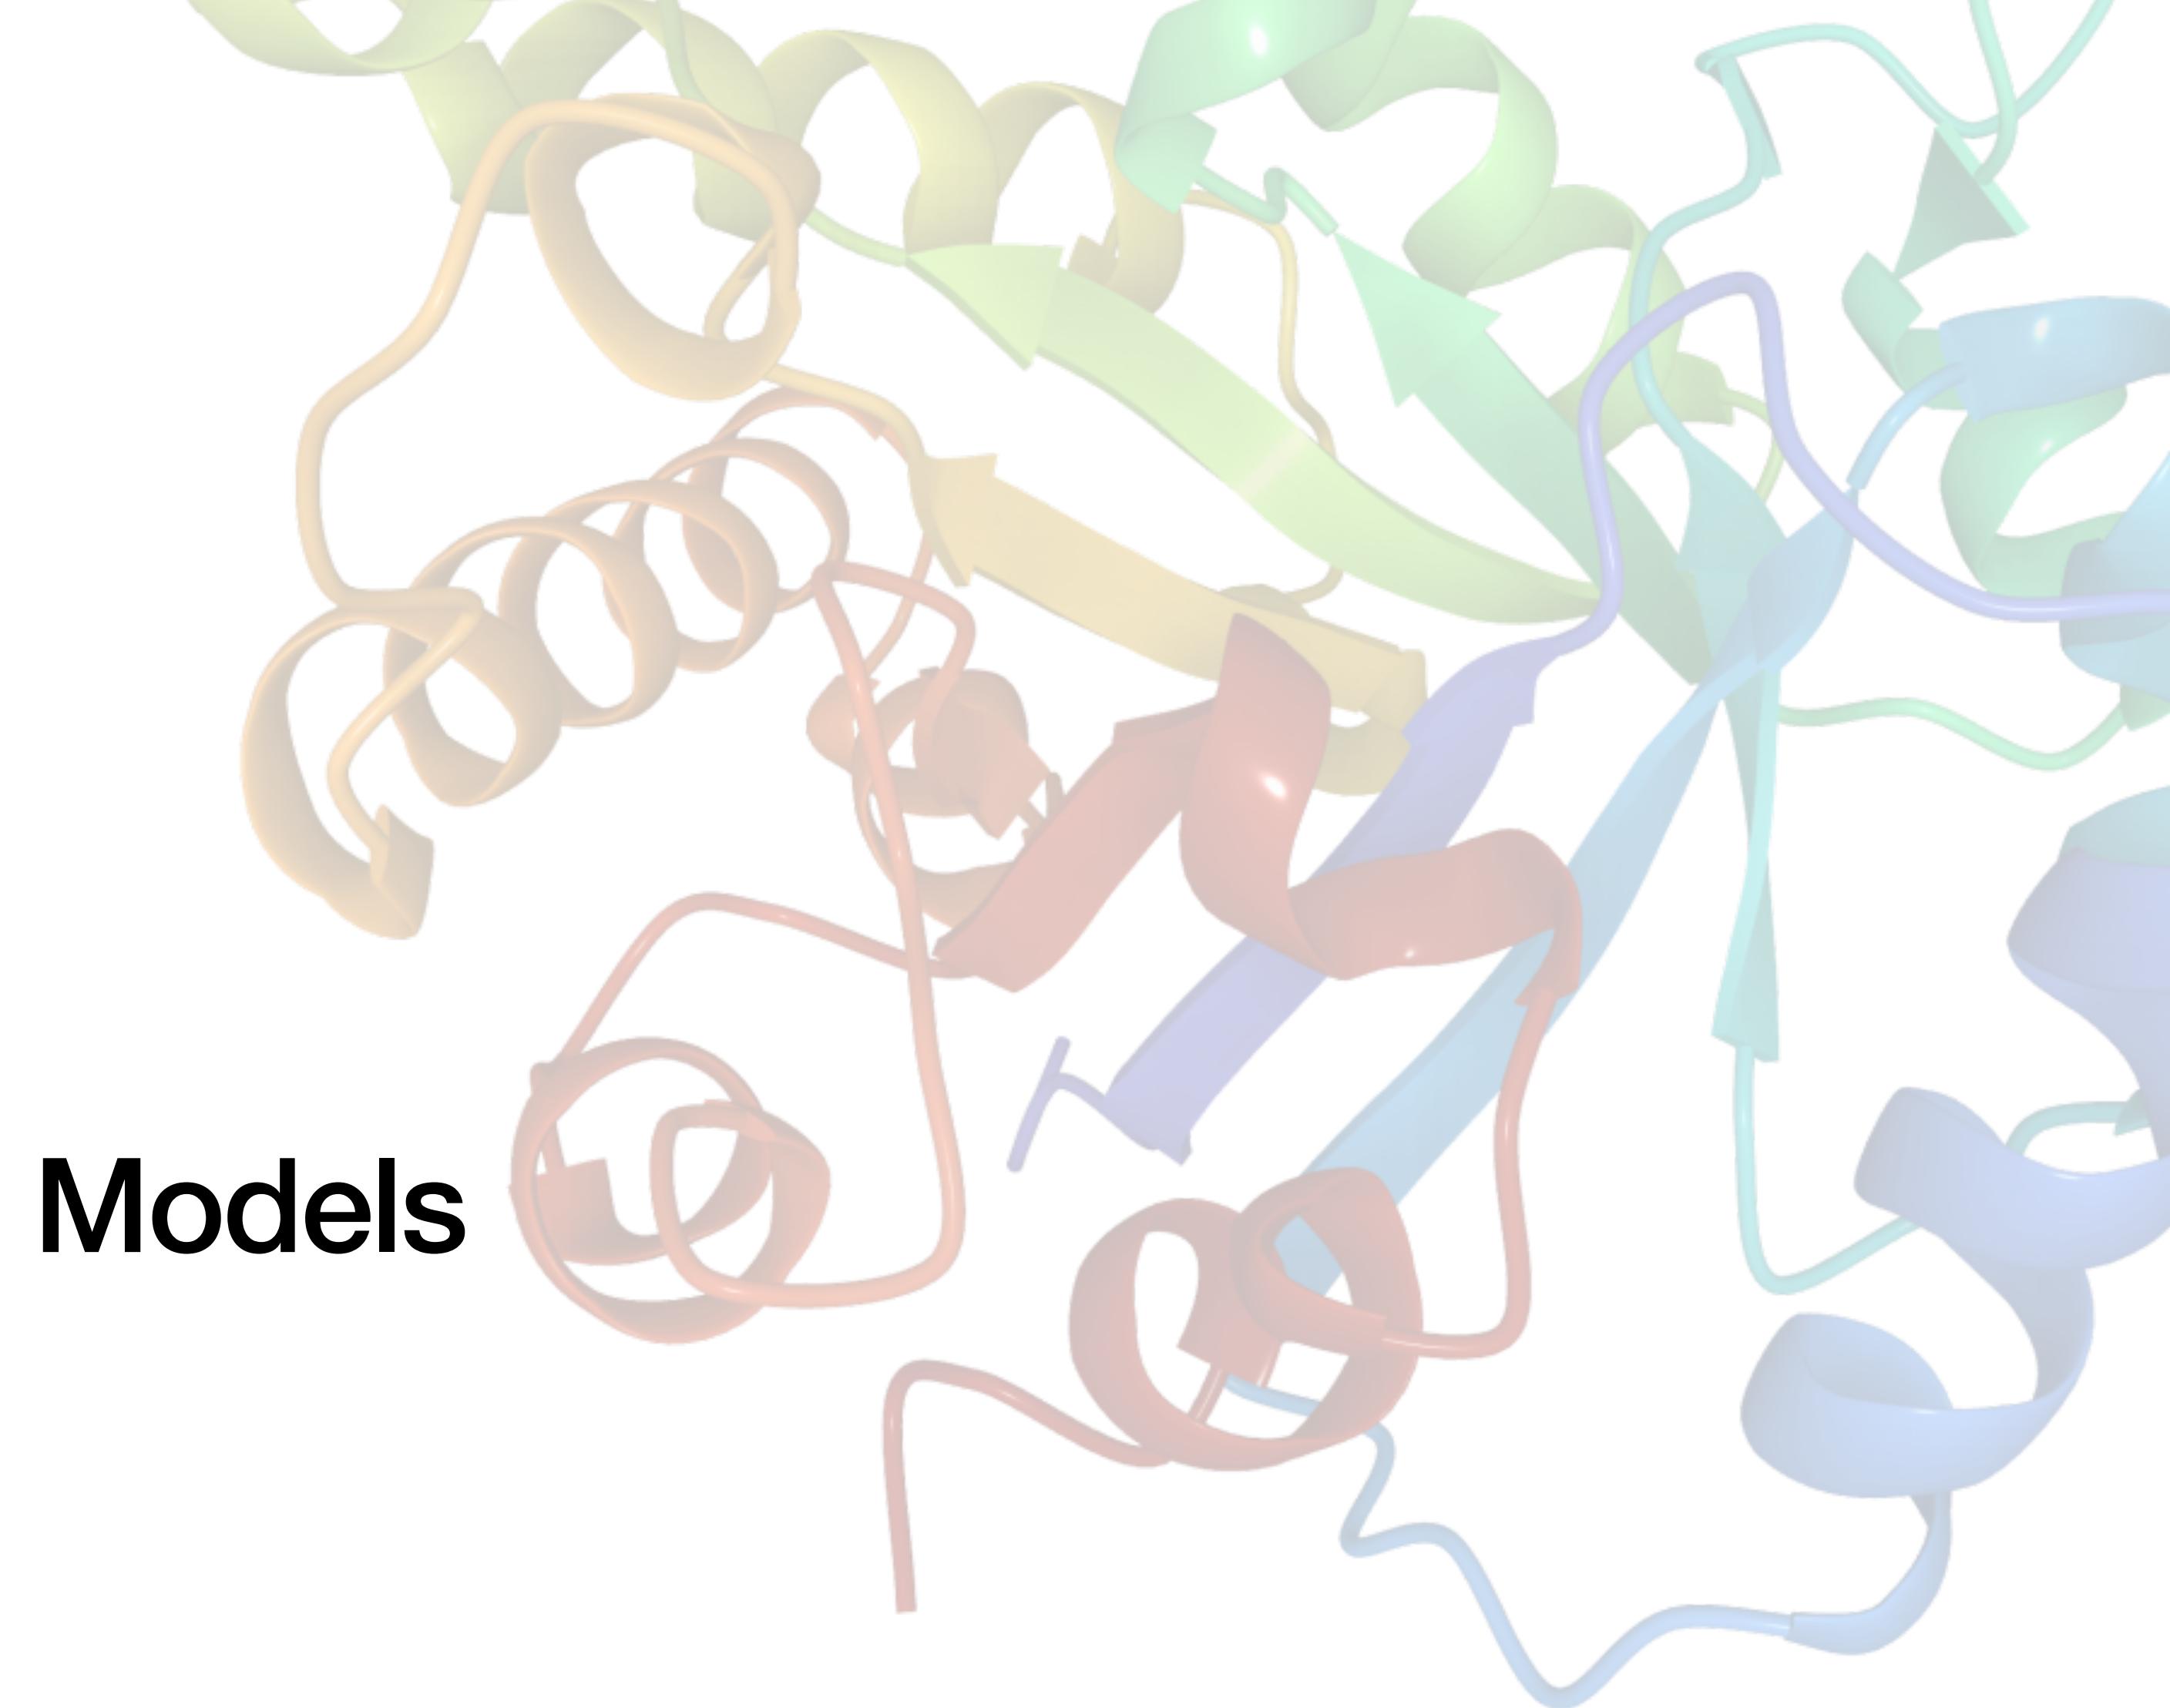
- Simulating x_t
- Backprop through simulation
- (Unbiased) estimator of $\operatorname{div}(u_t)$
- Can compute $\log p(x)$

Maximum Likelihood Objective

$$D_{\text{KL}}(q \parallel p_1) = -\mathbb{E}_{x \sim q} \log p_1(x) + c$$

Part II:

Simulation Free Generative Models



Flow Matching

FLOW MATCHING FOR GENERATIVE MODELING

Yaron Lipman^{1,2} Ricky T. Q. Chen¹ Heli Ben-Hamu² Maximilian Nickel¹ Matt Le¹

¹Meta AI (FAIR) ²Weizmann Institute of Science

ICLR 2023

BUILDING NORMALIZING FLOWS WITH STOCHASTIC INTERPOLANTS

Michael S. Albergo

Center for Cosmology and Particle Physics
New York University
New York, NY 10003, USA
albergo@nyu.edu

Eric Vanden-Eijnden

Courant Institute of Mathematical Sciences
New York University
New York, NY 10012, USA
eve2@cims.nyu.edu

ICLR 2023

FLOW STRAIGHT AND FAST: LEARNING TO GENERATE AND TRANSFER DATA WITH RECTIFIED FLOW

Xingchao Liu*, Chengyue Gong*, Qiang Liu

Department of Computer Science
University of Texas at Austin
{xcliu, cygong, lqiang}@cs.utexas.edu

ICLR 2023

Flow Matching

NON-DENOISING FORWARD-TIME DIFFUSIONS

Stefano Peluchetti
Date: 18 Nov 2021

Rejected ICLR 2022

FLOW MATCHING FOR GENERATIVE MODELING

Yaron Lipman^{1,2} Ricky T. Q. Chen¹ Heli Ben-Hamu² Maximilian Nickel¹ Matt Le¹
¹Meta AI (FAIR) ²Weizmann Institute of Science

ICLR 2023

BUILDING NORMALIZING FLOWS WITH STOCHASTIC INTERPOLANTS

Michael S. Albergo
Center for Cosmology and Particle Physics
New York University
New York, NY 10003, USA
albergo@nyu.edu

Eric Vanden-Eijnden
Courant Institute of Mathematical Sciences
New York University
New York, NY 10012, USA
eve2@cims.nyu.edu

ICLR 2023

FLOW STRAIGHT AND FAST: LEARNING TO GENERATE AND TRANSFER DATA WITH RECTIFIED FLOW

Xingchao Liu*, Chengyue Gong*, Qiang Liu
Department of Computer Science
University of Texas at Austin
{xcliu, cygong, lqiang}@cs.utexas.edu

ICLR 2023

Flow Matching

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$

Core Principle:

u_t generates p_t
iff they satisfy the continuity equation

Construct:

- Target probability path p_t s.t. $p_0 = p$, $p_1 \approx q$
- Generating velocity field u_t

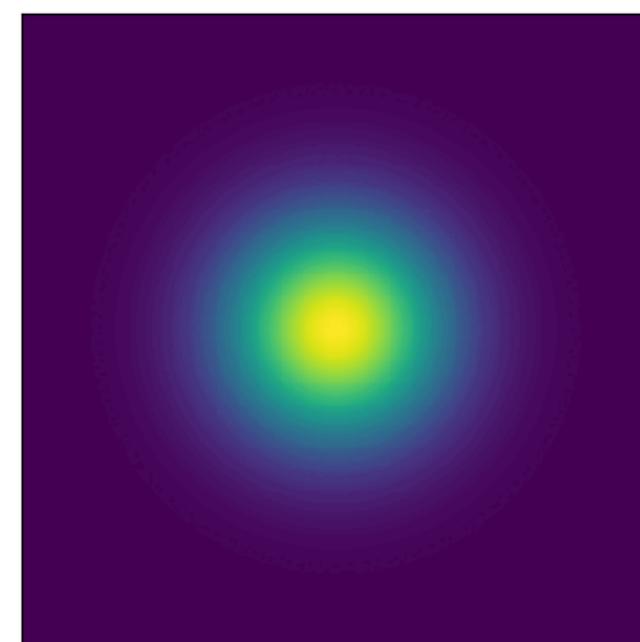
Find a tractable optimization objective

Conditional Probability Paths

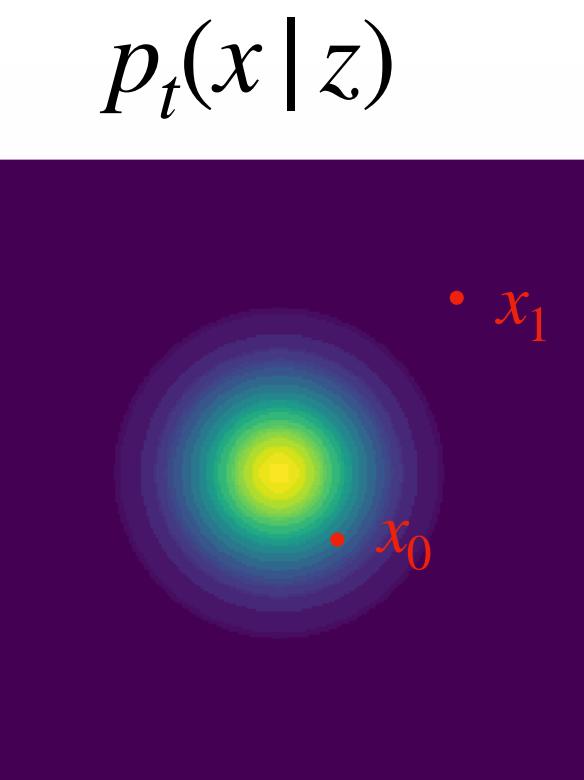
Marginal path

Law of total probability $z := (x_0, x_1)$

$$p_t(x) = \int p_t(x | z) q(z) dz$$



Conditional path

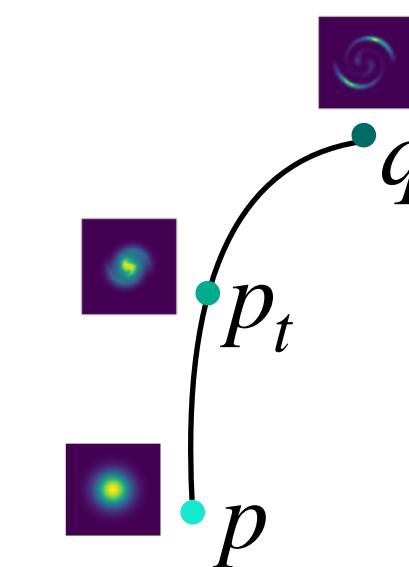


Note! This is the simplest case, we will consider other z in the future

Boundary conditions:

$$p_0 = p$$

$$p_1 = q$$

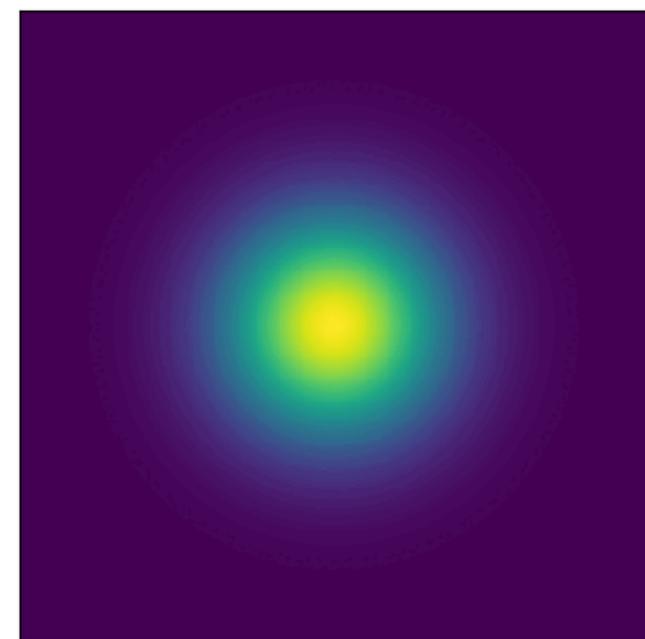


$$p_0(\cdot | z) = \delta_{x_0}$$

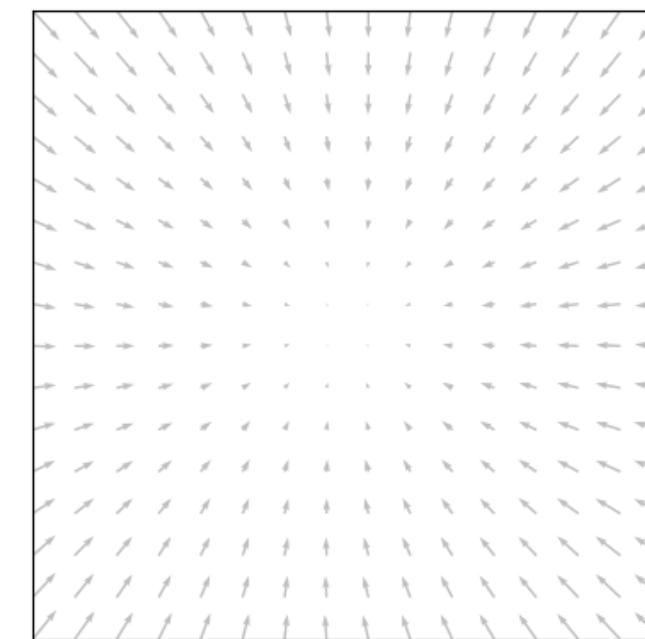
$$p_1(\cdot | z) = \delta_{x_1}$$

The marginalization “trick”

$$p_t(x) = \int p_t(x | z) q(z) dz$$

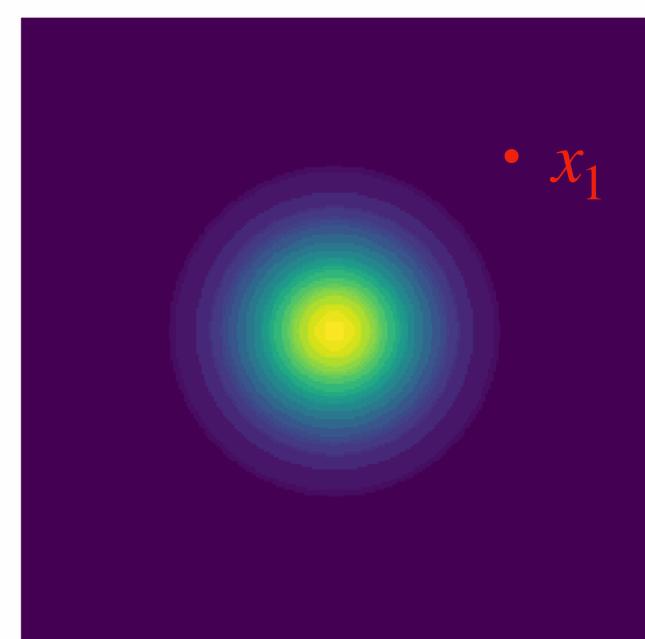


$$u_t(x) = \int u_t(x | z) \frac{p_t(x | z) q(z)}{p_t(x)} dz$$



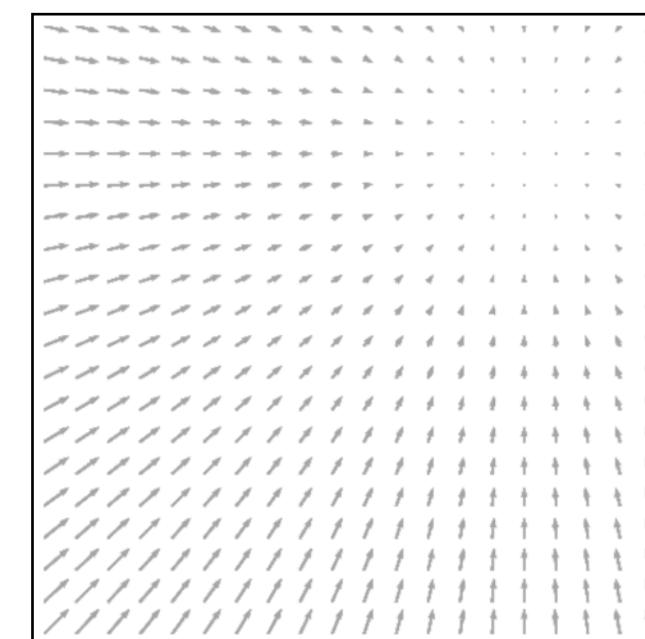
Marginal path

$$p_t(x | z)$$



Conditional path

$$u_t(x | z)$$



Continuity Equation PDE

$$\partial_t p_t = - \operatorname{div}(p_t u_t)$$

The marginalization “trick”

Claim:

$$u_t(x) = \int u_t(x | z) \frac{p_t(x | z)q(z)}{p_t(x)} dz$$

generates

$$p_t(x) = \int p_t(x | z)q(z)dz$$

Proof:

$$\begin{aligned} \operatorname{div}(p_t(x)u_t(x)) &= \\ &= \operatorname{div}\left(\int u_t(x | z)p_t(x | z)q(z)dz\right) \\ &= \int \operatorname{div}(u_t(x | z)p_t(x | z))q(z)dz \\ &= \int -\frac{d}{dt}p_t(x | z)q(z)dz \\ &= -\frac{d}{dt} \int p_t(x | z)q(z)dz \\ &= -\frac{d}{dt}p_t(x) \end{aligned}$$

Flow Matching

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$

$$u_t(x) = \int u_t(x|z) \frac{p_t(x|z)q(z)}{p_t(x)} dz$$

Construct:

- Target probability path p_t s.t. $p_0 = p$, $p_1 \approx q$
- Generating velocity field u_t



Find a tractable optimization objective

Flow Matching

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$

$$u_t(x) = \int u_t(x|z) \frac{p_t(x|z)q(z)}{p_t(x)} dz$$

Useful example:
 $z = (x_0, x_1)$

Construct:

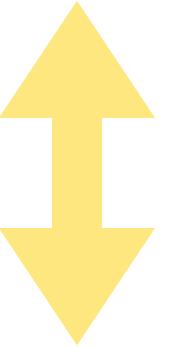
- Target probability path p_t s.t. $p_0 = p$, $p_1 \approx q$
- Generating velocity field u_t



Find a tractable optimization objective

Conditional Flow Matching Loss

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$



$$L_{\text{CFM}}(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x) - u_t(x | z)\|^2$$

The gradients of losses coincide:

$$\nabla_\theta L_{\text{FM}} = \nabla_\theta L_{\text{CFM}}$$

Conditional Flow Matching Loss

$$L_{\text{FM}}(q \| p_1) = \min \mathbb{E}_{t, p_t(x)} \|v_t(\mathbf{x}) - u_t(x)\|^2 \quad L_{\text{CFM}}(q \| p_1) = \min \mathbb{E}_{t, q(x_1), p_t(x|x_1)} \|v_t(\mathbf{x}) - u_t(x|z)\|^2$$

$$\|v_t(x) - u_t(x)\|^2 = \underline{\|v_t(x)\|^2} - 2\langle v_t(x), u_t(x) \rangle + \cancel{\|u_t(x)\|^2}$$

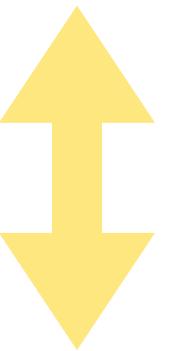
Do not depend on θ

$$\|v_t(x) - u_t(x|z)\|^2 = \underline{\|v_t(x)\|^2} - 2\langle v_t(x), u_t(x|z) \rangle + \cancel{\|u_t(x|z)\|^2}$$

Law of total expectation $\mathbb{E}_{t, p_t(x)} \|v_t(x)\|^2 = \mathbb{E}_{t, q(x_1), p_t(x|z)} \|v_t(x)\|^2$

Flow Matching

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$



$$L_{\text{CFM}}(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x) - u_t(x | z)\|^2$$

Construct:

- Target probability path p_t s.t. $p_0 = p$, $p_1 \approx q$
- Generating velocity field u_t

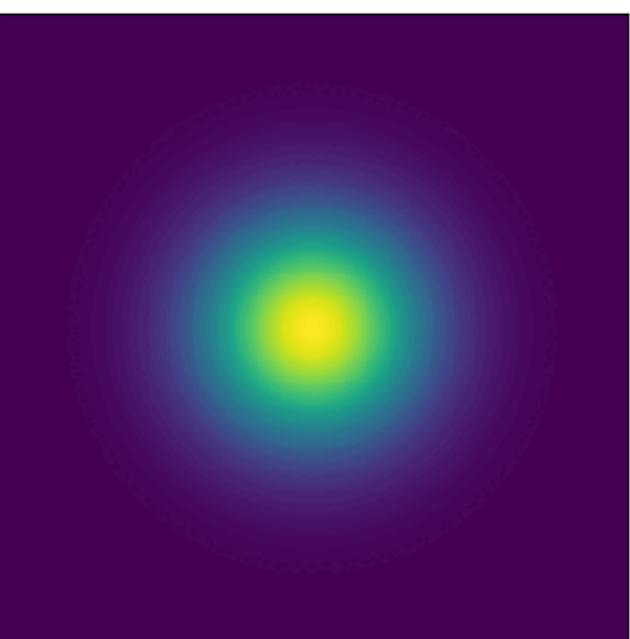


Find a tractable optimization objective



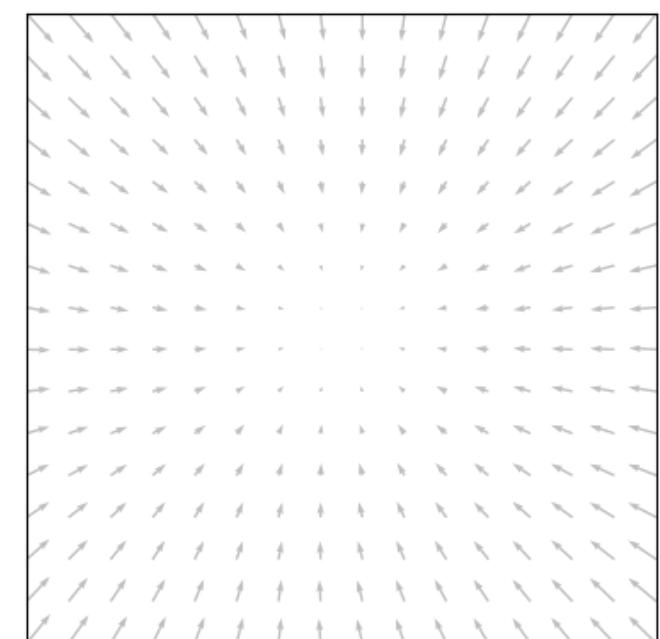
Conditional Flows

$$p_t(x) = \int p_t(x | z) q(z) dz$$



Marginal path

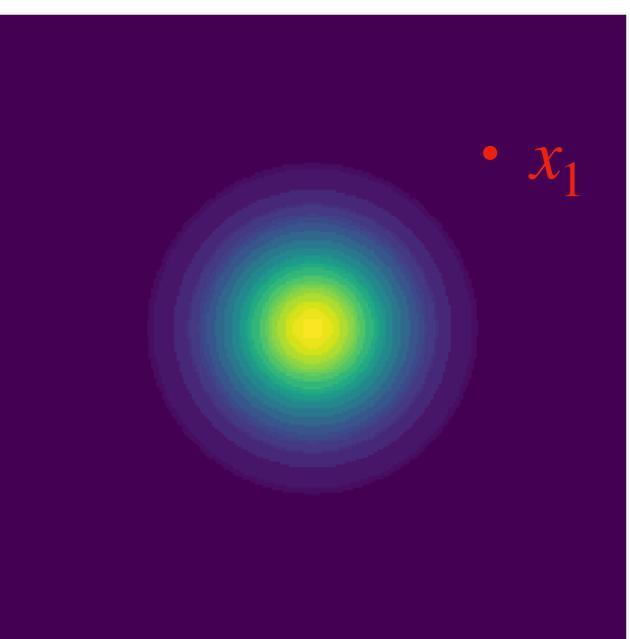
$$u_t(x) = \int u_t(x | z) \frac{p_t(x | z) q(z)}{p_t(x)} dz$$



Construct a conditional flow
s.t.

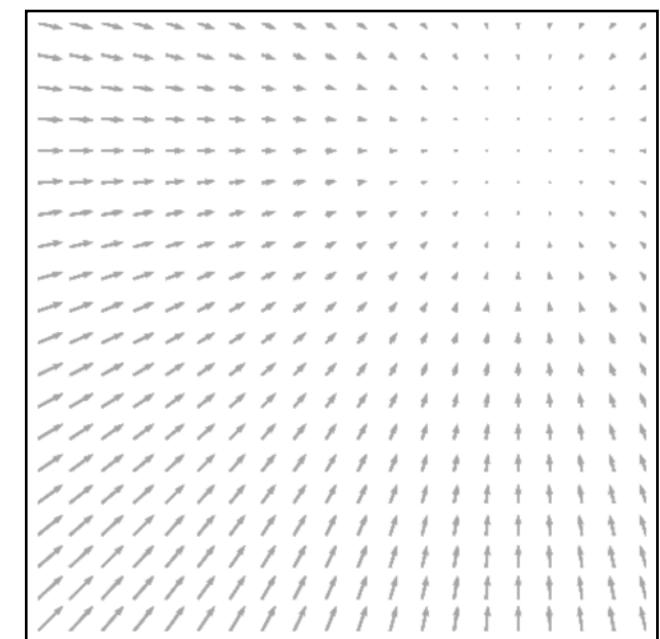
$$\psi_0(x | z) = x , \psi_1(x | z) = x_1$$

$$p_t(x | x_1)$$

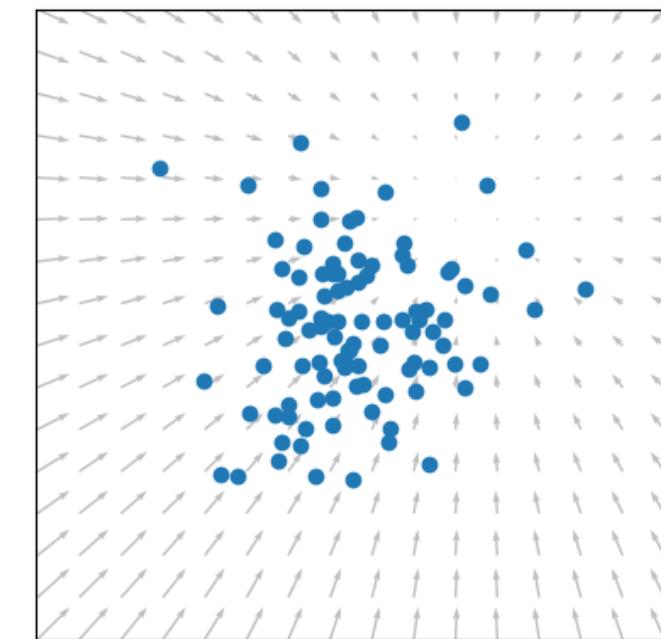


Conditional path

$$u_t(x | x_1)$$



$$\psi_t(x_0 | x_1)$$

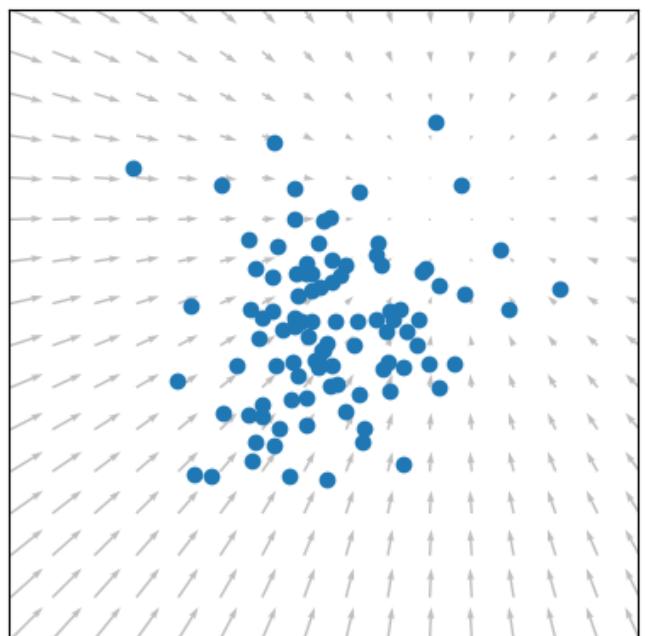


Affine Conditional Flows

Construct a conditional flow
s.t.

$$\psi_0(x|x_1) = x \quad , \quad \psi_1(x|x_1) = x_1$$

$$\psi_t(x_0|x_1)$$



Affine conditional flow:

$$\psi_t(x_0|x_1) = x_t = \alpha_t x_1 + \sigma_t x_0$$

Conditional velocity field:

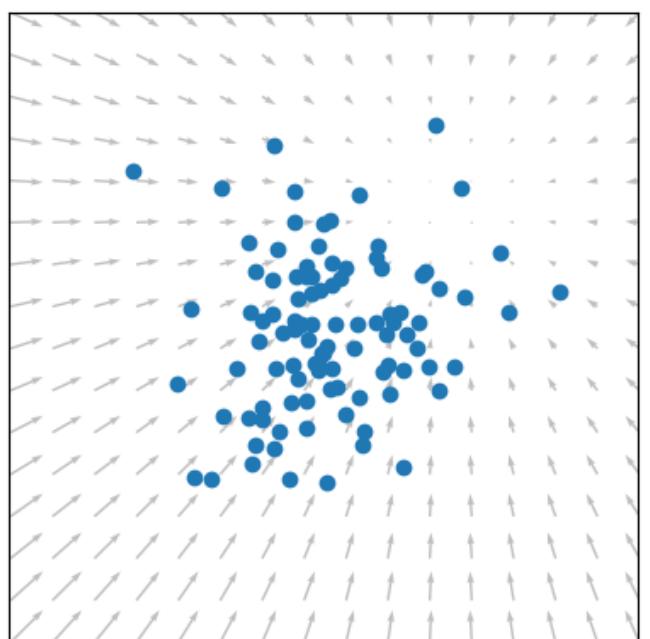
$$\frac{d}{dt} \psi_t(x_0|x_1) = u_t(\psi_t(x_0|x_1)|x_1)$$

$$u_t(x_t|x_1) = \dot{\alpha}_t x_1 + \dot{\sigma}_t x_0$$

$$u_t(x|x_1) = \frac{\dot{\sigma}_t}{\sigma_t} x + (\dot{\alpha}_t - \alpha_t \frac{\dot{\sigma}_t}{\sigma_t}) x_1$$

Affine Conditional Flows

Construct a conditional flow
s.t.
 $\psi_0(x | x_1) = x$, $\psi_1(x | x_1) = x_1$



$$\psi_t(x_0 | x_1)$$

Affine conditional flow:

$$\psi_t(x_0 | x_1) = x_t = \alpha_t x_1 + \sigma_t x_0$$

Conditional velocity field:

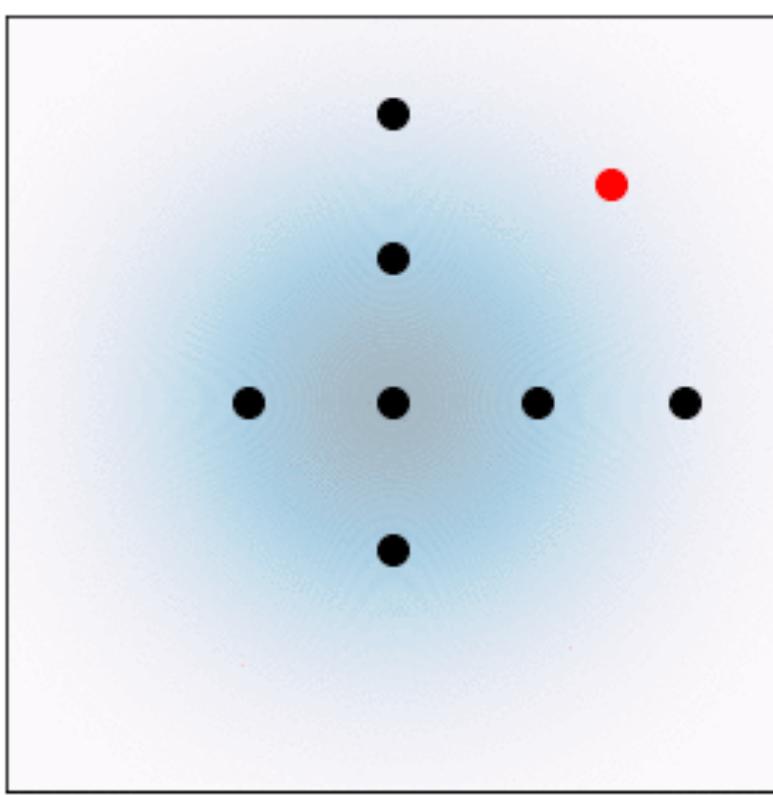
$$\frac{d}{dt}x_t = u_t(x_t | x_1)$$

$$u_t(x_t | x_1) = \dot{\alpha}_t x_1 + \dot{\sigma}_t x_0$$

$$u_t(x_t | x_1) = \frac{\dot{\sigma}_t}{\sigma_t}x + (\dot{\alpha}_t - \alpha_t \frac{\dot{\sigma}_t}{\sigma_t})x_1$$

The default setting

Construct a conditional flow
s.t.
 $\psi_0(x|x_1) = x$, $\psi_1(x|x_1) = x_1$



flow coefficients:

$$\alpha_t = t \quad , \quad \sigma_t = 1 - t$$

flow:

$$\psi_t(x_0|x_1) = x_t = tx_1 + (1 - t)x_0$$

$$u_t(x_t|x_1) = x_1 - x_0$$

$$u_t(x|x_1) = \frac{x_1 - x}{1 - t}$$

Equivalent to a
diffusion loss when

$$\sigma_t = t!$$

Gaussian Affine Conditional Flows

Construct a conditional flow
s.t.
 $\psi_0(x|x_1) = x$, $\psi_1(x|x_1) = x_1$

Gaussian source:

$$x_0 \sim \mathcal{N}(0, I)$$

Affine conditional flow:

$$\psi_t(x_0 | x_1) = x_t = \alpha_t x_1 + \sigma_t x_0$$



$$x_t \sim \mathcal{N}(\alpha_t x_1, \sigma_t^2 I)$$

Popular Diffusion Paths:

Variance Exploding:

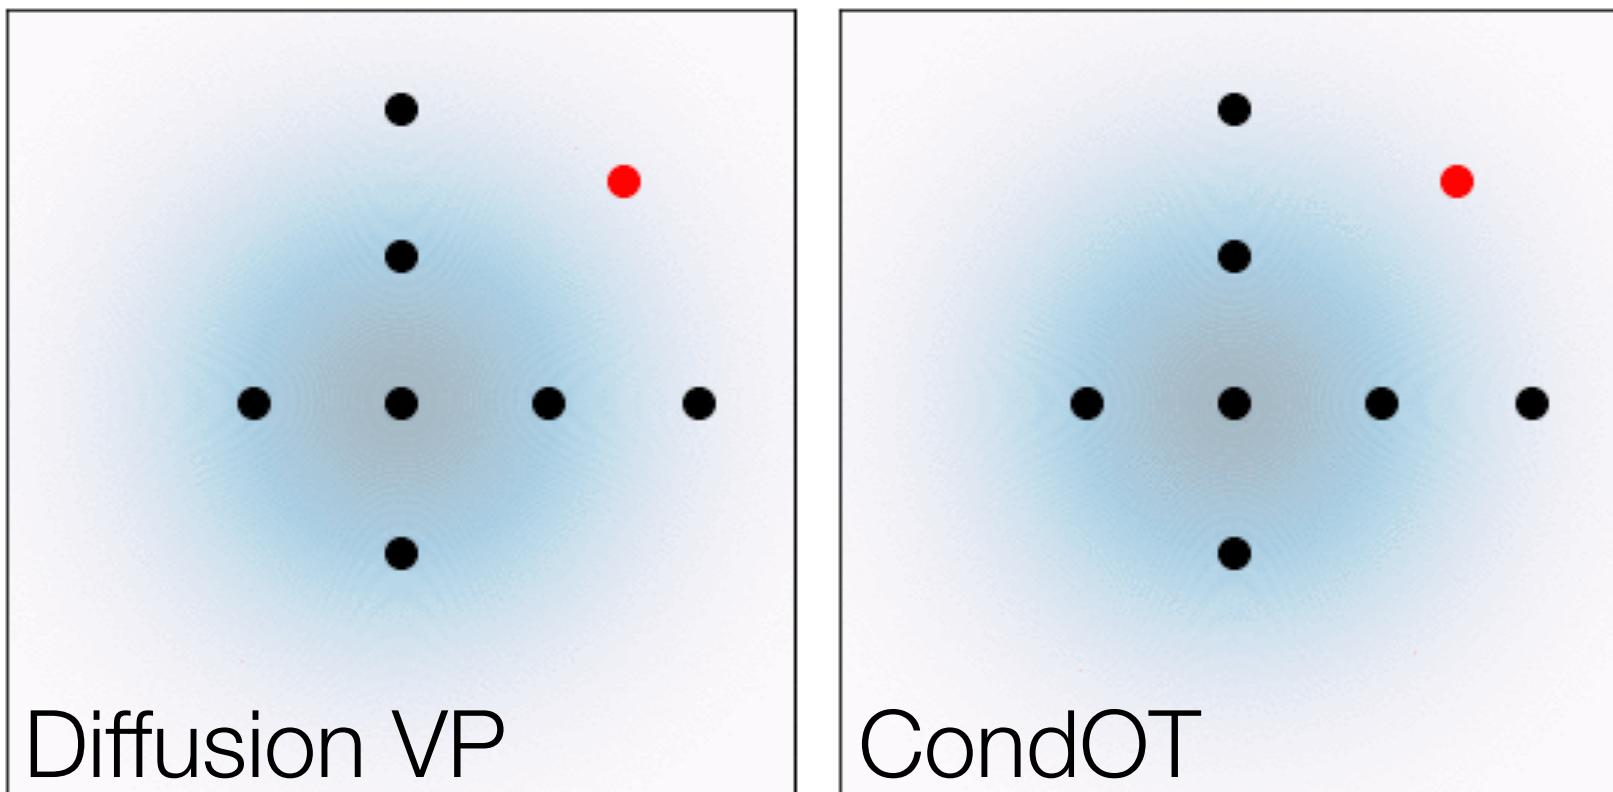
$$x_t \sim \mathcal{N}(x_1, \sigma_t^2 I)$$

Variance Preserving:

$$x_t \sim \mathcal{N}(\alpha_t x_1, (1 - \alpha_t)_t I)$$

Gaussian Affine Conditional Flows

Construct a conditional flow
s.t.
 $\psi_0(x|x_1) = x$, $\psi_1(x|x_1) = x_1$

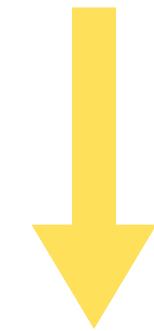


Gaussian source:

$$x_0 \sim \mathcal{N}(0, I)$$

Affine conditional flow:

$$\psi_t(x_0 | x_1) = x_t = \alpha_t x_1 + \sigma_t x_0$$



$$x_t \sim \mathcal{N}(\alpha_t x_1, \sigma_t^2 I)$$

Popular Diffusion Paths:

Variance Exploding:

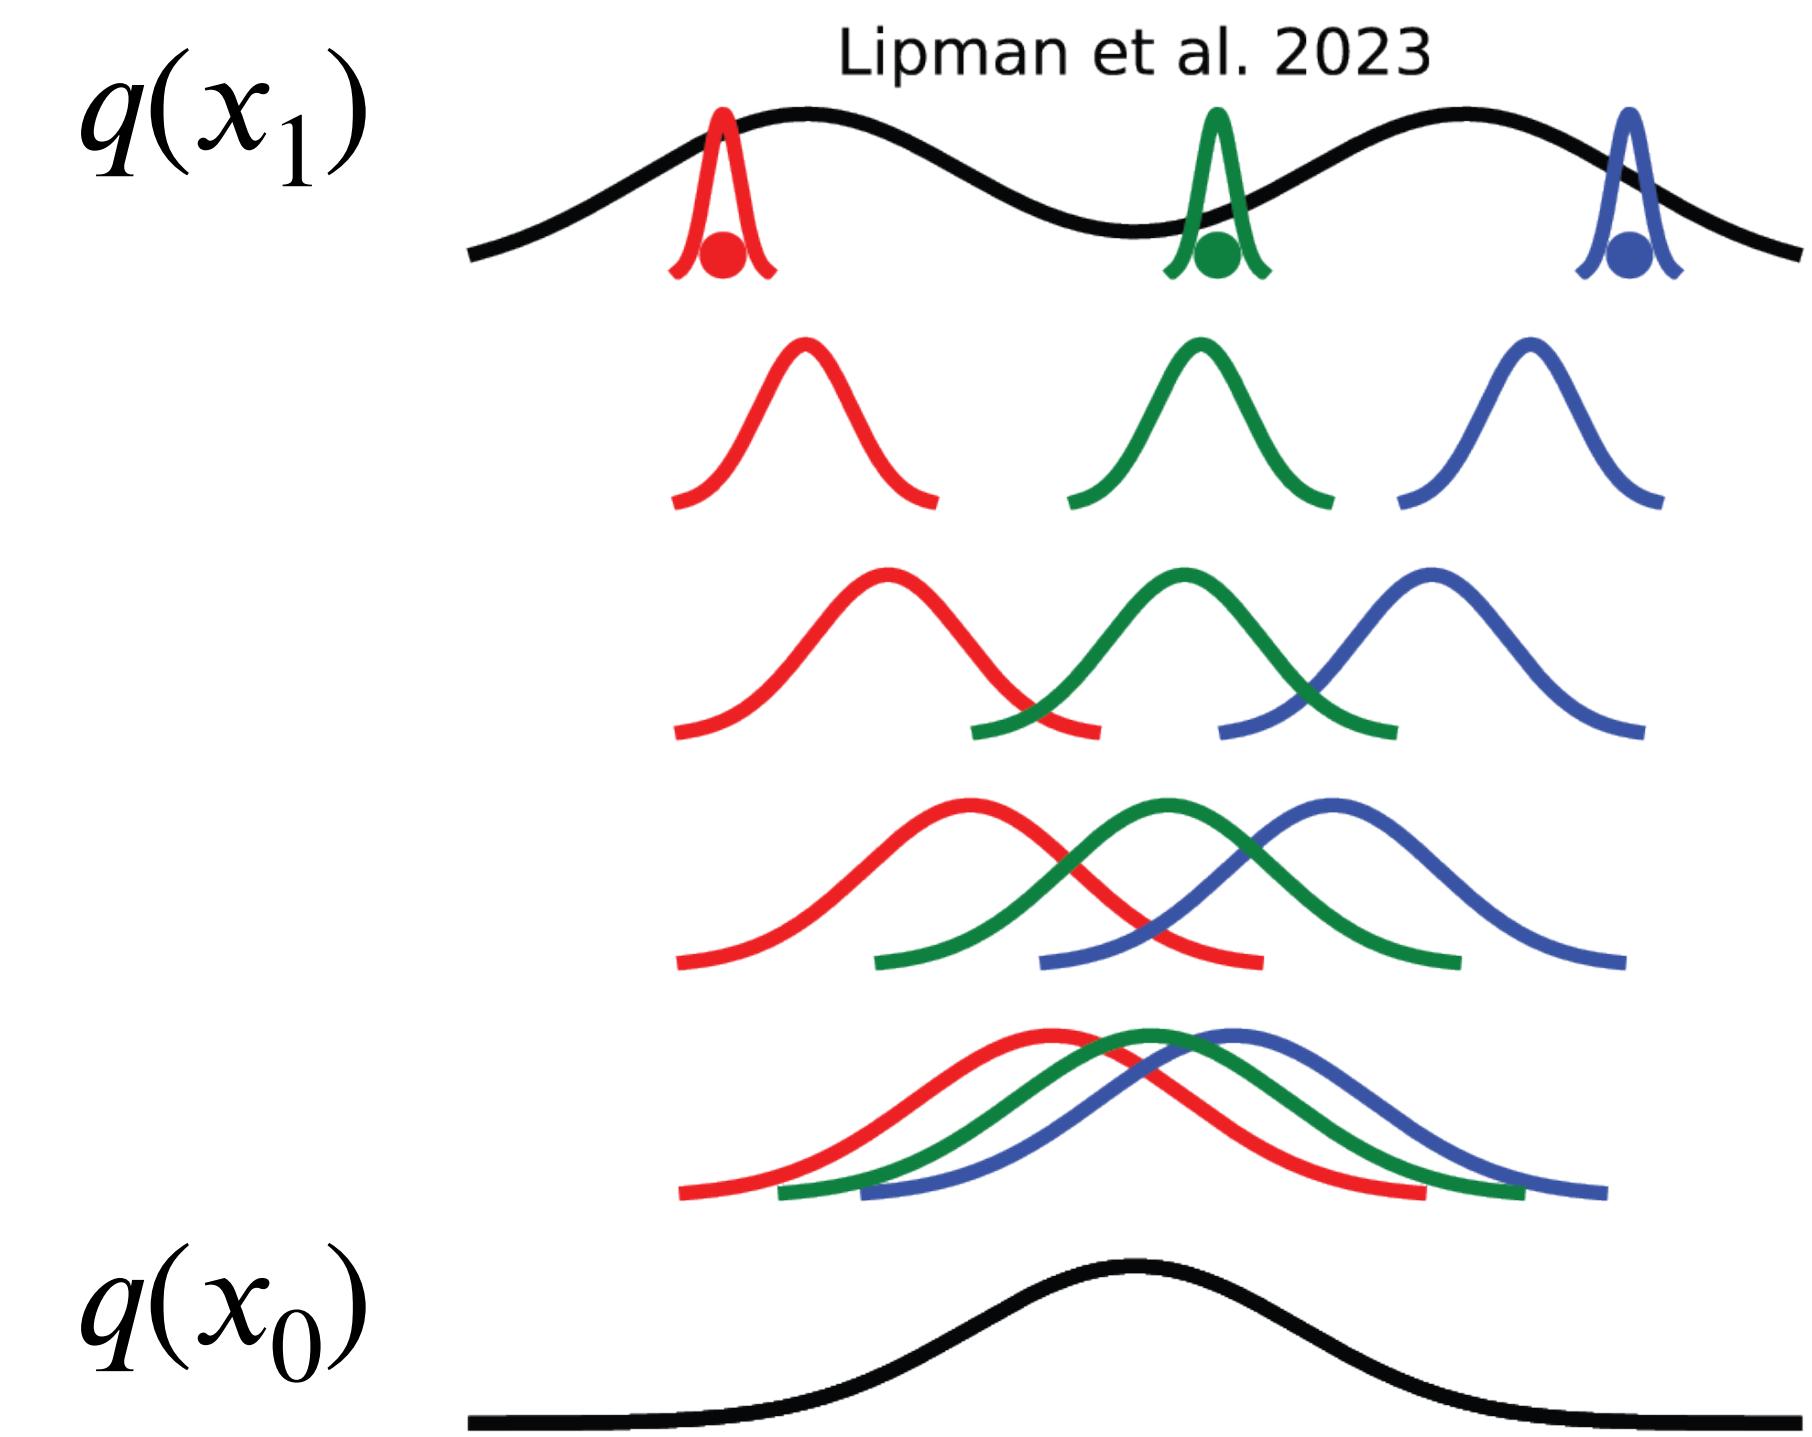
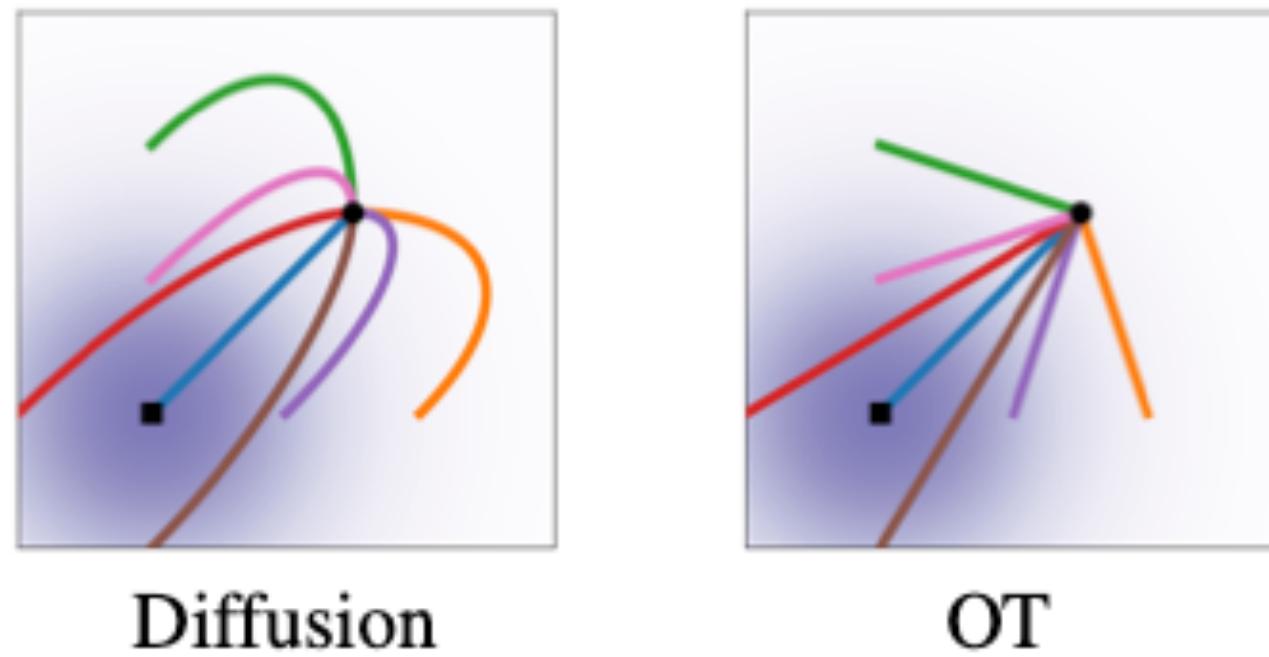
$$x_t \sim \mathcal{N}(x_1, \sigma_t^2 I)$$

Variance Preserving:

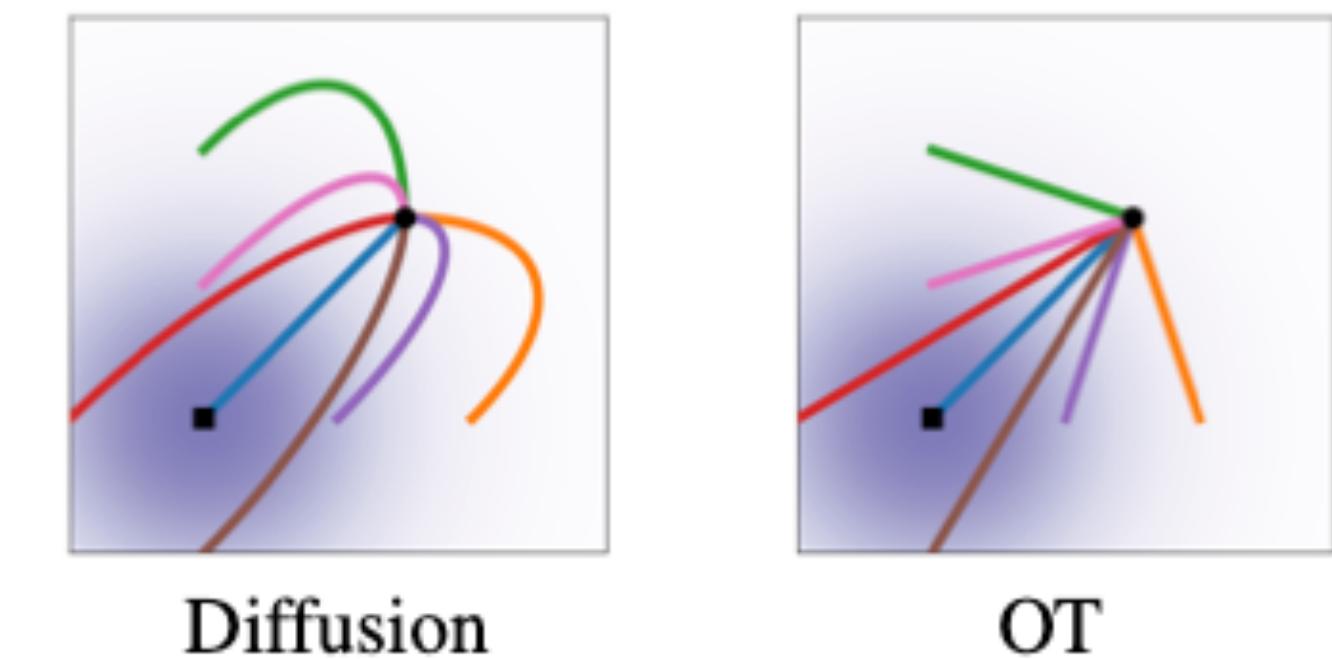
$$x_t \sim \mathcal{N}(\alpha_t x_1, (1 - \alpha_t)_t I)$$

A Geometric / Algorithmic point of view

Flow Matching

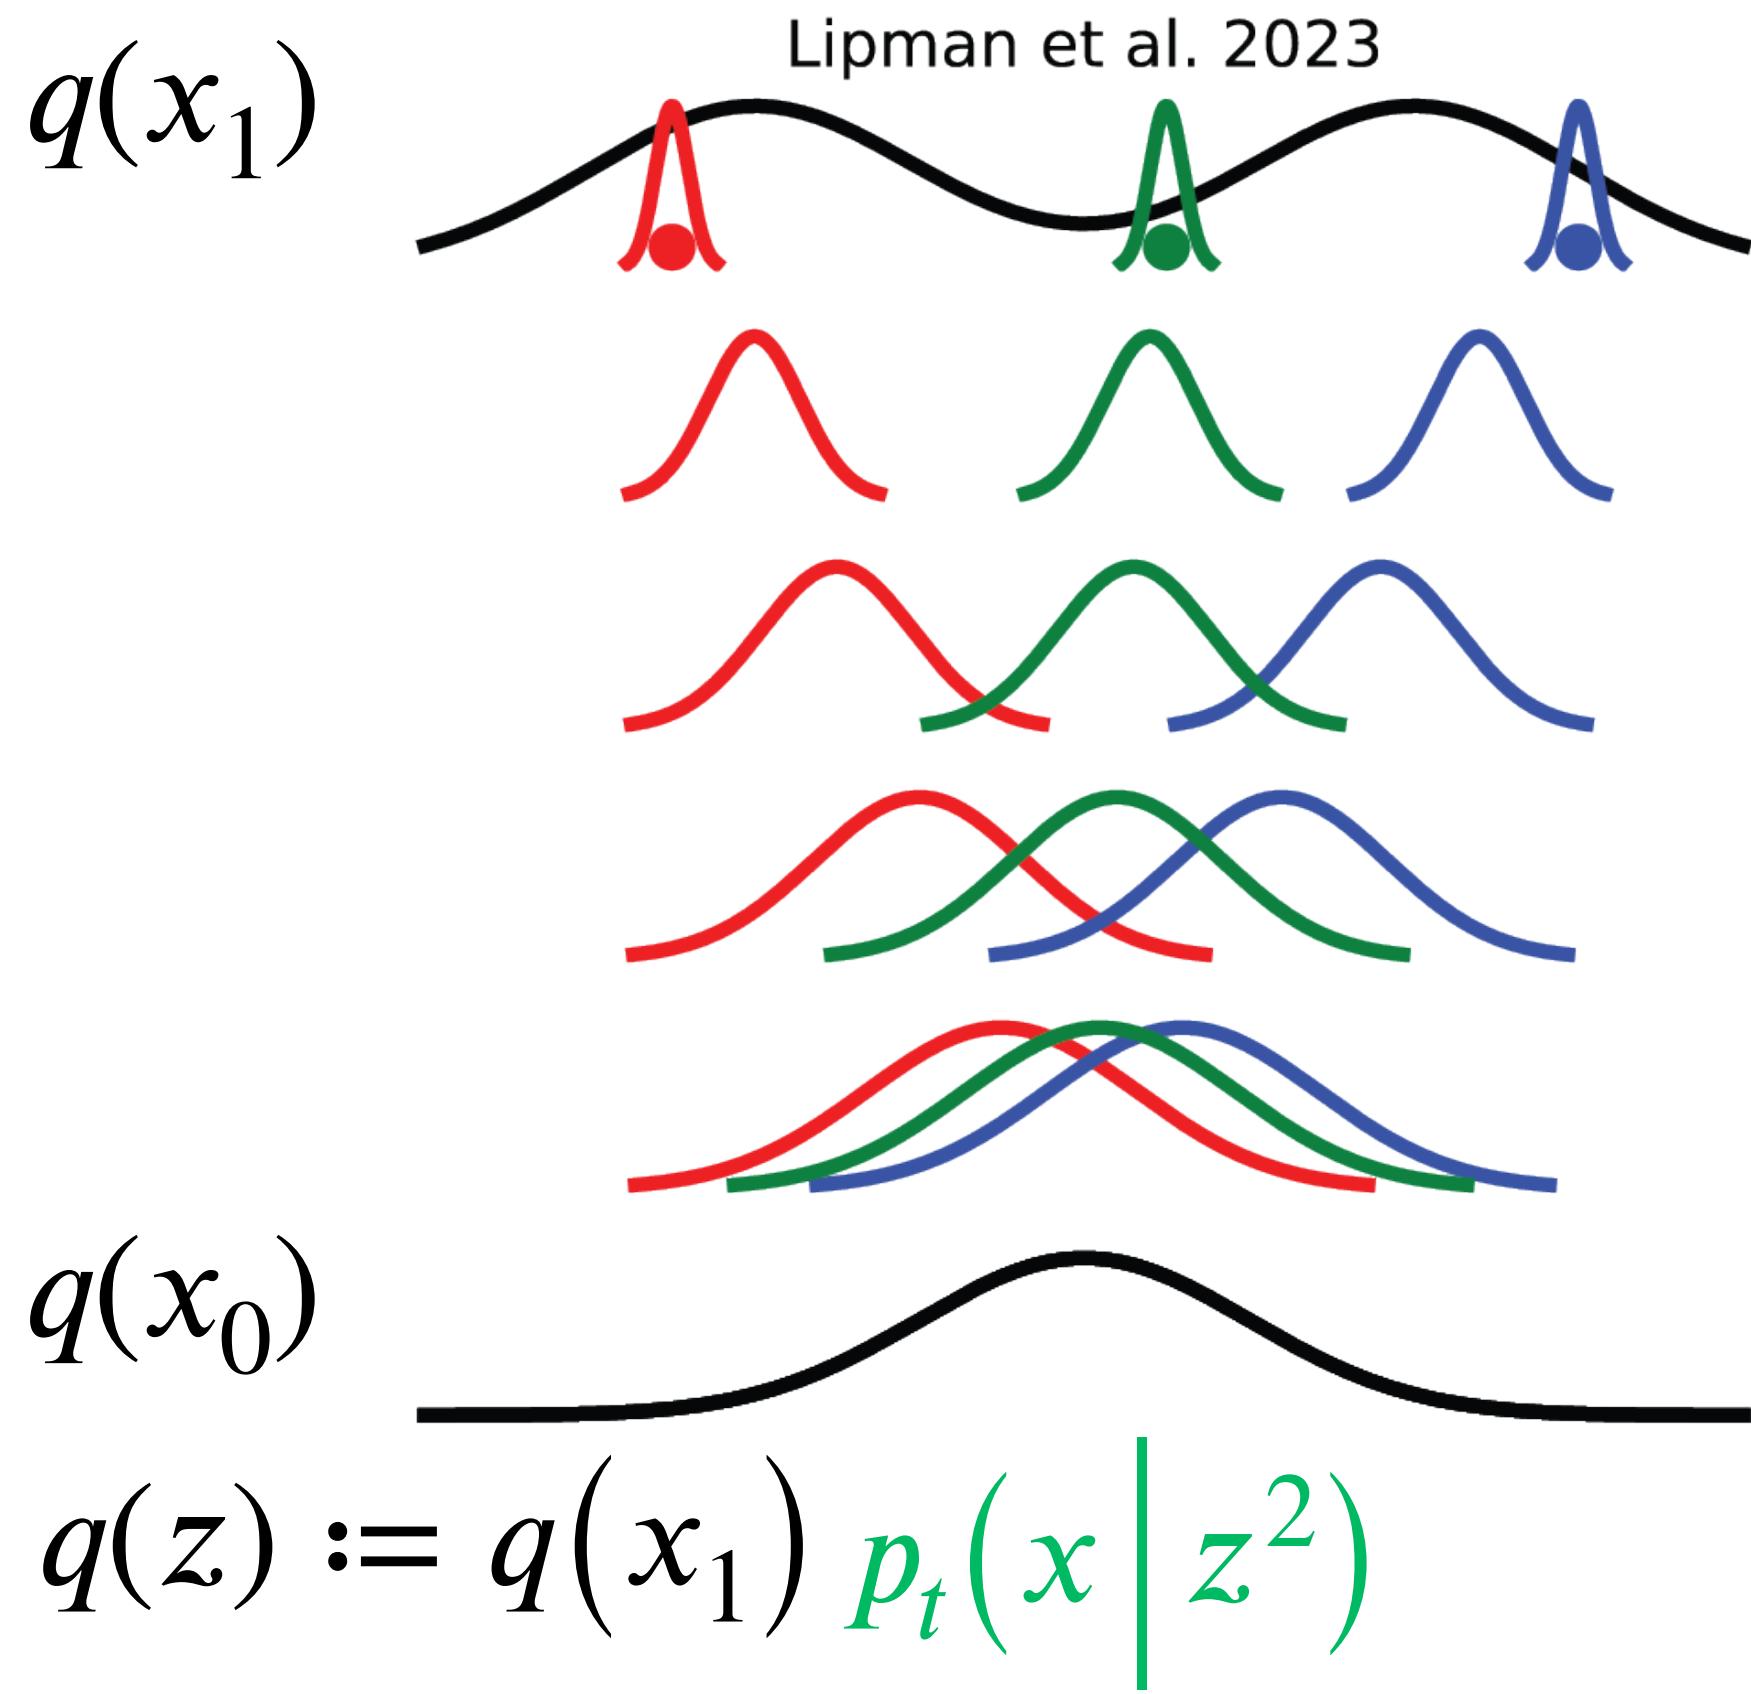


Flow Matching

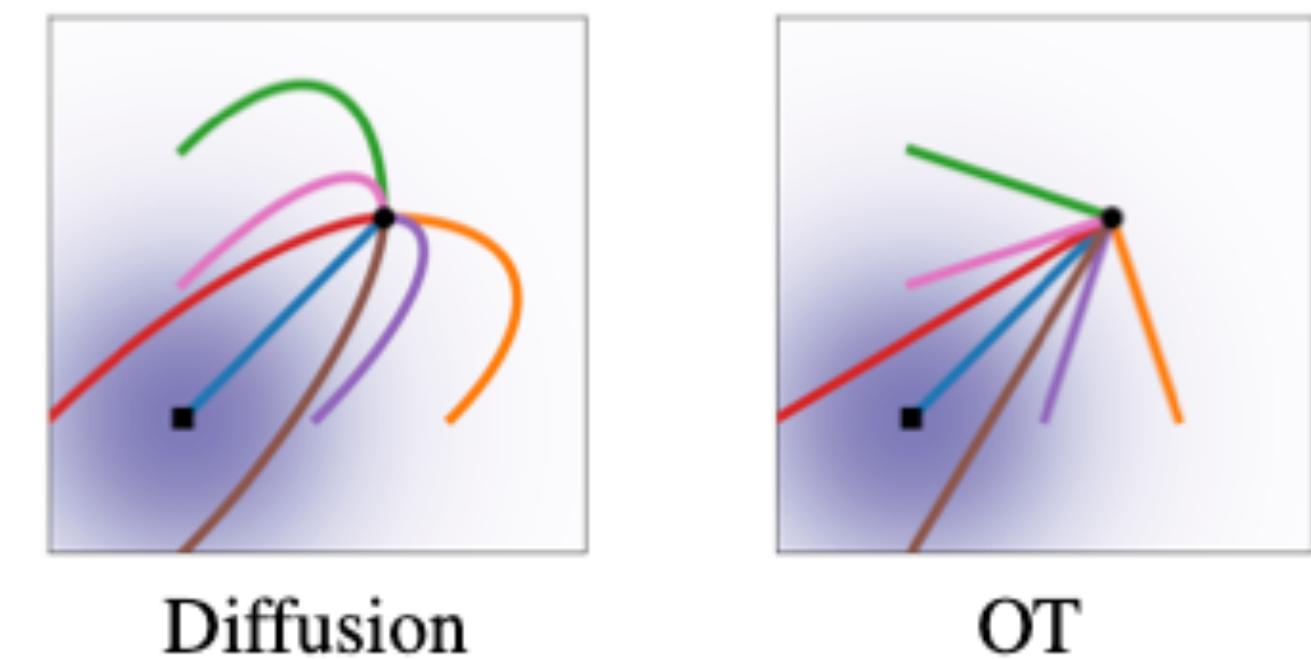


$$p_t(x \mid z^1)$$

$$p_t(x \mid z^3)$$

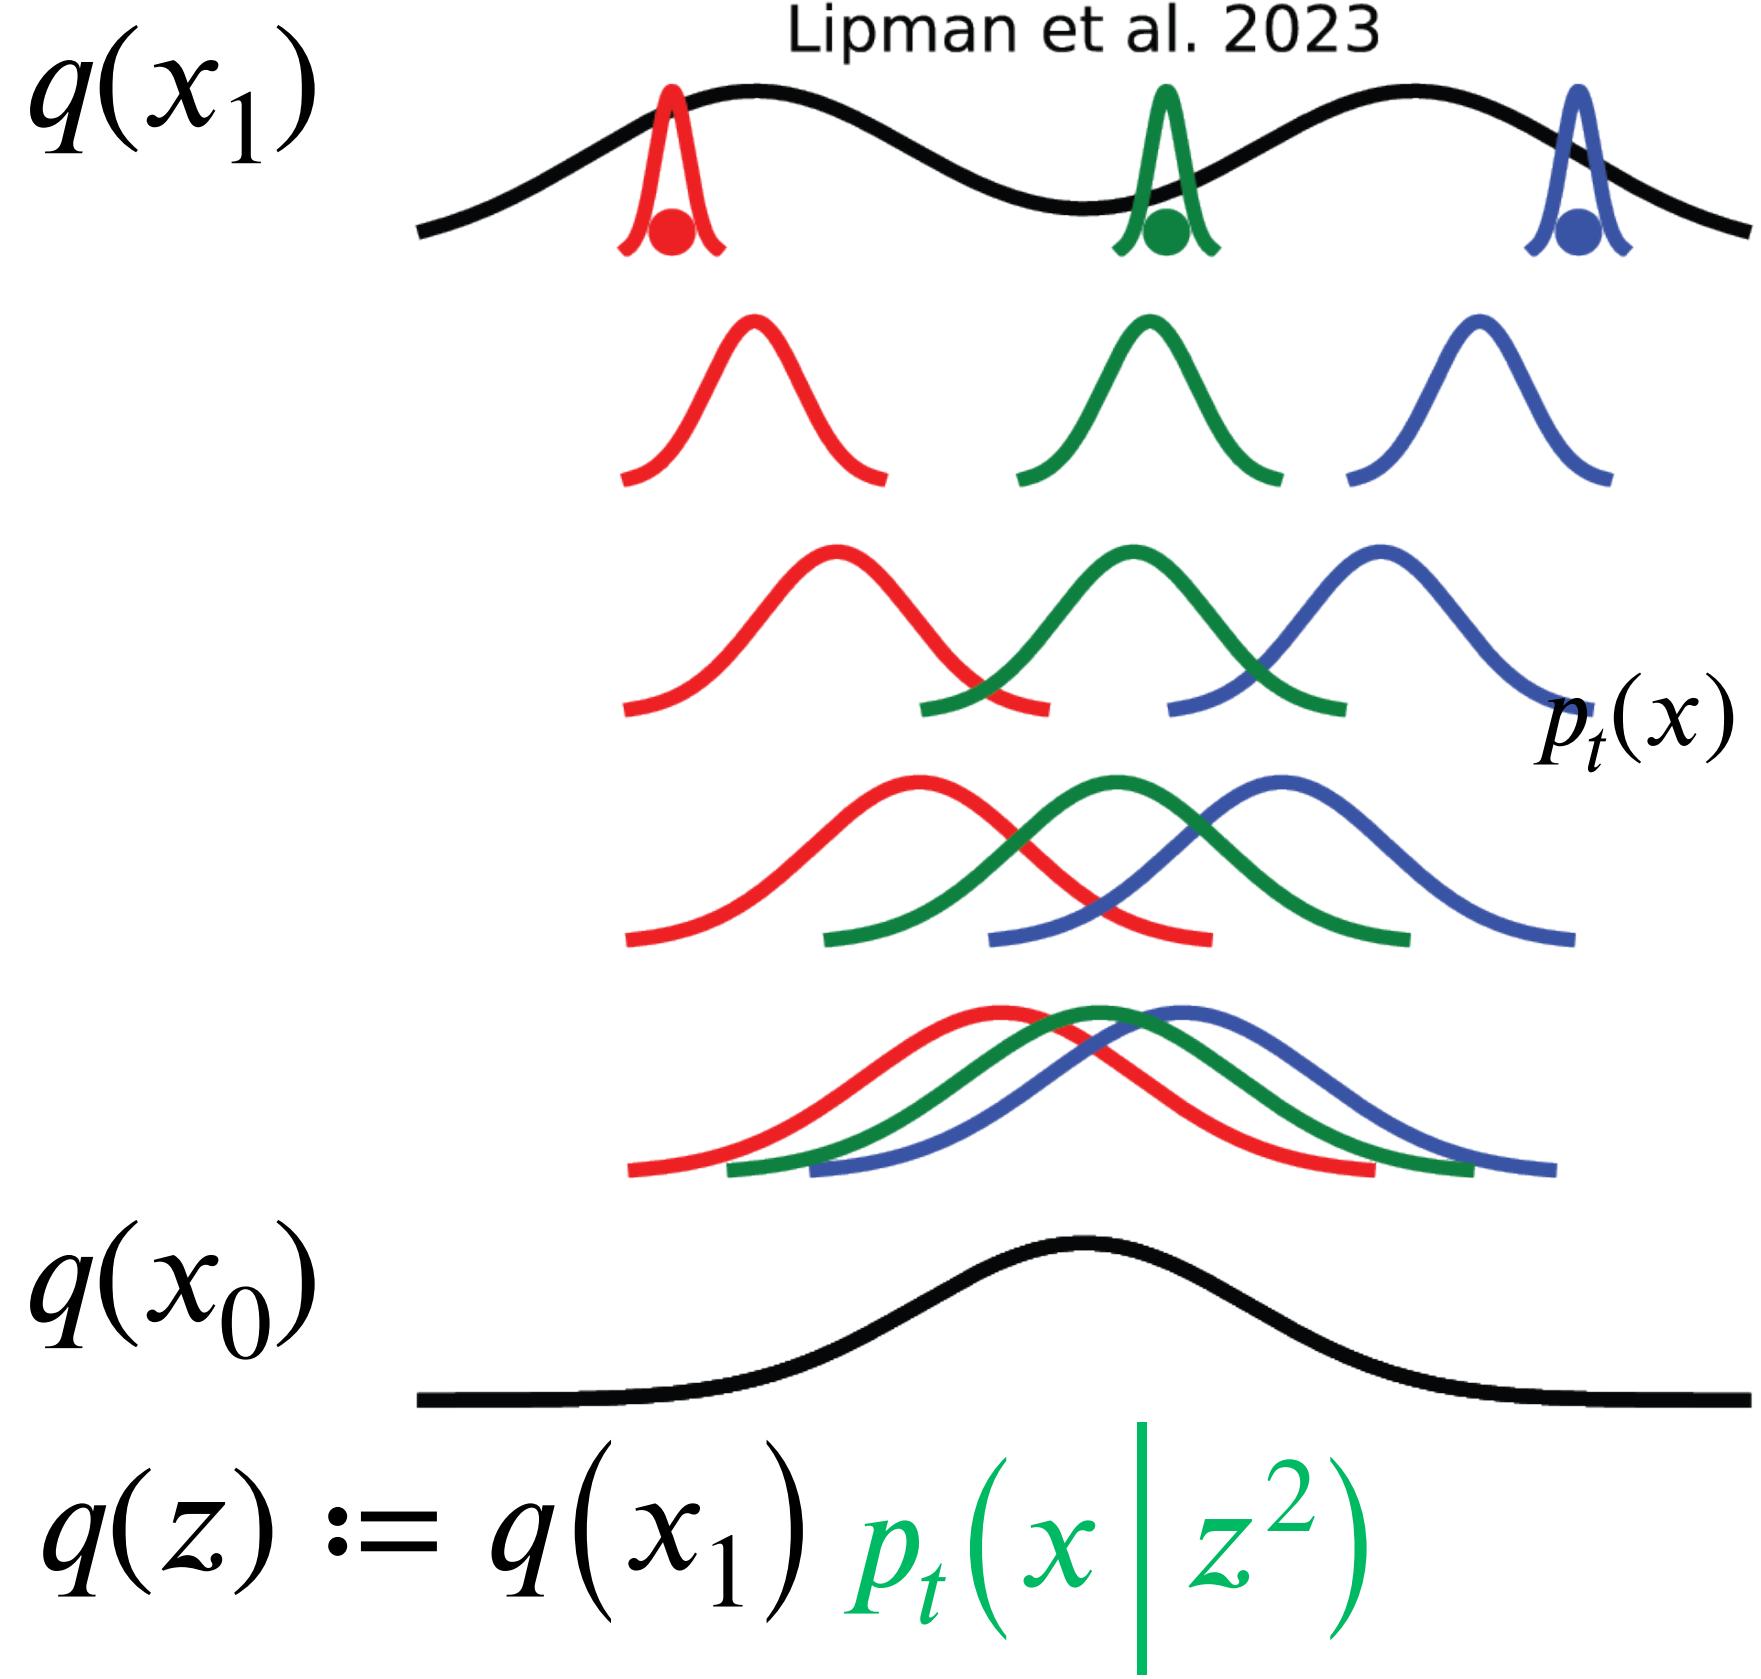


Flow Matching

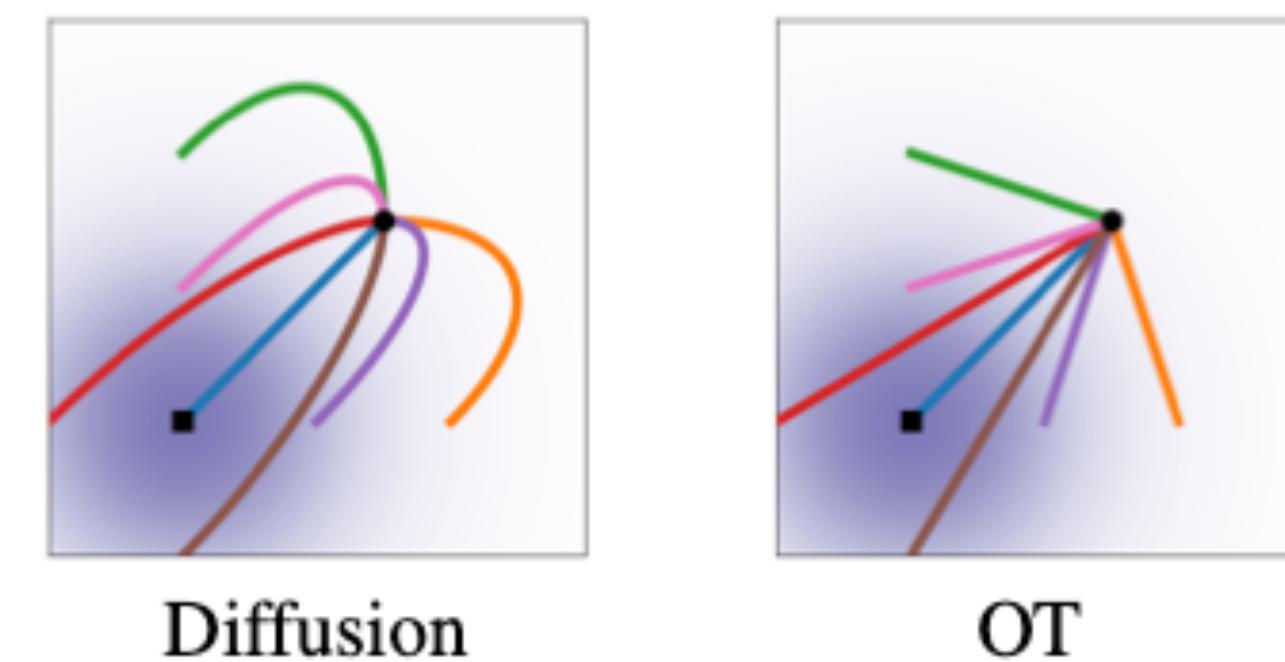


$$p_t(x | z^1)$$

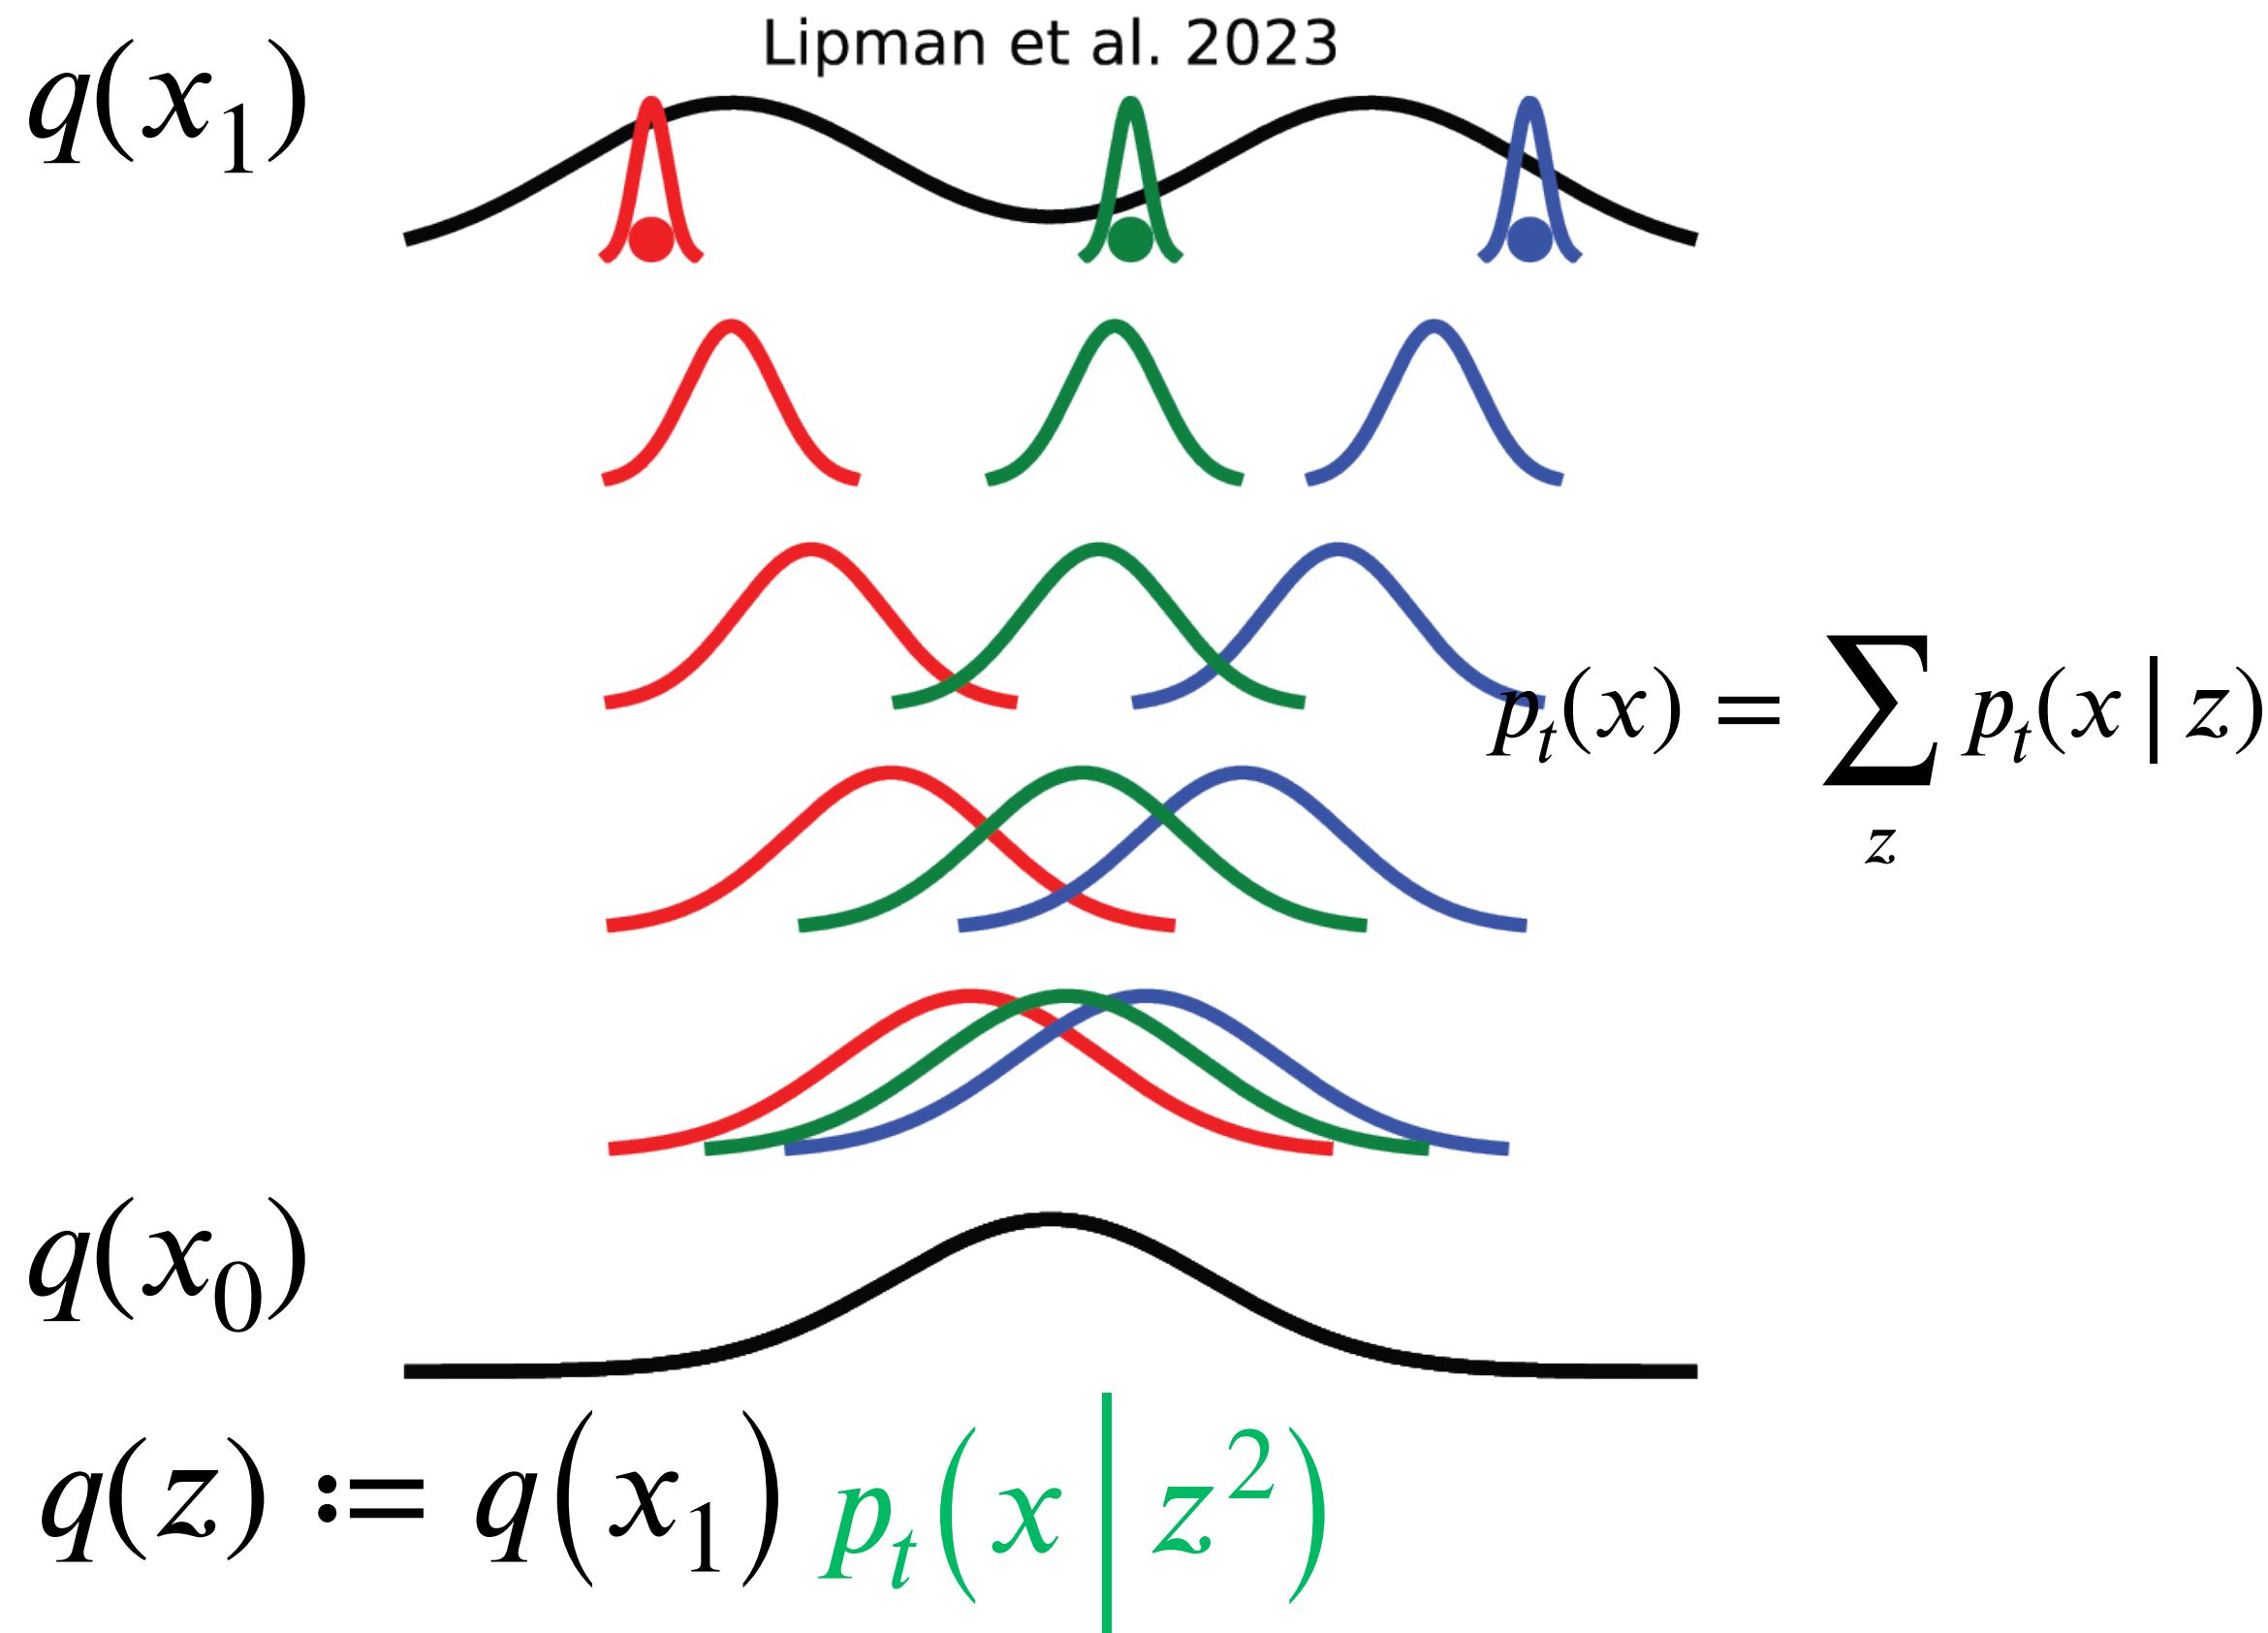
$$p_t(x | z^3)$$



Flow Matching



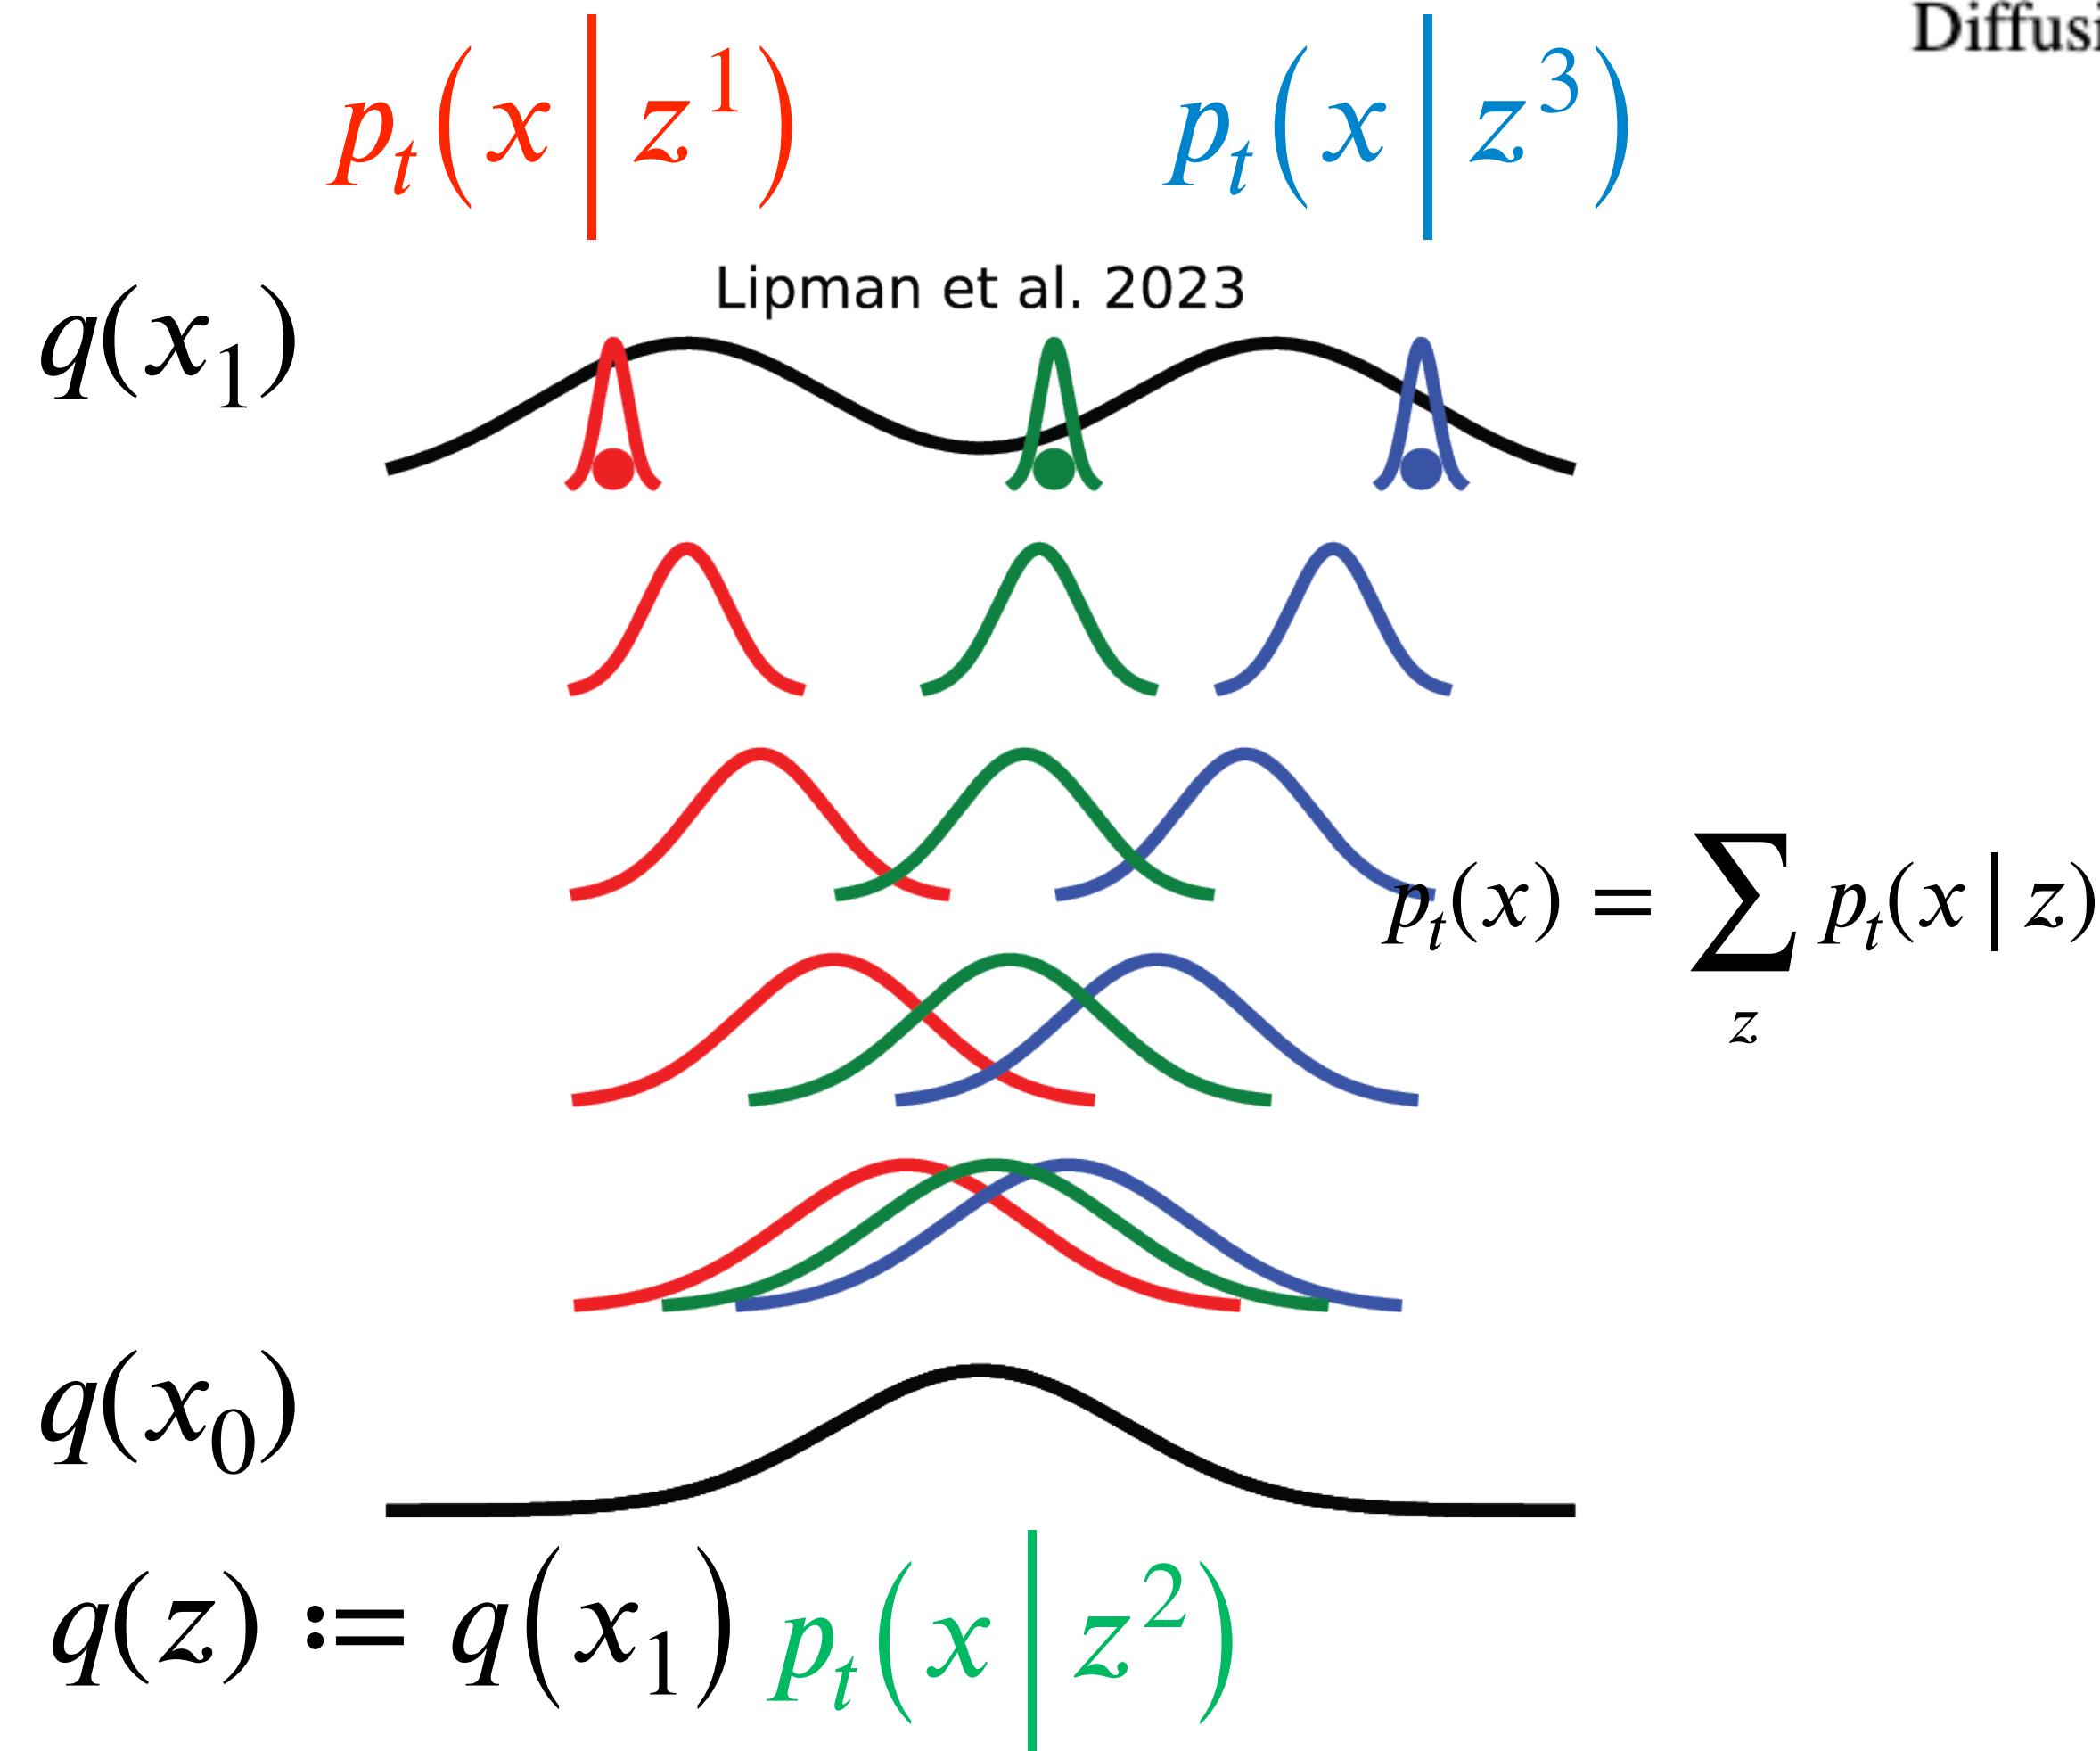
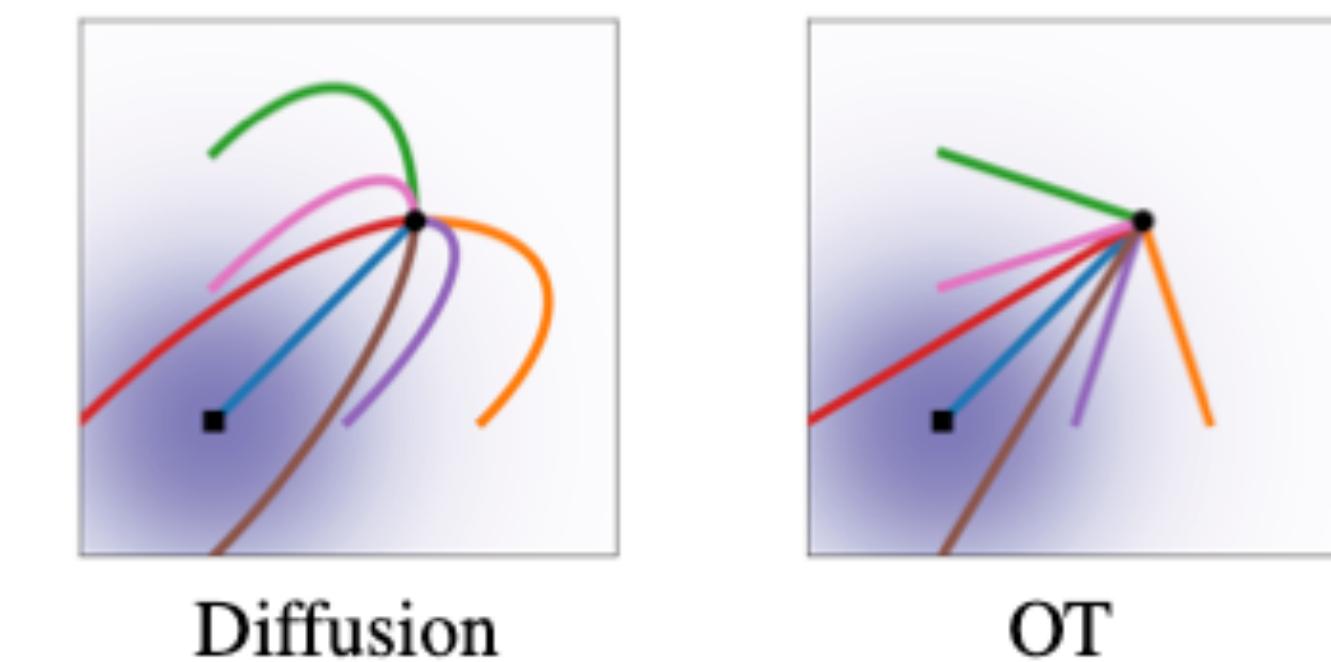
$$p_t(x | z^1) \quad p_t(x | z^3)$$



Closed form Gaussian Conditional Flow! $p_t(x | z) = N(x | \mu_t, \sigma_t)$ then

$$u_t(x | z) = \frac{\sigma'_t(z)}{\sigma_t(z)}(x - \mu_t(z)) + \mu'_t(z)$$

Flow Matching

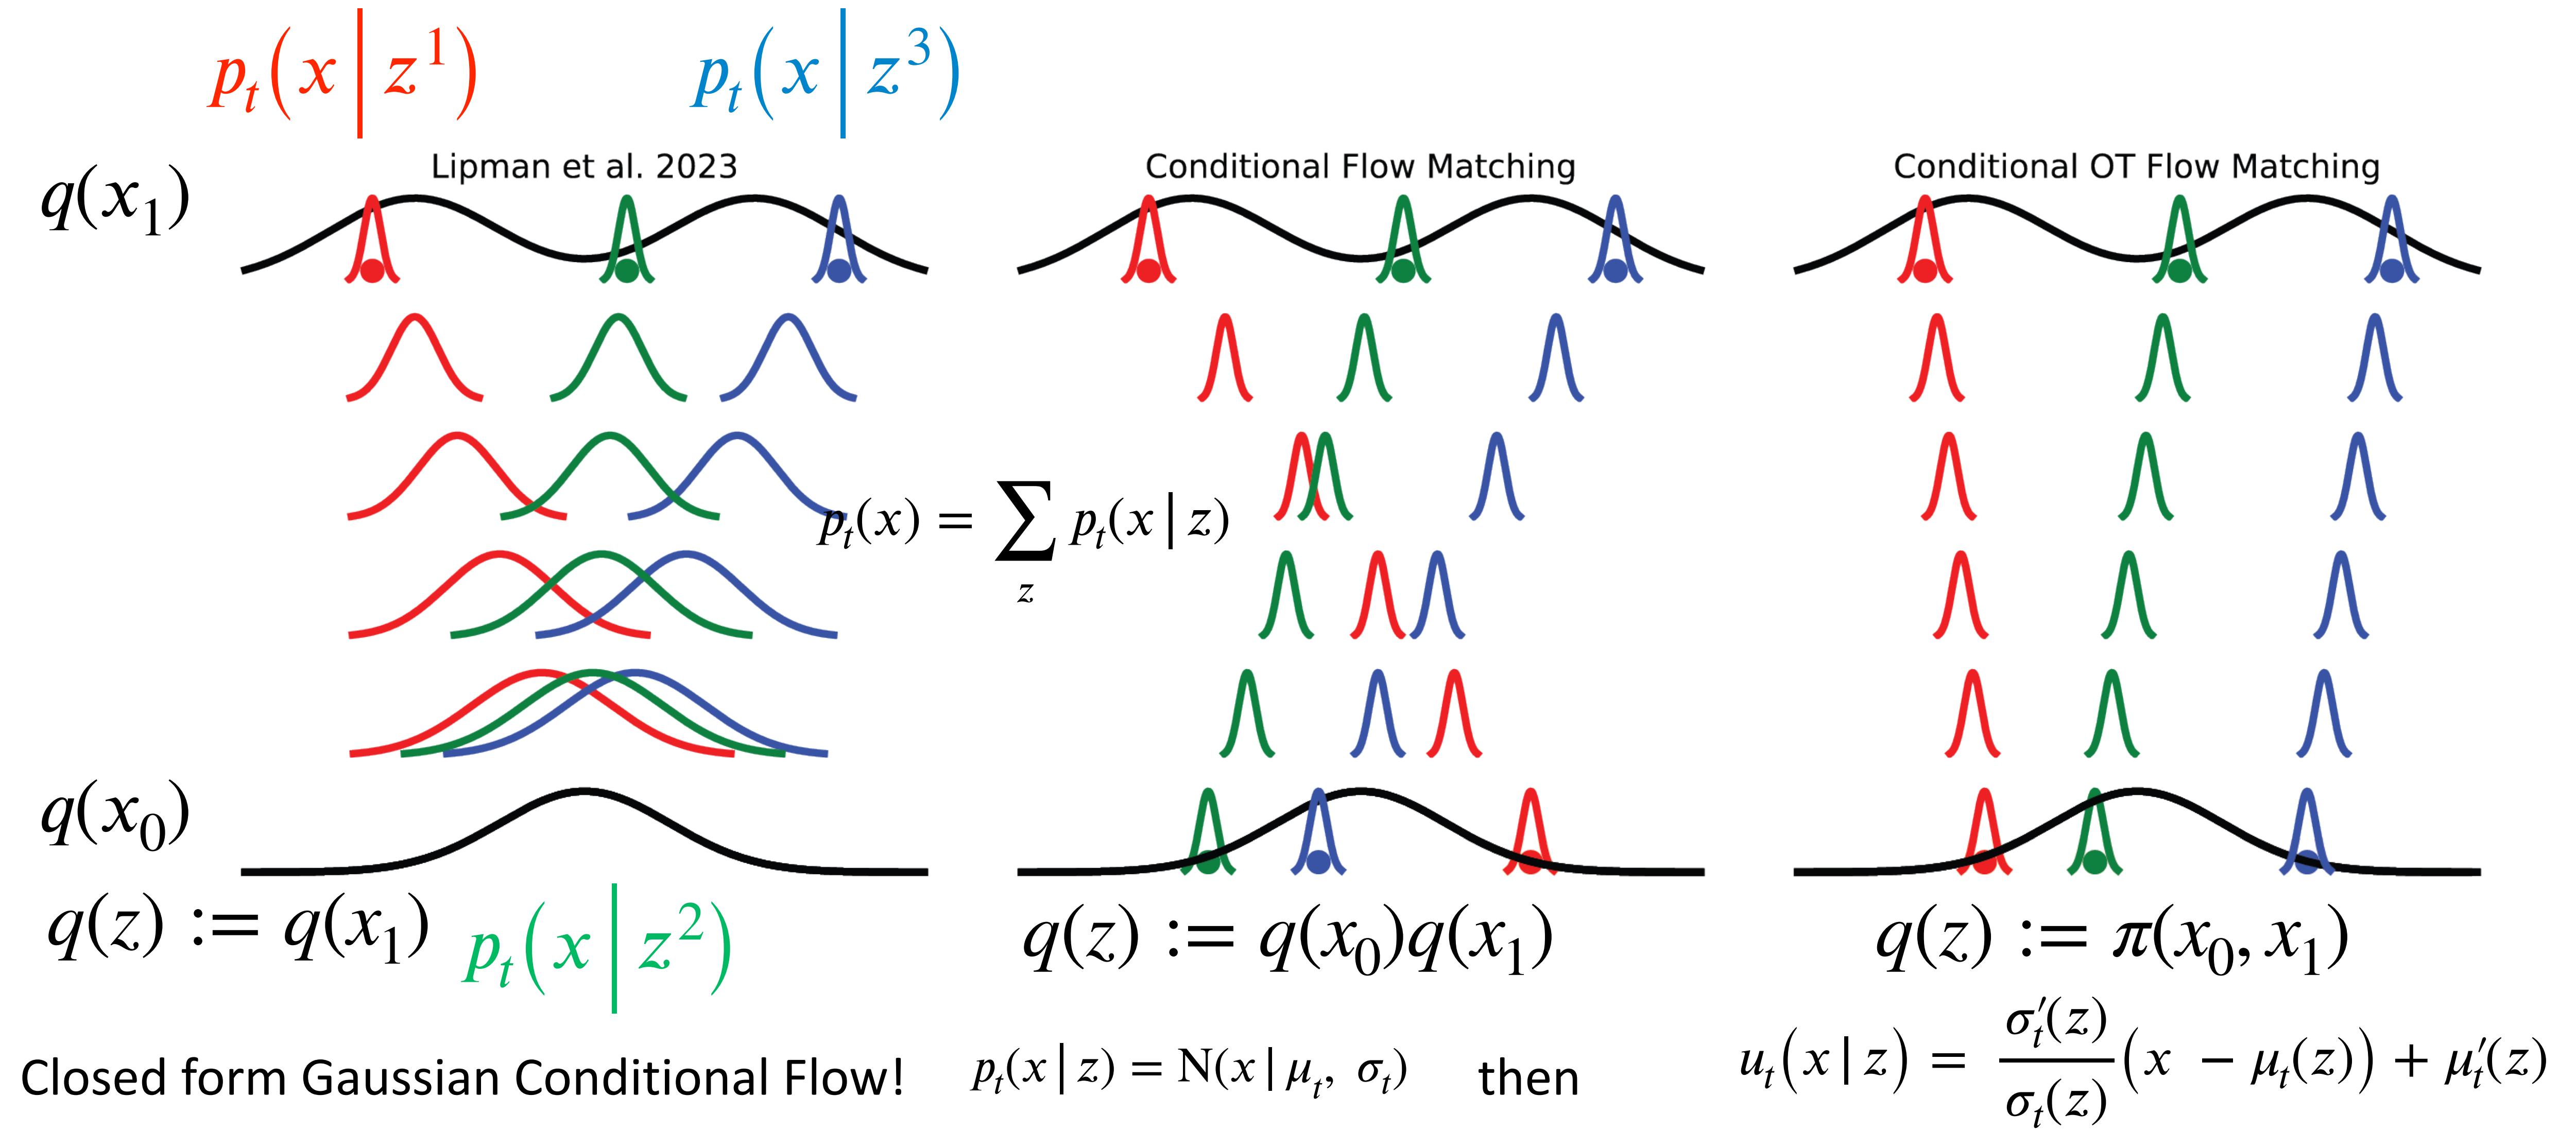


Marginalization for $u_t(x)$ that generates $p_t(x)$ given $q(z)$

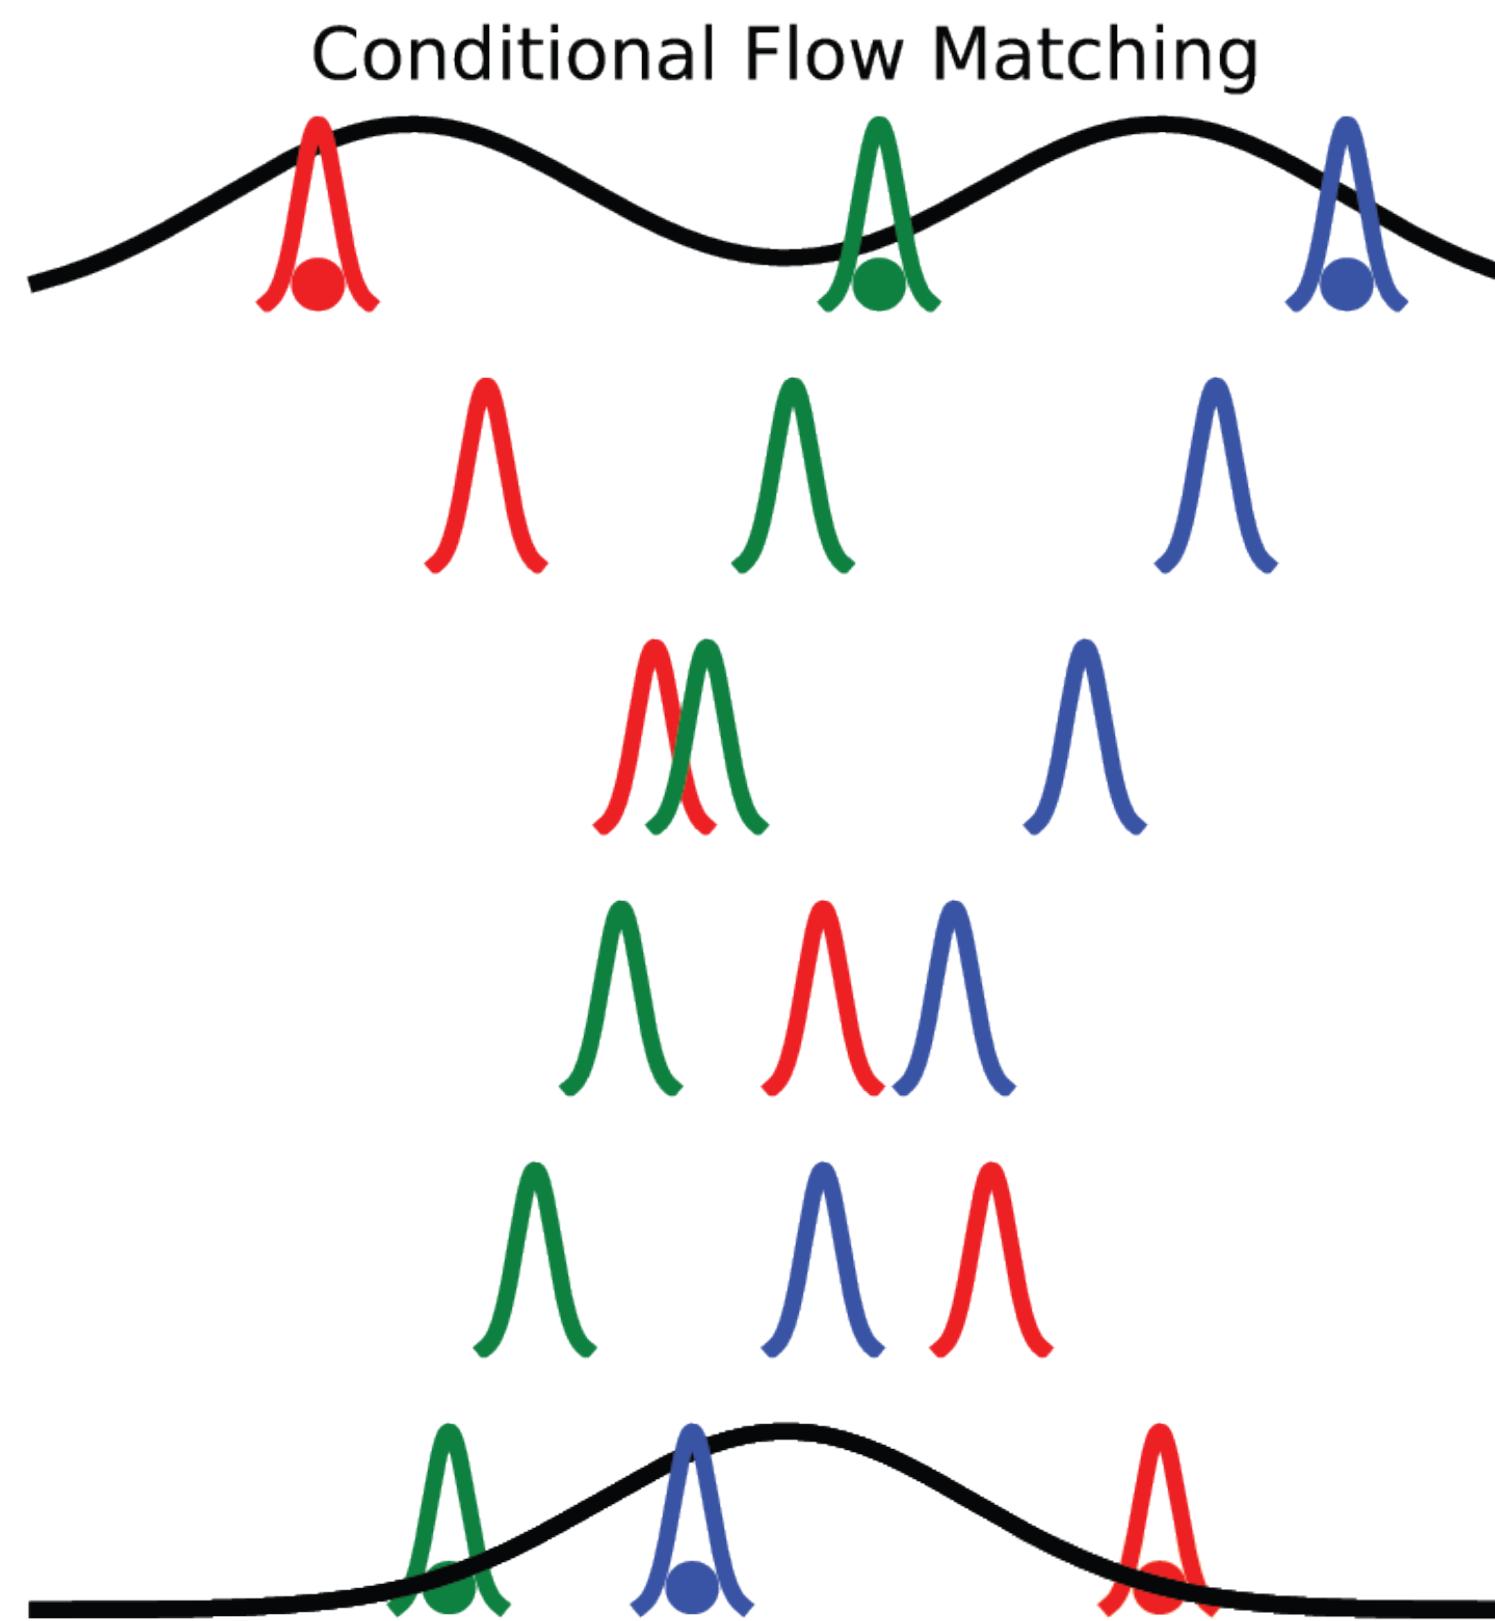
$$u_t(x) = \int \frac{u_t(x | z)p_t(x | z)q(z)}{p_t(x)} dz$$

Closed form Gaussian Conditional Flow! $p_t(x | z) = N(x | \mu_t, \sigma_t)$ then $u_t(x | z) = \frac{\sigma'_t(z)}{\sigma_t(z)}(x - \mu_t(z)) + \mu'_t(z)$

Conditional Flow Matching



Objective: $L_{CFM}(\theta) = \mathbb{E}_{t, q(z), p_t(x|z)} \|v_\theta(t, x) - u_t(x|z)\|_2^2$



Algorithm 1 Simplified Conditional Flow Matching

Input: Sample-able distributions $\mathbf{X}_0, \mathbf{X}_1$, bandwidth σ , batchsize b , initial network v_θ .

while Training **do**

/* Sample batches of size b i.i.d. from the datasets */

$\mathbf{x}_0 \sim \mathbf{X}_0; \quad \mathbf{x}_1 \sim \mathbf{X}_1$

$t \sim \mathcal{U}(0, 1)$

$\mu_t \leftarrow t\mathbf{x}_1 + (1 - t)\mathbf{x}_0$

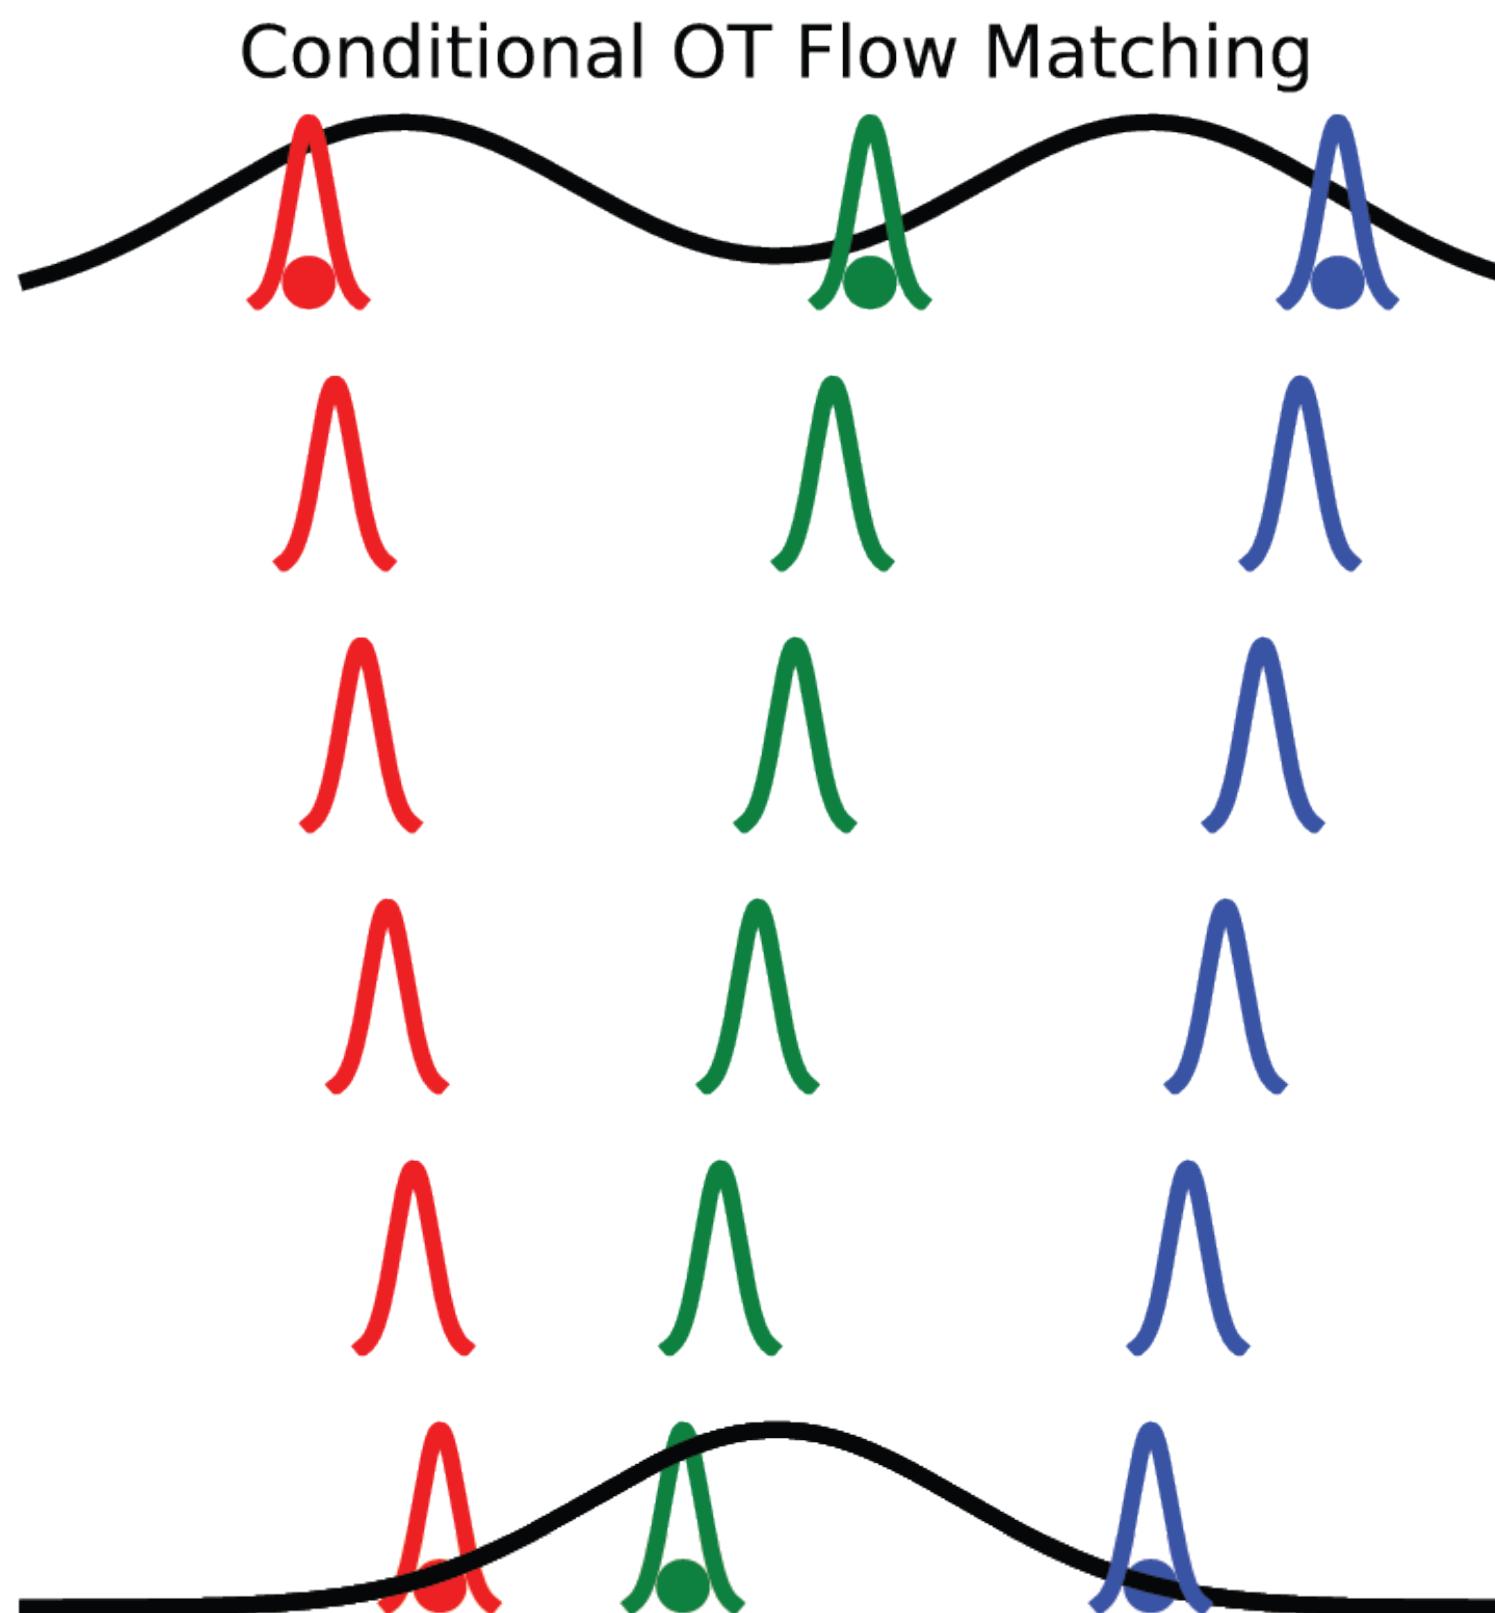
$\mathbf{x} \sim \mathcal{N}(\mu_t, \sigma^2 I)$

$L_{CFM}(\theta) \leftarrow \|v_\theta(t, x) - (\mathbf{x}_1 - \mathbf{x}_0)\|^2$

$\theta \leftarrow \text{Update}(\theta, \nabla_\theta L_{CFM}(\theta))$

return v_θ

+ Static (mini-batch) OT Resampling Step



Algorithm 2 Minibatch OT Conditional Flow Matching

Input: Sample-able distributions $\mathbf{X}_0, \mathbf{X}_1$, bandwidth σ , batch size b , initial network v_θ .

while Training **do**

/* Sample batches of size b i.i.d. from the datasets */

$$\mathbf{x}_0 \sim \mathbf{X}_0; \quad \mathbf{x}_1 \sim \mathbf{X}_1$$

$$\pi \leftarrow \text{OT}(\mathbf{x}_1, \mathbf{x}_0)$$

$$(\mathbf{x}_0, \mathbf{x}_1) \sim \pi$$

$$t \sim \mathcal{U}(0, 1)$$

$$\mu_t \leftarrow t\mathbf{x}_1 + (1 - t)\mathbf{x}_0$$

$$\mathbf{x} \sim \mathcal{N}(\mu_t, \sigma^2 I)$$

$$L_{CFM}(\theta) \leftarrow \|v_\theta(t, \mathbf{x}) - (\mathbf{x}_1 - \mathbf{x}_0)\|^2$$

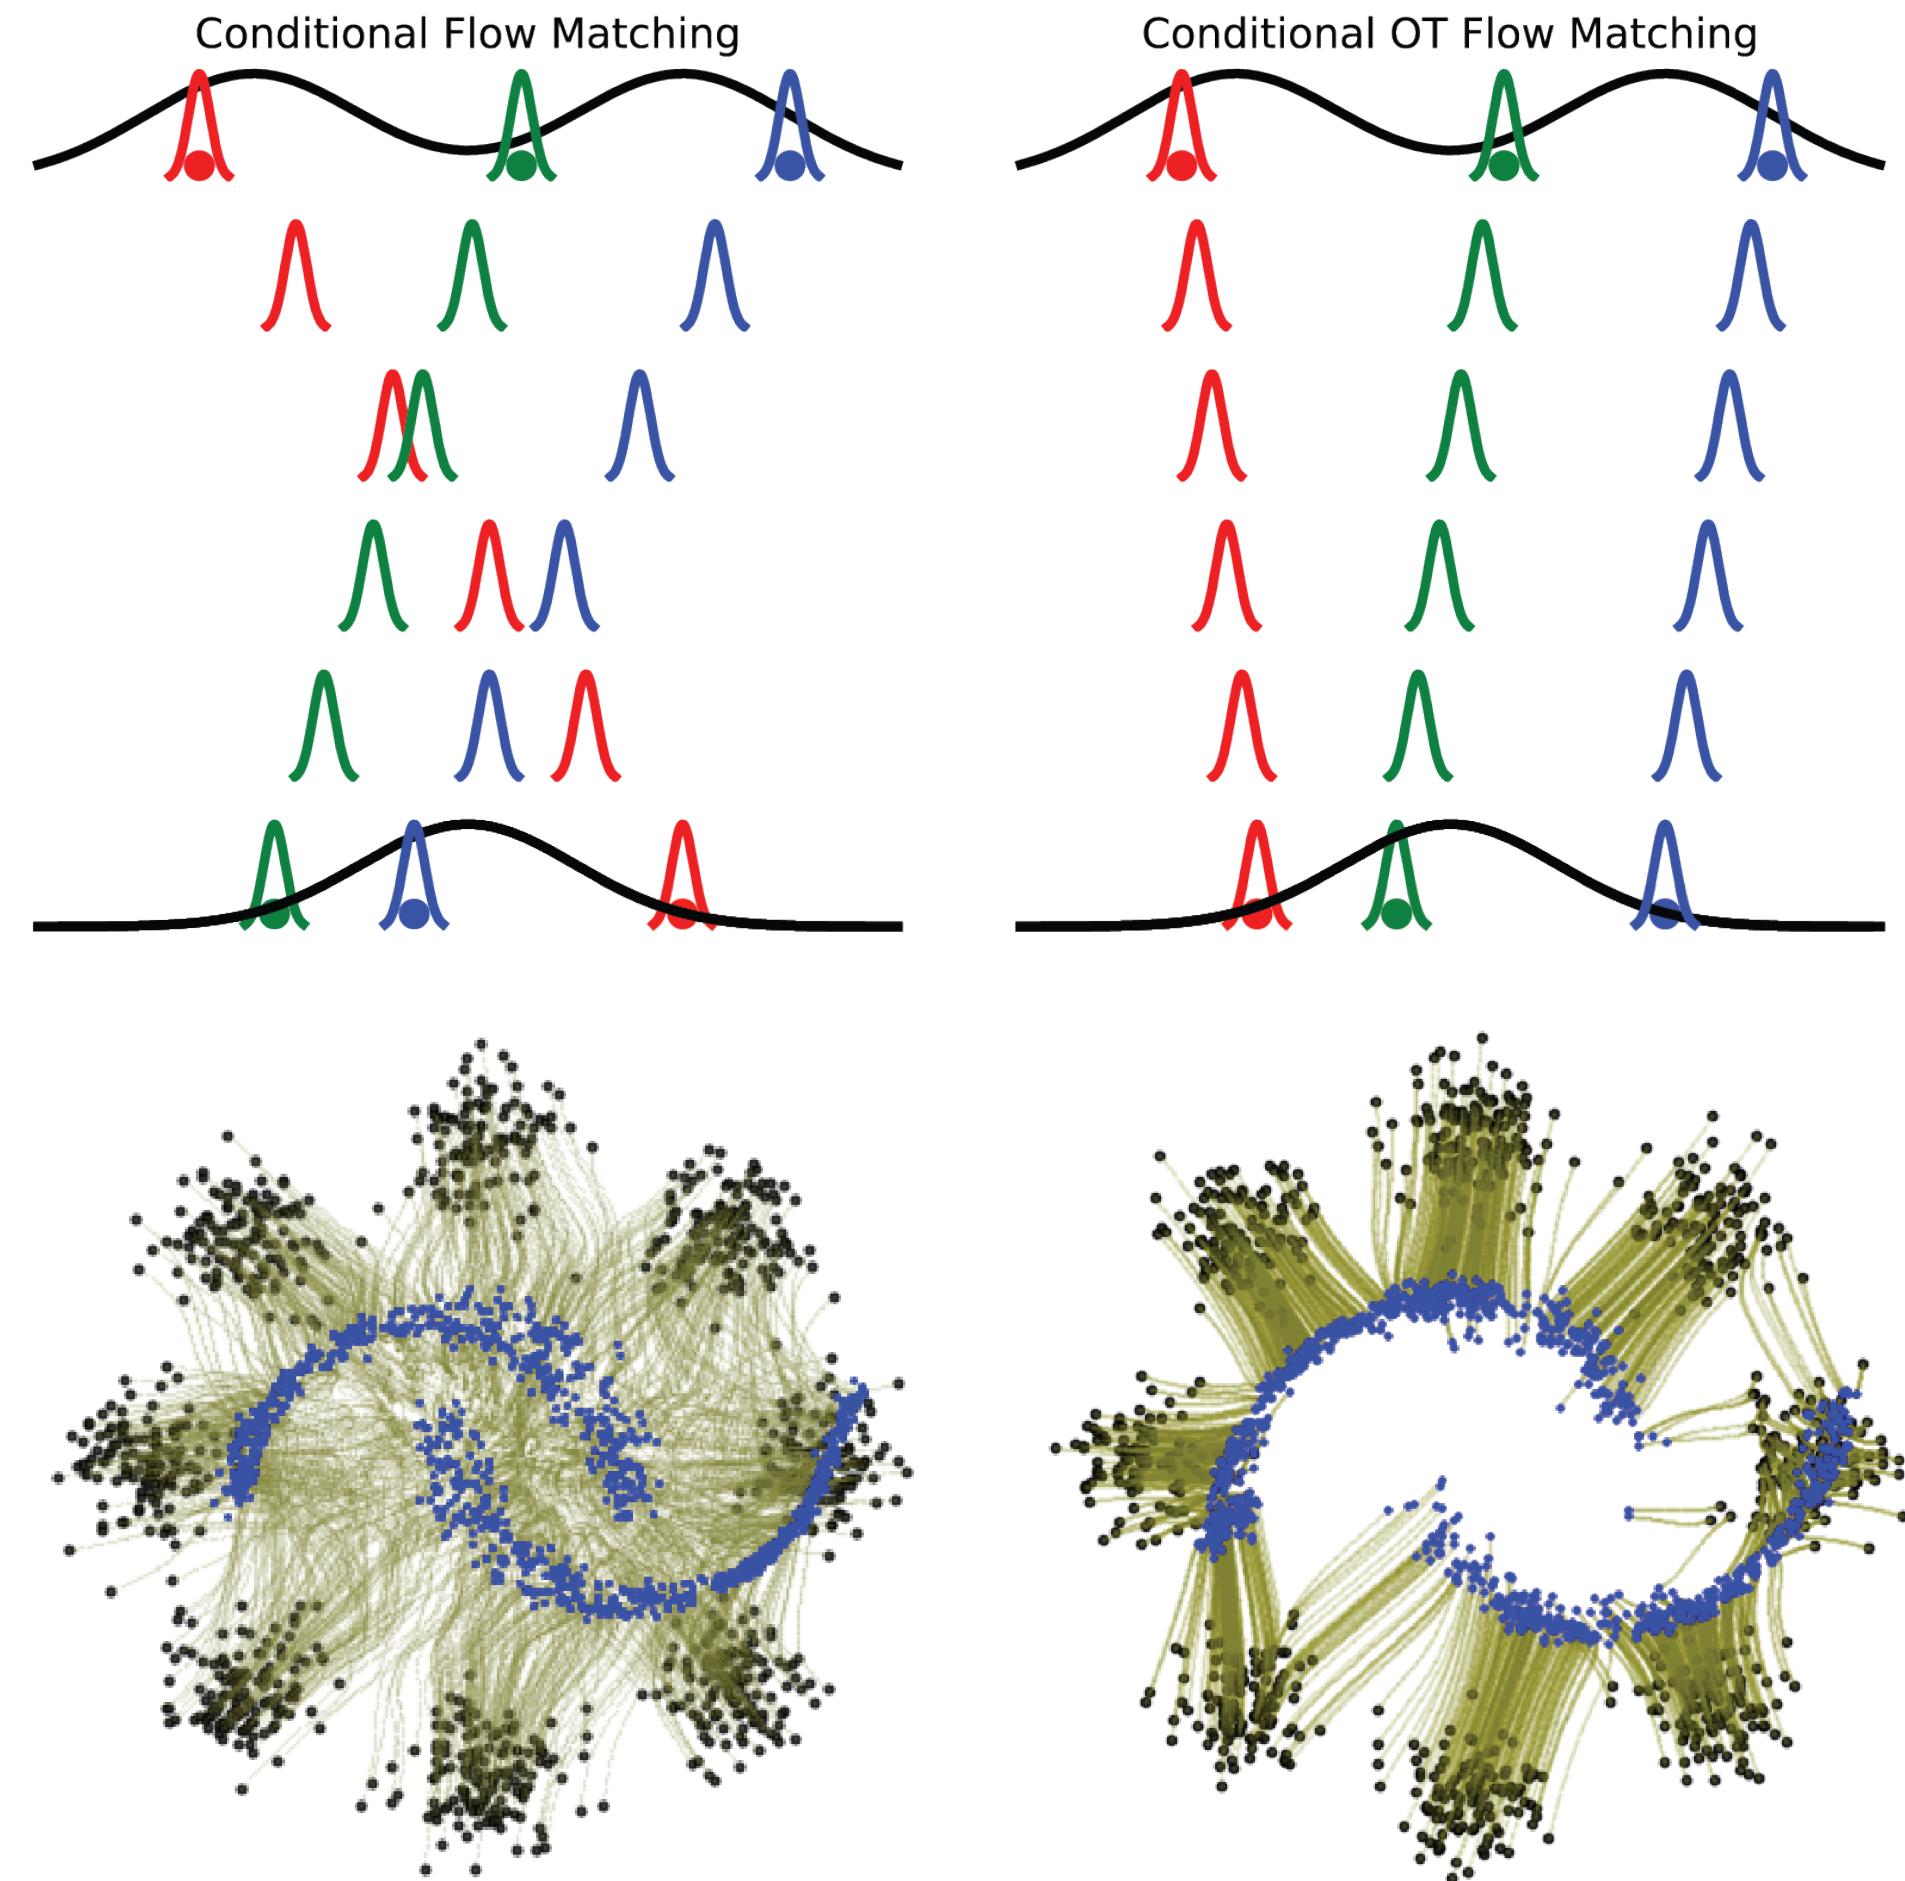
$$\theta \leftarrow \text{Update}(\theta, \nabla_\theta L_{CFM}(\theta))$$

return v_θ

Why use optimal transport in flow matching?

More general framework:

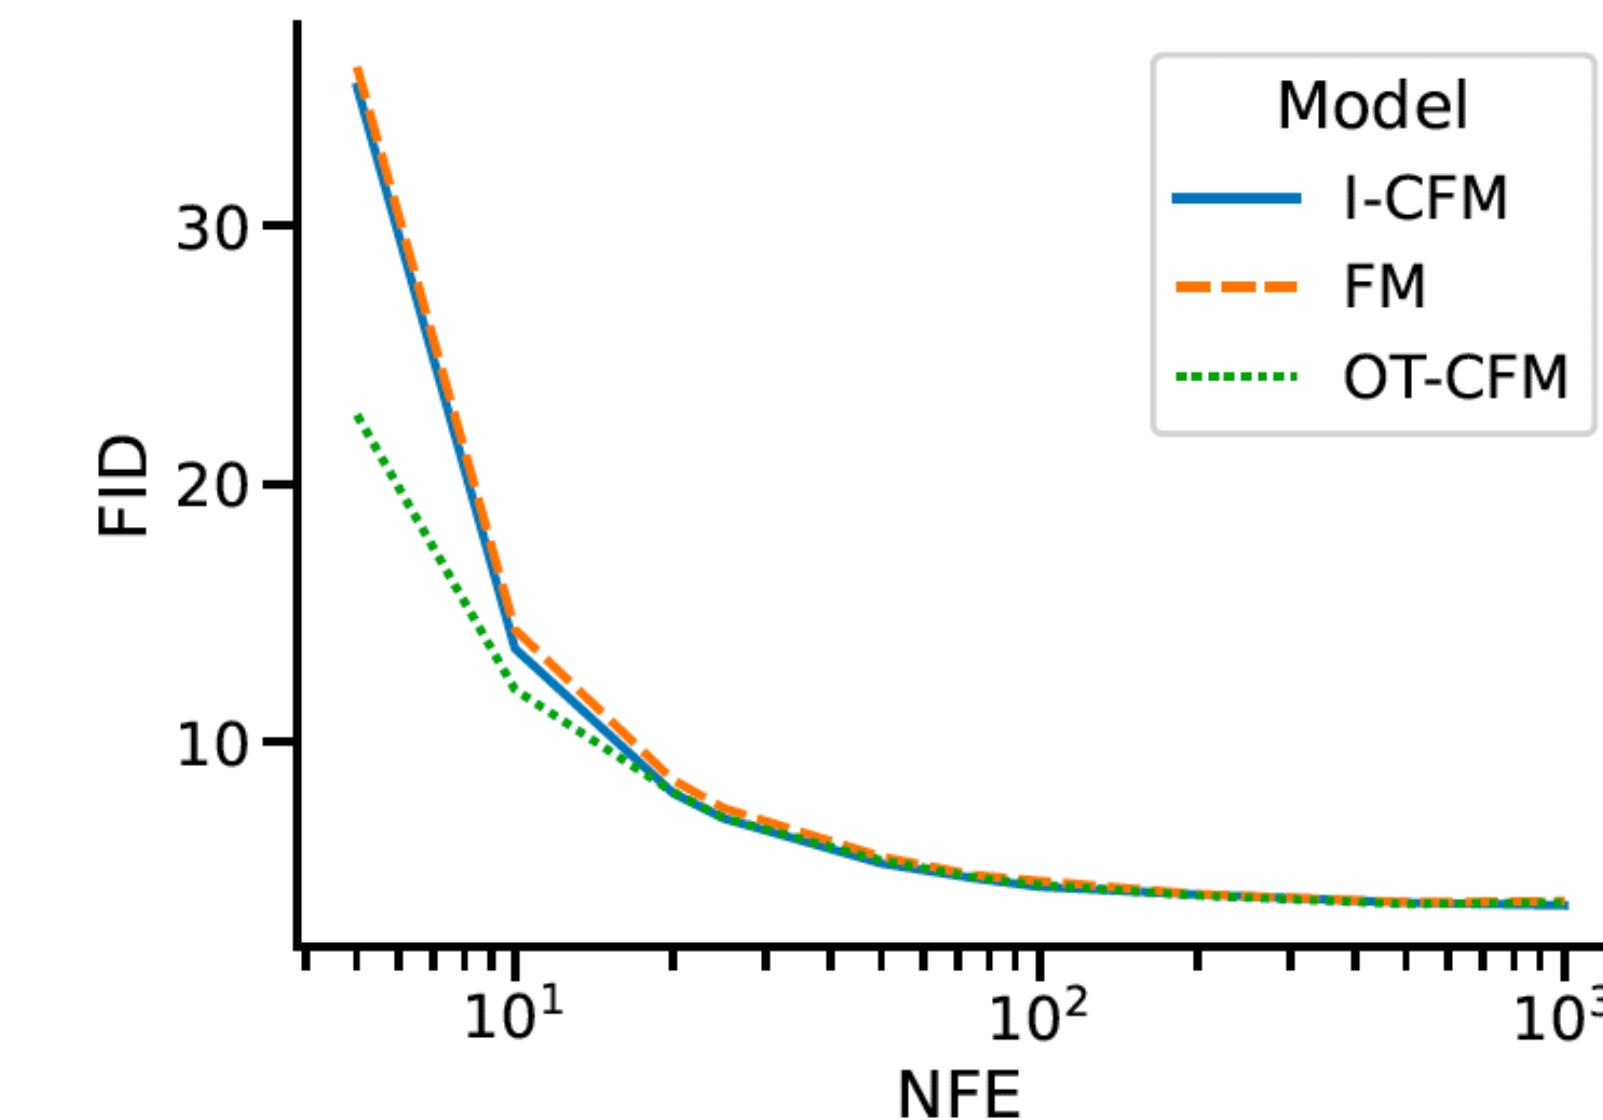
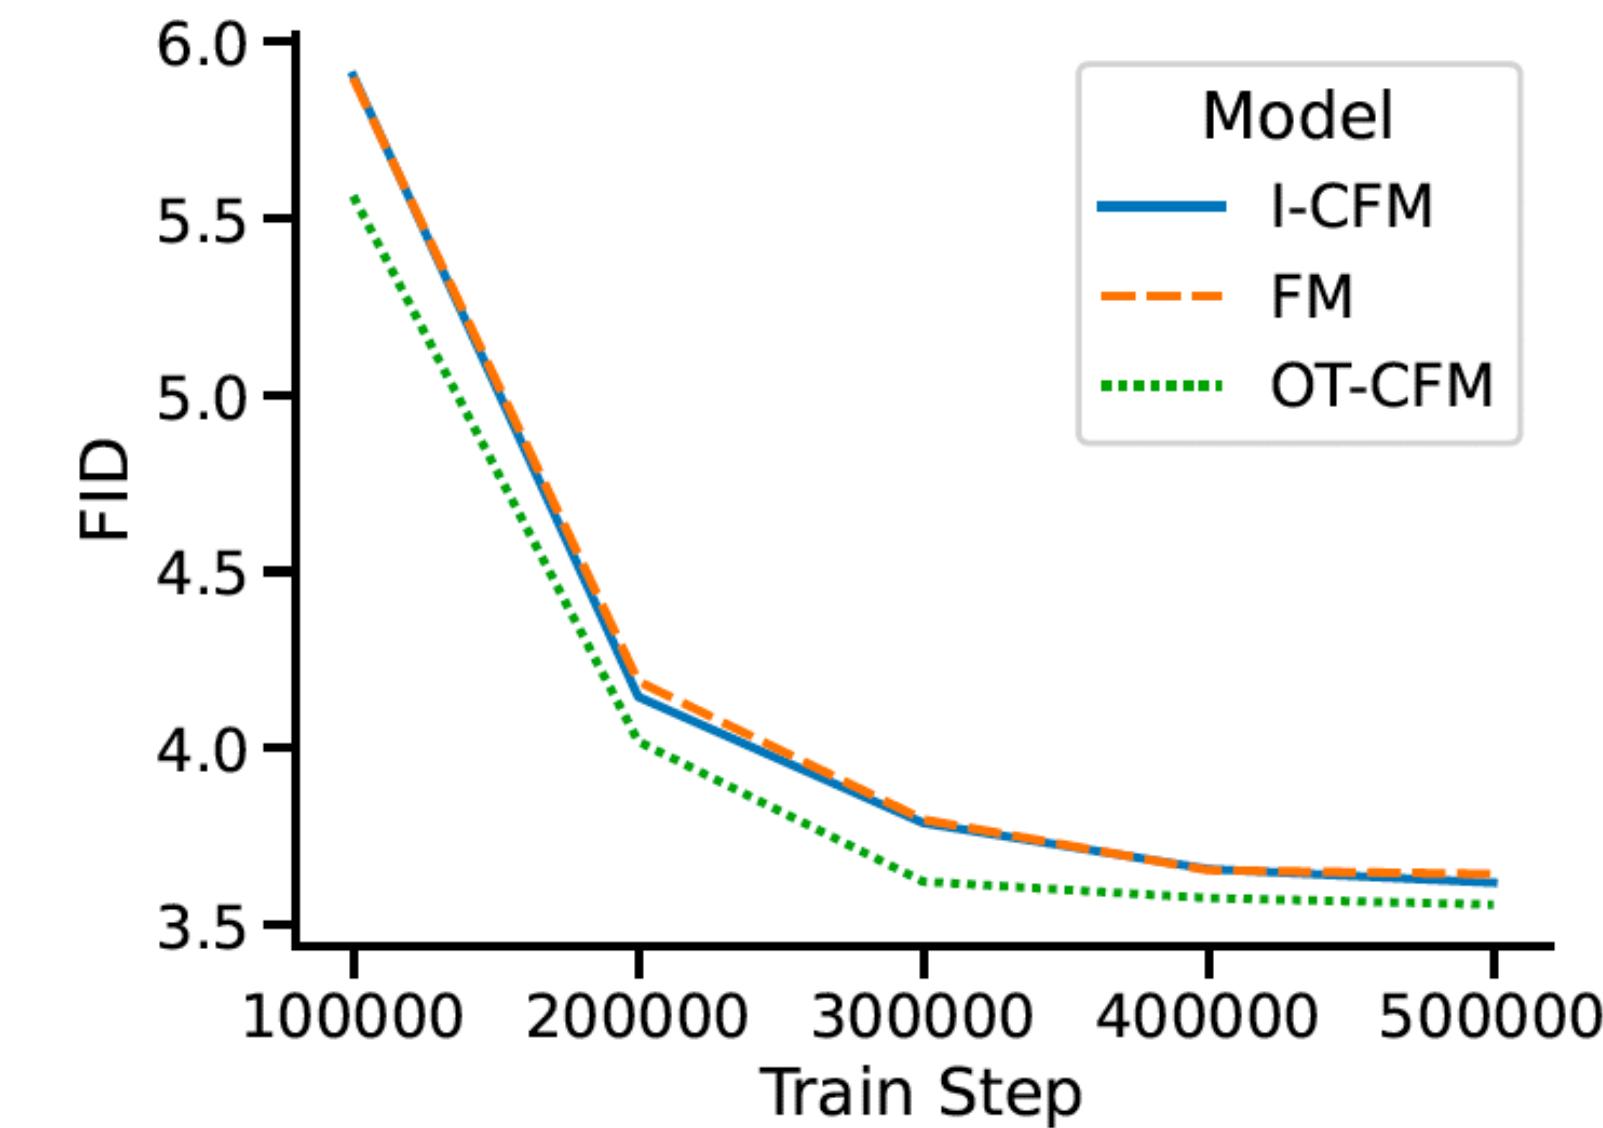
- Reduced variance in the objective via optimal transport leads to **faster training**
- Straighter inference paths via optimal transport leads to **faster inference**
- Can be applied to new problems where we care about the paths



Why use optimal transport in flow matching?

More general framework:

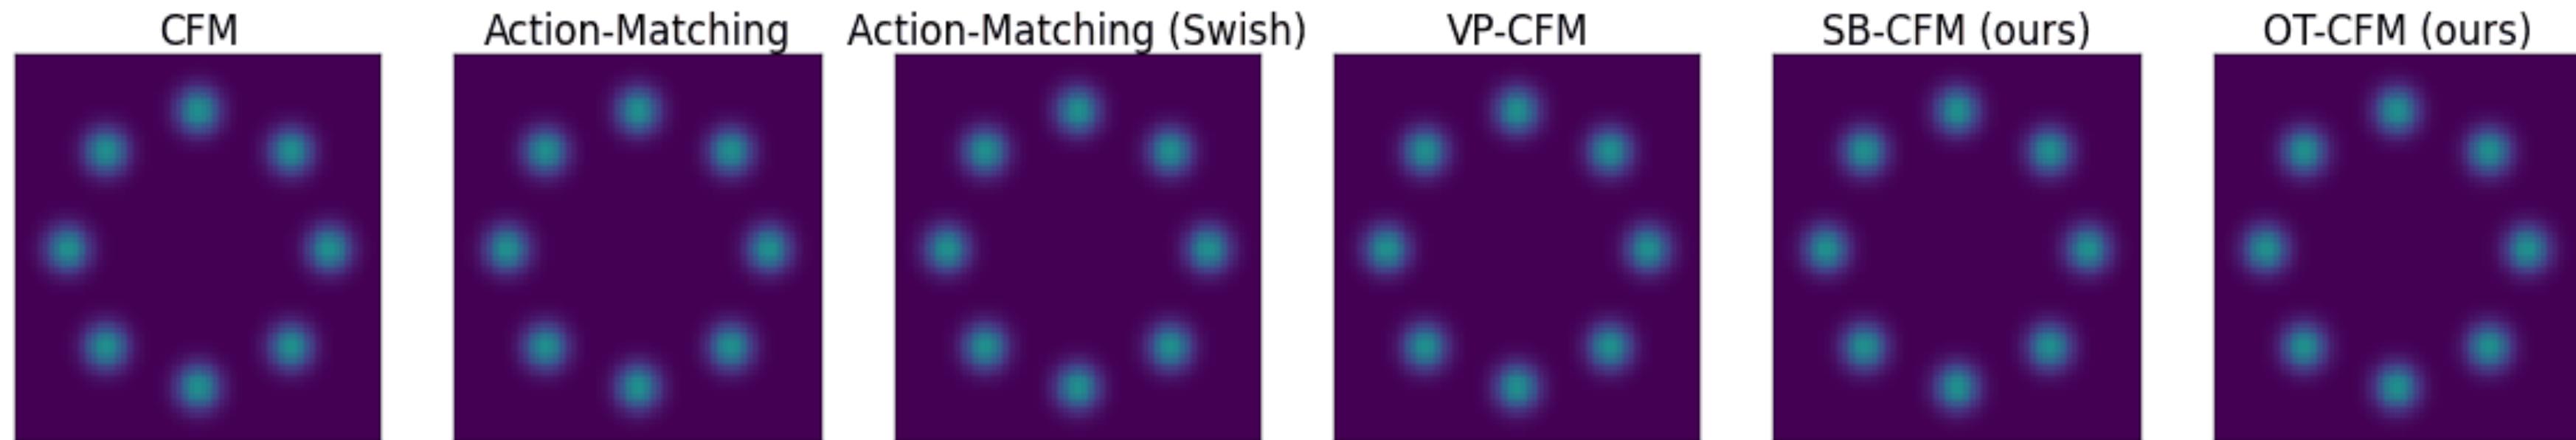
- Reduced variance in the objective via optimal transport leads to **faster training**
- Straighter inference paths via optimal transport leads to **faster inference**
- Can be applied to new problems where we care about the paths



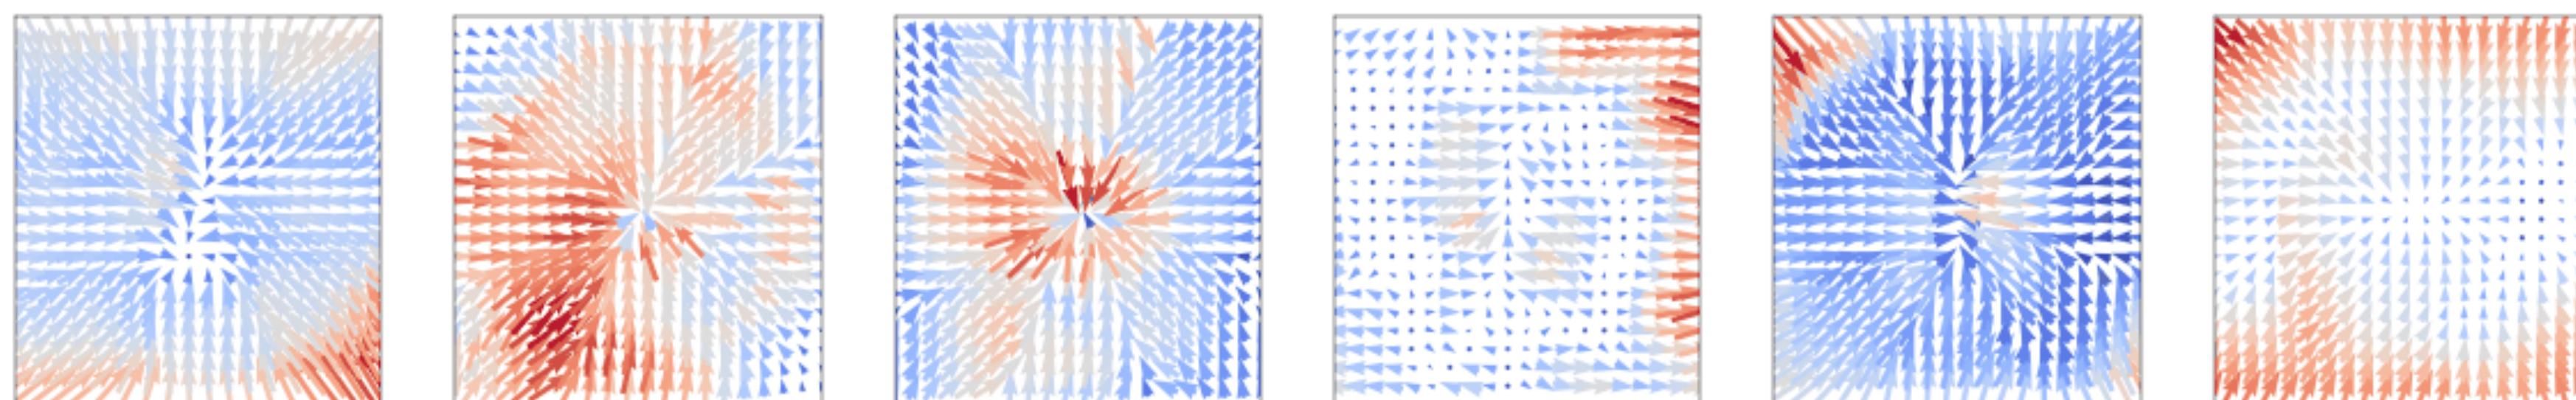
Comparing choices of $u_t(x | z)$, $p_t(x | z)$, and $q_t(z)$

8gaussians to Moons T=0.00

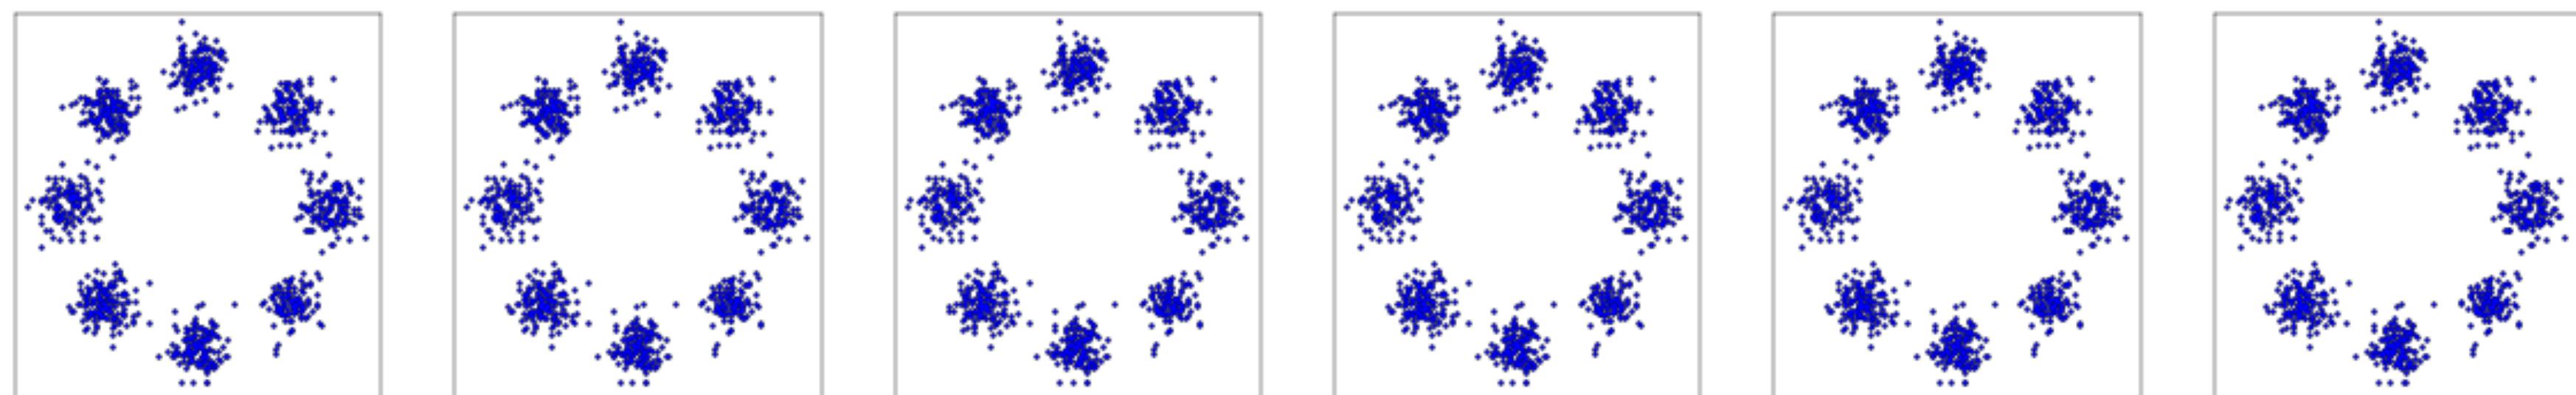
Marginal Probability
 $p_t(x)$



Marginal field $u_t(x)$



Time dependent flow $\psi_t(x_0)$



Flow Matching

$$L_{\text{FM}}(\theta) = \min \mathbb{E}_{t, p_t(x)} \|u_t^\theta(x) - u_t(x)\|^2$$

Add slide on types of prediction



$$L_{\text{CFM}}(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x) - u_t(x | z)\|^2$$

Construct:

- Target probability path p_t s.t. $p_0 = p$, $p_1 \approx q$
- Generating velocity field u_t

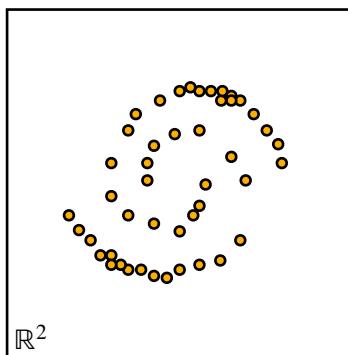


Find a tractable optimization objective



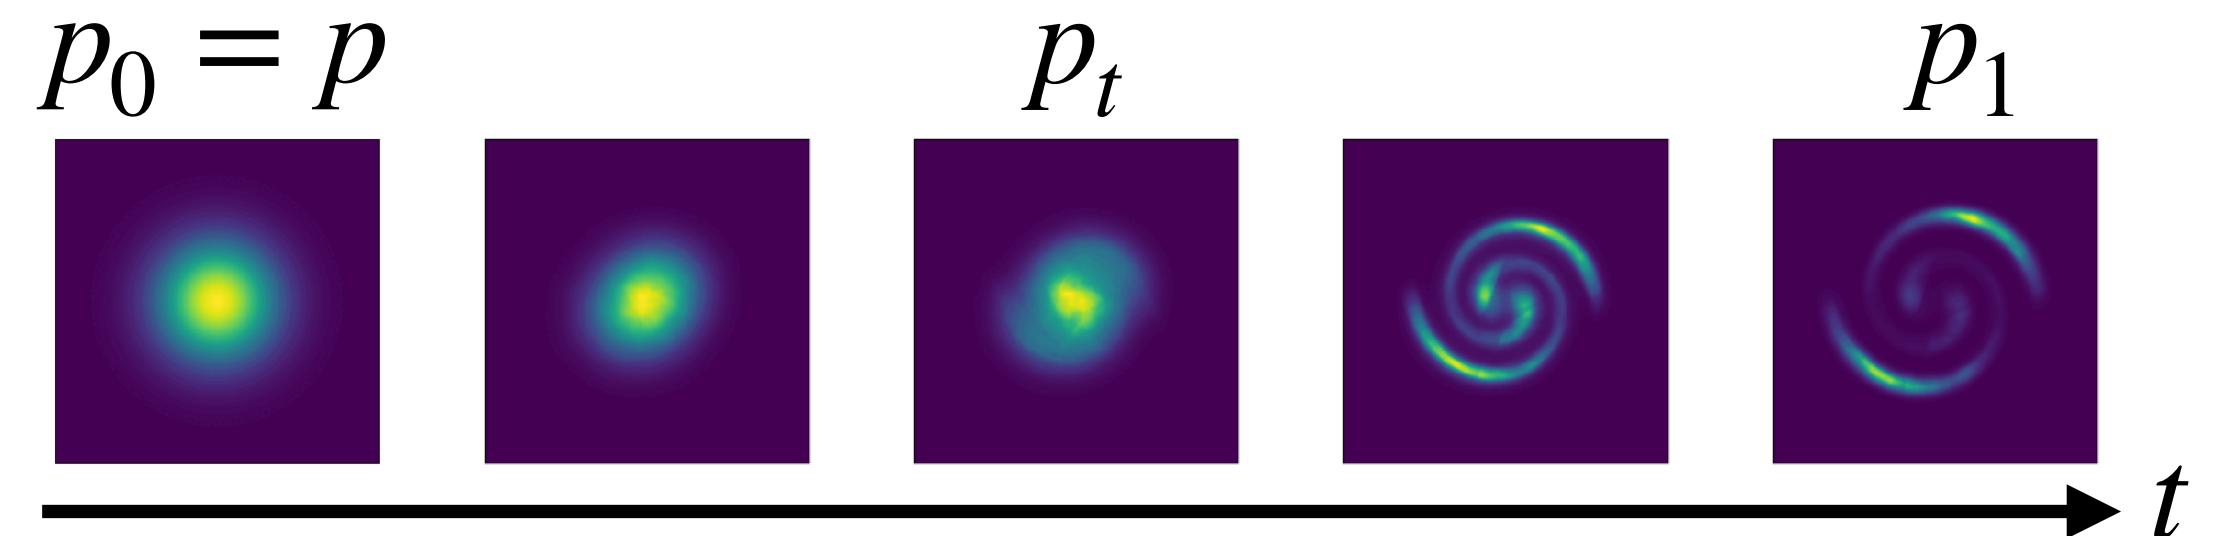
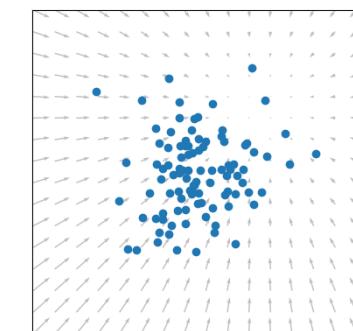
Recipe: Flow Matching

- Given: samples $x_1 \sim q$



- Construct: p_t s.t. $p_0 = p$, $p_1 \approx q$

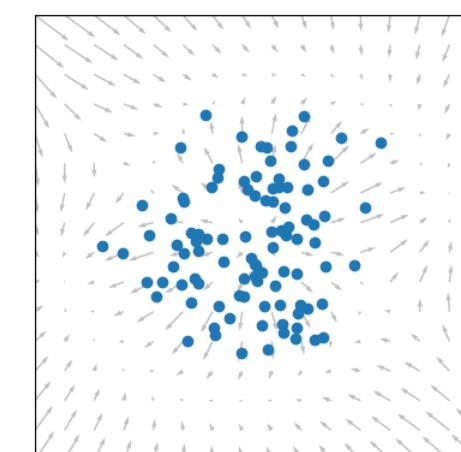
via conditional flows $\psi_t(x | z)$



- Learn: velocity field u_t with CFM loss

s.t. $\psi_t(x_0) \sim p_t$ where $x_0 \sim p$

$$L_{\text{CFM}}(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x) - u_t(x|z)\|^2$$



$$\psi_t : [0,1] \times \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

Recipe: Flow Matching

- Given: samples $x_1 \sim q$
- Design Choices:
 - Choose a source distribution p
 - Construct p_t s.t. $p_0 = p, p_1 \approx q$
 - Choose conditional flow $\psi_t(x | z)$
- Learn: velocity field u_t s.t. $\psi_t(x_0) \sim p_t$ where $x_0 \sim p$
 - We learn velocities by square regressing using L_2 -norm

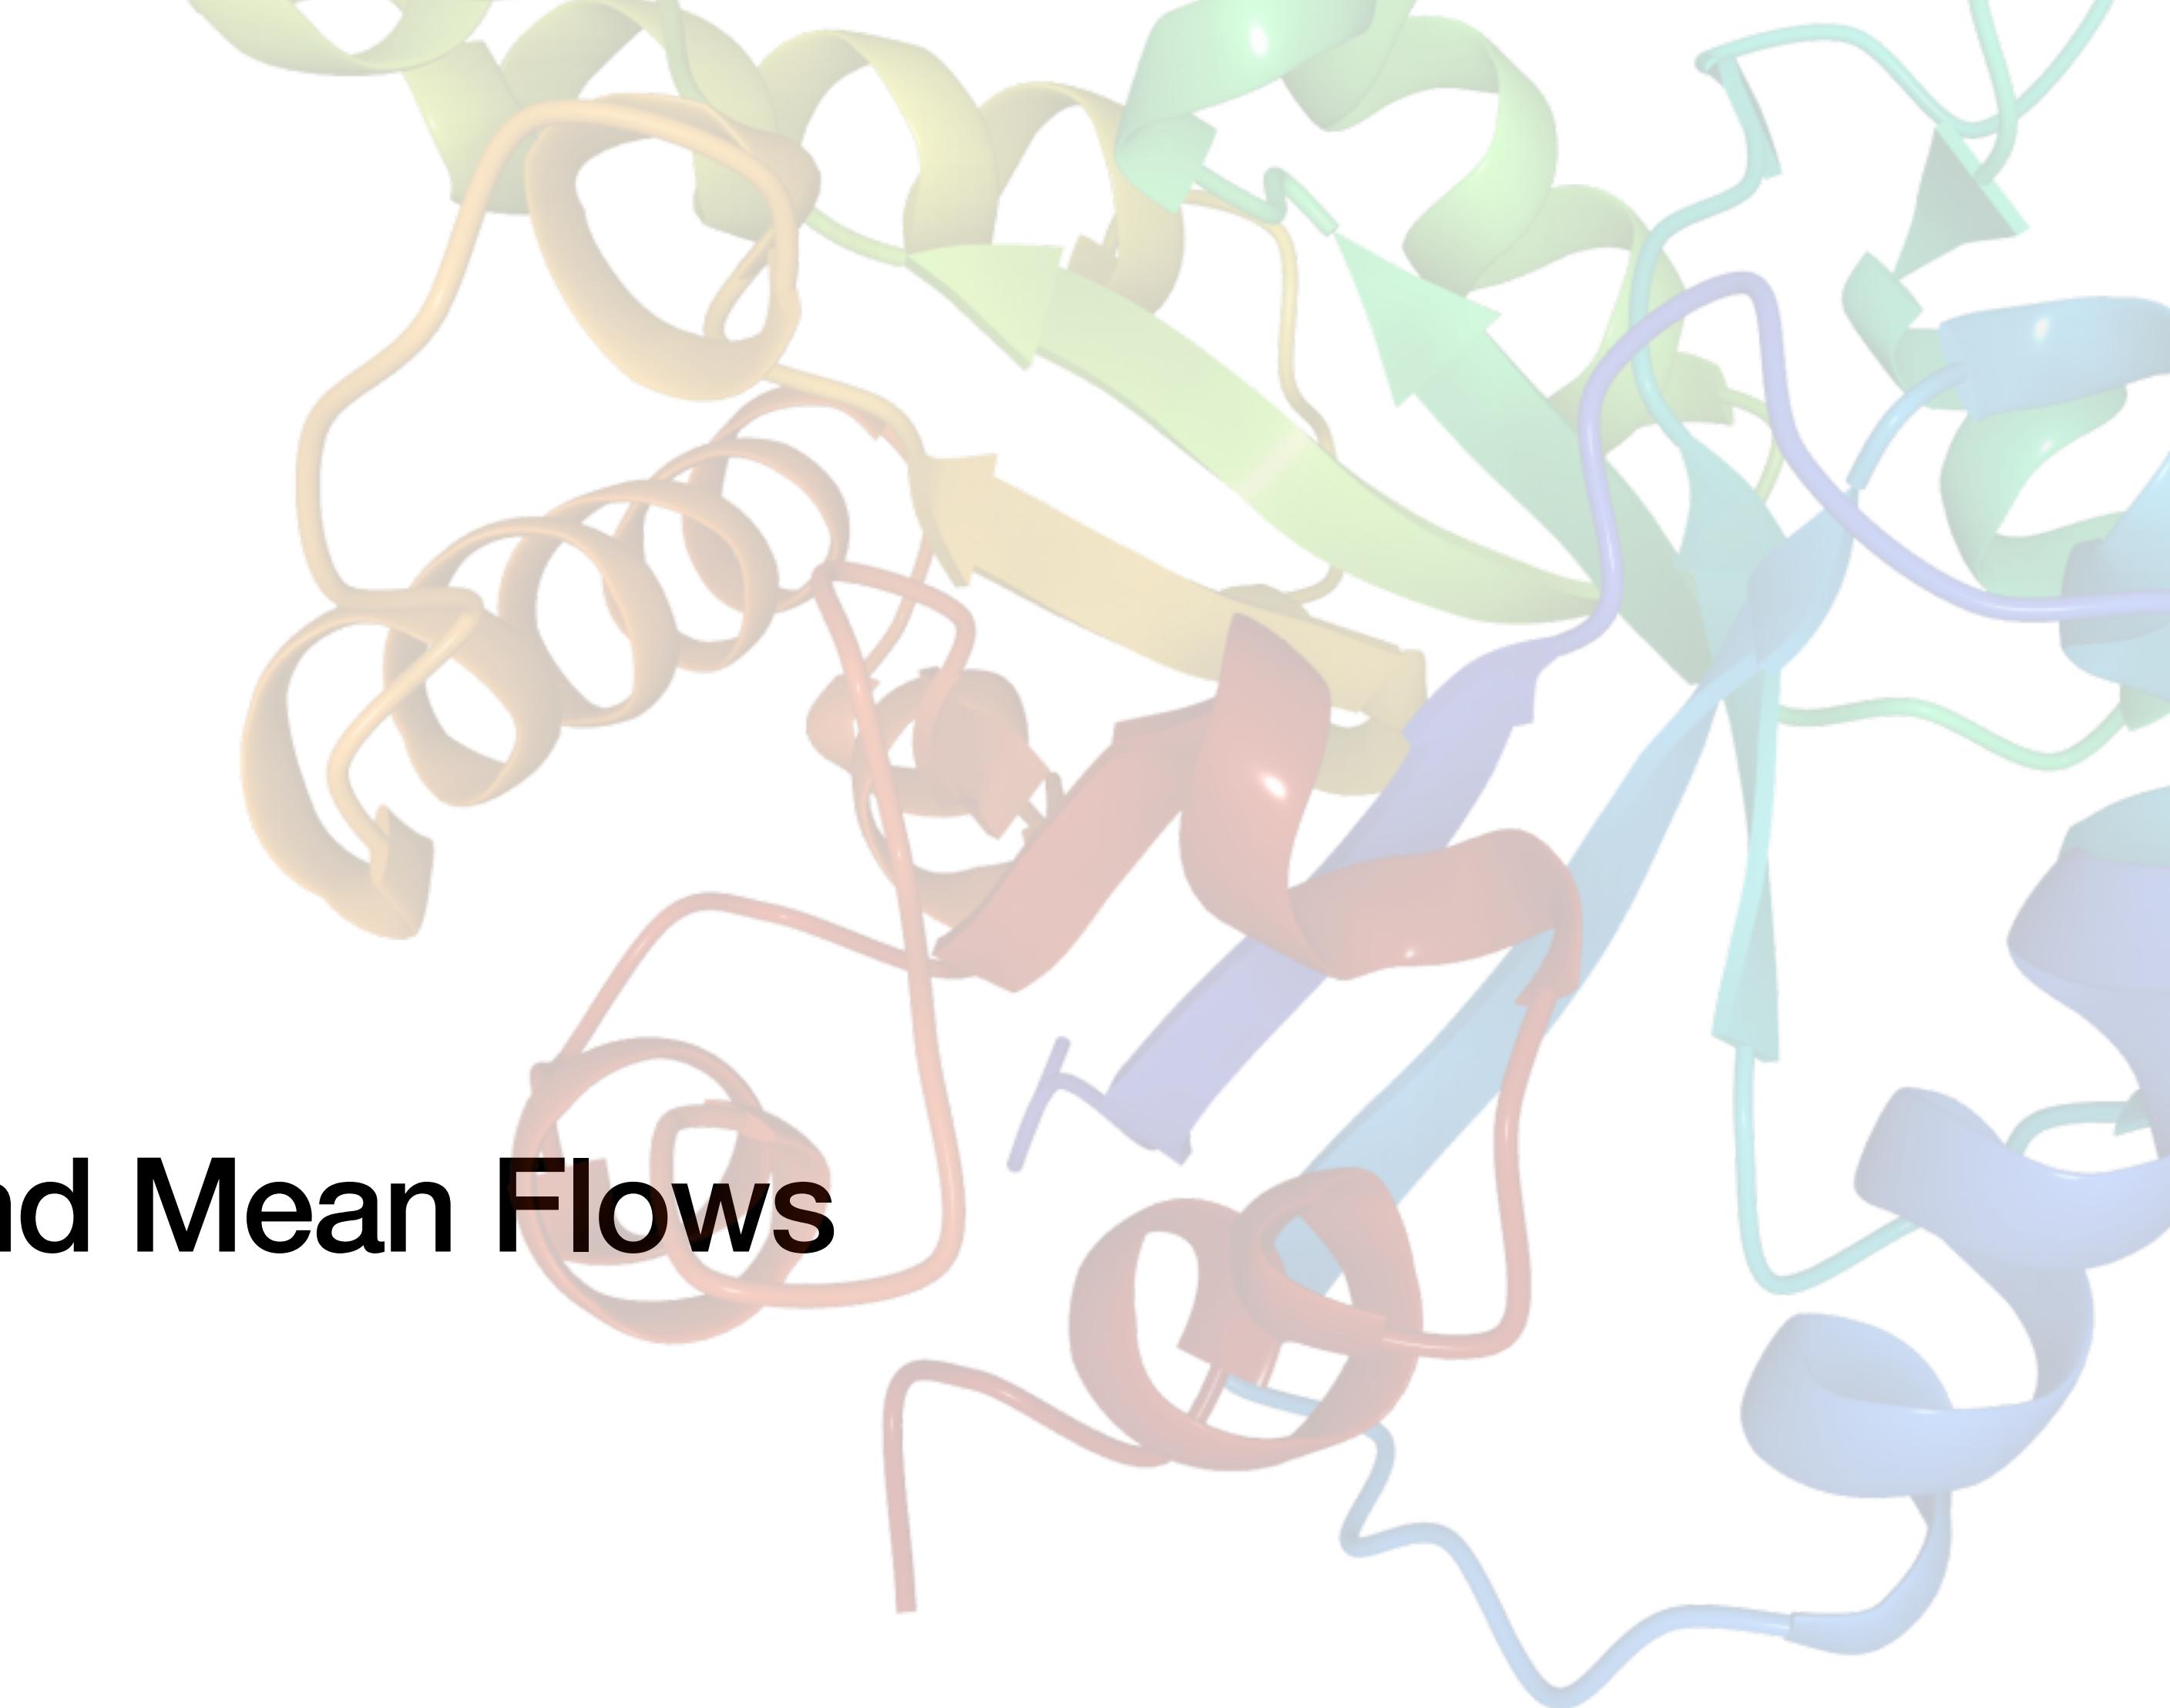
Exercise



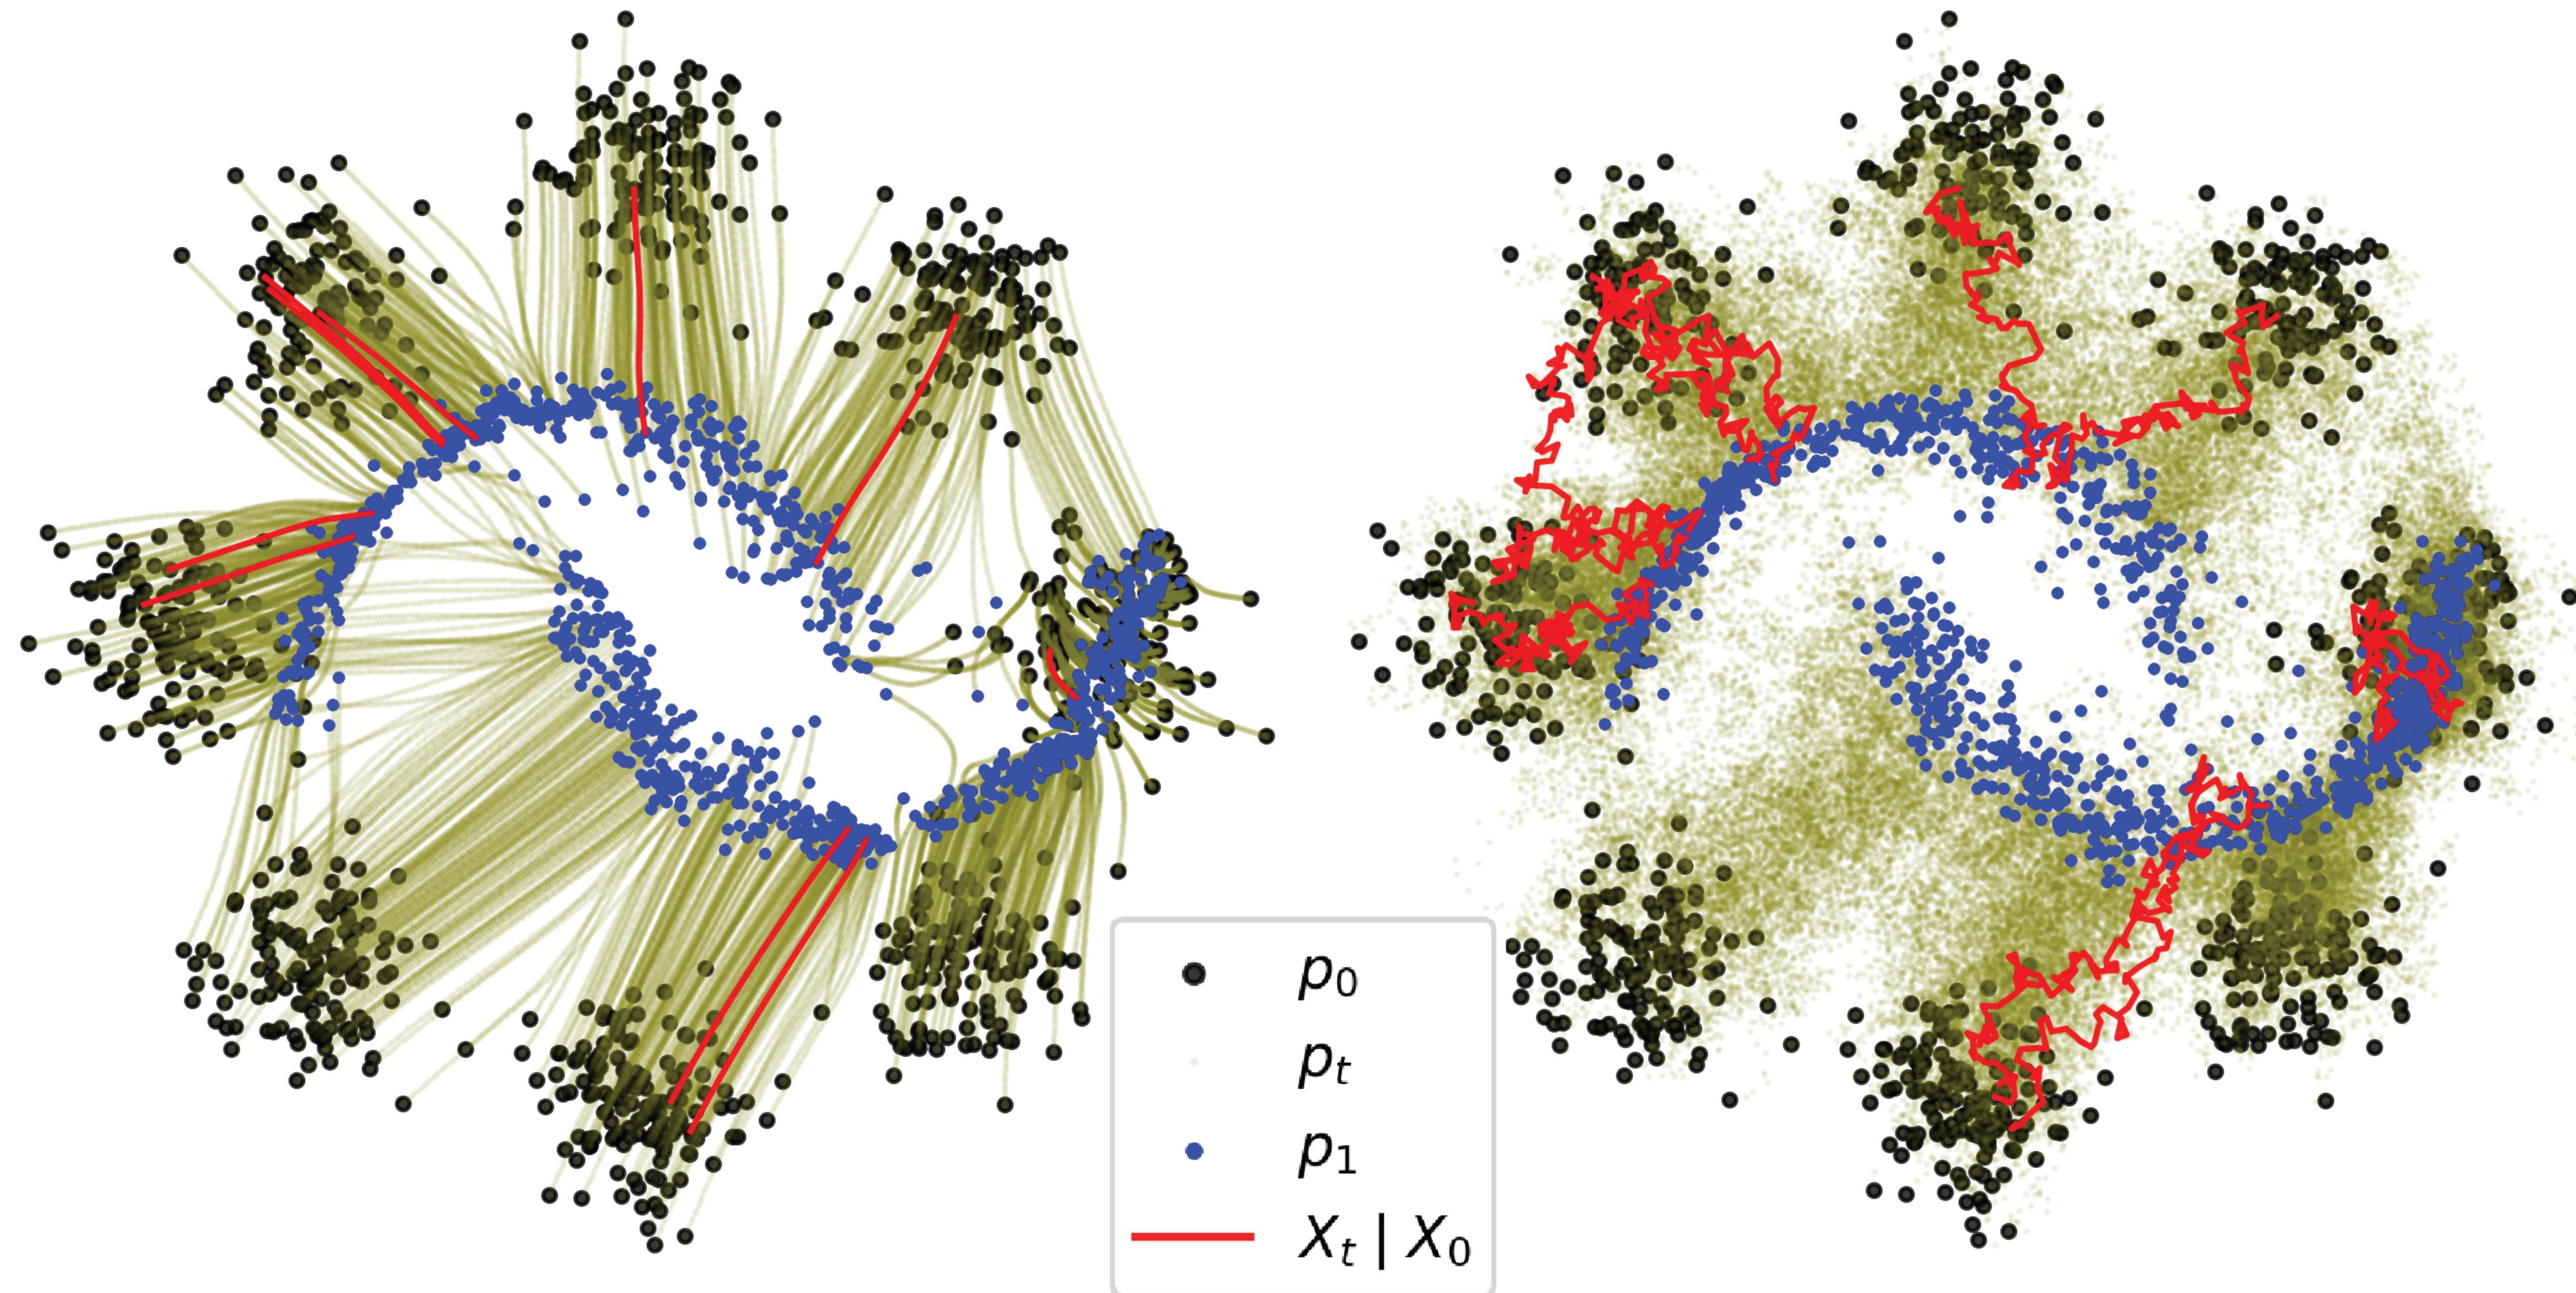
- Implement basic flow matching training and inference
- Implement optimal transport conditional flow matching
- Implement variance preserving variant
- How does it compare to MLE training?

<https://github.com/atong01/ProbAI-2025-flow-matching/>

Part III: Bridge Matching and Mean Flows



Bridge Matching



Let's first recall diffusion!

Diffusion and Score Models

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data → Noise

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t$$

Noise → Data

Learn the score by regressing to conditional scores:

$$\min_{\theta} \mathbb{E}_{p_{data}, p_t(x|x_{data})} [\|s_t^{\theta}(x) - \nabla \log p_t(x|x_{data})\|^2]$$

Simulation-free

Known SDEs: Variance Exploding
Variance Preserving

Different Integration

Forward
SDE

$$dx_t = f_t(x_t)dt + g_t dw_t$$

Data → Noise

Reverse
SDE

$$d\bar{x}_t = (f_t(x_t) - g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t$$

Noise → Data

Forward
Probability
Flow ODE

$$dx_t = (f_t(x_t) - \frac{1}{2}g_t^2 \nabla \log p_t)dt$$

Data → Noise

Reverse
Probability
Flow ODE

$$d\bar{x}_t = - (f_t(x_t) - \frac{1}{2}g_t^2 \nabla \log p_t)dt$$

Noise → Data

Different Integration with deterministic drift

$$u_t = f_t - \frac{1}{2}g_t^2 \nabla \log p_t$$

Forward
SDE

$$dx_t = (u_t + \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t dw_t \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
SDE

$$d\bar{x}_t = (u_t - \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t \quad \text{Noise} \rightarrow \text{Data}$$

Forward
Probability
Flow ODE

$$dx_t = u_t(x_t)dt \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
Probability
Flow ODE

$$d\bar{x}_t = -u_t(x_t)dt \quad \text{Noise} \rightarrow \text{Data}$$

Different Integration with deterministic drift

$$u_t = f_t - \frac{1}{2}g_t^2 \nabla \log p_t$$

Bridge Matching!

Forward
SDE

$$dx_t = \left(u_t + \frac{1}{2}g_t^2 \nabla \log p_t \right) dt + g_t dw_t \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
SDE

$$d\bar{x}_t = \left(u_t - \frac{1}{2}g_t^2 \nabla \log p_t \right) dt + g_t d\bar{w}_t \quad \text{Noise} \rightarrow \text{Data}$$

Forward
Probability
Flow ODE

$$dx_t = u_t(x_t)dt \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
Probability
Flow ODE

$$d\bar{x}_t = -u_t(x_t)dt \quad \text{Noise} \rightarrow \text{Data}$$

Different Integration with deterministic drift

$$u_t = f_t - \frac{1}{2}g_t^2 \nabla \log p_t$$

Forward
SDE

$$dx_t = (u_t + \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t dw_t \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
SDE

$$d\bar{x}_t = (u_t - \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t \quad \text{Noise} \rightarrow \text{Data}$$

Forward
Probability
Flow ODE

$$dx_t = u_t(x_t)dt$$

Flow Matching!

Reverse
Probability
Flow ODE

$$d\bar{x}_t = -u_t(x_t)dt$$

Data → Noise

Noise → Data

Different Integration with deterministic drift

$$u_t = f_t - \frac{1}{2}g_t^2 \nabla \log p_t$$

Forward
SDE

$$dx_t = (u_t + \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t dw_t \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
SDE

$$d\bar{x}_t = (u_t - \frac{1}{2}g_t^2 \nabla \log p_t)dt + g_t d\bar{w}_t \quad \text{Noise} \rightarrow \text{Data}$$

Forward
Probability
Flow ODE

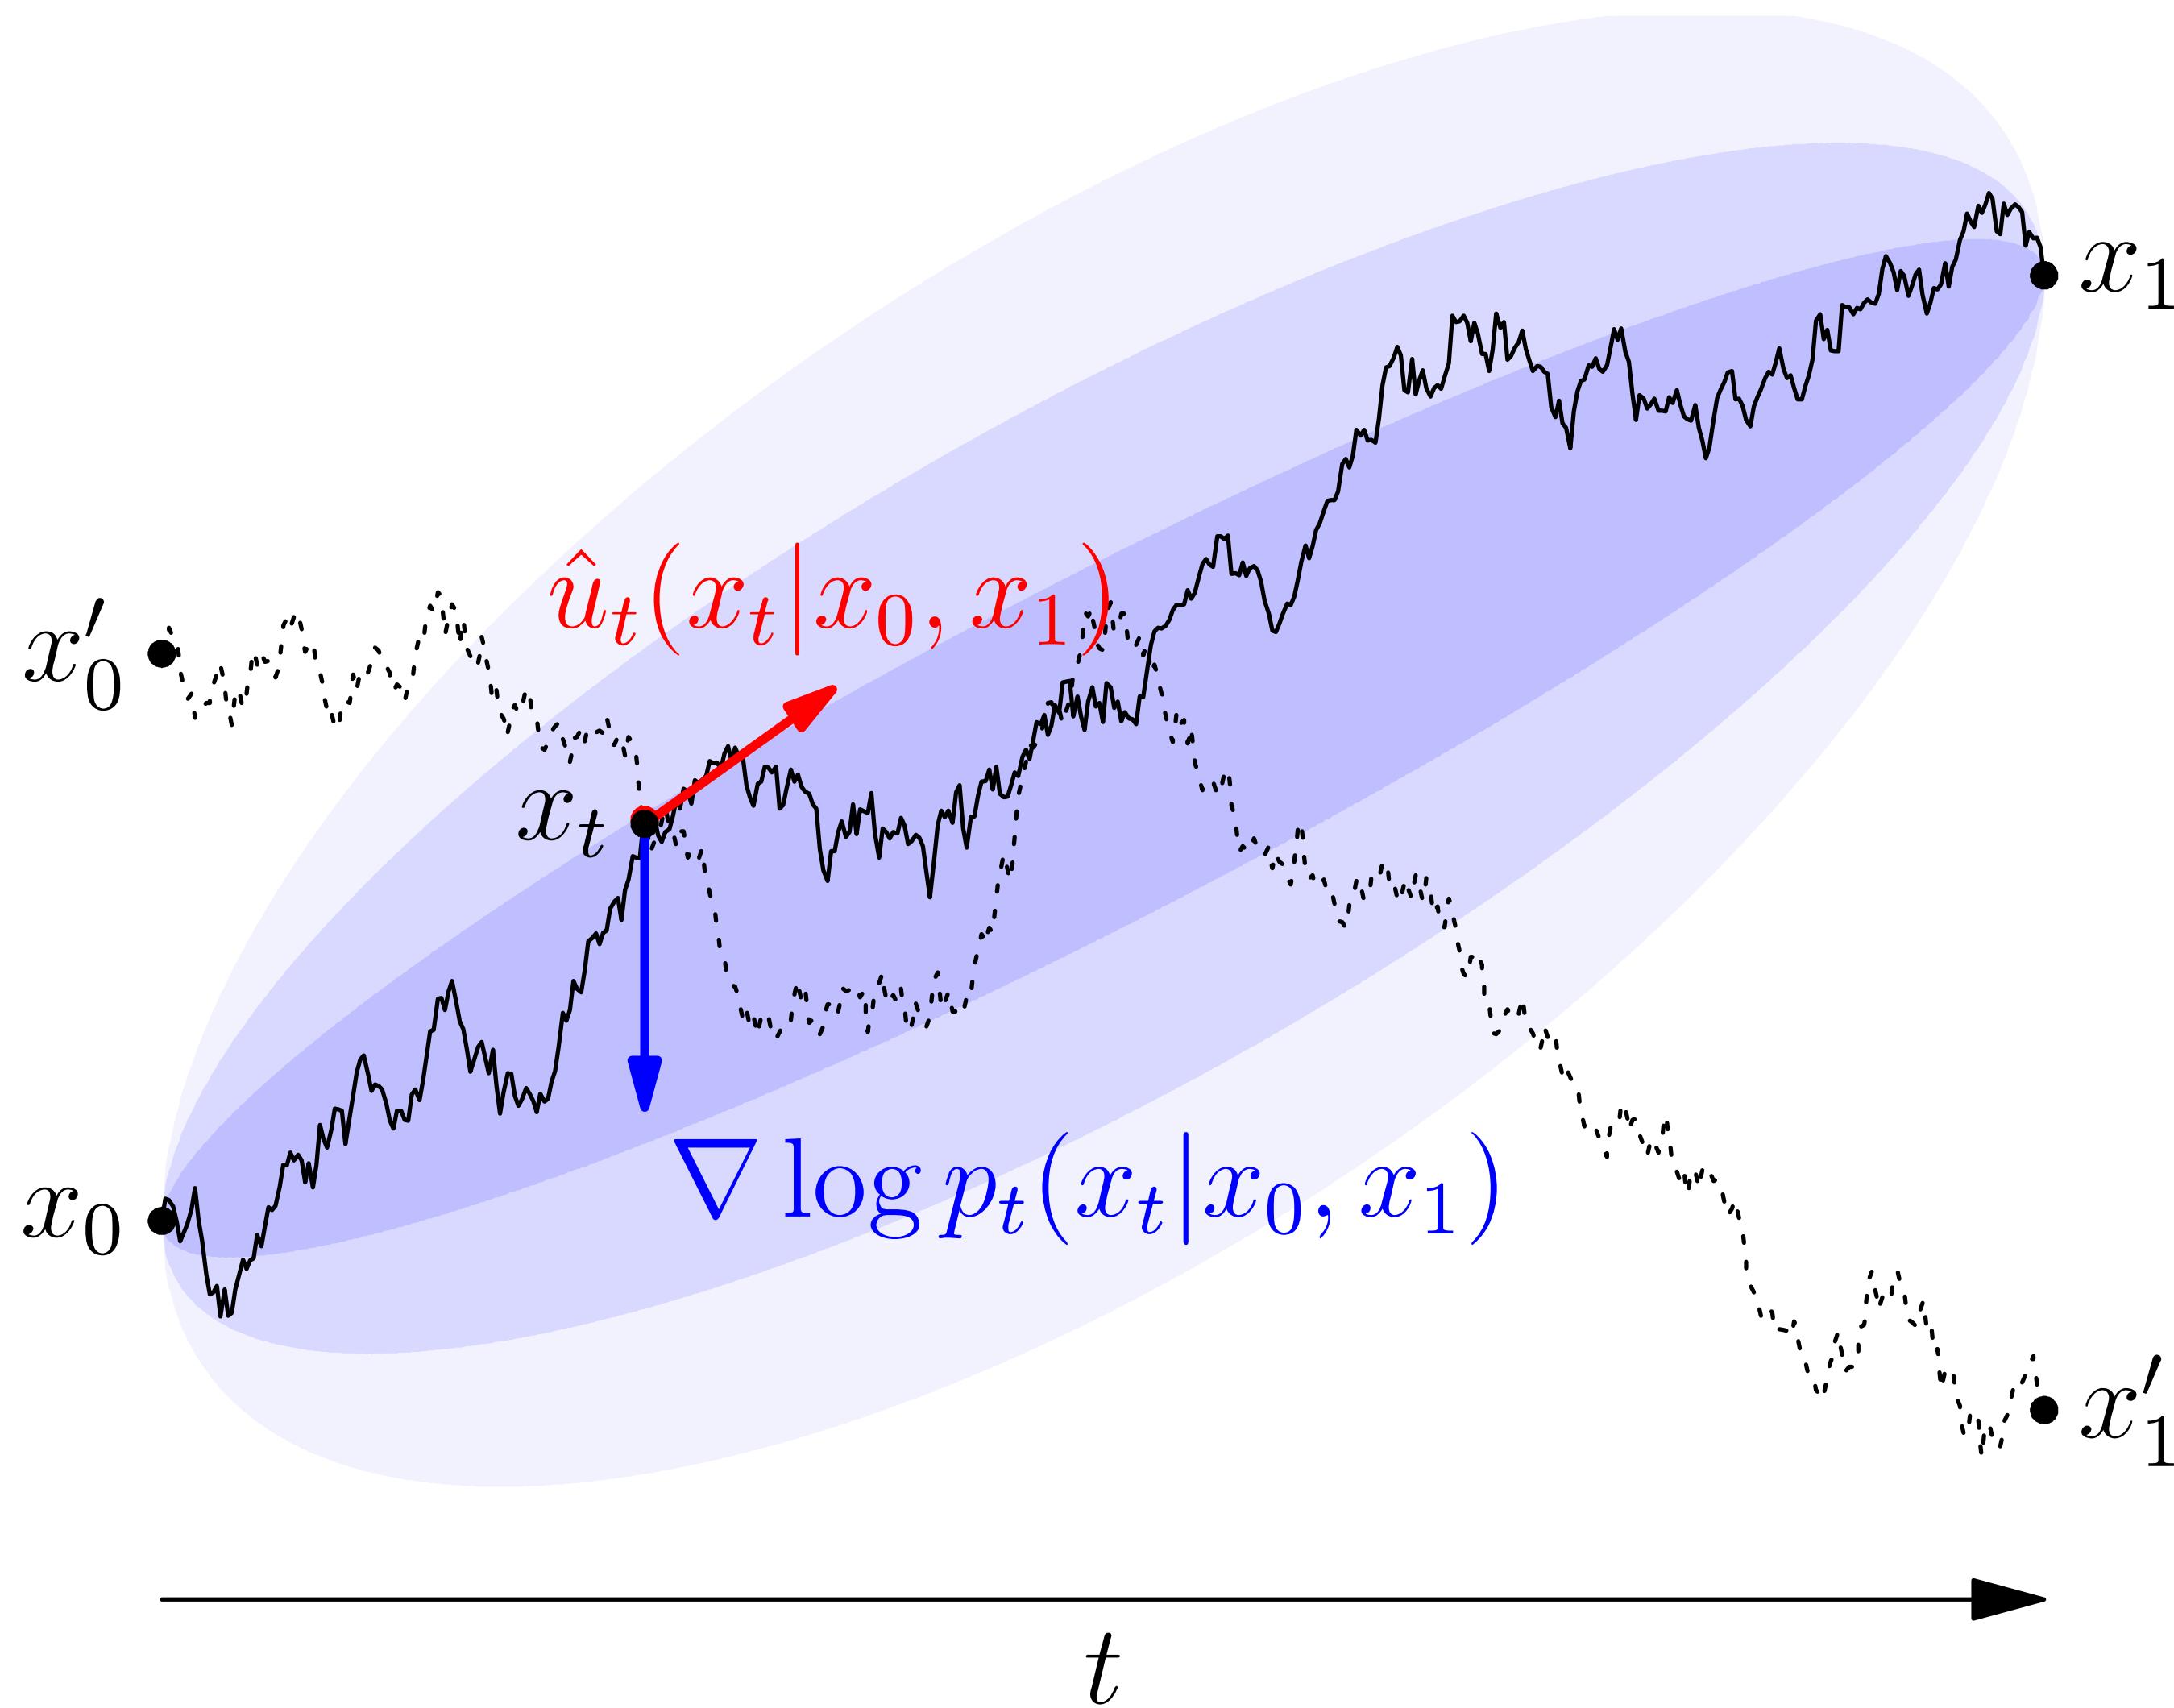
$$dx_t = u_t(x_t)dt \quad \text{Data} \rightarrow \text{Noise}$$

Reverse
Probability
Flow ODE

$$d\bar{x}_t = -u_t(x_t)dt \quad \text{Noise} \rightarrow \text{Data}$$

Score + Flow Matching!

A Geometric Intuition



Schrödinger Bridges

Problem statement

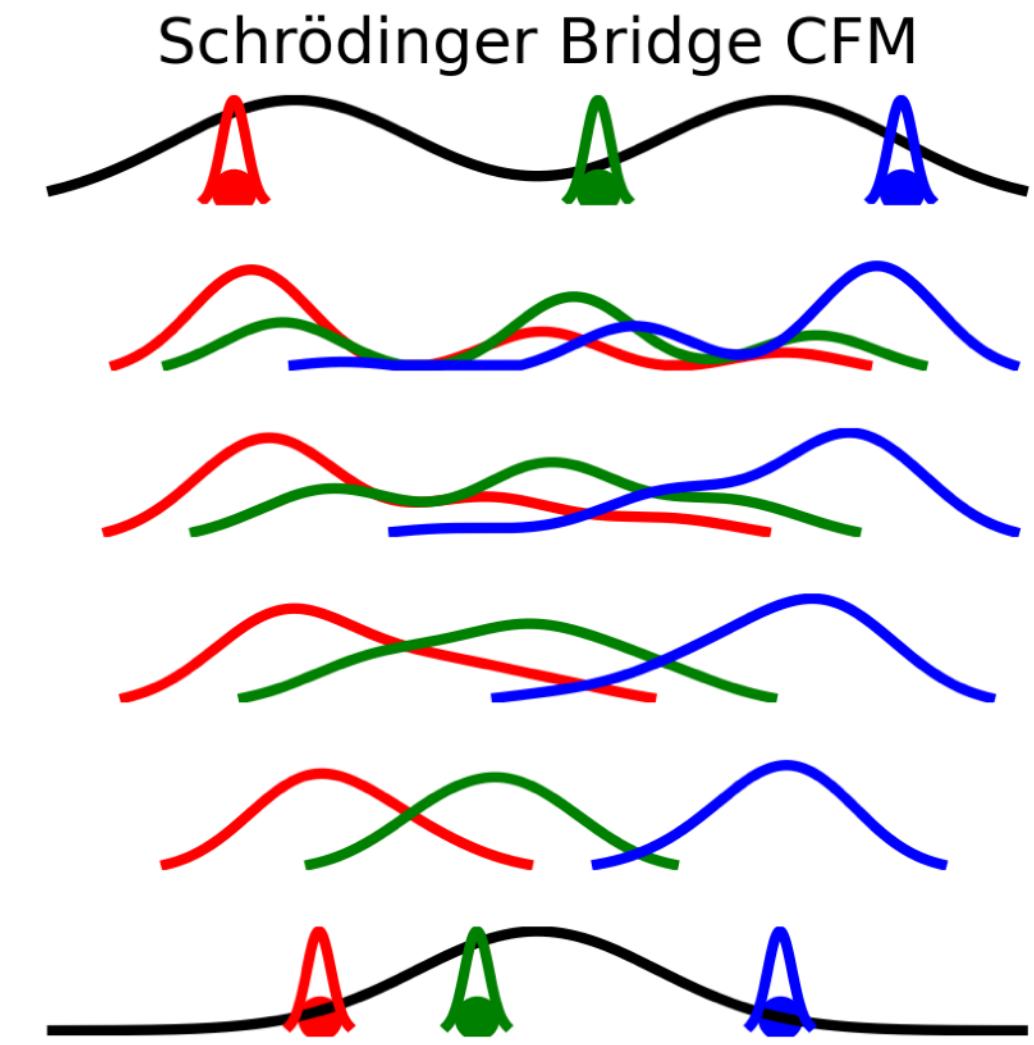
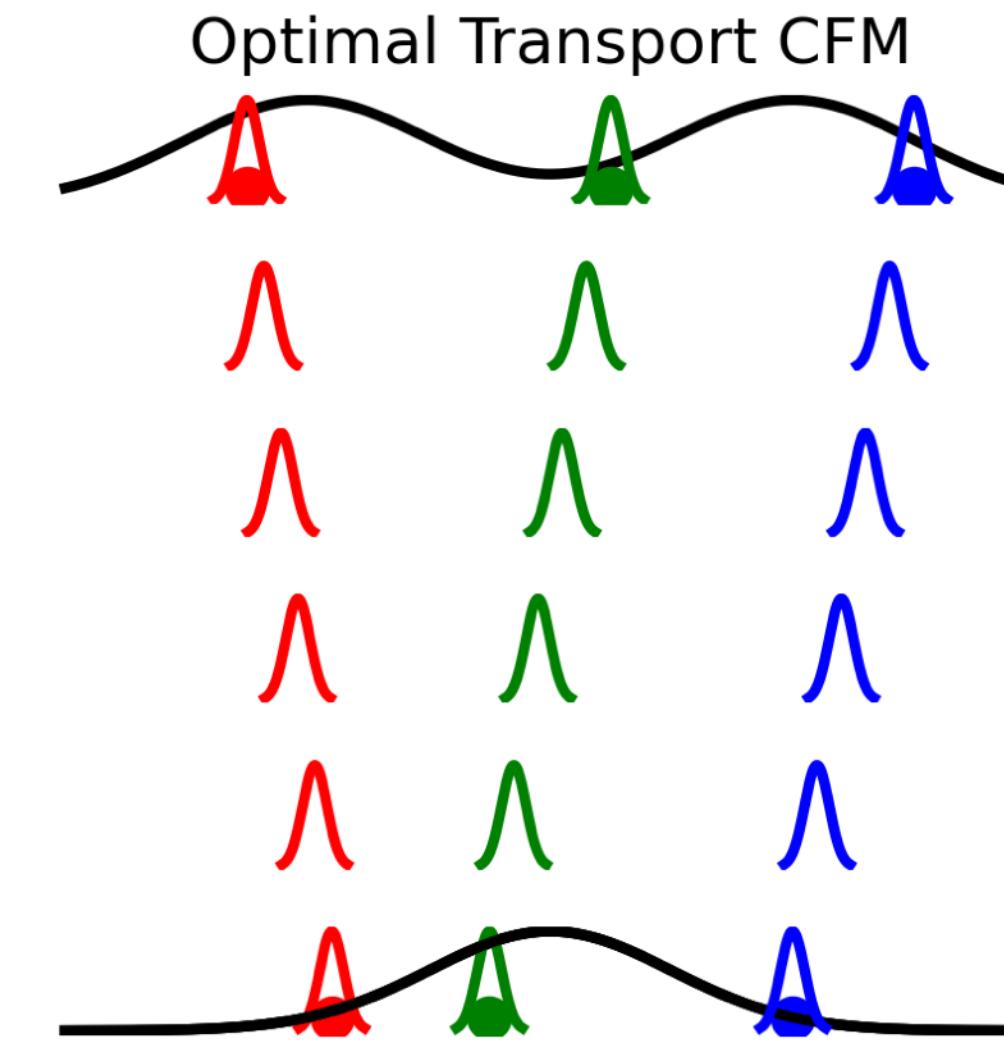
$$\mathbb{P}^* = \min_{\mathbb{P}: p_0 = q_0, p_1 = q_1} \text{KL}(\mathbb{P} \parallel \mathbb{Q})$$

“Most likely stochastic process under observation”

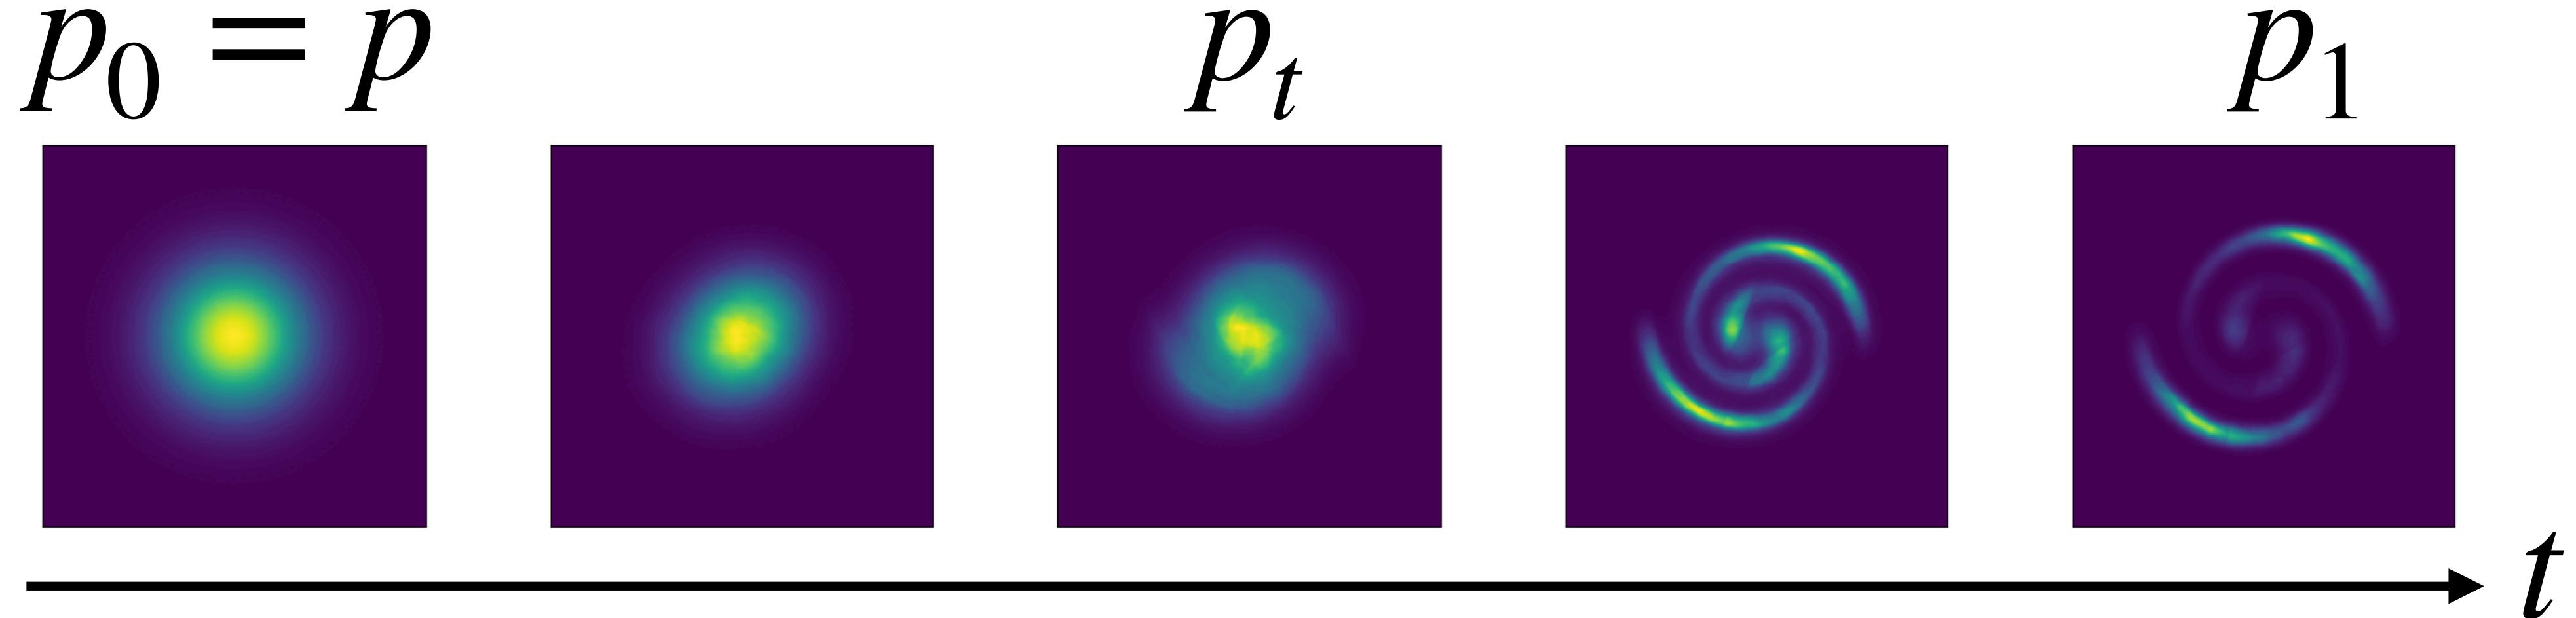
Diffusion Schrödinger Bridges as mixtures of Brownian bridges

$$p_t(x) = \int p_t(x|x_0, x_1) d\pi_{2\sigma^2}^*(x_0, x_1)$$

$$p_t(x|x_0, x_1) = \mathcal{N}(x; (1-t)x_0 + tx_1, \sigma^2 t(1-t))$$



How do we speed up Inference?



Consistency Models

Yang Song¹ Prafulla Dhariwal¹ Mark Chen¹ Ilya Sutskever¹

ONE STEP DIFFUSION VIA SHORTCUT MODELS

Kevin Frans
UC Berkeley
kvfrans@berkeley.edu

Danijar Hafner
UC Berkeley

Sergey Levine
UC Berkeley

Pieter Abbeel
UC Berkeley

Flow Straight and Fast:
Learning to Generate and Transfer Data with Rectified Flow

Xingchao Liu*
University of Texas at Austin
xcliu@utexas.edu

Chengyue Gong*
University of Texas at Austin
cygong@cs.utexas.edu

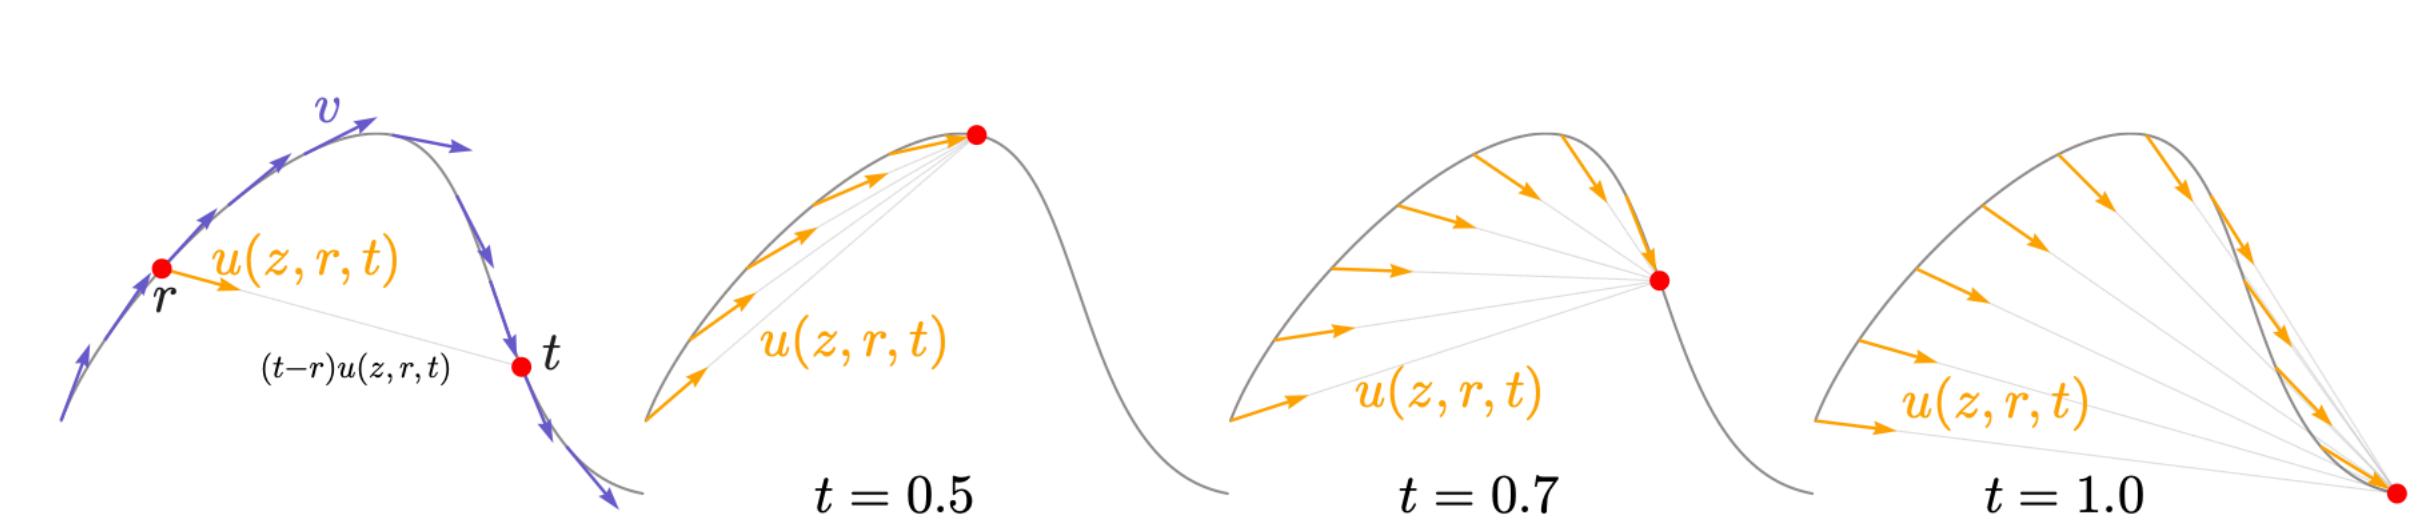
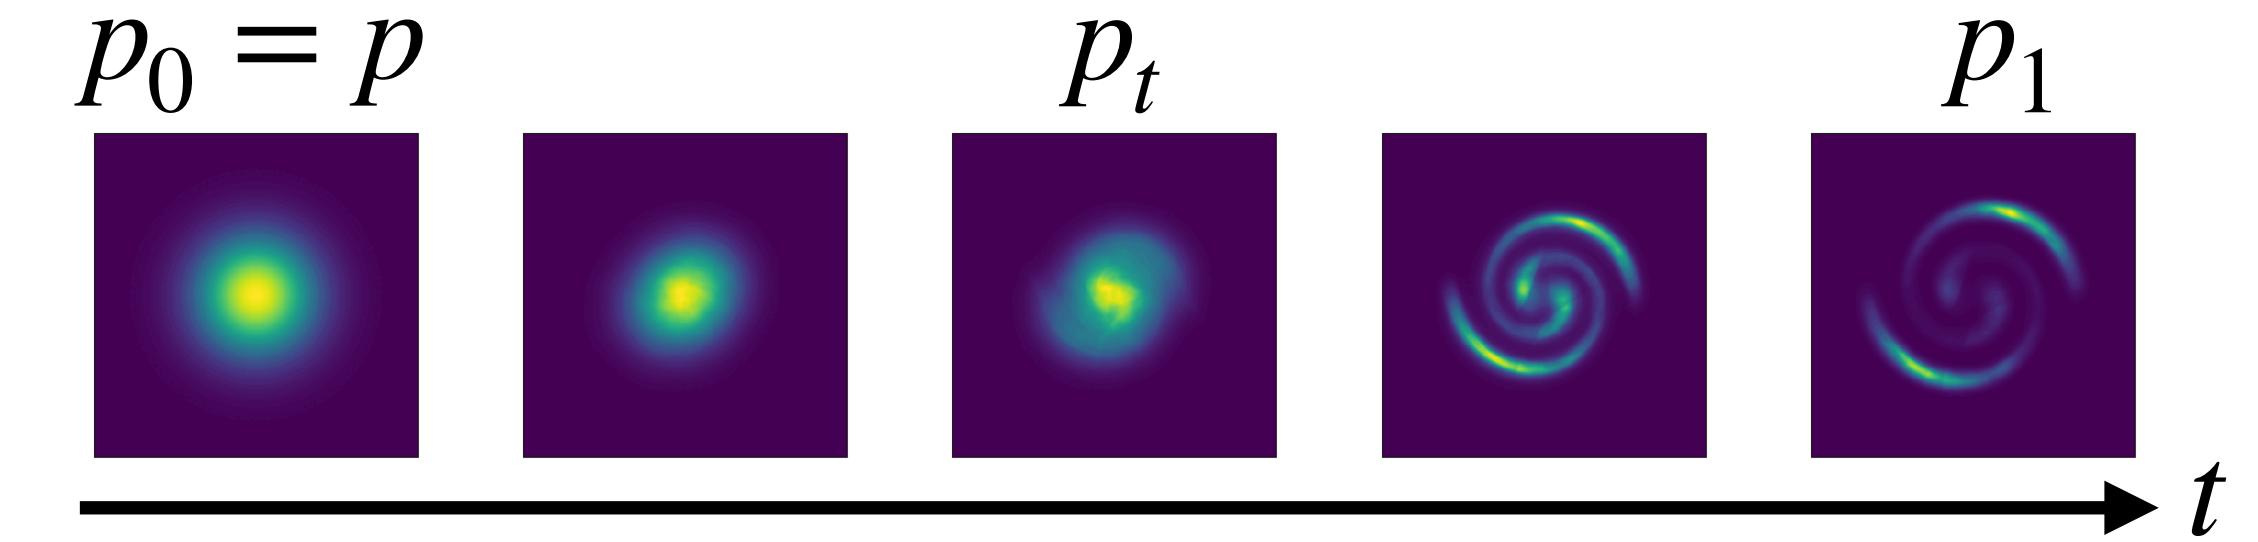
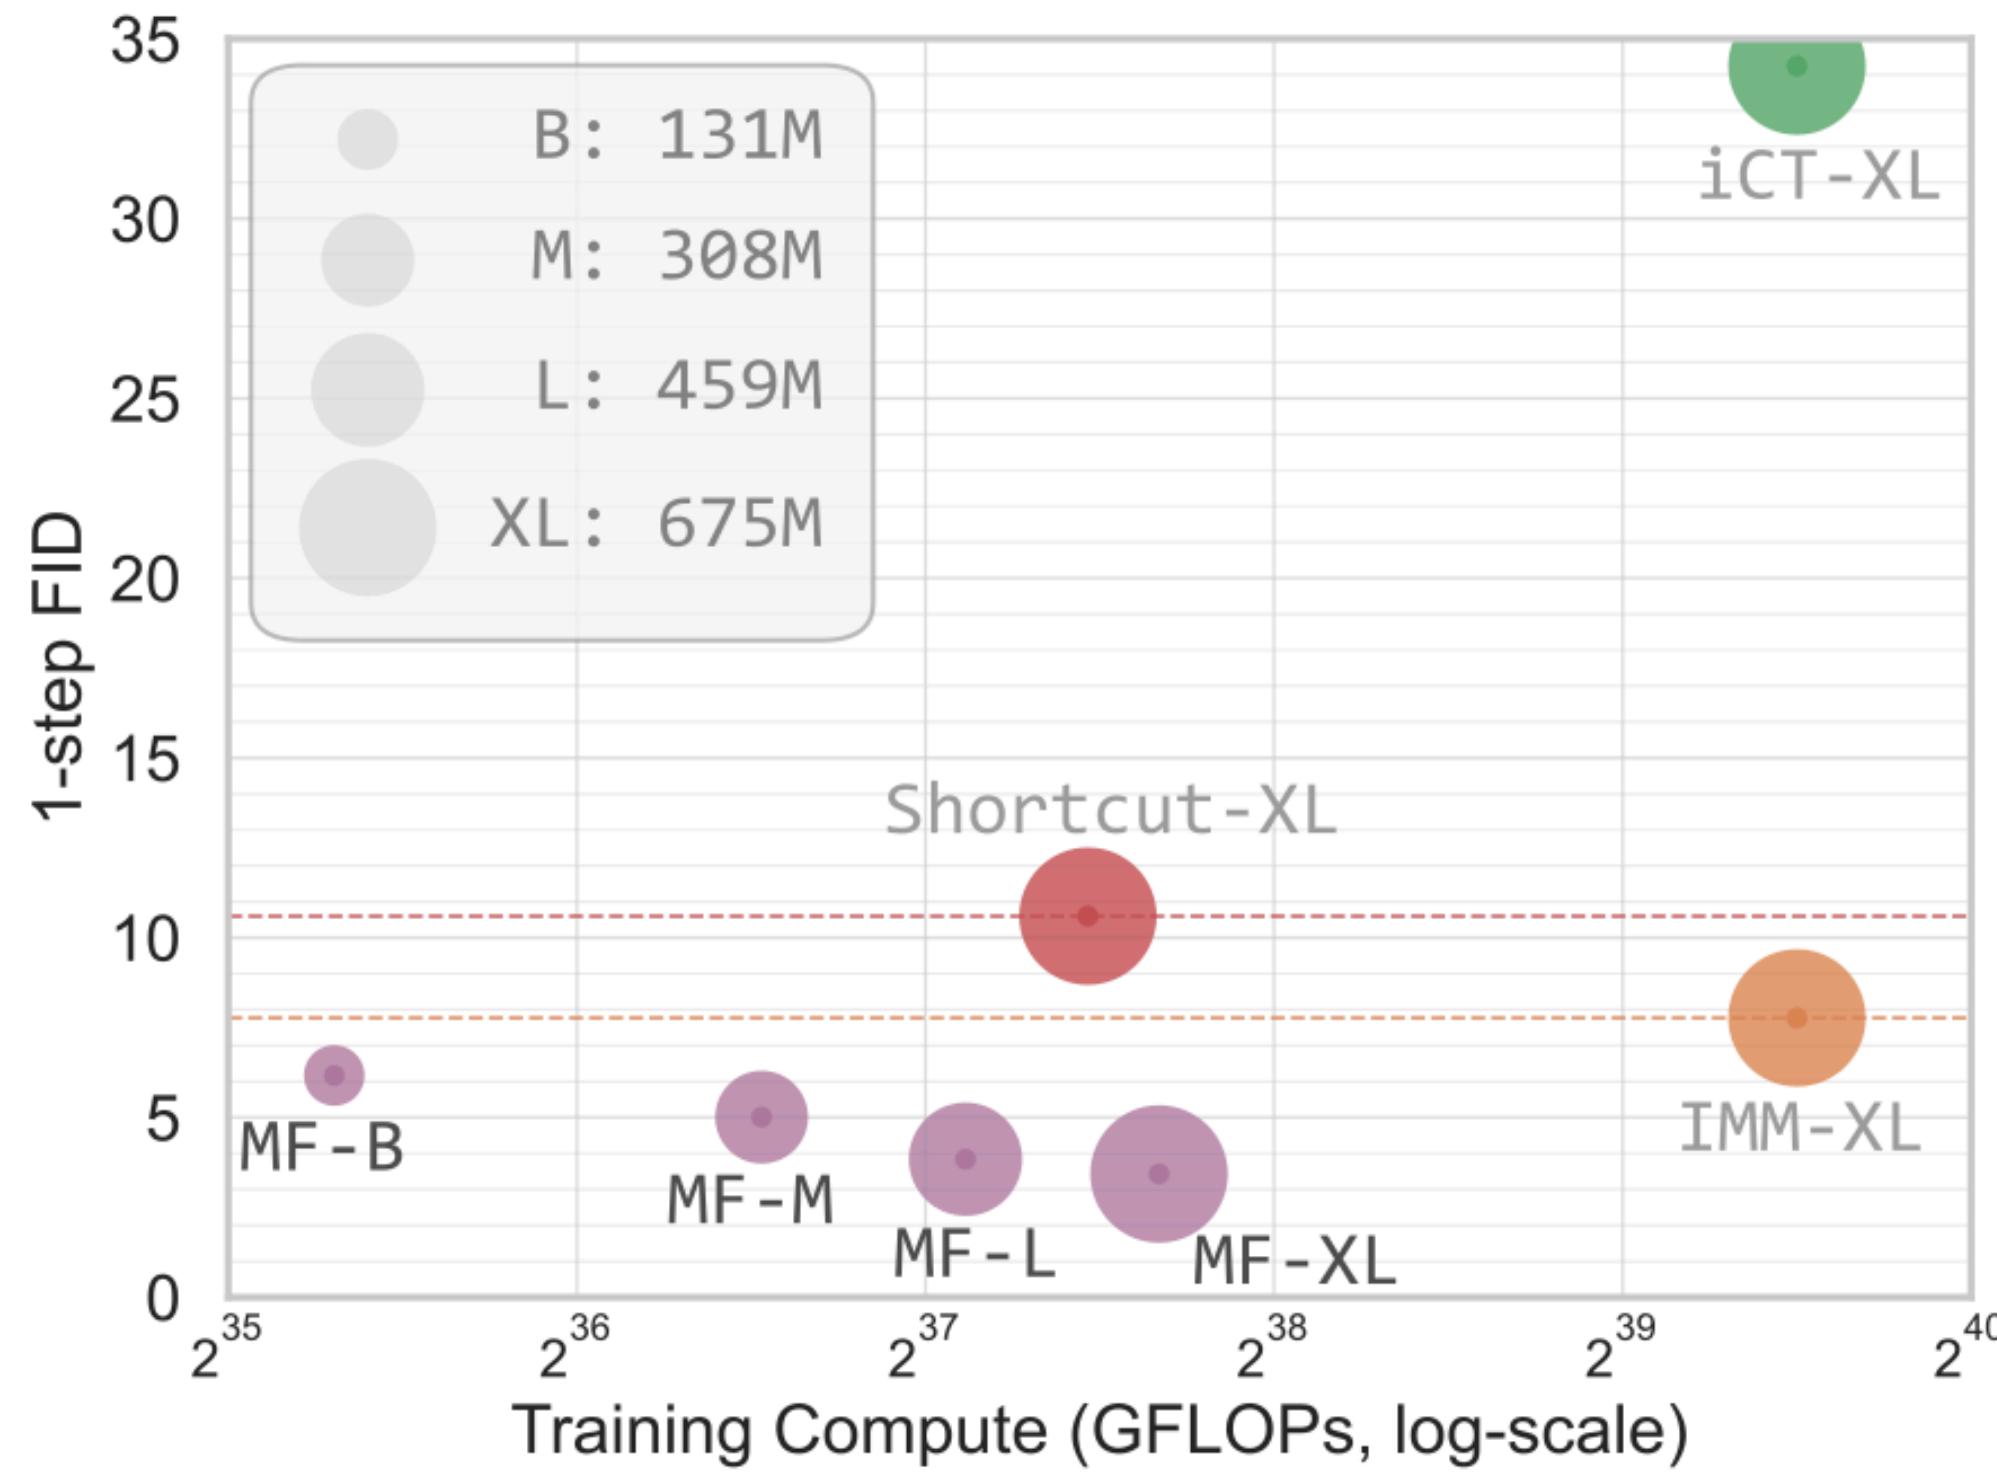
Qiang Liu
University of Texas at Austin
lqiang@cs.utexas.edu

How do we speed up Inference? – Mean Flows

Mean Flows for One-step Generative Modeling

Zhengyang Geng^{1*} Mingyang Deng² Xingjian Bai² J. Zico Kolter¹ Kaiming He²

¹CMU ²MIT



Idea:
**Condition on current
and target time!**

$u_\theta(x_t, t)$
↓
 $u_\theta(x_t, r, t)$

How do we learn $u_\theta(x_t, r, t)$?

$$(t - r)u(x_t, r, t) = \int_r^t v(x_t, \tau) d\tau$$

Derivative w.r.t. t

$$u(x_t, r, t) + (t - r)\frac{d}{dt}u(x_t, r, t) = v(x_t, t)$$

Rearrange

$$u_{target}(x_t, r, t) = v(x_t, t) - (t - r)\frac{d}{dt}u(x_t, r, t)$$

$$L(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x, r, t) - stopgrad(u_{target})\|^2$$

How do we learn $u_\theta(x_t, r, t)$?

$$(t - r)u(x_t, r, t) = \int_r^t v(x_t, \tau) d\tau$$

Derivative w.r.t. t

$$u(x_t, r, t) + (t - r)\frac{d}{dt}u(x_t, r, t) = v(x_t, t)$$

Rearrange

$$u_{target}(x_t, r, t) = v(x_t, t) - (t - r)$$

$$L(\theta) = \min \mathbb{E}_{t, q(z), p_t(x|z)} \|u_t^\theta(x, r, t) - u_{target}(x_t, r, t)\|$$

Algorithm 1 MeanFlow: Training.

Note: in PyTorch and JAX, jvp returns the function output and JVP.

```
# fn(z, r, t): function to predict u
# x: training batch
```

```
t, r = sample_t_r()
e = randn_like(x)
```

```
z = (1 - t) * x + t * e
v = e - x
```

```
u, dudt = jvp(fn, (z, r, t), (v, 0, 1))
```

```
u_tgt = v - (t - r) * dudt
error = u - stopgrad(u_tgt)
```

```
loss = metric(error)
```

Algorithm 2 MeanFlow: 1-step Sampling

```
e = randn(x_shape)
x = e - fn(e, r=0, t=1)
```

Practical Considerations

Three practical considerations

- Model conditioning

EDM Keras et al. 2022

	VP [49]	VE [49]	iDDPM [37] + DDIM [47]	Ours (“EDM”)
Sampling (Section 3)				
ODE solver	Euler	Euler	Euler	2^{nd} order Heun
Time steps	$t_{i < N}$	$1 + \frac{i}{N-1}(\epsilon_s - 1)$	$\sigma_{\max}^2 (\sigma_{\min}^2 / \sigma_{\max}^2)^{\frac{i}{N-1}}$	$u_{\lfloor j_0 + \frac{M-1-j_0}{N-1} i + \frac{1}{2} \rfloor}$, where $u_M = 0$ $u_{j-1} = \sqrt{\frac{u_j^2 + 1}{\max(\bar{\alpha}_{j-1}/\bar{\alpha}_j, C_1)} - 1}$
Schedule	$\sigma(t)$	$\sqrt{e^{\frac{1}{2}\beta_d t^2 + \beta_{\min} t} - 1}$	\sqrt{t}	t
Scaling	$s(t)$	$1 / \sqrt{e^{\frac{1}{2}\beta_d t^2 + \beta_{\min} t}}$	1	1
Network and preconditioning (Section 5)				
Architecture of F_θ	DDPM++	NCSN++	DDPM	(any)
Skip scaling $c_{\text{skip}}(\sigma)$	1	1	1	$\sigma_{\text{data}}^2 / (\sigma^2 + \sigma_{\text{data}}^2)$
Output scaling $c_{\text{out}}(\sigma)$	$-\sigma$	σ	$-\sigma$	$\sigma \cdot \sigma_{\text{data}} / \sqrt{\sigma_{\text{data}}^2 + \sigma^2}$
Input scaling $c_{\text{in}}(\sigma)$	$1 / \sqrt{\sigma^2 + 1}$	1	$1 / \sqrt{\sigma^2 + 1}$	$1 / \sqrt{\sigma^2 + \sigma_{\text{data}}^2}$
Noise cond. $c_{\text{noise}}(\sigma)$	$(M-1) \sigma^{-1}(\sigma)$	$\ln(\frac{1}{2}\sigma)$	$M-1 - \arg \min_j u_j - \sigma $	$\frac{1}{4} \ln(\sigma)$
Training (Section 5)				
Noise distribution	$\sigma^{-1}(\sigma) \sim \mathcal{U}(\epsilon_t, 1)$	$\ln(\sigma) \sim \mathcal{U}(\ln(\sigma_{\min}), \ln(\sigma_{\max}))$	$\sigma = u_j, j \sim \mathcal{U}\{0, M-1\}$	$\ln(\sigma) \sim \mathcal{N}(P_{\text{mean}}, P_{\text{std}}^2)$
Loss weighting	$\lambda(\sigma)$	$1/\sigma^2$	$1/\sigma^2$	$(\sigma^2 + \sigma_{\text{data}}^2) / (\sigma \cdot \sigma_{\text{data}})^2$
Parameters				
	$\beta_d = 19.9, \beta_{\min} = 0.1$	$\sigma_{\min} = 0.02$	$\bar{\alpha}_j = \sin^2(\frac{\pi}{2} \frac{j}{M(C_2+1)})$	$\sigma_{\min} = 0.002, \sigma_{\max} = 80$
	$\epsilon_s = 10^{-3}, \epsilon_t = 10^{-5}$	$\sigma_{\max} = 100$	$C_1 = 0.001, C_2 = 0.008$	$\sigma_{\text{data}} = 0.5, \rho = 7$
	$M = 1000$		$M = 1000, j_0 = 8^\dagger$	$P_{\text{mean}} = -1.2, P_{\text{std}} = 1.2$

* iDDPM also employs a second loss term L_{vlb} † In our tests, $j_0 = 8$ yielded better FID than $j_0 = 0$ used by iDDPM

$\sigma(t)$	$\sqrt{e^{\frac{1}{2}\beta_d t^2 + \beta_{\min} t} - 1}$	\sqrt{t}	t	t
$s(t)$	$1/\sqrt{e^{\frac{1}{2}\beta_d t^2 + \beta_{\min} t}}$	1	1	1

and preconditioning (Section 5)

Component	DDPM++	NCSN++	DDPM	(any)
$c_{\text{skip}}(\sigma)$	1	1	1	$\text{EDM Keras et al. } 2022 + \sigma_{\text{data}}^2$
$c_{\text{out}}(\sigma)$	$-\sigma$	σ	$-\sigma$	$\sigma \cdot \sigma_{\text{data}} / \sqrt{\sigma_{\text{data}}^2 + \sigma^2}$
$c_{\text{in}}(\sigma)$	$1/\sqrt{\sigma^2 + 1}$	1	$1/\sqrt{\sigma^2 + 1}$	$1/\sqrt{\sigma^2 + \sigma_{\text{data}}^2}$
$c_{\text{noise}}(\sigma)$	$(M - 1) \sigma^{-1}(\sigma)$	$\ln(\frac{1}{2}\sigma)$	$M - 1 - \arg \min_j u_j - \sigma $	$\frac{1}{4} \ln(\sigma)$

Section 5)

Parameter	DDPM++	NCSN++	DDPM	(any)
Initial distribution	$\sigma^{-1}(\sigma) \sim \mathcal{U}(\epsilon_t, 1)$	$\ln(\sigma) \sim \mathcal{U}(\ln(\sigma_{\min}), \ln(\sigma_{\max}))$	$\sigma = u_j, \quad j \sim \mathcal{U}\{0, M-1\}$	$\ln(\sigma) \sim \mathcal{N}(P_{\text{mean}}, P_{\text{std}}^2)$
Scaling	$1/\sigma^2$	$1/\sigma^2$	$1/\sigma^2$ (note: *)	$(\sigma^2 + \sigma_{\text{data}}^2) / (\sigma \cdot \sigma_{\text{data}})^2$
Hyperparameters	$\beta_d = 19.9, \beta_{\min} = 0.1$ $\epsilon_s = 10^{-3}, \epsilon_t = 10^{-5}$ $M = 1000$	$\sigma_{\min} = 0.02$ $\sigma_{\max} = 100$	$\bar{\alpha}_j = \sin^2(\frac{\pi}{2} \frac{j}{M(C_2+1)})$ $C_1 = 0.001, C_2 = 0.008$ $M = 1000, j_0 = 8^\dagger$	$\sigma_{\min} = 0.002, \sigma_{\max} = 80$ $\sigma_{\text{data}} = 0.5, \rho = 7$ $P_{\text{mean}} = -1.2, P_{\text{std}} = 1.2$

DDPM also employs a second loss term L_{vlb}

^{*} In our tests, $j_0 = 8$ yielded better FID than $j_0 = 0$ used by iDDPM

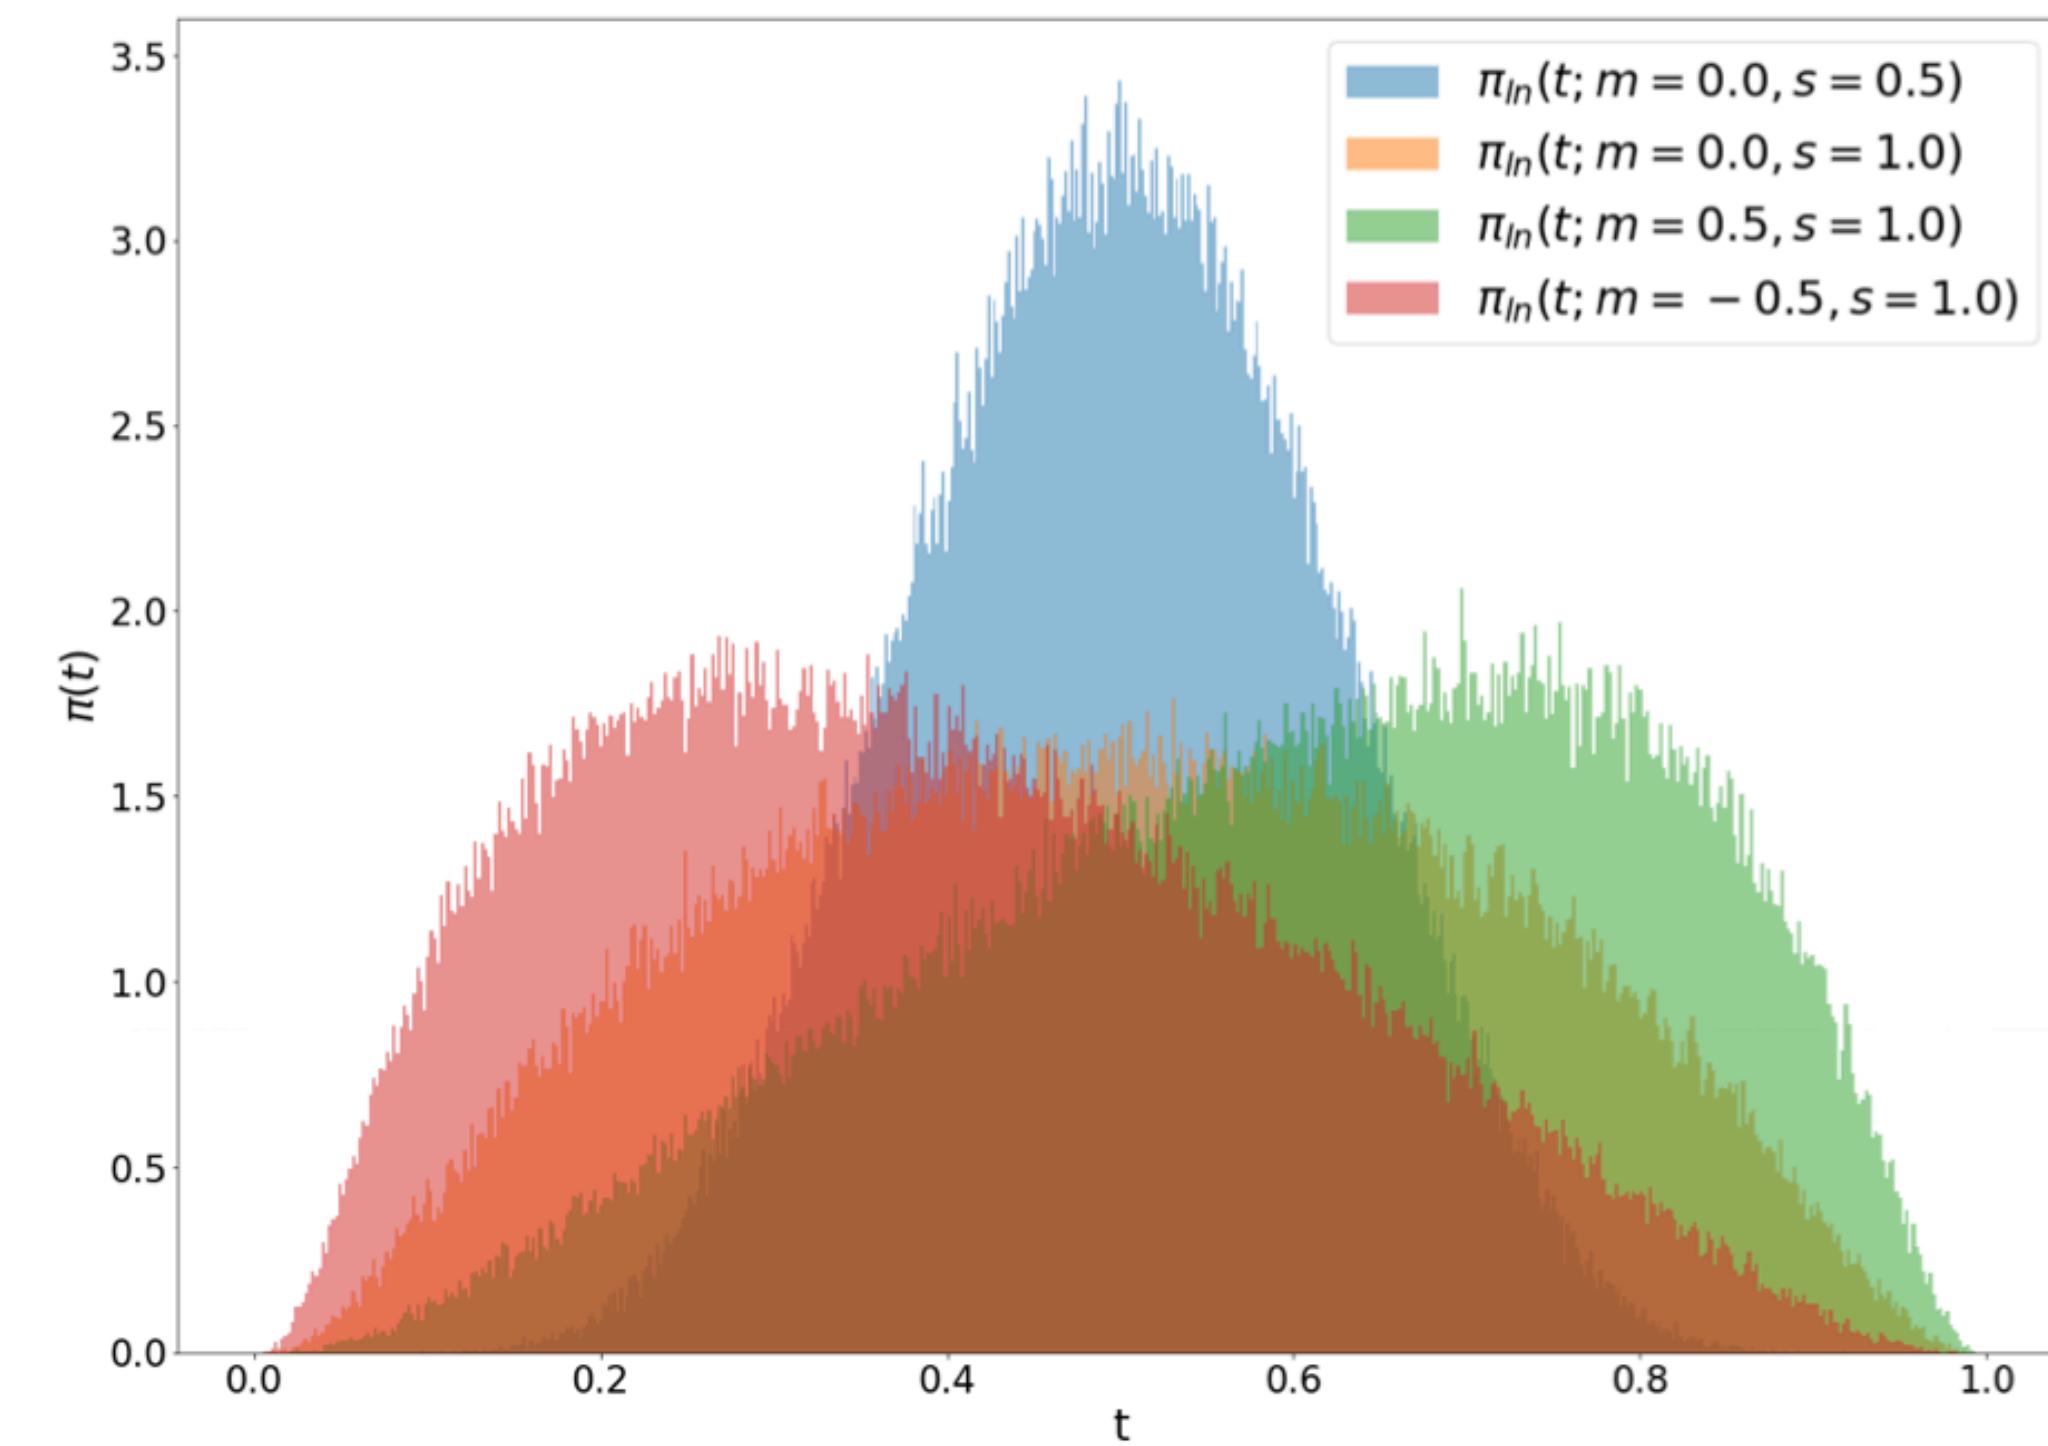
Actually is equivalent to a flow matching model!

Practical Considerations

Three practical considerations

- Model conditioning
- Sampling t

Stable diffusion 3 found orange was best



Practical Considerations

Three practical considerations

- Model conditioning
- Sampling t
- Exponential Moving Average on weights

$$\mathbf{x}_0^{\text{EMA}} = \mathbf{x}_0, \text{ and } \mathbf{x}_{t+1}^{\text{EMA}} = \alpha \mathbf{x}_t^{\text{EMA}} + (1 - \alpha) \mathbf{x}_{t+1}.$$

Practical Considerations

Three practical considerations

- Model conditioning
- Sampling t
- Exponential Moving Average on weights

Two unintuitive things / open questions

- Gaussian prior works way too well
- OT doesn't seem to help beyond a certain batch size

In theory we would think choosing our prior p_0 close to p_1 would be great. In practice $p_0 = \mathcal{N}(0,1)$ is almost always better

Practical Considerations

Three practical considerations

- Model conditioning
- Sampling t
- Exponential Moving Average on weights

Two unintuitive things / open questions

- Gaussian prior works way too well
- OT doesn't seem to help beyond a certain batch size

SIMPLE REFLOW: IMPROVED TECHNIQUES FOR FAST FLOW MODELS

Beomsu Kim
Apple and KAIST

Yu-Guan Hsieh
Apple

Michal Klein
Apple

Marco Cuturi
Apple

Jong Chul Ye
KAIST

Bahjat Kawar
Apple

James Thornton
Apple

Tested batch sizes up to 50k,
didn't see improvement beyond
2-4k.

Method	CIFAR10			AFHQv2			FFHQ			ImageNet (cond.)			Reference
	NFE	FID	STN	NFE	FID	STN	NFE	FID	STN	NFE	FID	STN	
DM ODE													
EDM	35 9	1.97 37.91	14.19 —	79 9	1.96 28.03	28.41 —	79 9	2.39 56.84	27.15 —	79 9	2.30 35.46	26.76 —	(Karras et al., 2022)
DPM-Solver	9	4.98	—	—	—	—	9	9.26	—	9	6.64	—	(Lu et al., 2022)
AMED-Solver	9	2.63	—	—	—	—	9	4.24	—	9	5.60	—	(Zhou et al., 2024)
FM ODE													
MinCurv	9	8.76	5.87	9	13.63	10.45	9	10.44	10.49	—	—	—	(Lee et al., 2023)
FM-OT	142	6.35	—	—	—	—	—	—	—	138	14.45	—	(Lipman et al., 2023)
OT-CFM	100	4.44	—	—	—	—	—	—	—	—	—	—	(Tong et al., 2023)
MOT-50	—	—	—	—	—	—	—	—	—	132	11.82	—	(Pooladian et al., 2023)
FM*	100	2.96	10.73	100	2.73	16.20	100	3.30	16.71	—	—	—	Baseline
MOT-512*	100	3.29	8.77	100	5.53	13.45	100	4.69	14.29	—	—	—	—
MOT-1024*	100	3.18	8.59	100	5.83	13.45	100	4.84	14.07	—	—	—	—
MOT-4096*	100	3.16	8.34	100	6.18	12.68	100	4.92	13.47	—	—	—	—
ReFlow	110	3.36	—	—	—	—	—	—	—	—	—	—	(Liu et al., 2022)
Simple ReFlow*	9	2.23	1.64	9	2.30	3.30	9	2.84	2.87	9	3.49	2.72	Ours
+ Guidance*	9	1.98	2.49	9	1.91	5.60	9	2.67	3.24	9	1.74	3.92	—

Additional Resources

TorchCFM: a Conditional Flow Matching library

[paper arxiv.2302.00482](#) [paper arxiv.2307.03672](#) PyTorch 1.8+ Lightning 1.6+

Config Hydra 1.2 Code Style Black Pre-commit enabled TorchCFM Tests passing

codecov 34% Code Quality Main failing License MIT Lightning-Hydra-Template

downloads 128k downloads/month 14k

Flow Matching Guide and Code

Yaron Lipman¹, Marton Havasi¹, Peter Holderith², Neta Shaul³, Matt Le¹, Brian Karrer¹, Ricky T. Q. Chen¹, David Lopez-Paz¹, Heli Ben-Hamu³, Itai Gat¹

¹FAIR at Meta, ²MIT CSAIL, ³Weizmann Institute of Science

Flow Matching (FM) is a recent framework for generative modeling that has achieved state-of-the-art performance across various domains, including image, video, audio, speech, and biological structures. This guide offers a comprehensive and self-contained review of FM, covering its mathematical foundations, design choices, and extensions. By also providing a PyTorch package featuring relevant examples (*e.g.*, image and text generation), this work aims to serve as a resource for both novice and experienced researchers interested in understanding, applying and further developing FM.

Date: December 10, 2024

Code: flow_matching library at https://github.com/facebookresearch/flow_matching



Elucidating the Design Space of Diffusion-Based Generative Models

Tero Karras
NVIDIA

Miika Aittala
NVIDIA

Timo Aila
NVIDIA

Samuli Laine
NVIDIA

ICLR Blogposts 2025

A Visual Dive into Conditional Flow Matching

Conditional flow matching (CFM) was introduced by three simultaneous papers at ICLR 2023, through different approaches (conditional matching, rectifying flows and stochastic interpolants).

The main part of this post, Section 2, explains CFM by using both visual intuitions and insights on its probabilistic formulations. Section 1 introduces normalizing flows; it can be skipped by reader familiar with the topic, or that wants to cover them later. Section 3 opens on the links between CFM and other approaches, and ends with a 'CFM playground'.

AUTHORS
Anne Gagneux
Sérgolène Martin

Rémi Emonet

Quentin Bertrand

Mathurin Massias

AFFILIATIONS
INRIA, ENS de Lyon
Technische Universität Berlin
Inria, Université Jean Monnet
Inria, Université Jean Monnet
Inria, ENS de Lyon

PUBLISHED
April 28, 2025

Thank you and Questions?

Flow/Diffusion Models

